

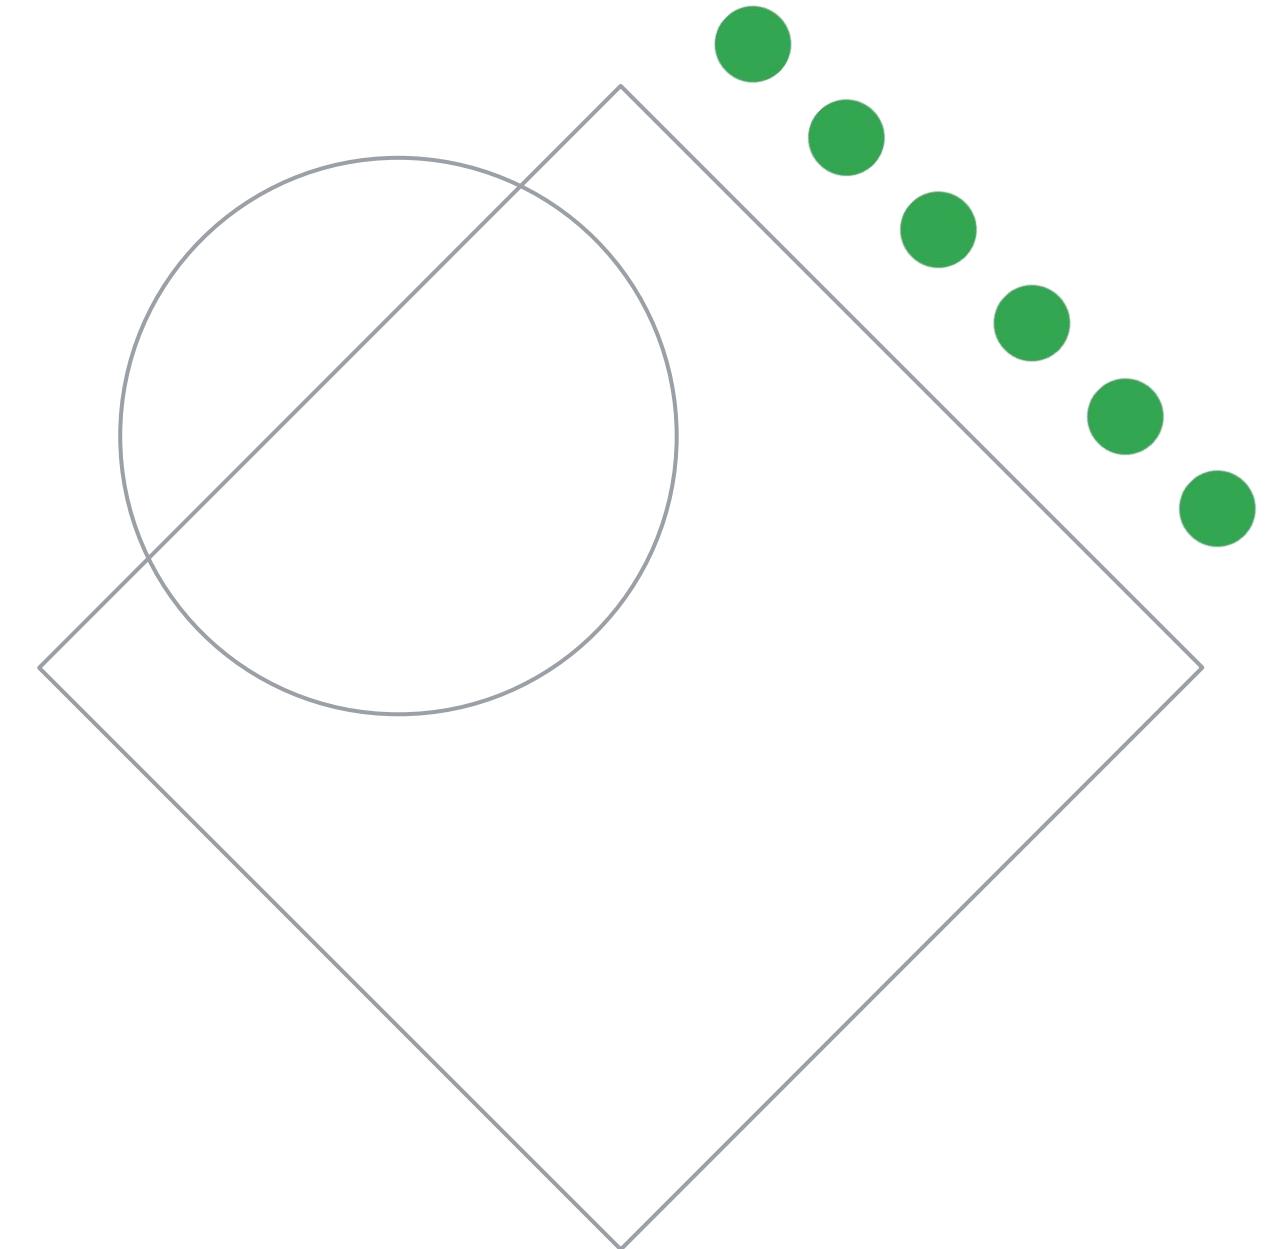
TensorFlow on Google Cloud

In this course, you learn to ...

- 01 Create TensorFlow 2.x and Keras machine learning models
- 02 Describe Tensorflow 2.x key components
- 03 Use the tf.data library to manipulate data and large datasets
- 04 Use the Keras Sequential and Functional APIs for simple and advanced model creation
- 05 Train, deploy, and productionalize ML models at scale with Vertex AI



Introduction to the TensorFlow Ecosystem



In this module, you learn to ...

01

Recall the TensorFlow API hierarchy

02

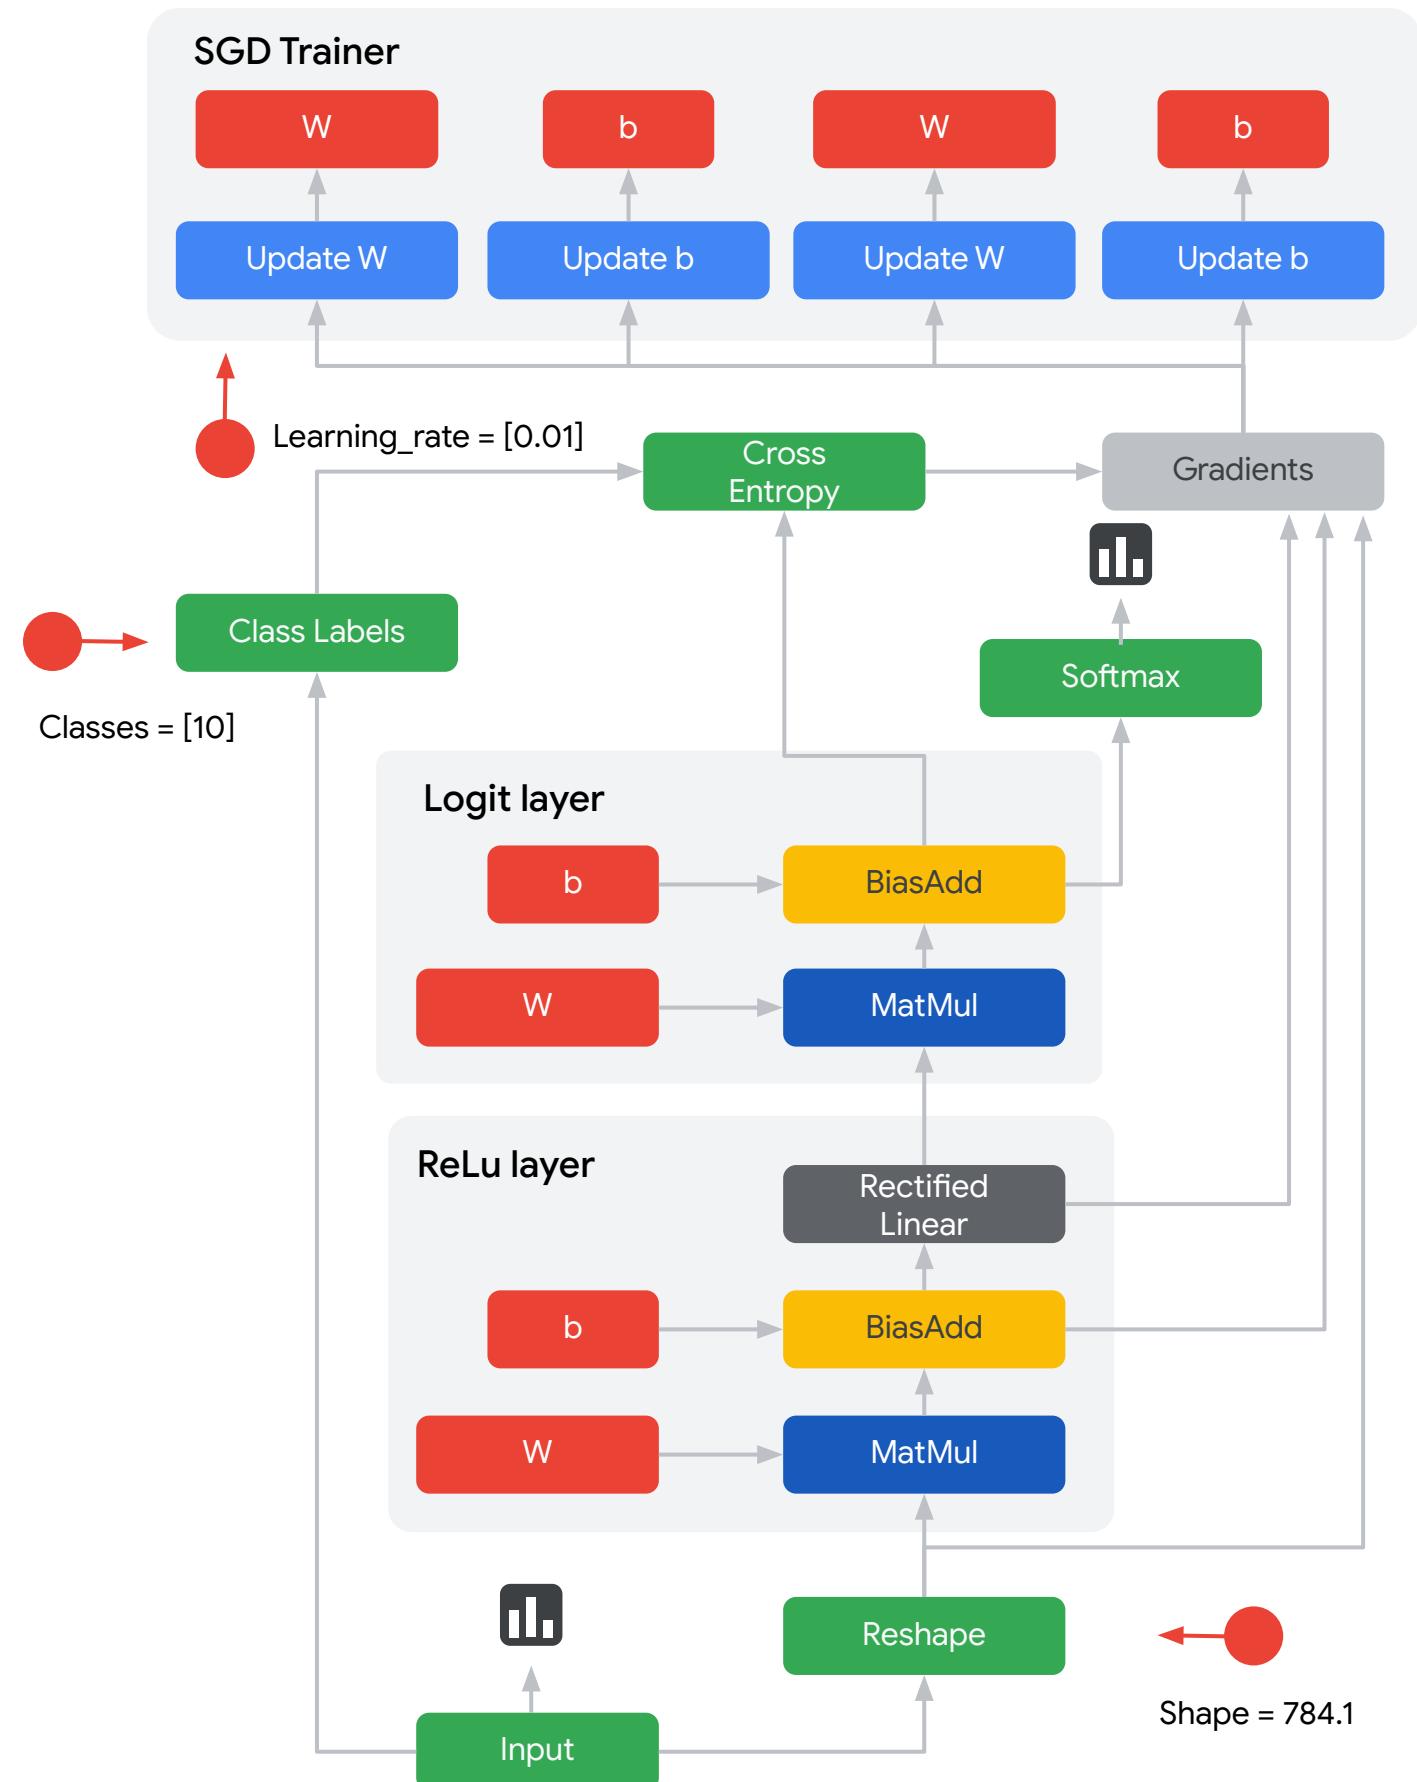
Understand TensorFlow's building blocks:
Tensors and operations

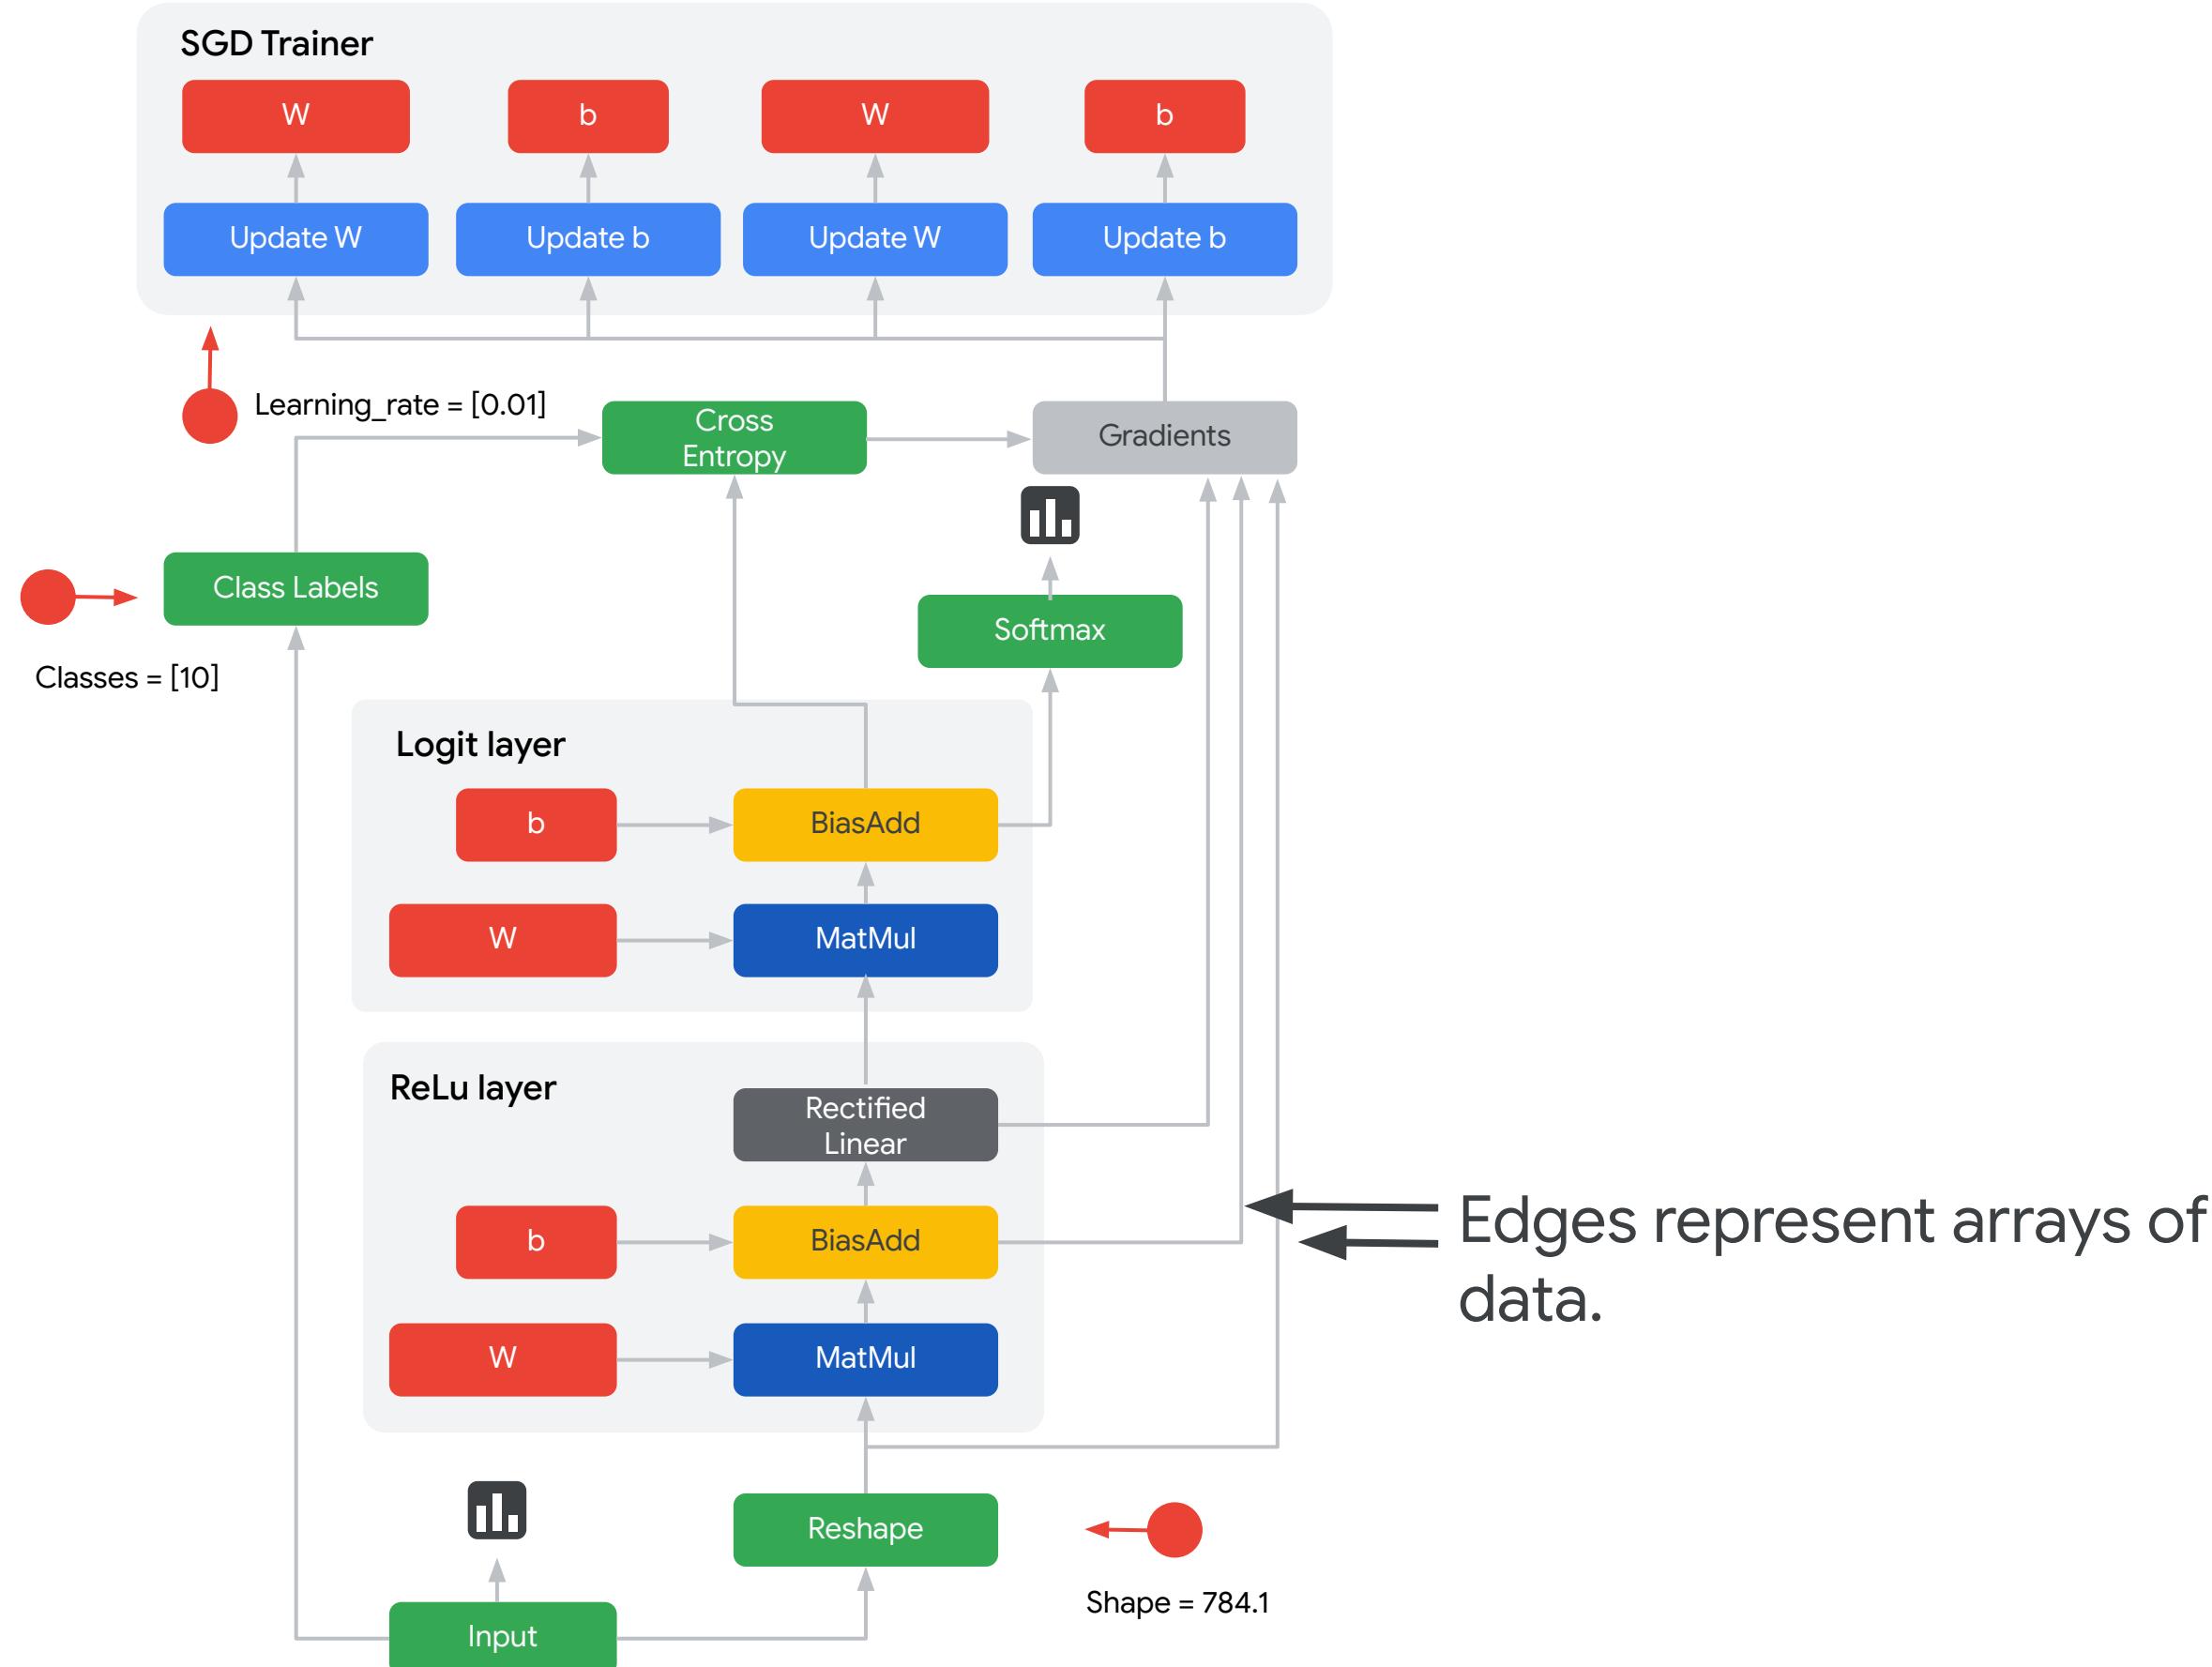
03

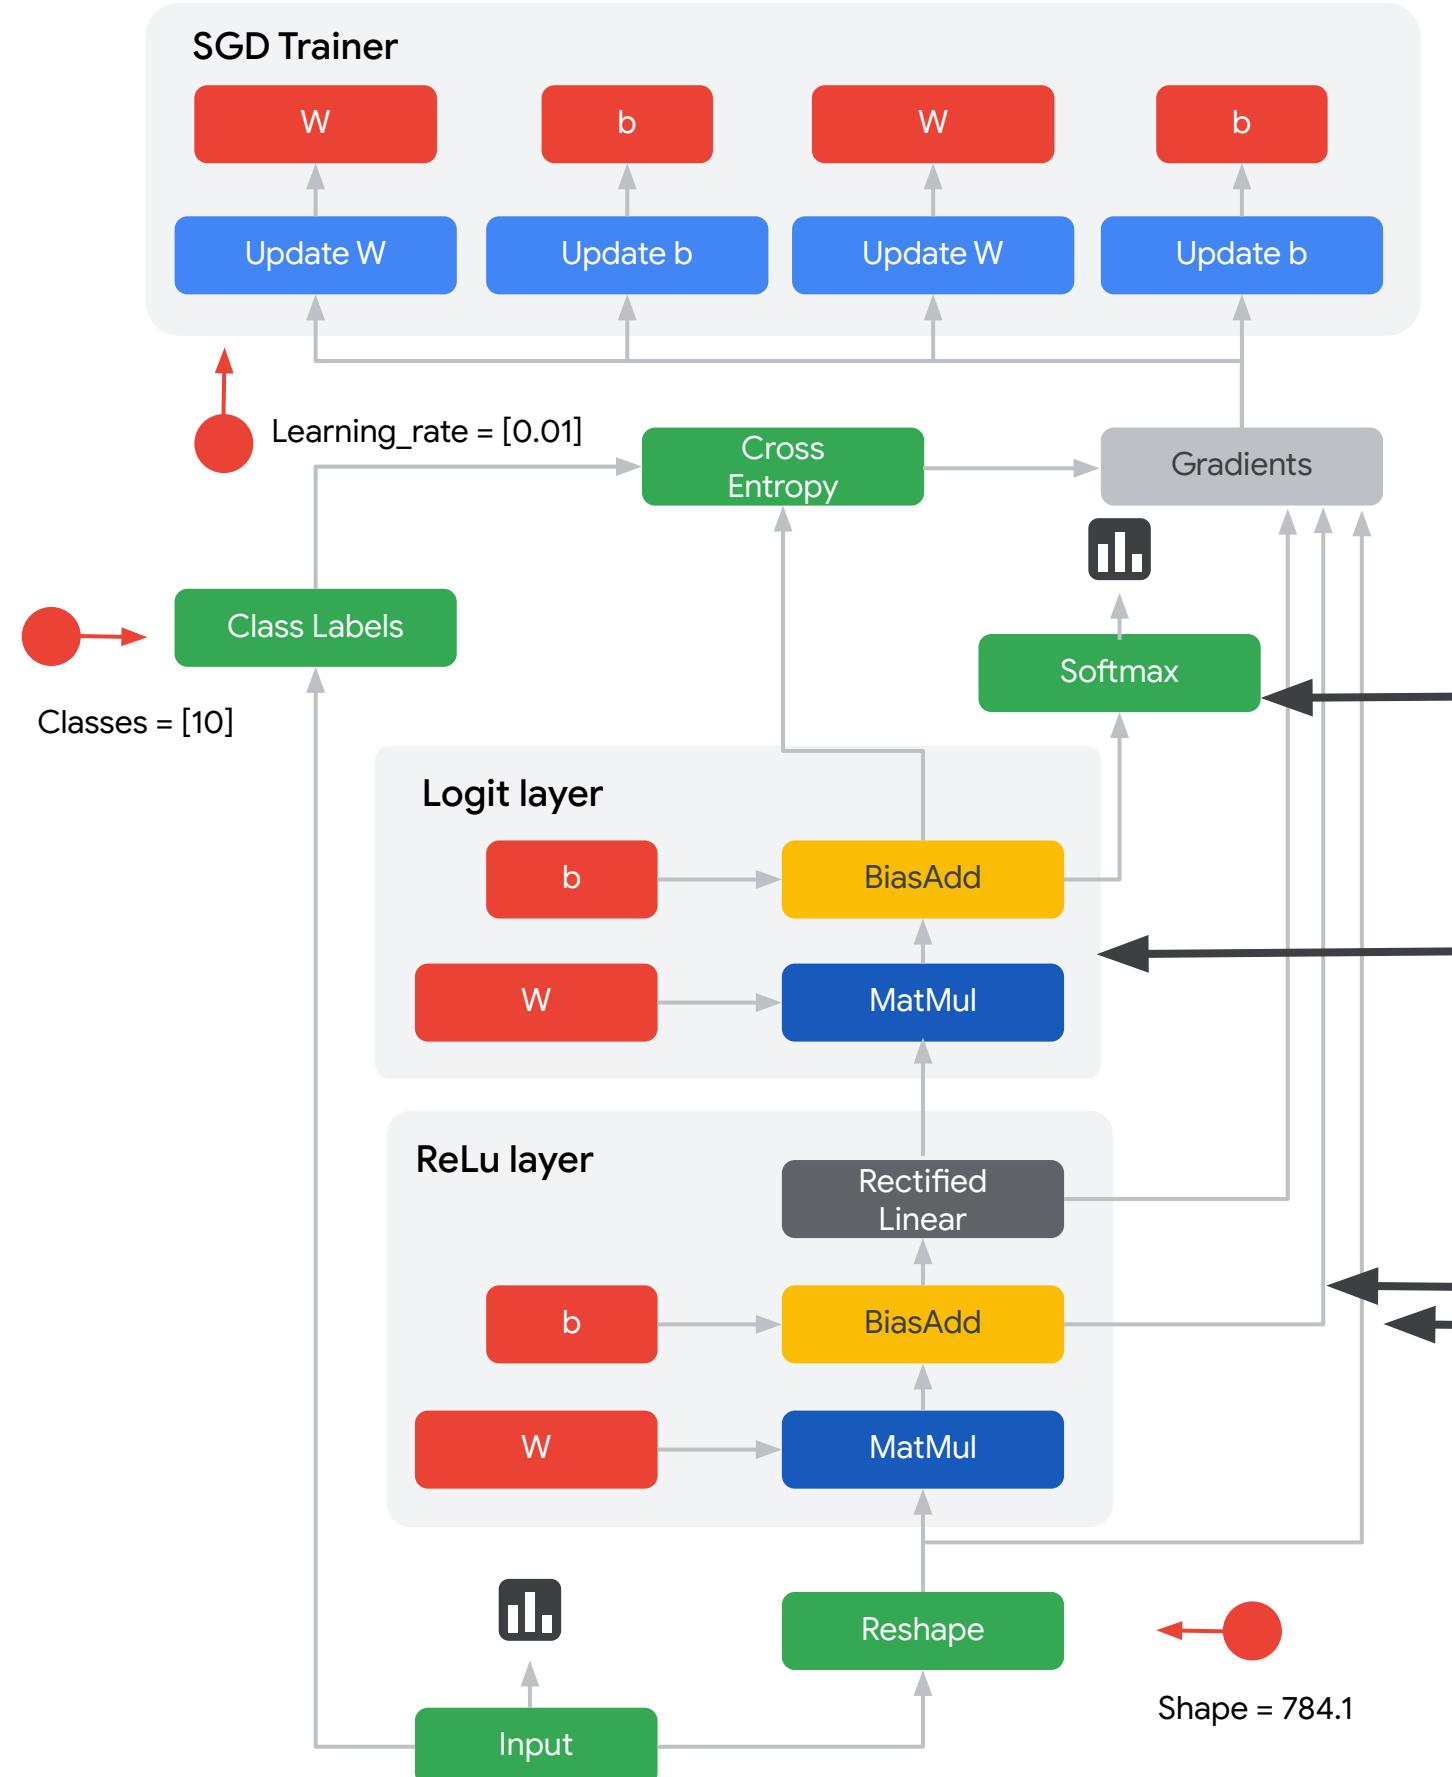
Write low-level TensorFlow programs



**TensorFlow is an open-source,
high-performance library for
numerical computation that
uses directed graphs**







Nodes represent mathematical operations.

Edges represent arrays of data.

A tensor is an N-dimensional array of data



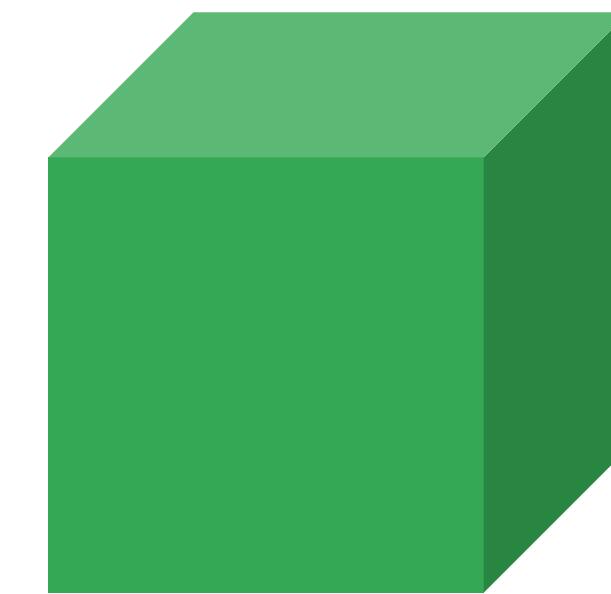
Rank 0
Tensor



Rank 1
Tensor



Rank 2
Tensor



Rank 3
Tensor



Rank 4
Tensor

A tensor is an N-dimensional array of data



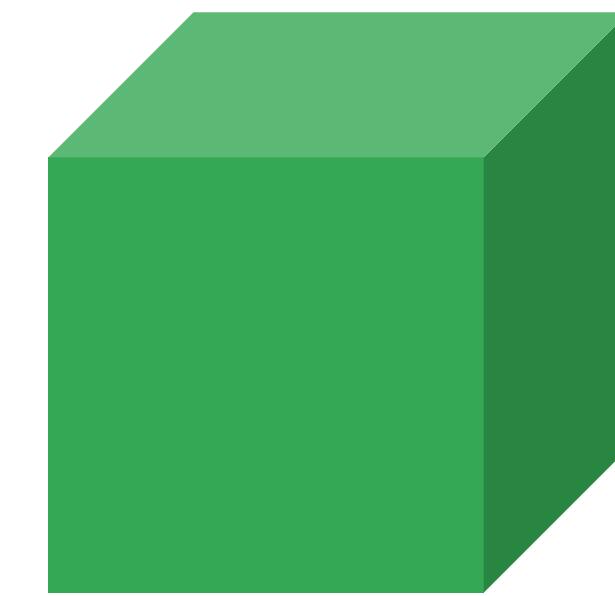
Rank 0
Tensor
scalar



Rank 1
Tensor



Rank 2
Tensor



Rank 3
Tensor



Rank 4
Tensor

A tensor is an N-dimensional array of data



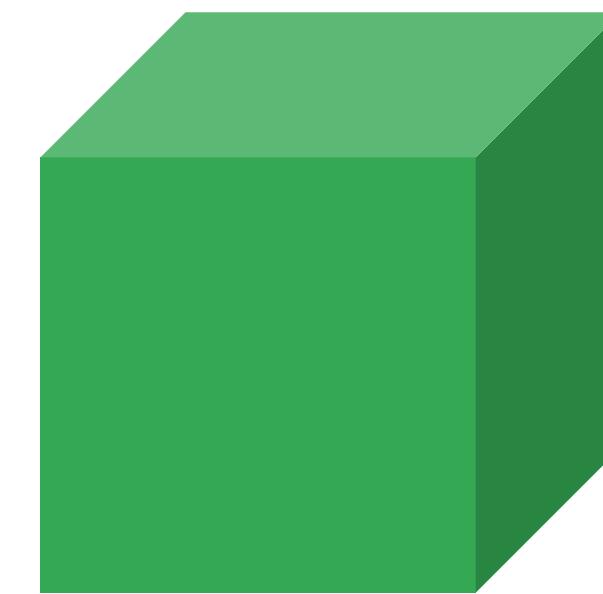
Rank 0
Tensor
scalar



Rank 1
Tensor
vector



Rank 2
Tensor



Rank 3
Tensor



Rank 4
Tensor

A tensor is an N-dimensional array of data



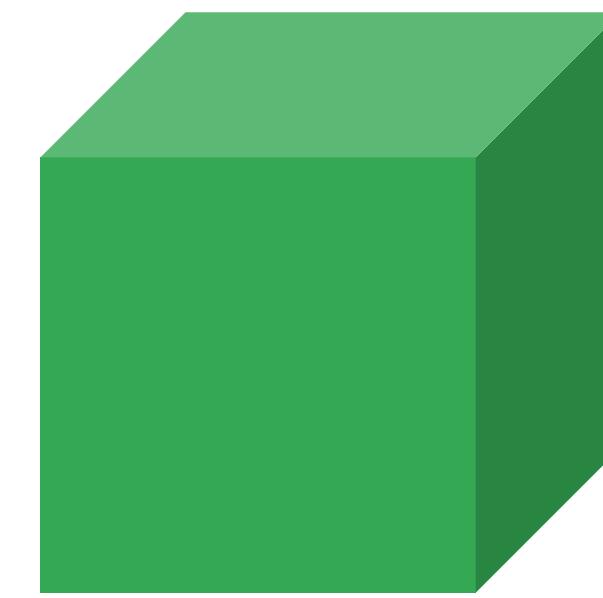
Rank 0
Tensor
scalar



Rank 1
Tensor
vector



Rank 2
Tensor
matrix

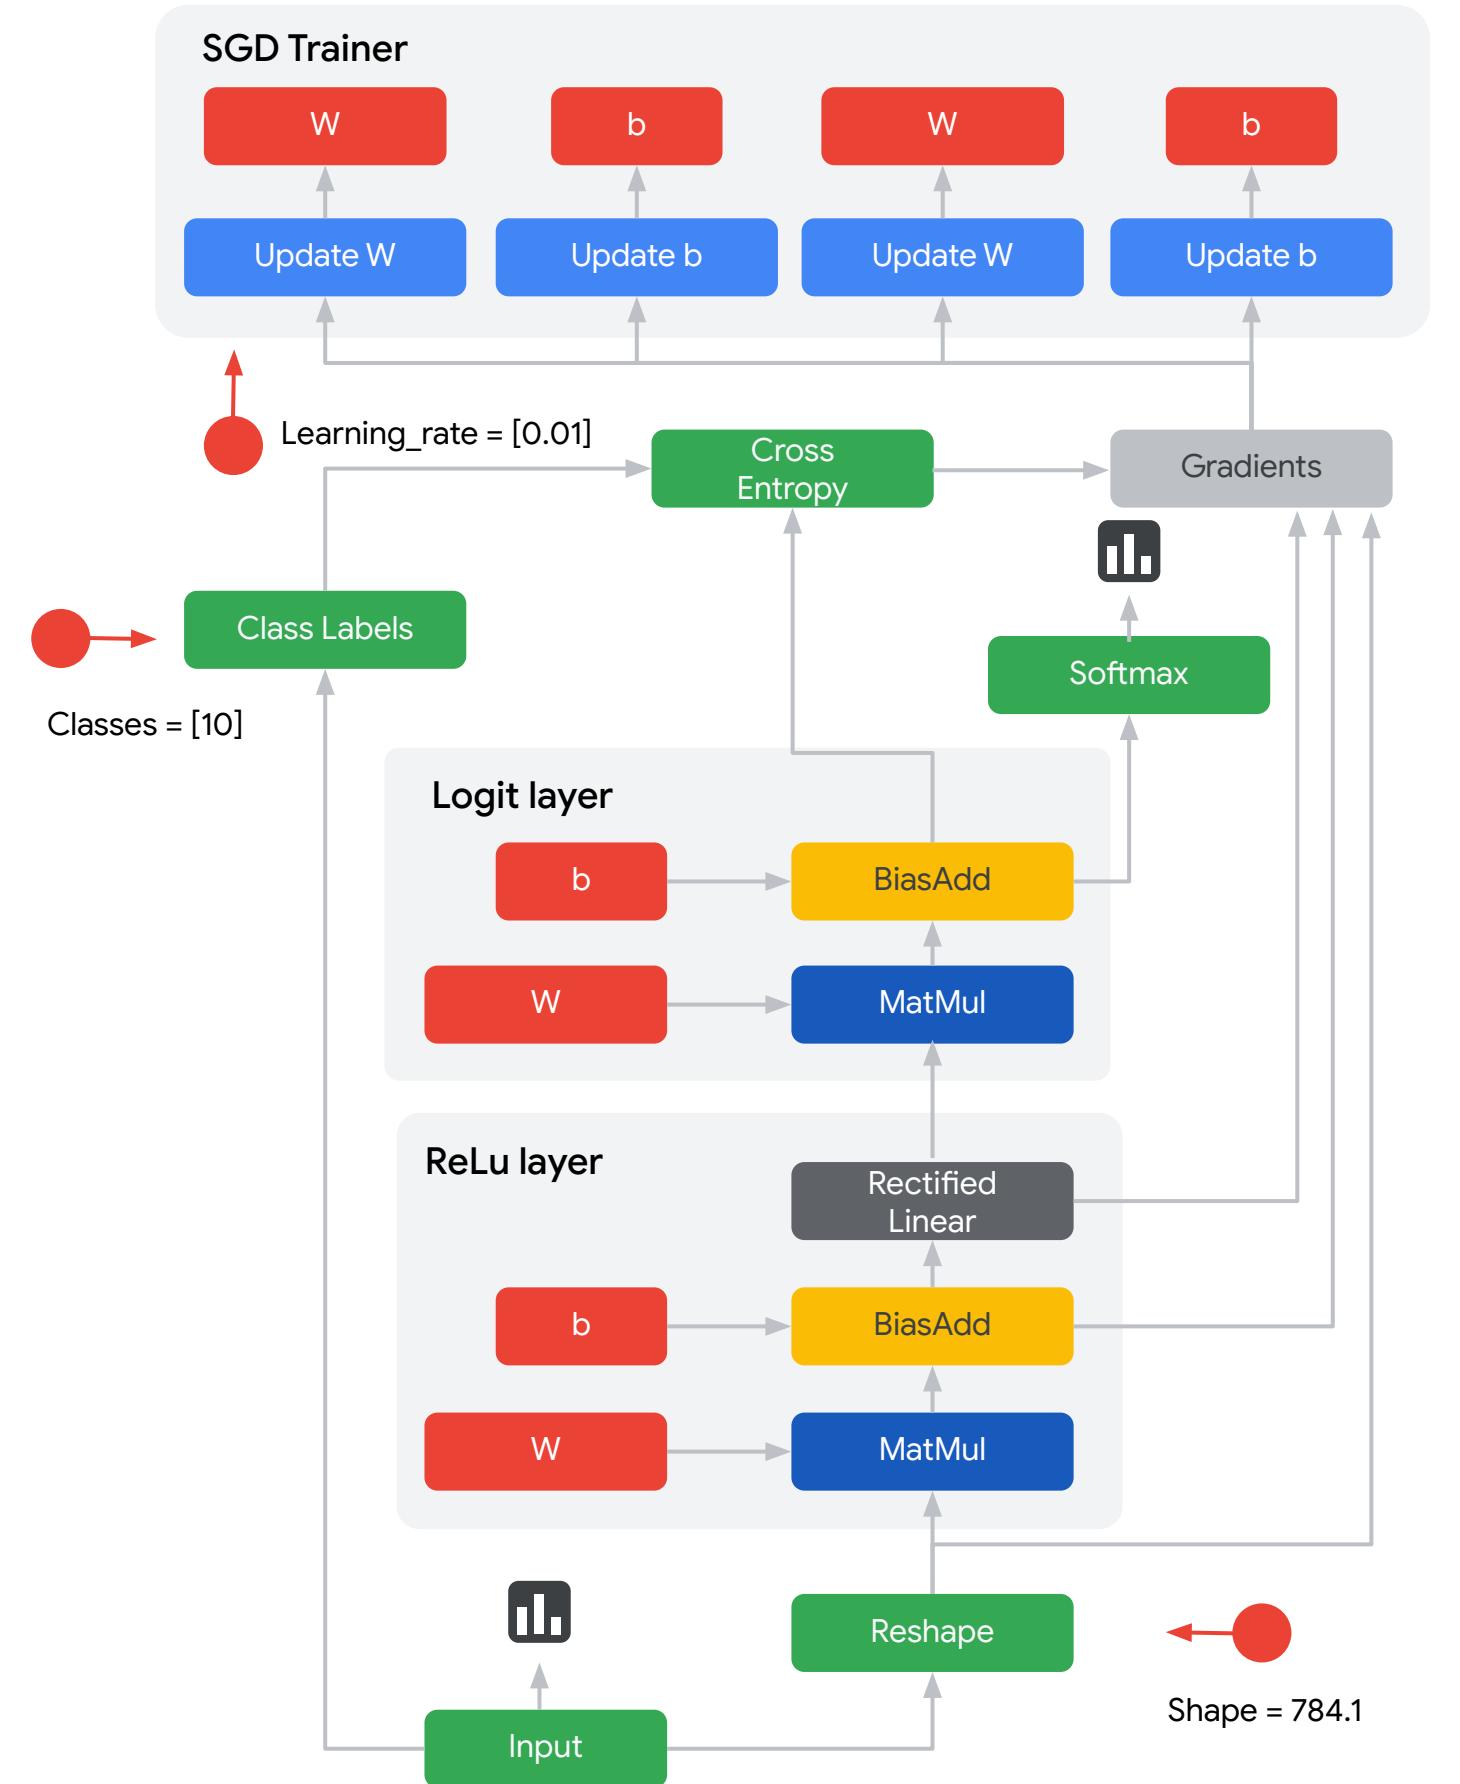


Rank 3
Tensor

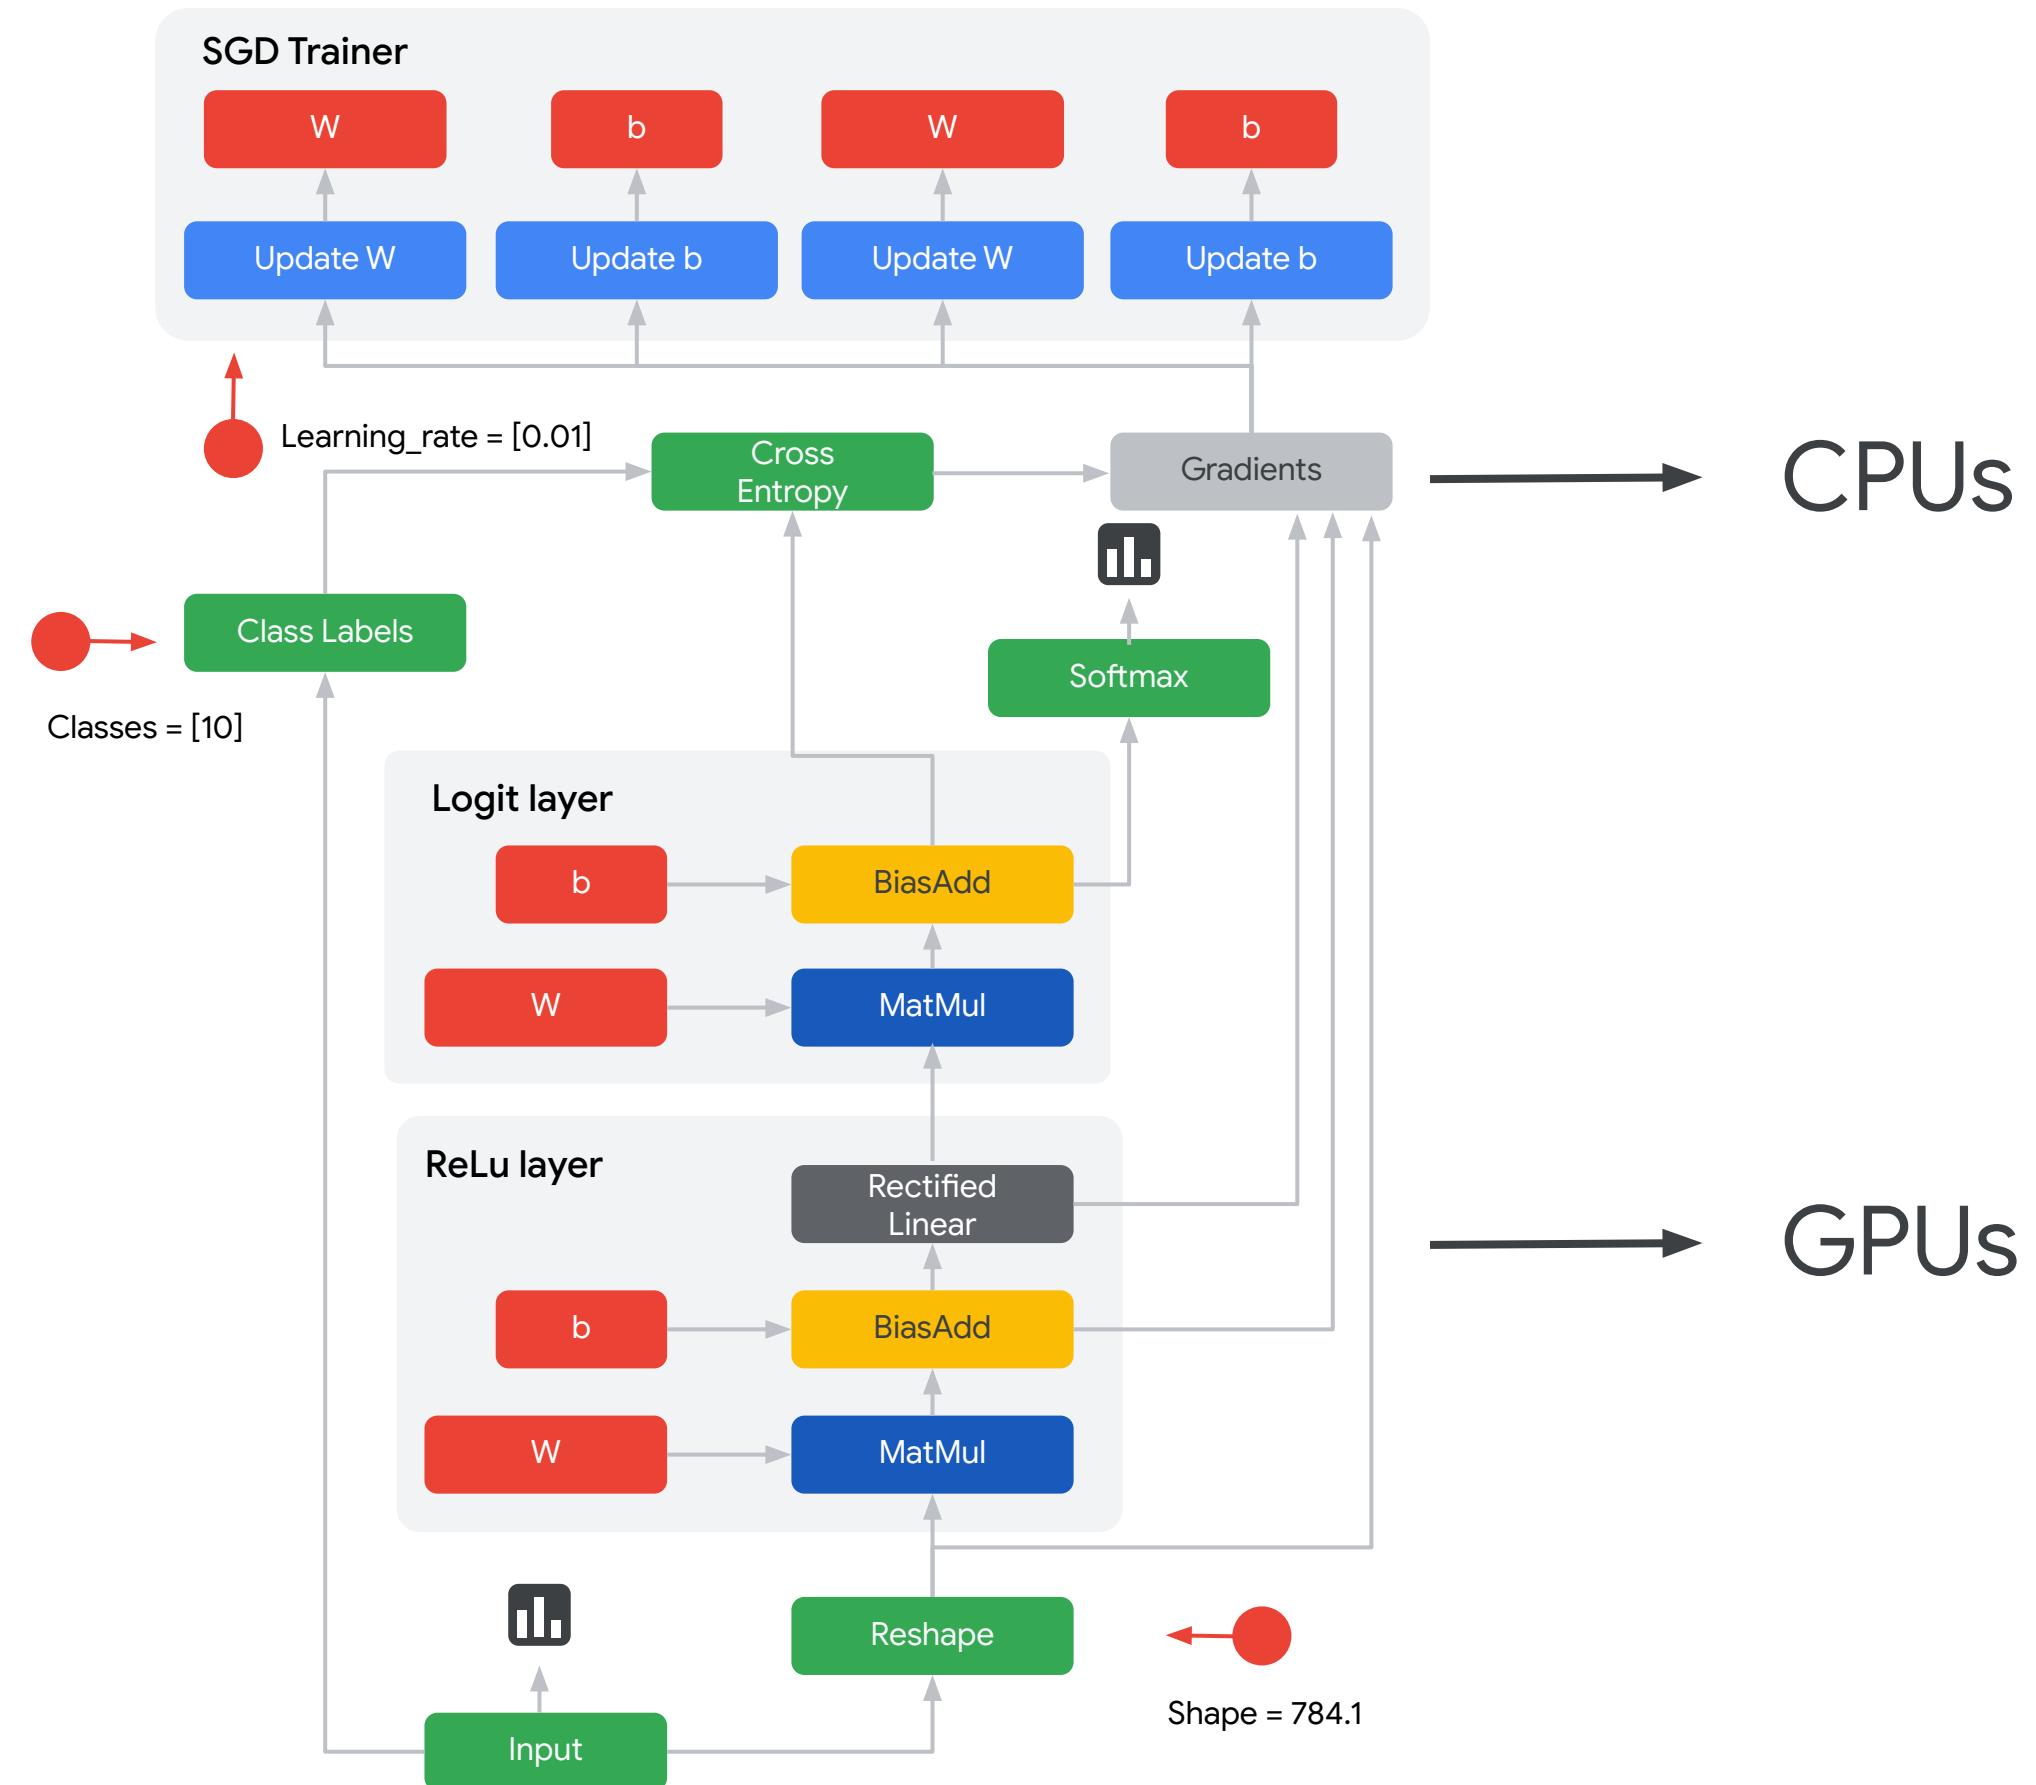


Rank 4
Tensor

TensorFlow graphs are portable between different devices



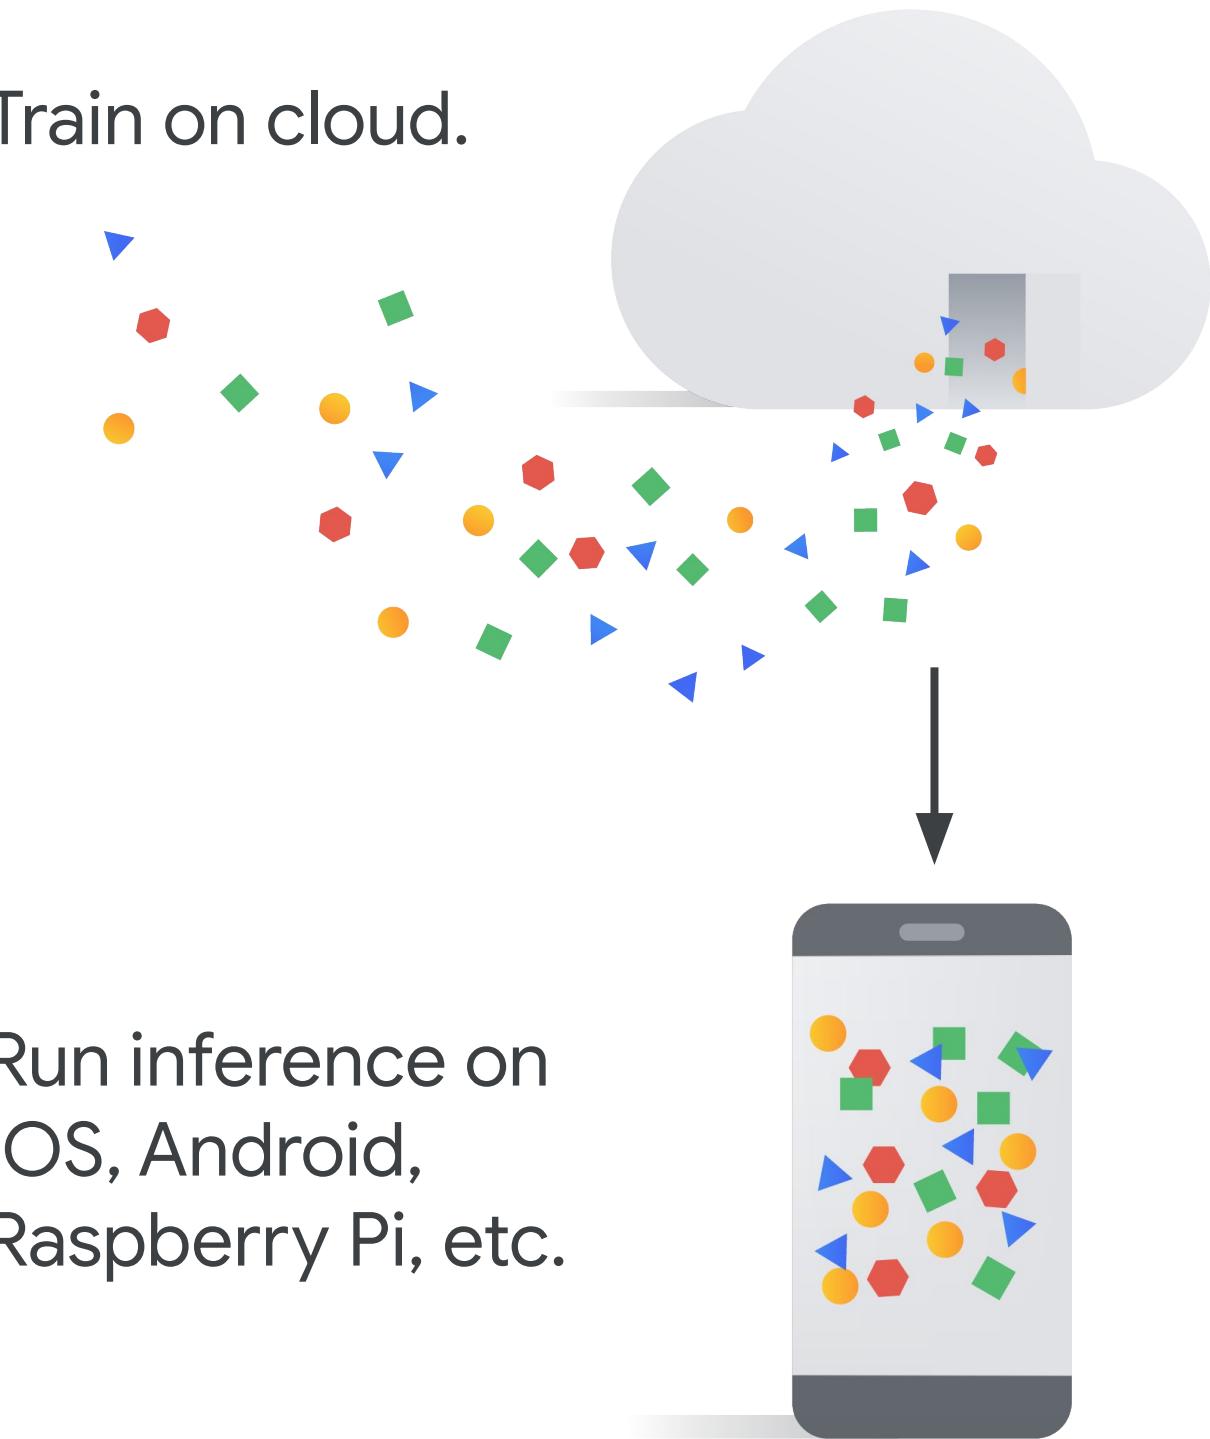
TensorFlow
graphs are
portable between
different devices

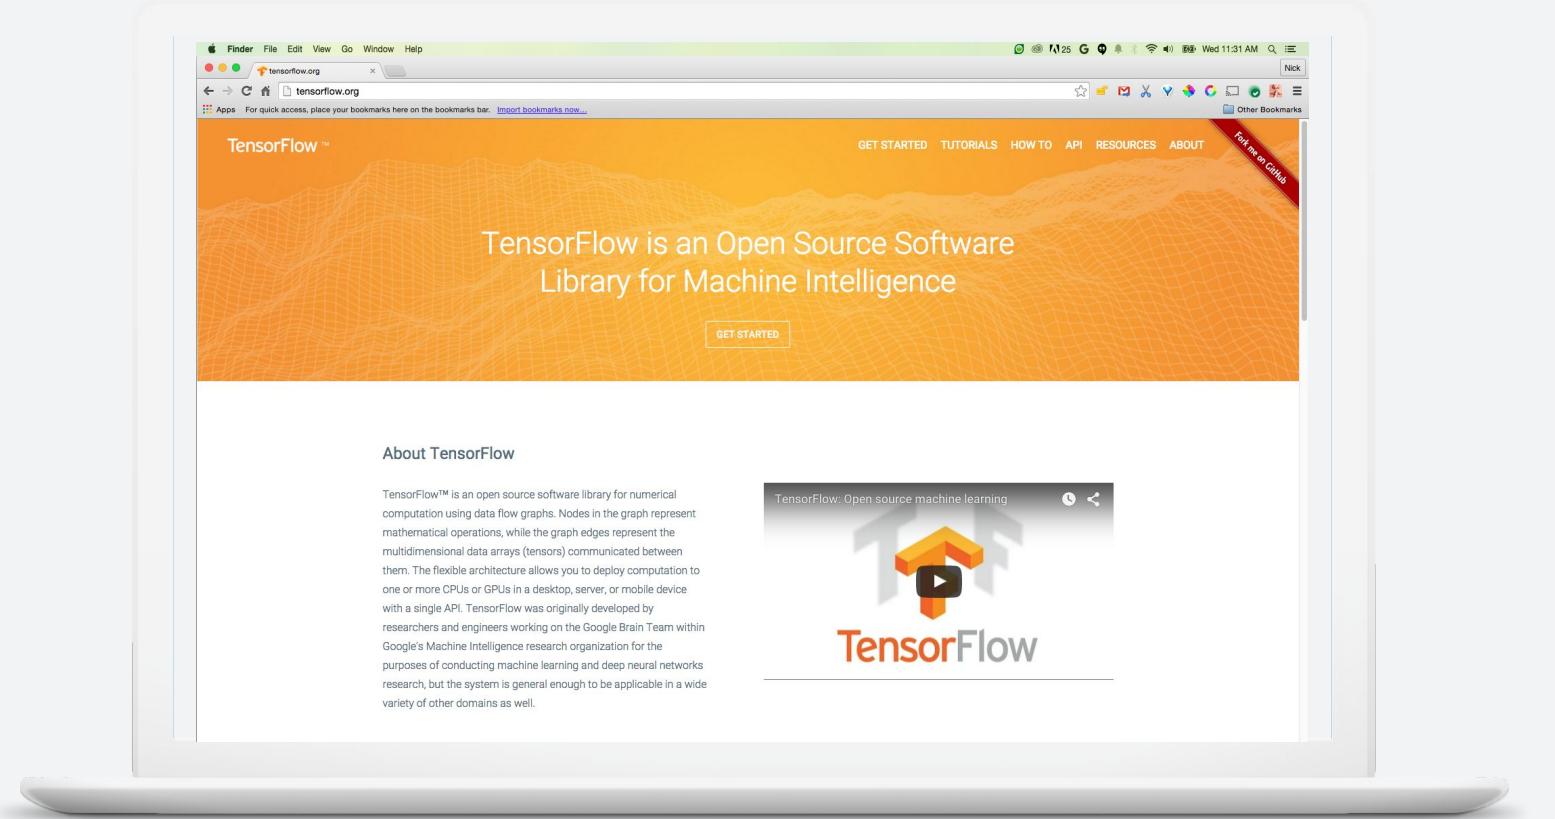


**TensorFlow Lite provides
on-device inference of
ML models on mobile
devices and is available
for a variety of hardware.**

Announcing TensorFlow Lite:
<https://developers.googleblog.com/2017/11/announcing-tensorflow-lite.html>

Train on cloud.





TensorFlow is popular among both deep learning researchers and machine learning engineers.

#1 repository for “machine learning” category on GitHub

TensorFlow API hierarchy

TensorFlow
contains
multiple
abstraction
layers.

CPU

GPU

TPU

Android

TF runs on different
hardware

TensorFlow
contains
multiple
abstraction
layers.

Core TensorFlow (C++)

C++ API is quite low level

CPU

GPU

TPU

Android

TF runs on different
hardware

TensorFlow
contains
multiple
abstraction
layers.

Core TensorFlow (Python)

Python API gives you full control

Core TensorFlow (C++)

C++ API is quite low level

CPU

GPU

TPU

Android

TF runs on different hardware

TensorFlow contains multiple abstraction layers.

`tf.losses`, `tf.metrics`, `tf.optimizers`, etc.

Components useful when building custom NN models

Core TensorFlow (Python)

Python API gives you full control

Core TensorFlow (C++)

C++ API is quite low level

CPU

GPU

TPU

Android

TF runs on different hardware

TensorFlow contains multiple abstraction layers.

`tf.estimator`, `tf.keras`, `tf.data`

High-level APIs for distributed training

`tf.losses`, `tf.metrics`, `tf.optimizers`, etc.

Components useful when building custom NN models

Core TensorFlow (Python)

Python API gives you full control

Core TensorFlow (C++)

C++ API is quite low level

CPU

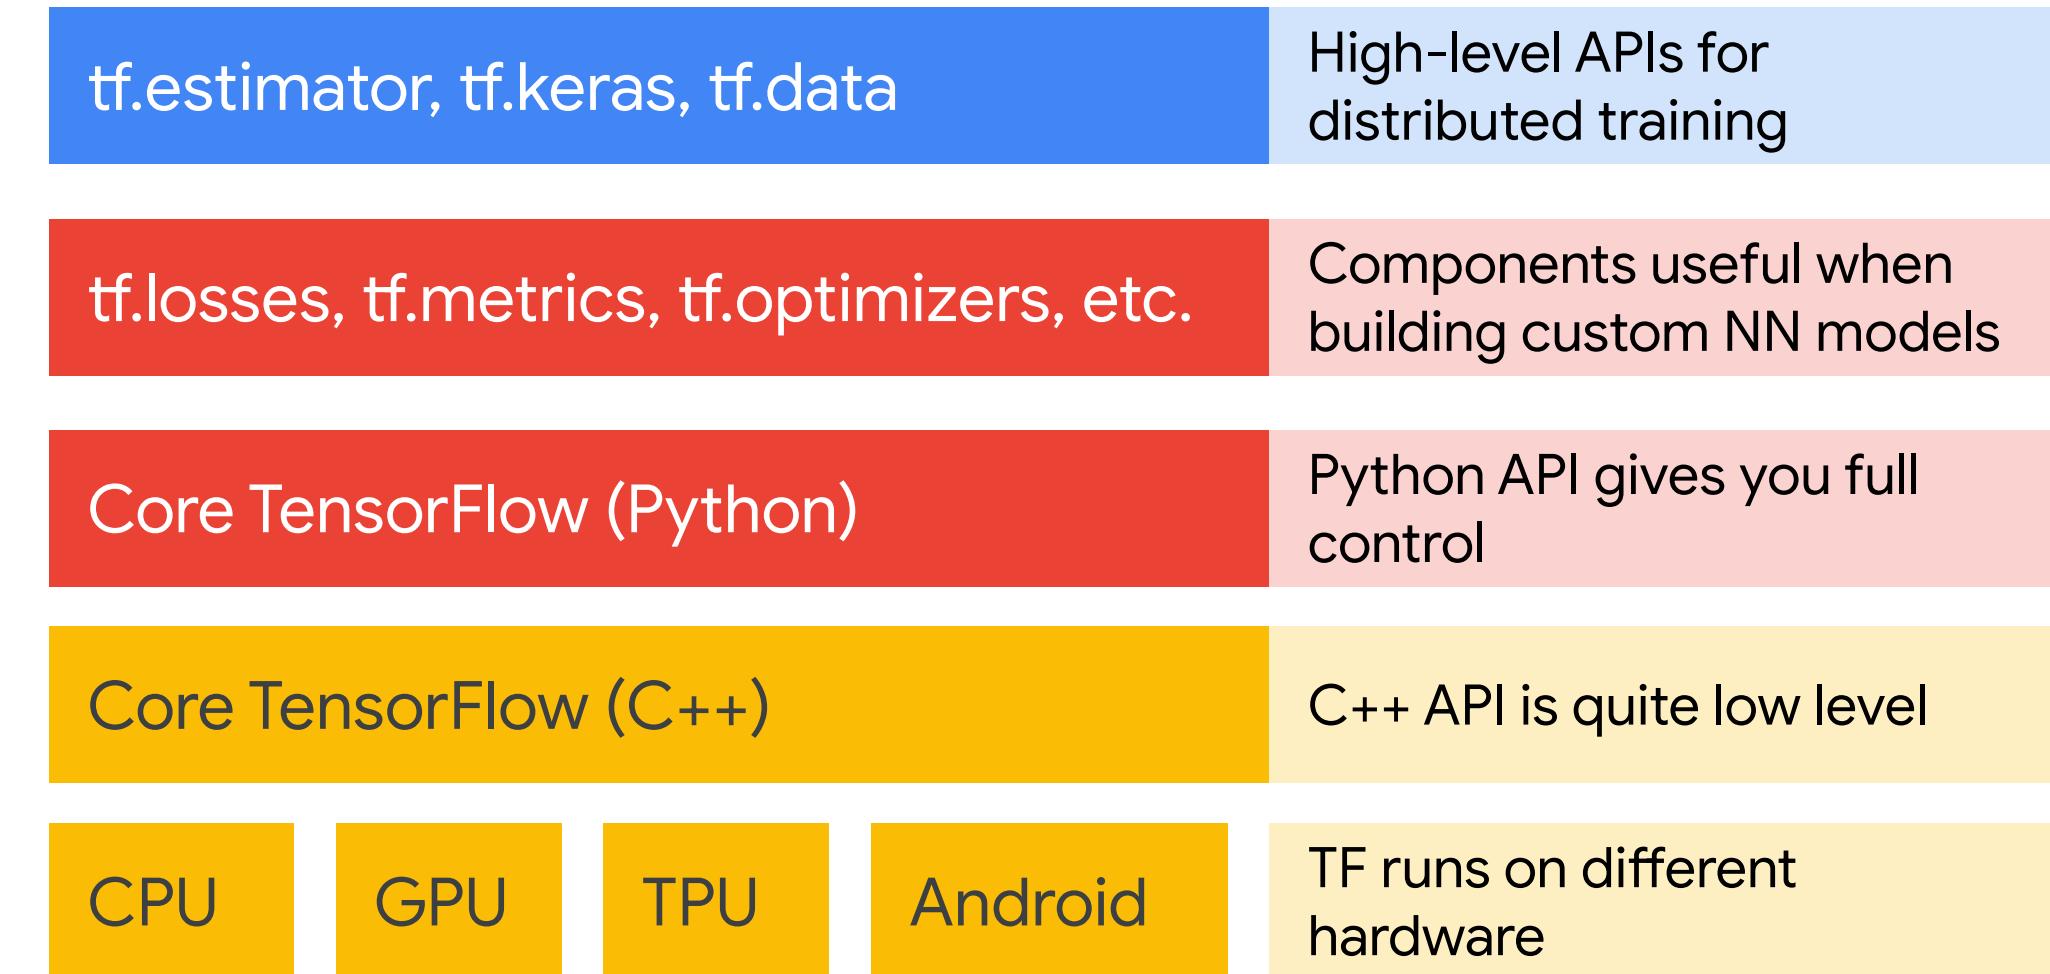
GPU

TPU

Android

TF runs on different hardware

TensorFlow contains multiple abstraction layers.

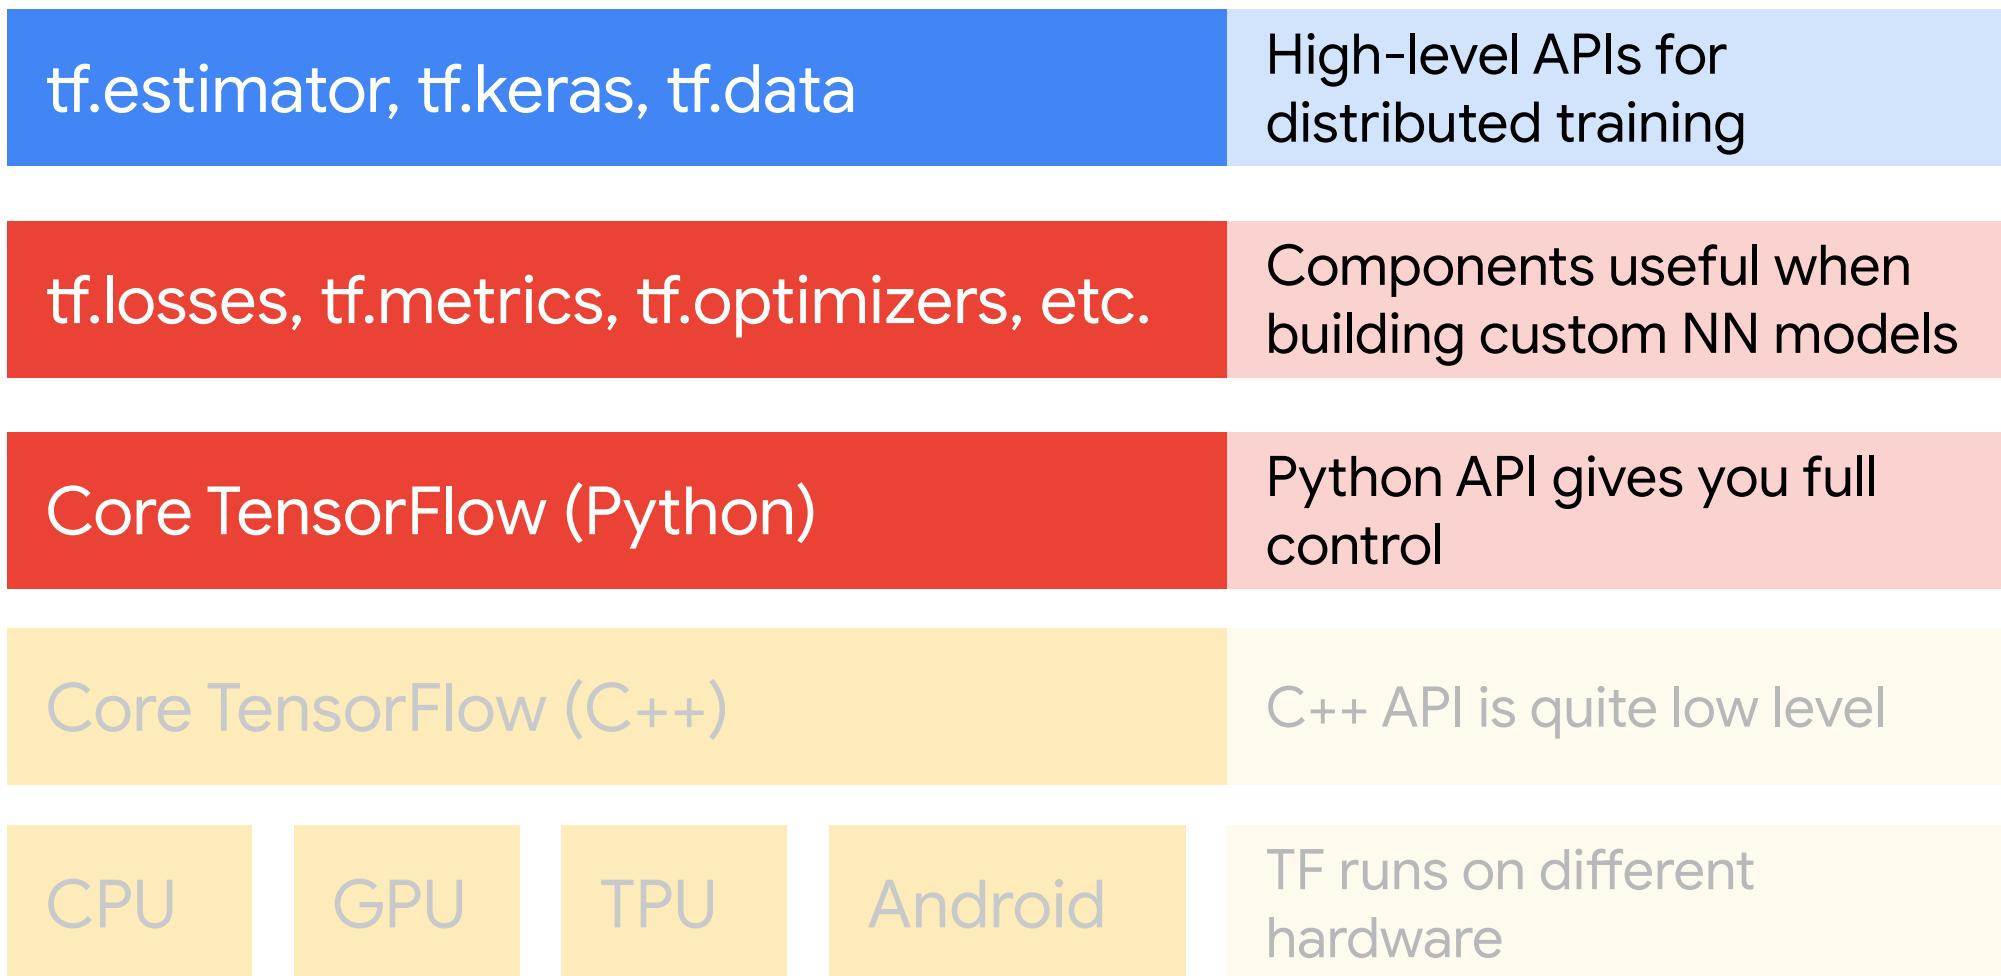


Vertex AI

Run TF at scale with AI Platform.

Note that AI Platform is now Vertex AI.

TensorFlow toolkit hierarchy



TensorFlow **tensors** and **variables**

A tensor is an N-dimensional array of data

Common name	Rank (Dimension)	Example	Shape of example
Scalar	0	<code>x = tf.constant(3)</code>	<code>()</code>

A tensor is an N-dimensional array of data

	Common name	Rank (Dimension)	Example	Shape of example
	Scalar	0	<code>x = tf.constant(3)</code>	(<code>)</code>
	Vector	1	<code>x = tf.constant([3, 5, 7])</code>	(<code>3,</code>)

A tensor is an N-dimensional array of data

	Common name	Rank (Dimension)	Example	Shape of example
	Scalar	0	<code>x = tf.constant(3)</code>	(<code>)</code>
	Vector	1	<code>x = tf.constant([3, 5, 7])</code>	(<code>3,</code>)
	Matrix	2	<code>x = tf.constant([[3, 5, 7], [4, 6, 8]])</code>	(<code>2, 3</code>)

A tensor is an N-dimensional array of data

	Common name	Rank (Dimension)	Example	Shape of example
	Scalar	0	<code>x = tf.constant(3)</code>	(<code>)</code>
	Vector	1	<code>x = tf.constant([3, 5, 7])</code>	(<code>3,</code>)
	Matrix	2	<code>x = tf.constant([[3, 5, 7], [4, 6, 8]])</code>	(<code>2, 3</code>)
	3D Tensor	3	<code>tf.constant([[[3, 5, 7], [4, 6, 8]], [[1, 2, 3], [4, 5, 6]]])</code>	(<code>2, 2, 3</code>)

A tensor is an N-dimensional array of data

	Common name	Rank (Dimension)	Example	Shape of example
	Scalar	0	<code>x = tf.constant(3)</code>	(<code>)</code>
	Vector	1	<code>x = tf.constant([3, 5, 7])</code>	(<code>3,</code>)
	Matrix	2	<code>x = tf.constant([[3, 5, 7], [4, 6, 8]])</code>	(<code>2, 3</code>)
	3D Tensor	3	<code>tf.constant([[[3, 5, 7], [4, 6, 8]], [[1, 2, 3], [4, 5, 6]]])</code>	(<code>2, 2, 3</code>)
	nD Tensor	n	<code>x1 = tf.constant([2, 3, 4])</code> <code>x2 = tf.stack([x1, x1])</code> <code>x3 = tf.stack([x2, x2, x2, x2])</code> <code>x4 = tf.stack([x3, x3])</code> ...	(<code>3, </code>) (<code>2, 3</code>) (<code>4, 2, 3</code>) (<code>2, 4, 2, 3</code>)

A tensor is an N-dimensional array of data

They behave like numpy n-dimensional arrays except that:

- `tf.constant` produces constant tensors
- `tf.Variable` produces tensors that can be modified

tf.constant produces constant tensors...

```
import tensorflow as tf

x = tf.constant([
    [3, 5, 7],
    [4, 6, 8]
])
```

Tensors can be sliced

```
import tensorflow as tf

x = tf.constant([
    [3, 5, 7],
    [4, 6, 8]
])

y = x[ :, 1]

[5, 6]
```

Tensors can be reshaped

```
import tensorflow as tf

x = tf.constant([
    [3, 5, 7],
    [4, 6, 8]
])

y = tf.reshape(x, [3, 2])

[[3 5]
 [7 4]
 [6 8]
]
```

A variable is a tensor whose value can be changed...

```
import tensorflow as tf

# x <- 2
x = tf.Variable(2.0, dtype=tf.float32,
name='my_variable')
```

A variable is a tensor whose value can be changed...

tf.Variable will typically hold model weights
that need to be updated in a training loop.

```
import tensorflow as tf

# x <- 2
x = tf.Variable(2.0, dtype=tf.float32, name='my_variable')

# x <- 48.5
x.assign(45.8)

# x <- x + 4
x.assign_add(4)

# x <- x - 3
x.assign_sub(3)
```

A variable is a tensor whose value can be changed...

```
import tensorflow as tf

# w * x
w = tf.Variable([[1., 2.]])
x = tf.constant([[3., 4.]])
tf.matmul(w, x)
```

GradientTape records operations for automatic differentiation

TensorFlow can compute the derivative of a function with respect to any parameter.

- the computation is recorded with [GradientTape](#)
- the function is expressed with [TensorFlow ops only!](#)

GradientTape records operations for automatic differentiation

```
def compute_gradients(X, Y, w0, w1):  
    with tf.GradientTape() as tape:  
        loss = loss_mse(X, Y, w0, w1)  
    return tape.gradient(loss, [w0, w1])
```

```
w0 = tf.Variable(0.0)  
w1 = tf.Variable(0.0)  
  
dw0, dw1 = compute_gradients(X, Y, w0, w1)
```

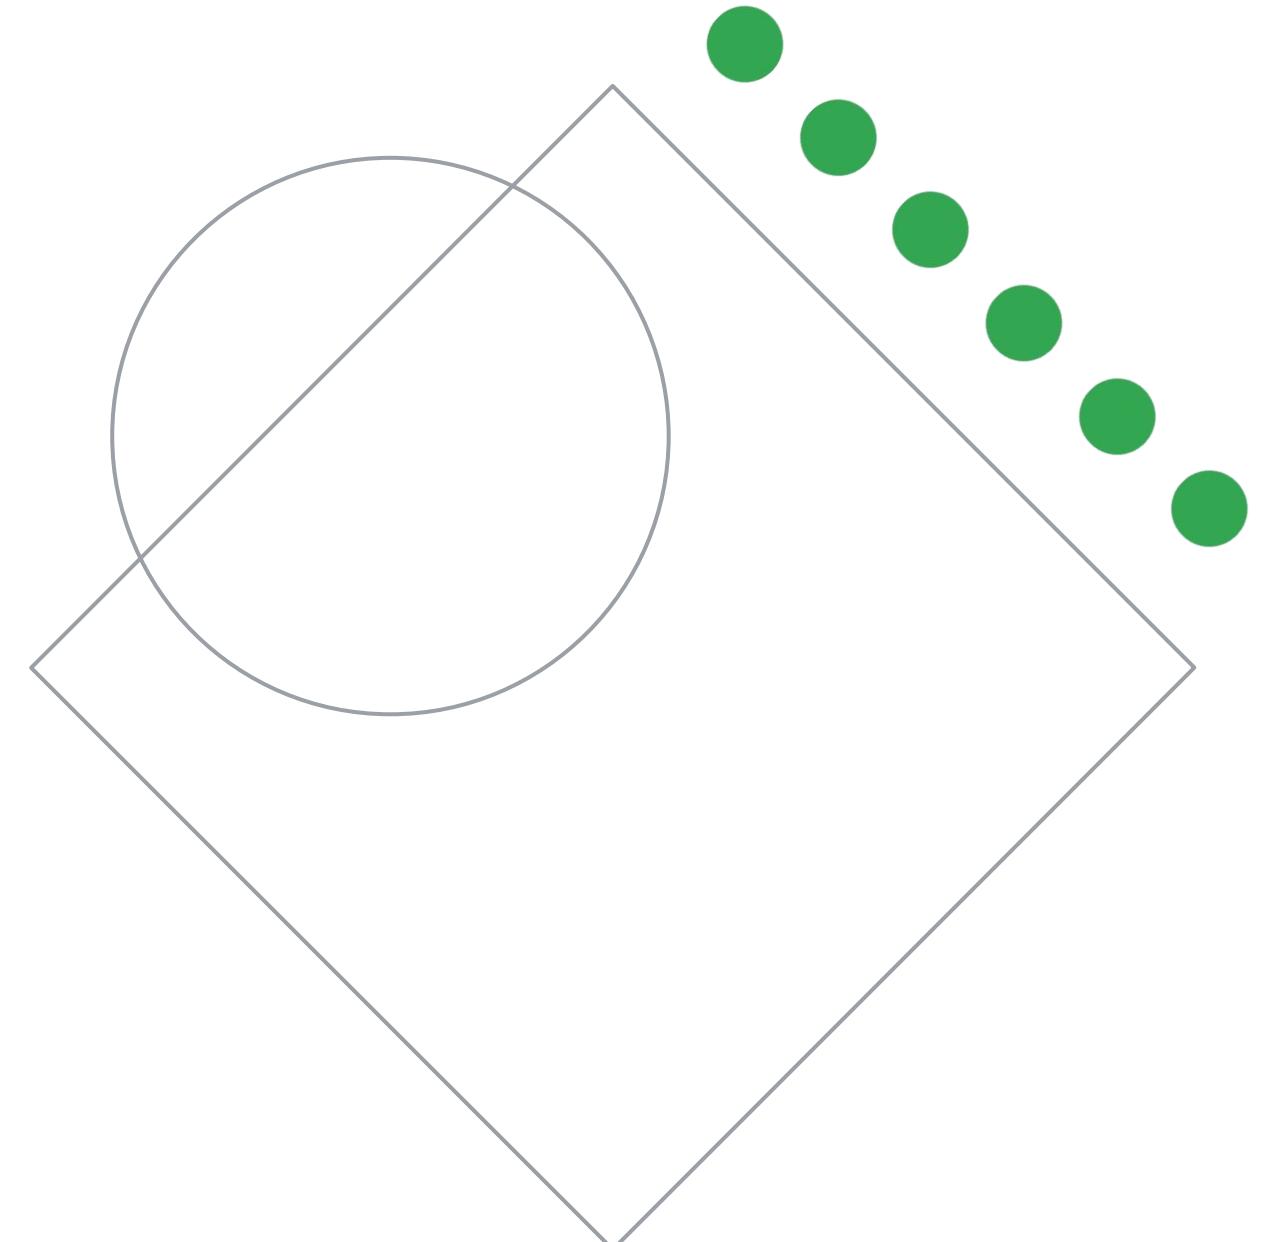
Record the computation with GradientTape when it's executed (not when it's defined!)

GradientTape records operations for automatic differentiation

```
def compute_gradients(X, Y, w0, w1):  
    with tf.GradientTape() as tape:  
        loss = loss_mse(X, Y, w0, w1)  
    return tape.gradient(loss, [w0, w1])  
  
w0 = tf.Variable(0.0)  
w1 = tf.Variable(0.0)  
  
dw0, dw1 = compute_gradients(X, Y, w0, w1)
```

Specify the function (loss) as well as the parameters you want to take the gradient with respect to ([w0, w1])

Design and Build an Input Data Pipeline



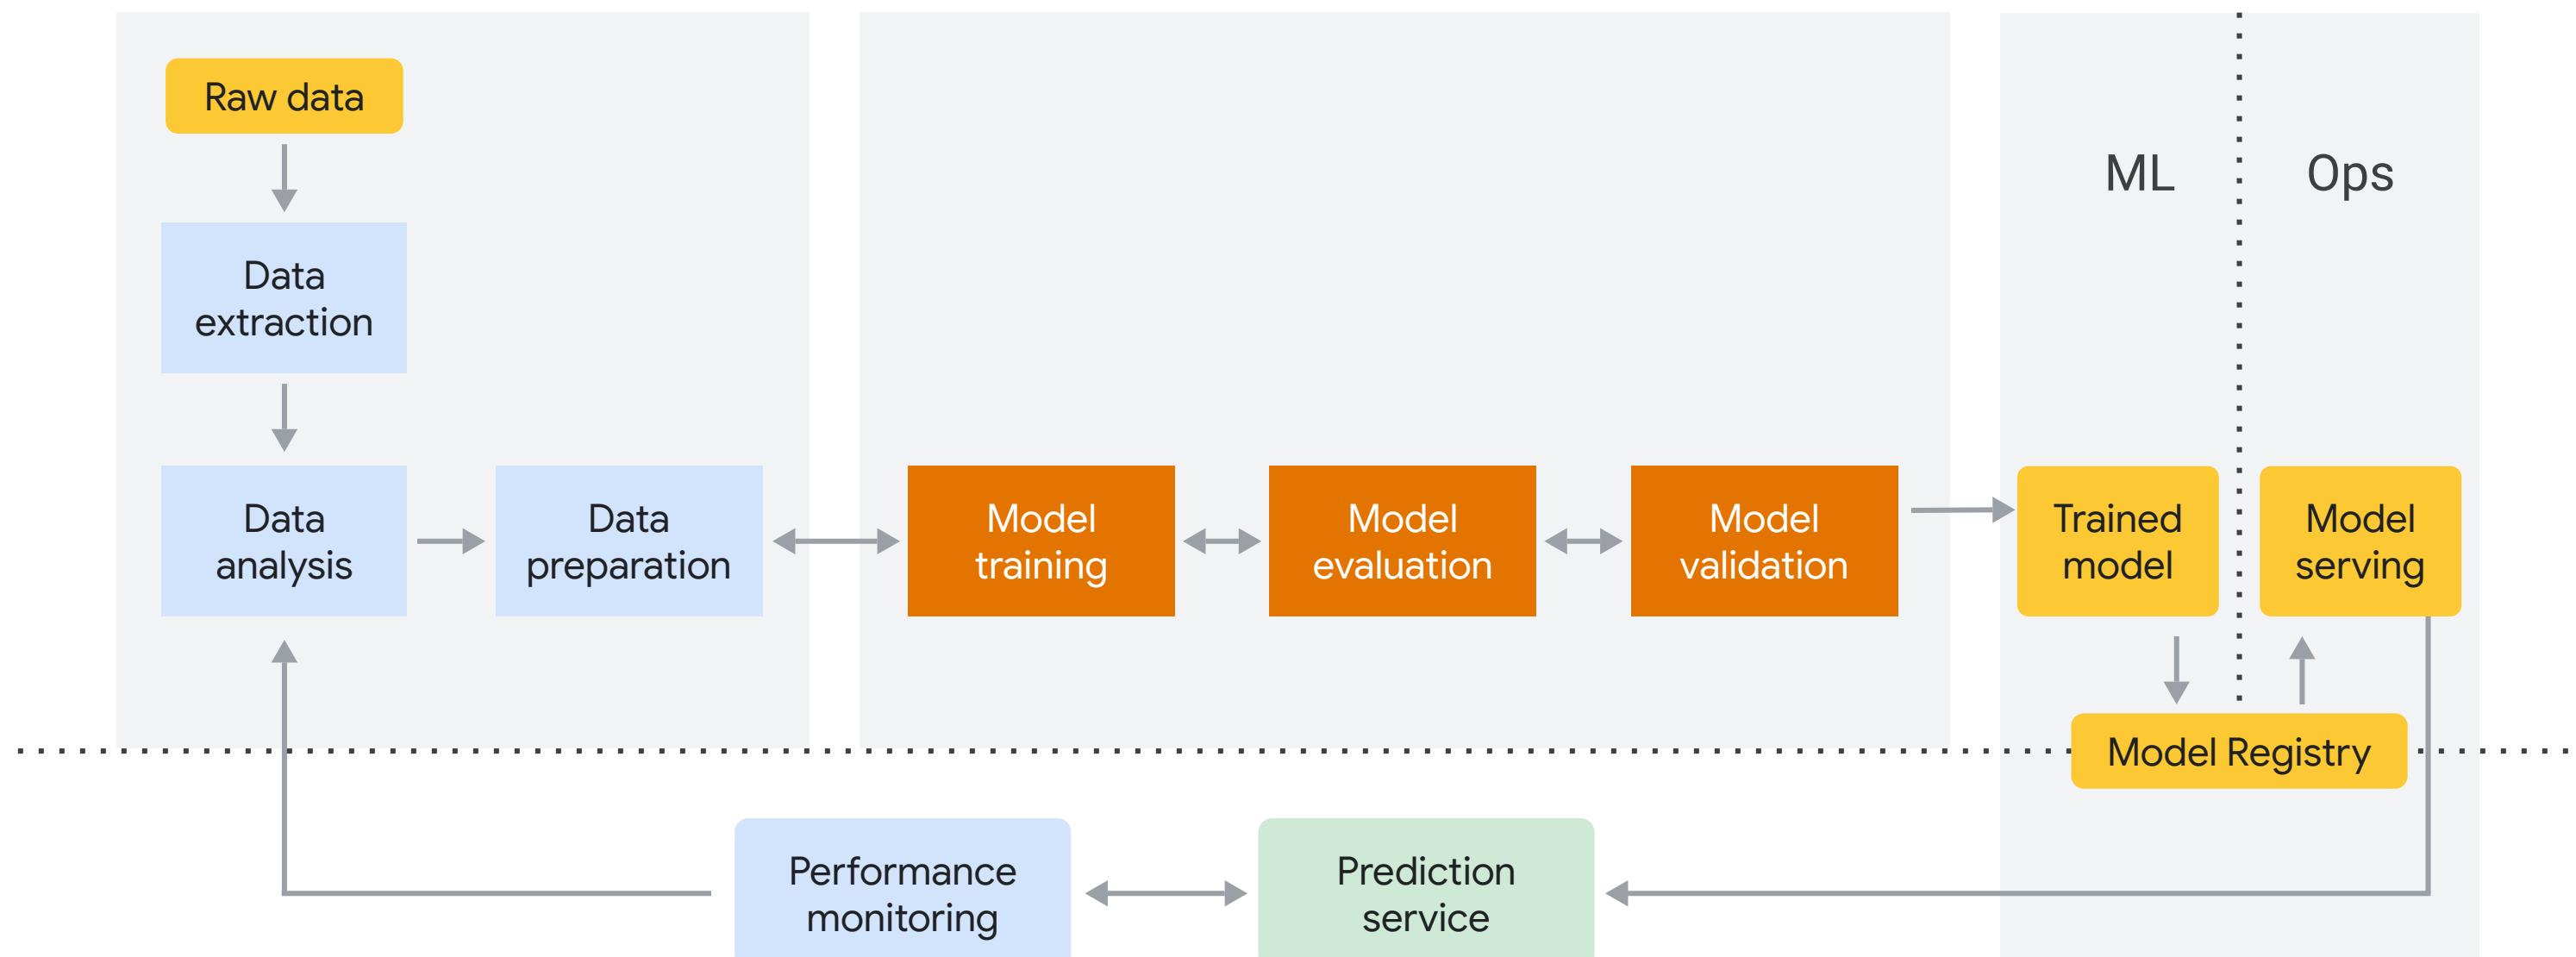
In this module, you learn to ...

- 01 Train on Large Datasets with tf.data
- 02 Work with in-memory files
- 03 Get the data ready for training
- 04 Describe embeddings
- 05 Understand scaling data with tf.Keras preprocessing layers



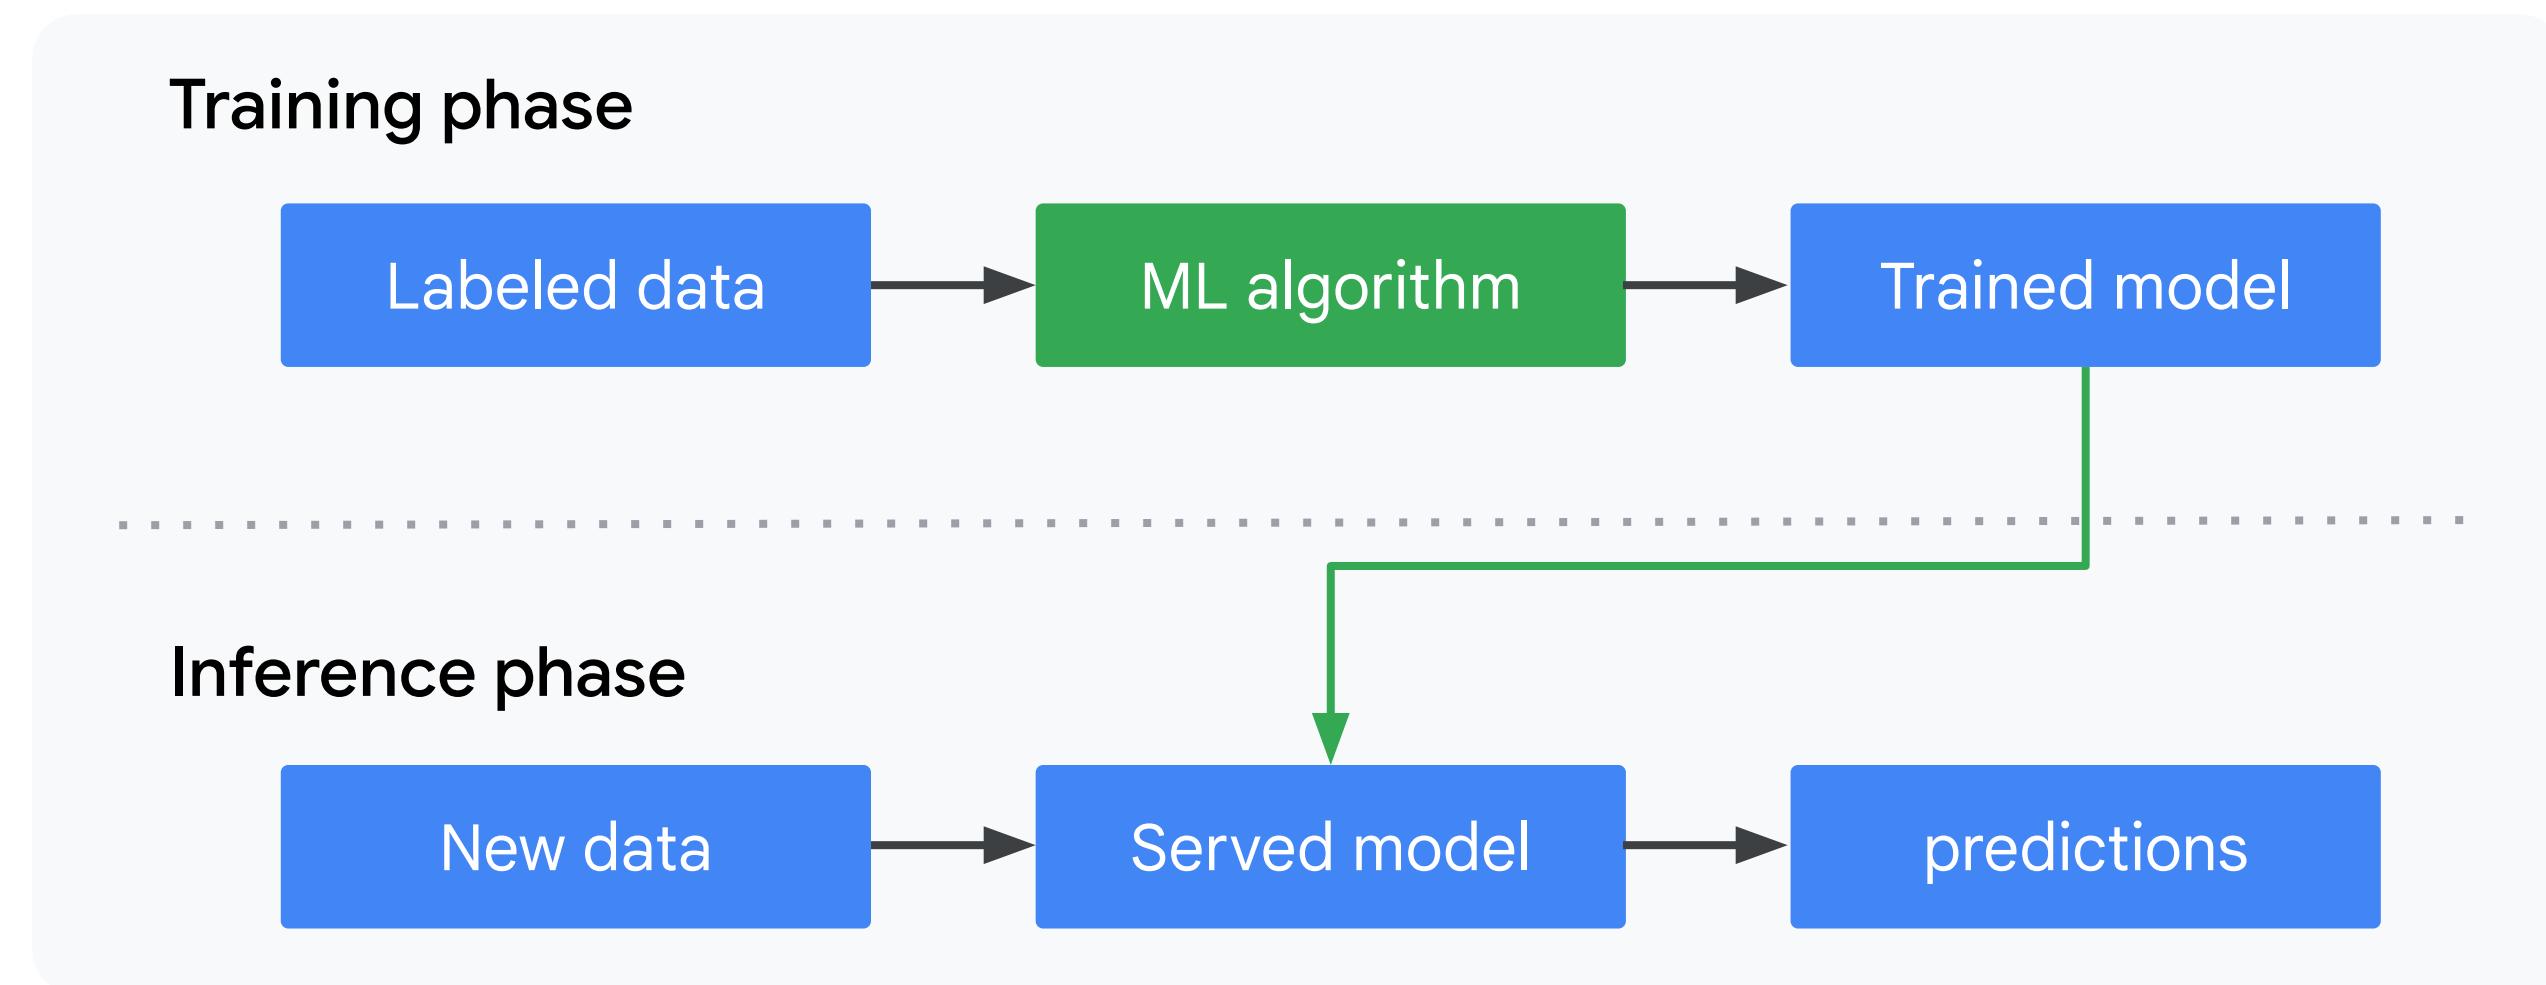
An ML recap

Staging/pre-production/production environments

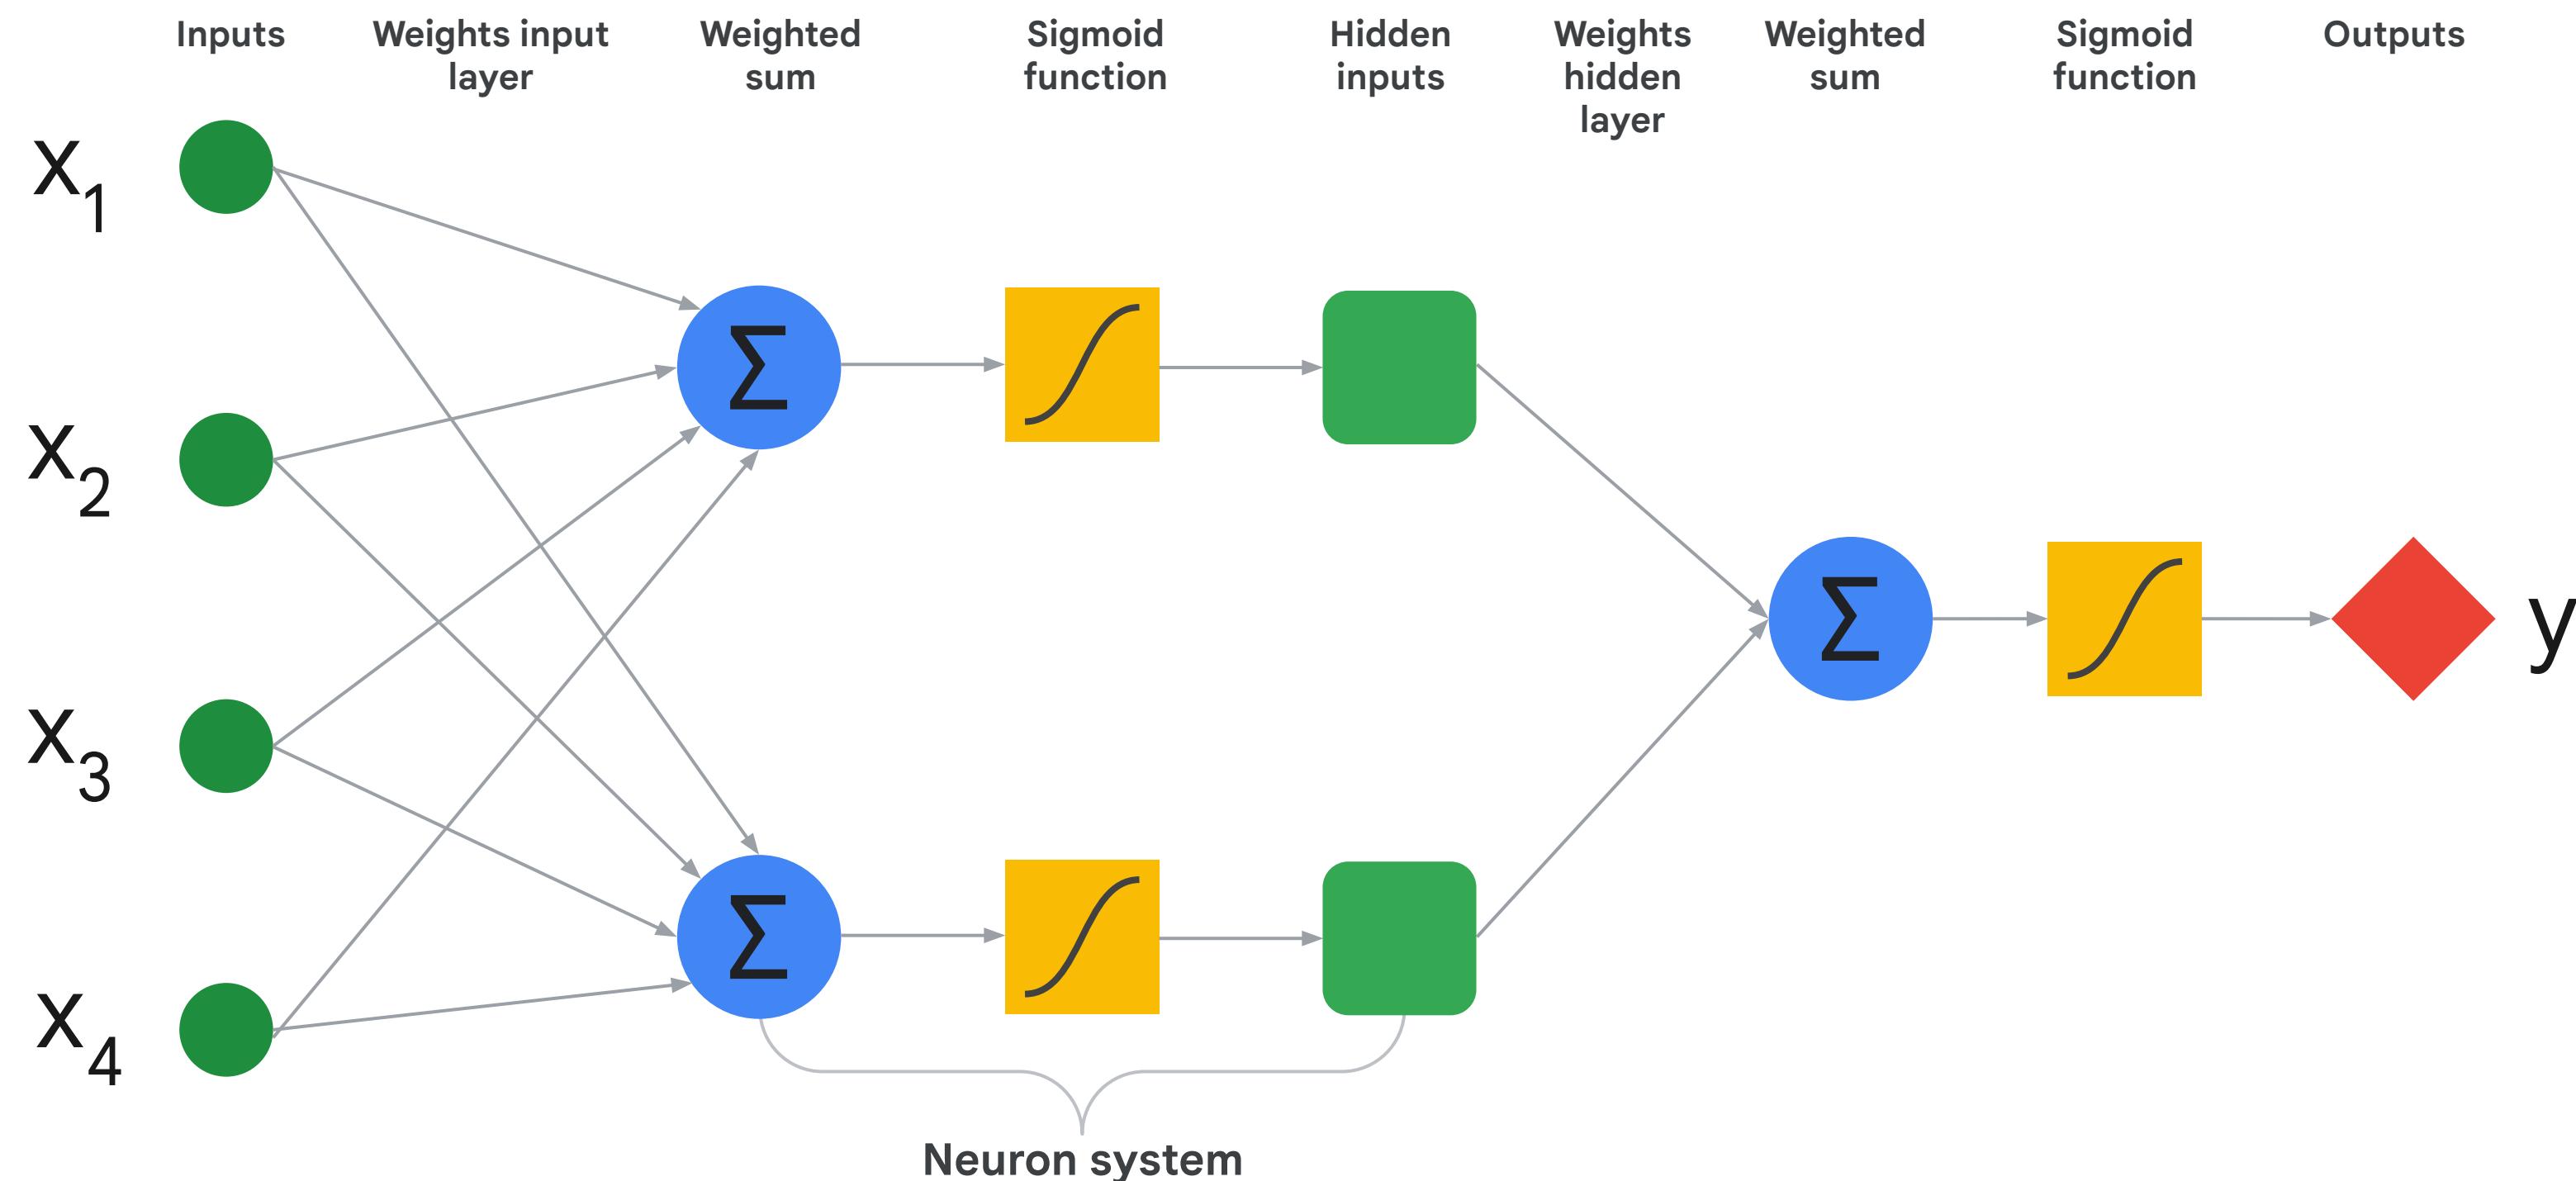


Experimentation/development/test environments

An ML recap



An ML recap



Performing a training step involves

- 1 Opening a file
(if it isn't open already)
- 2 Fetching a data entry
from the file
- 3 Using the data for
training

tf.data API



1

Build complex input
pipelines from simple,
reusable pieces.

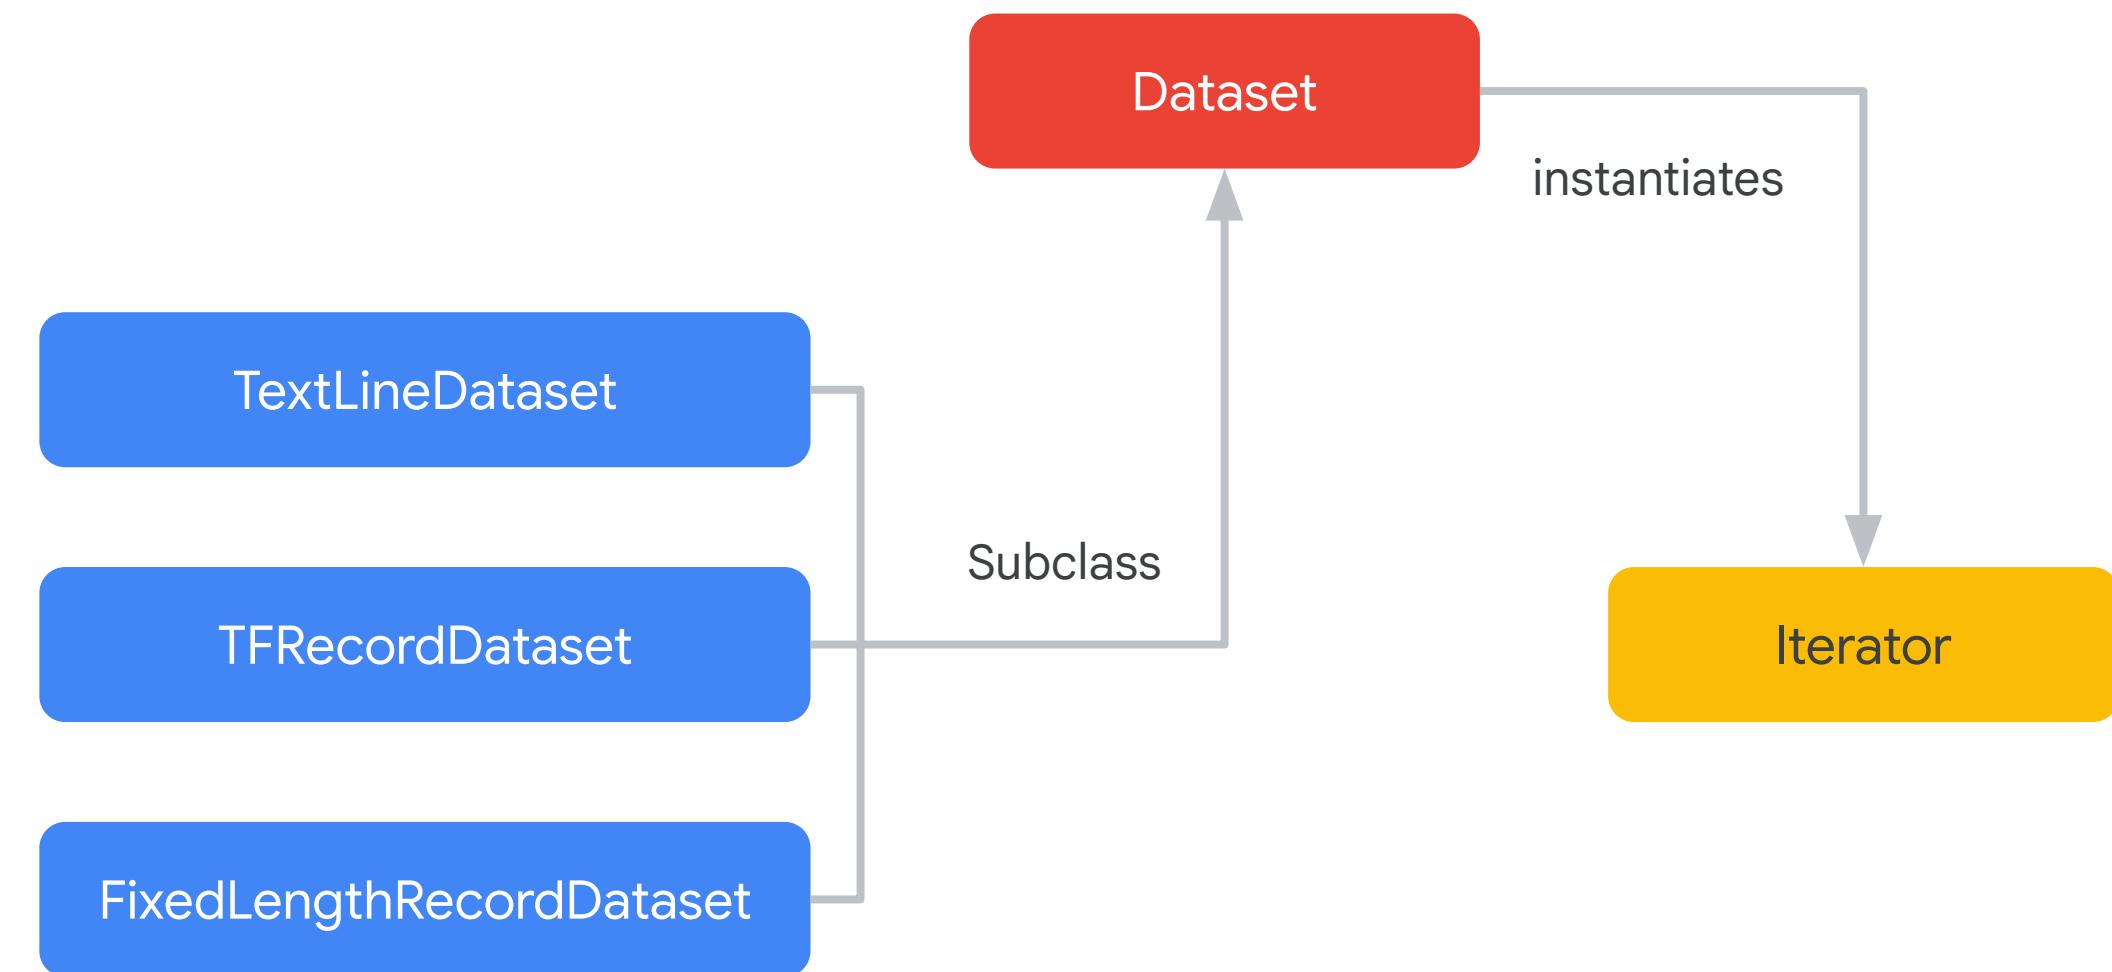
2

Build pipelines for
multiple data types

3

Handle large amounts of
data; perform complex
transformations

Multiple ways to feed TensorFlow models **with data**



A `tf.data.Dataset` allows you to

- Create data pipelines from
 - in-memory dictionary and lists of tensors
 - out-of-memory sharded data files
- Preprocess data in parallel (and cache result of costly operations)
`dataset = dataset.map(preproc_fun).cache()`

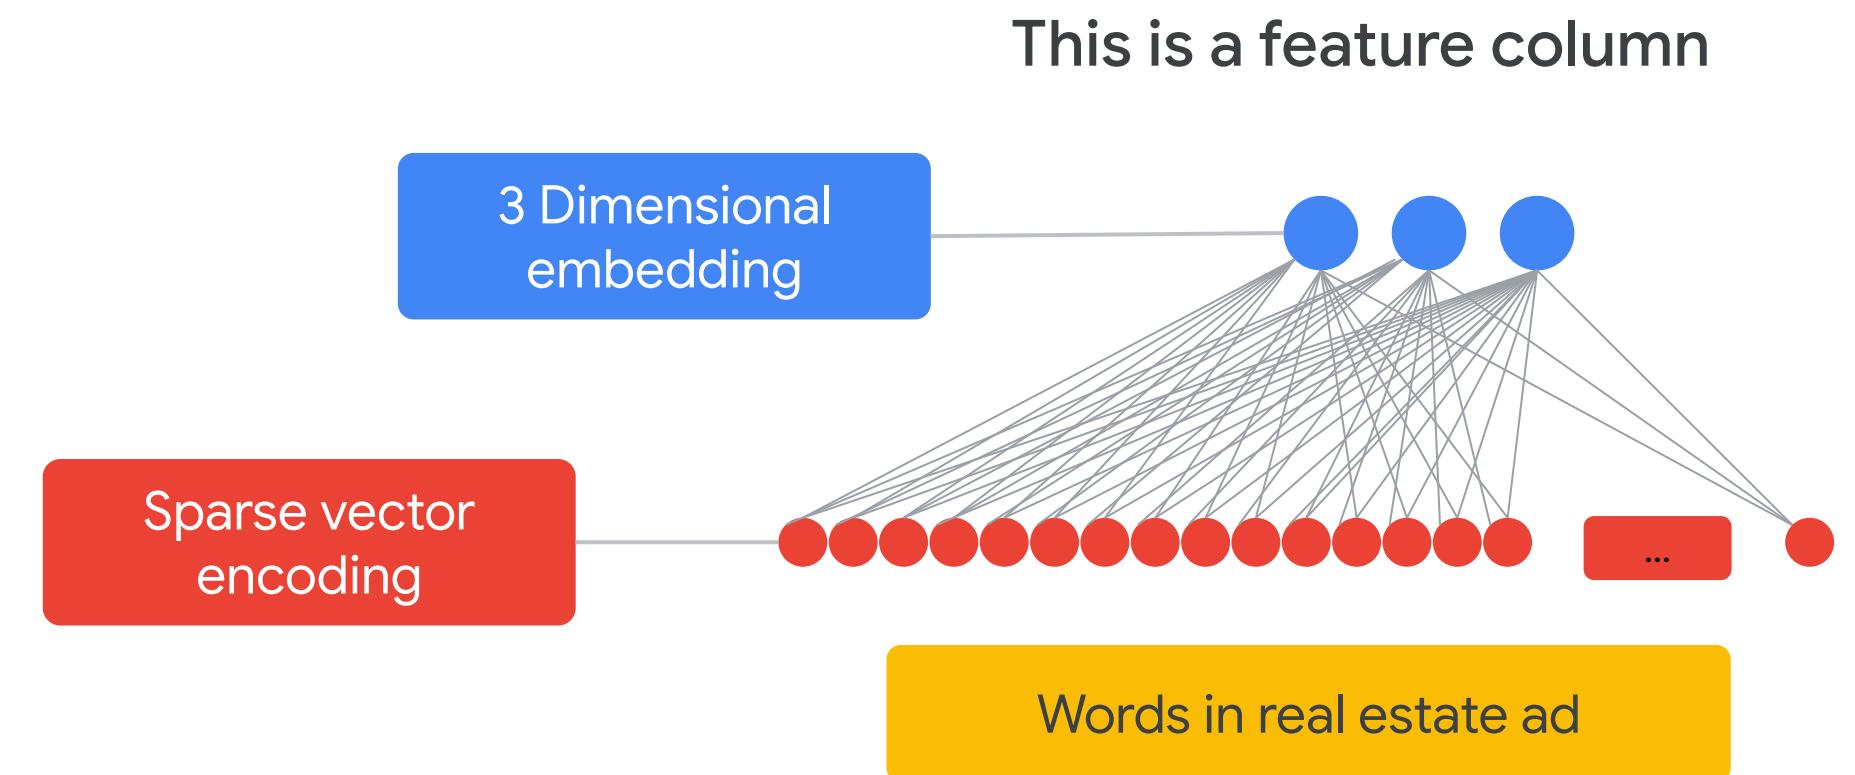
- Configure the way the data is fed into a model with a number of chaining methods

```
dataset = dataset.shuffle(1000).repeat(epochs).batch(batch_size,  
drop_remainder=True)
```

in a easy and very compact way

Embeddings are feature columns that function like layers

```
import tensorflow as tf  
  
sparse_word =  
    tf.feature_column.categorical_column_with_vocabulary_list  
    ('word', vocabulary_list=englishWords)  
  
embedded_word =  
    tf.feature_column.embedding_column(sparse_word, 3)
```



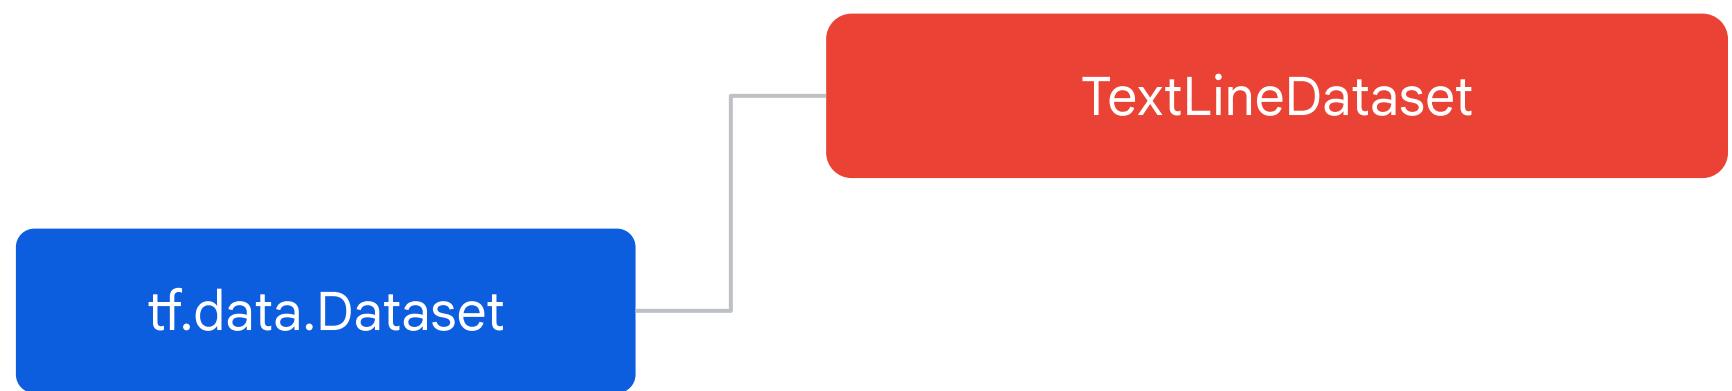
What about out-of-memory sharded datasets?

train.csv-00000-of-00011	9.23 MB
train.csv-00001-of-00011	16.82 MB
train.csv-00002-of-00011	44.18 MB
train.csv-00003-of-00011	14.63 MB
train.csv-00004-of-00011	
train.csv-00005-of-00011	
train.csv-00006-of-00011	
train.csv-00007-of-00011	
valid.csv-00000-of-00001	2.31 MB
valid.csv-00000-of-00009	19.47 MB
valid.csv-00001-of-00009	11.6 MB
valid.csv-00002-of-00009	9.5 MB
valid.csv-00003-of-00009	18.29 MB

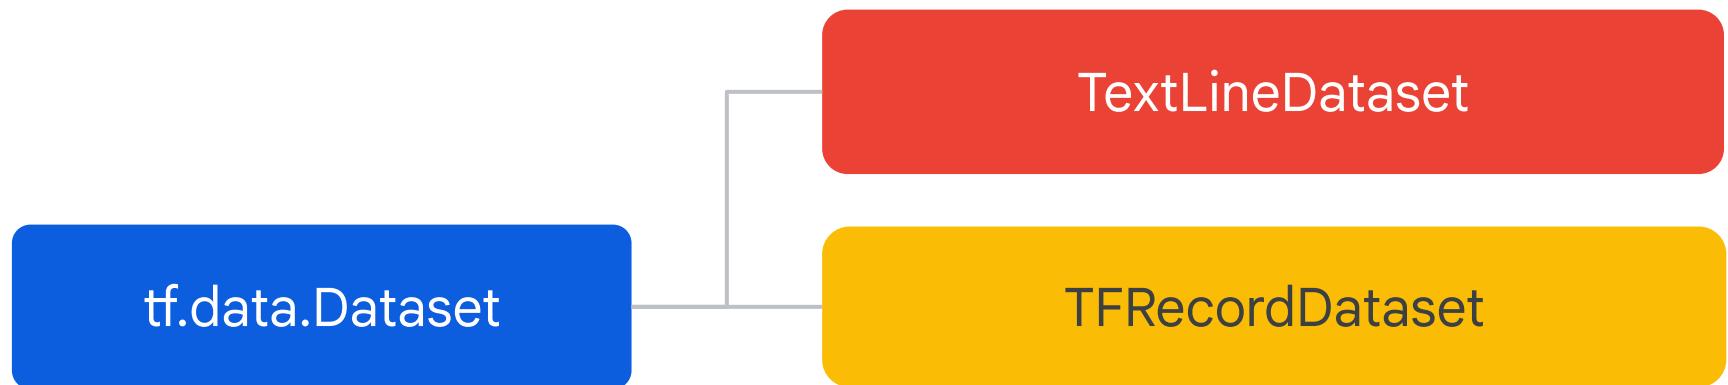
Datasets can be
created from
different **file formats**

`tf.data.Dataset`

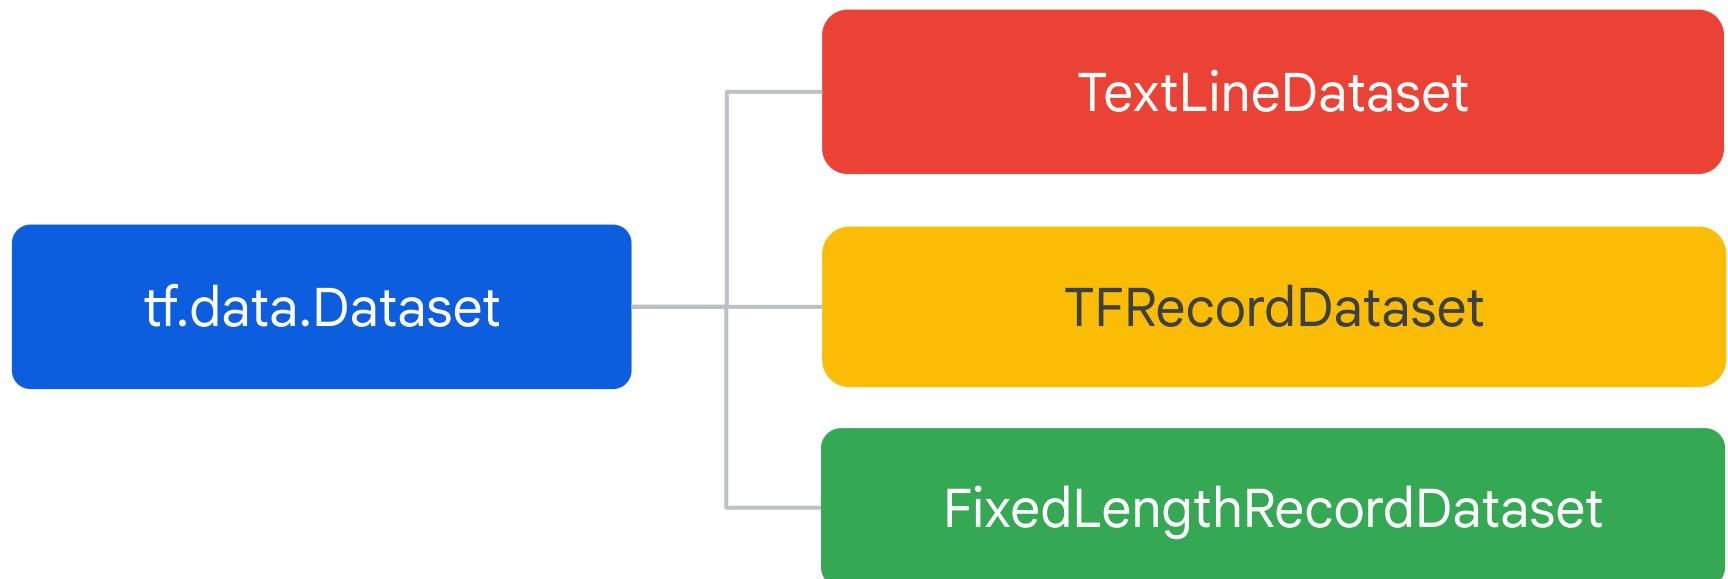
Datasets can be
created from
different **file formats**



Datasets can be
created from
different **file formats**

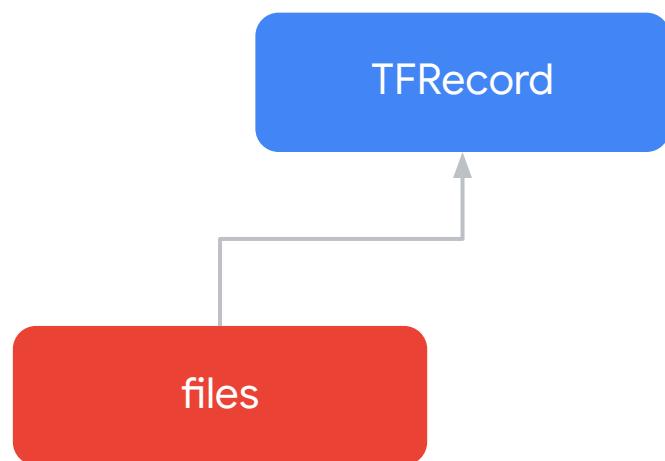


Datasets can be
created from
different **file formats**



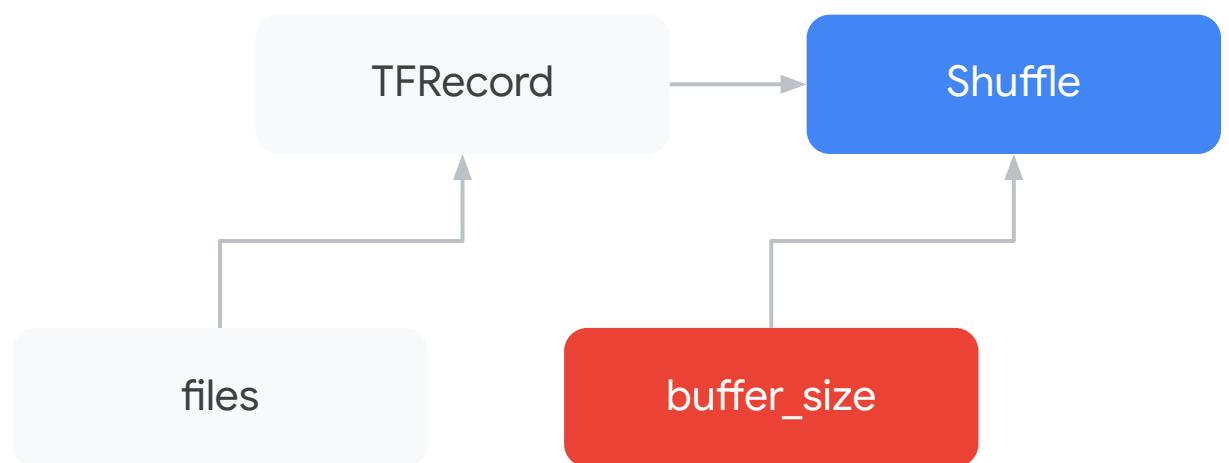
TFRecordDataset example

```
dataset = tf.data.TFRecordDataset(files)
```



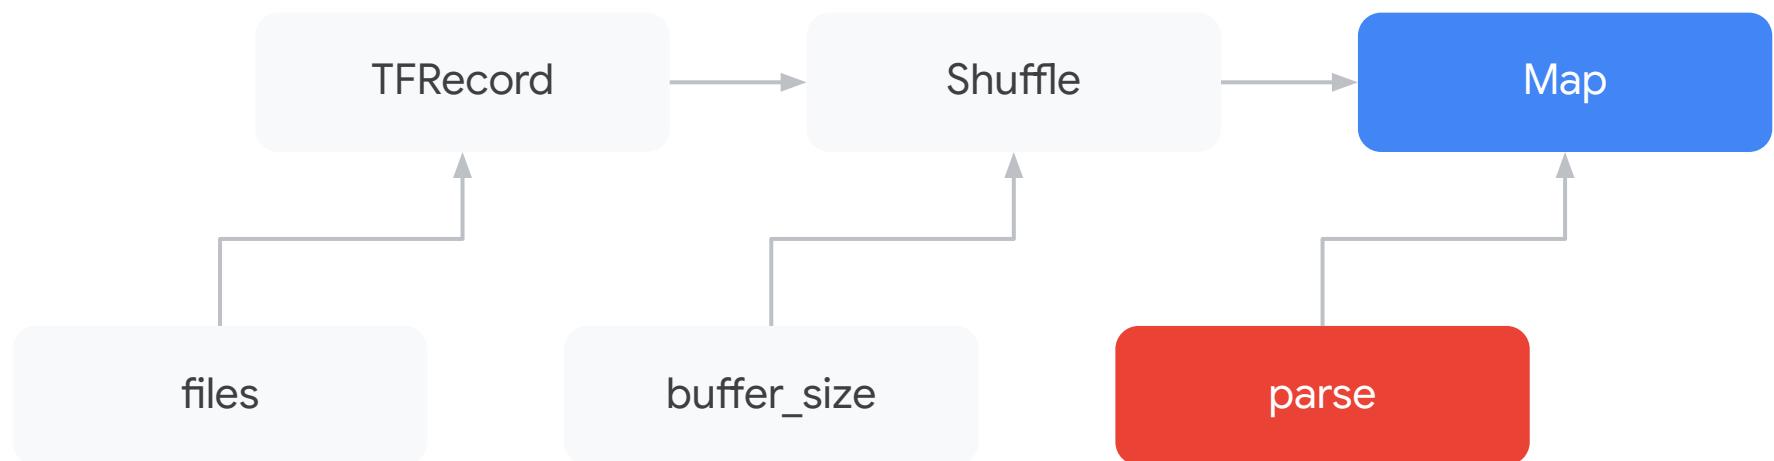
TFRecordDataset example

```
dataset = tf.data.TFRecordDataset(files)
dataset = dataset.shuffle(buffer_size=X)
```



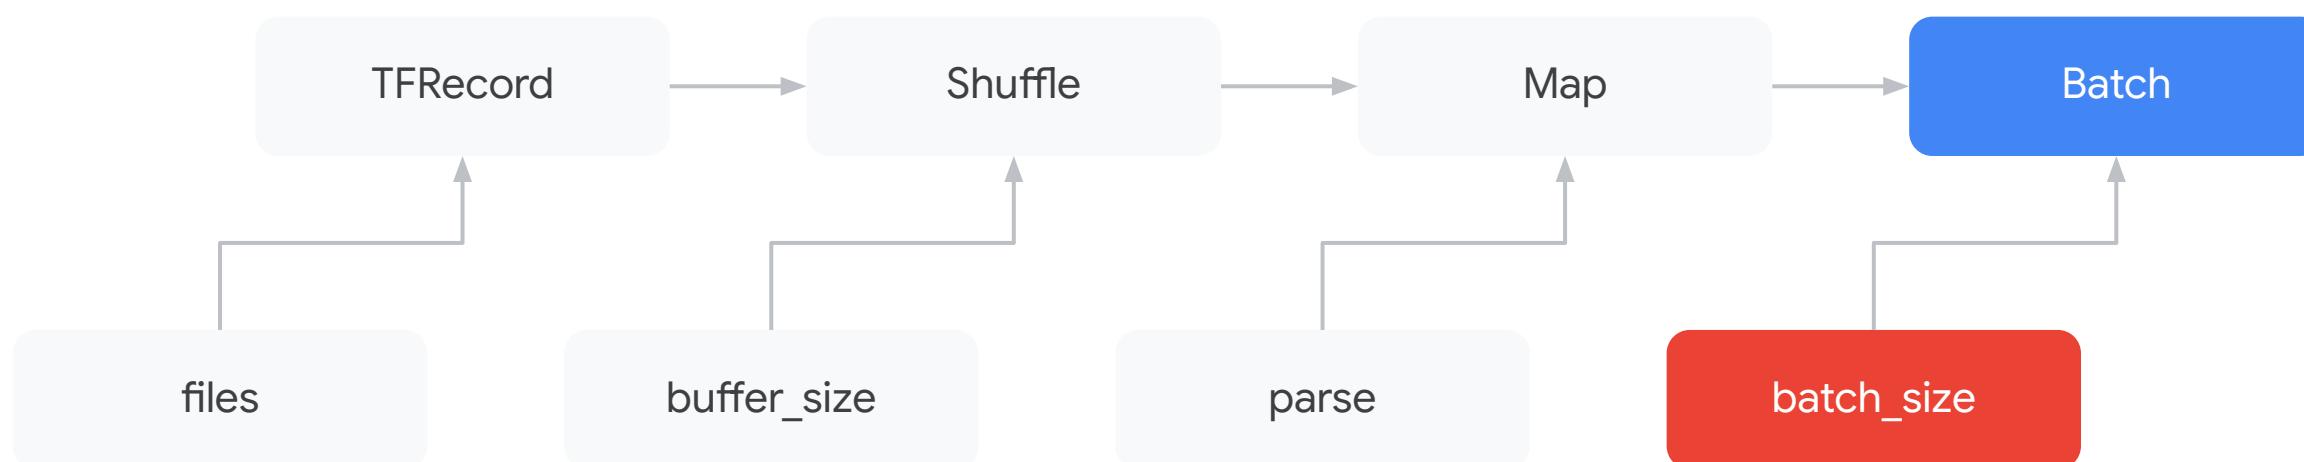
TFRecordDataset example

```
dataset = tf.data.TFRecordDataset(files)
dataset = dataset.shuffle(buffer_size=X)
dataset = dataset.map(lambda record: parse(record))
```



TFRecordDataset example

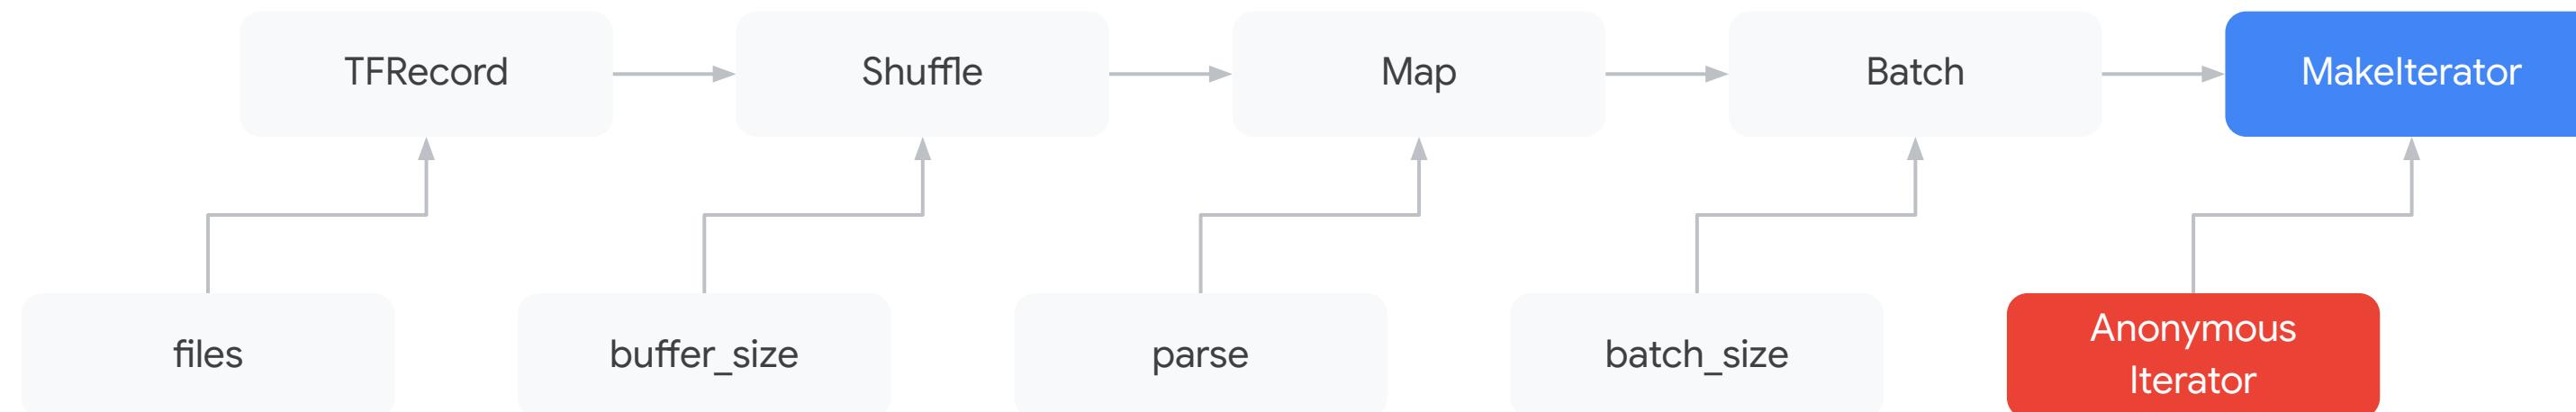
```
dataset = tf.data.TFRecordDataset(files)
dataset = dataset.shuffle(buffer_size=X)
dataset = dataset.map(lambda record: parse(record))
dataset = dataset.batch(batch_size=Y)
```



TFRecordDataset example

```
dataset = tf.data.TFRecordDataset(files)
dataset = dataset.shuffle(buffer_size=X)
dataset = dataset.map(lambda record: parse(record))
dataset = dataset.batch(batch_size=Y)
```

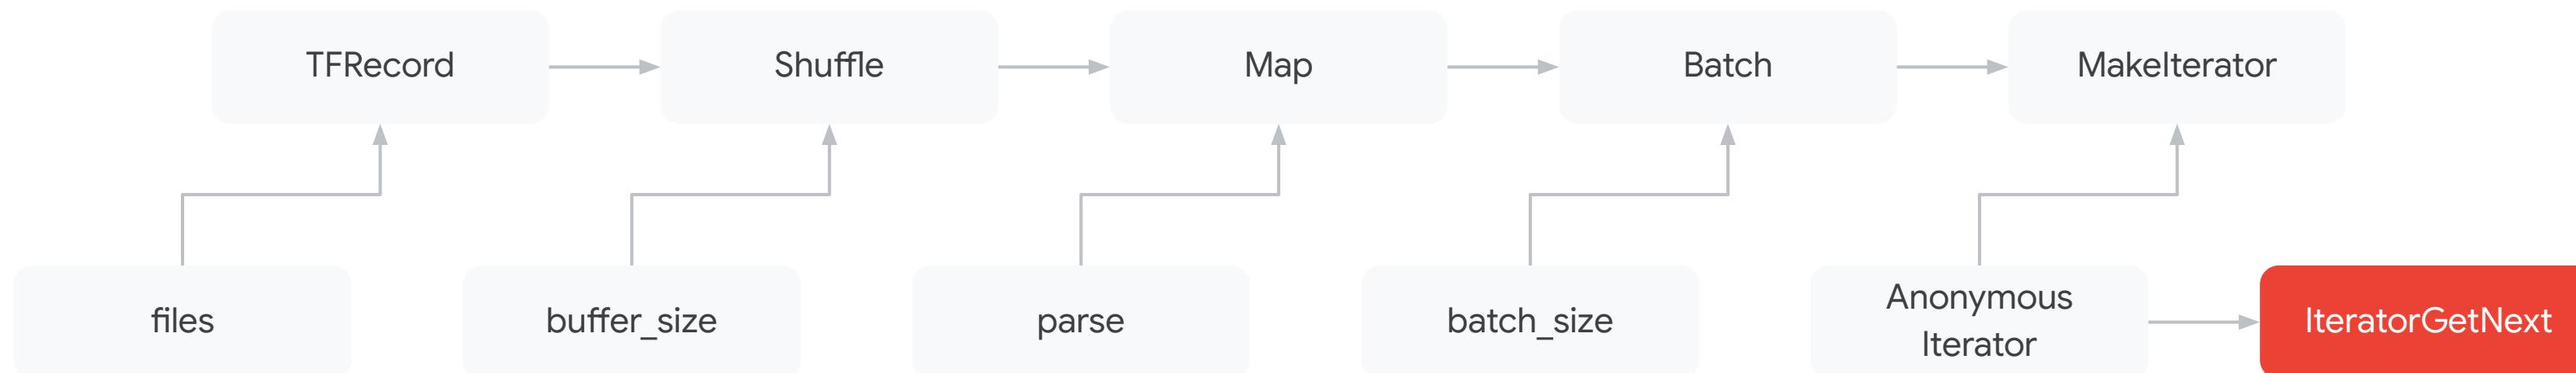
```
for element in dataset: # iter() is called
    ...
```



TFRecordDataset example

```
dataset = tf.data.TFRecordDataset(files)
dataset = dataset.shuffle(buffer_size=X)
dataset = dataset.map(lambda record: parse(record))
dataset = dataset.batch(batch_size=Y)

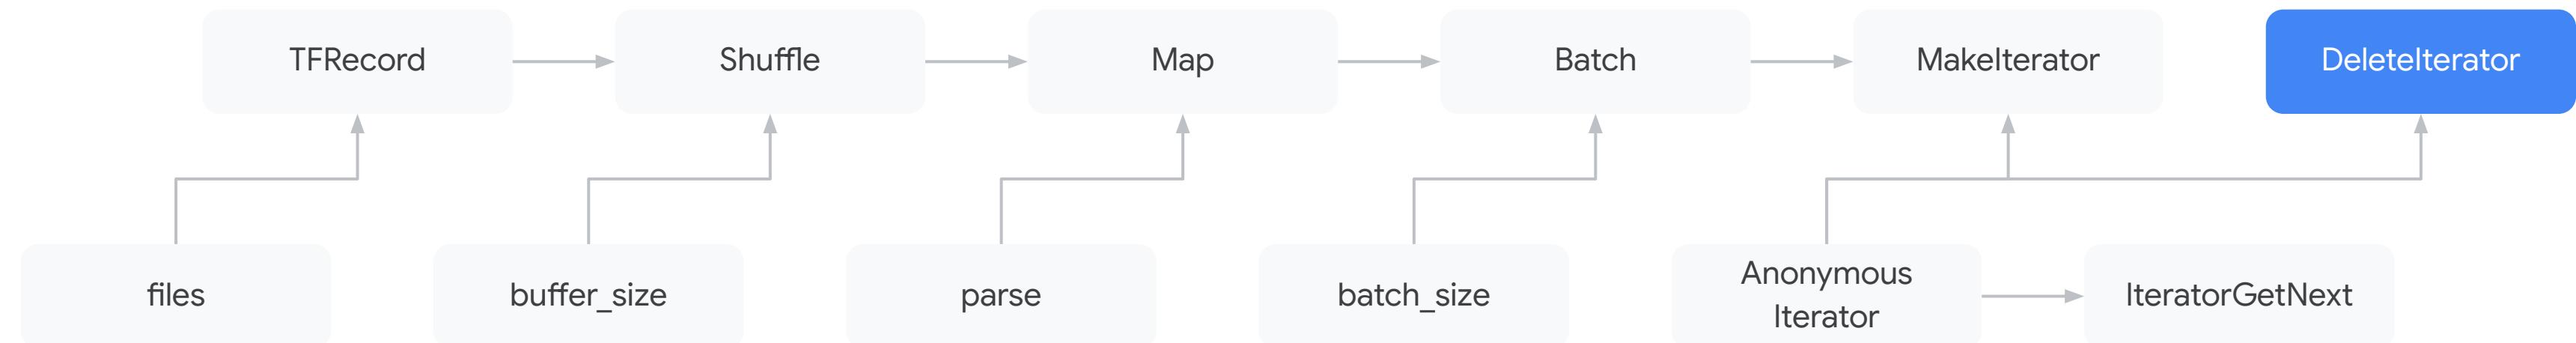
for element in dataset:
    ... # next() is called
```



TFRecordDataset example

```
dataset = tf.data.TFRecordDataset(files)
dataset = dataset.shuffle(buffer_size=X)
dataset = dataset.map(lambda record: parse(record))
dataset = dataset.batch(batch_size=Y)

for element in dataset:
    ... # iterator goes out of scope
```



Creating a dataset from in-memory tensors

```
def create_dataset(X, Y, epochs, batch_size):  
  
    dataset = tf.data.Dataset.from_tensor_slices((X, Y))  
  
    dataset = dataset.repeat(epochs).batch(batch_size, drop_remainder=True)  
  
    return dataset
```

X = [x_0, x_1, ..., x_n] Y = [y_0, y_1, ..., y_n]

The dataset is made of slices of (X, Y) along the 1st axis

Use `from_tensors()` or `from_tensor_slices()`

```
t = tf.constant([[4, 2], [5, 3]])
ds = tf.data.Dataset.from_tensors(t) # [[4, 2], [5, 3]]
```

```
t = tf.constant([[4, 2], [5, 3]])
ds = tf.data.Dataset.from_tensor_slices(t) # [4, 2], [5, 3]
```

Read one CSV file using TextLineDataset

```
def parse_row(records):
    cols = tf.decode_csv(records, record_defaults=[[0], ['house'], [0]])
    features = {'sq_footage': cols[0], 'type': cols[1]}
    label = cols[2]
    return features, label

dataset = "[line1, line2, etc.]"

def create_dataset(csv_file_path):
    dataset = tf.data.TextLineDataset(csv_file_path)
    dataset = dataset.map(parse_row)
    dataset = dataset.shuffle(1000).repeat(15).batch(128)
    return dataset
```

dataset = "[parse_row(line1),
parse_row(line2), etc.]"

property type

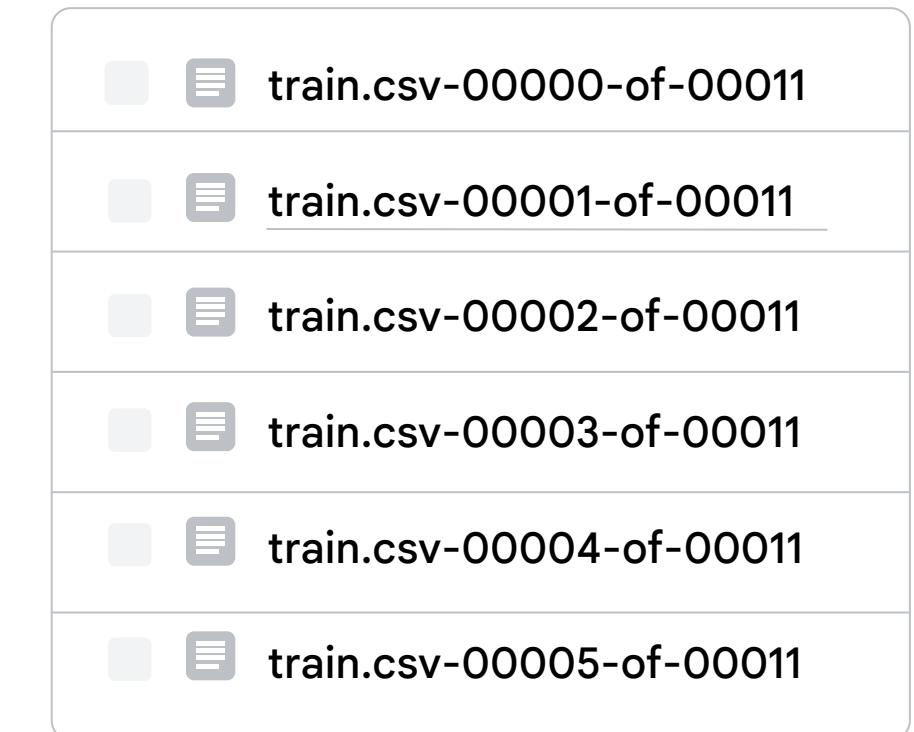
sq_footage	PRICE in K\$
1001, house,	501
2001, house,	1001
3001, house,	1501
1001, apt,	701
2001, apt,	1301
3001, apt,	1901
1101, house,	526
2101, house,	1026

Read a set of sharded CSV files using TextLineDataset

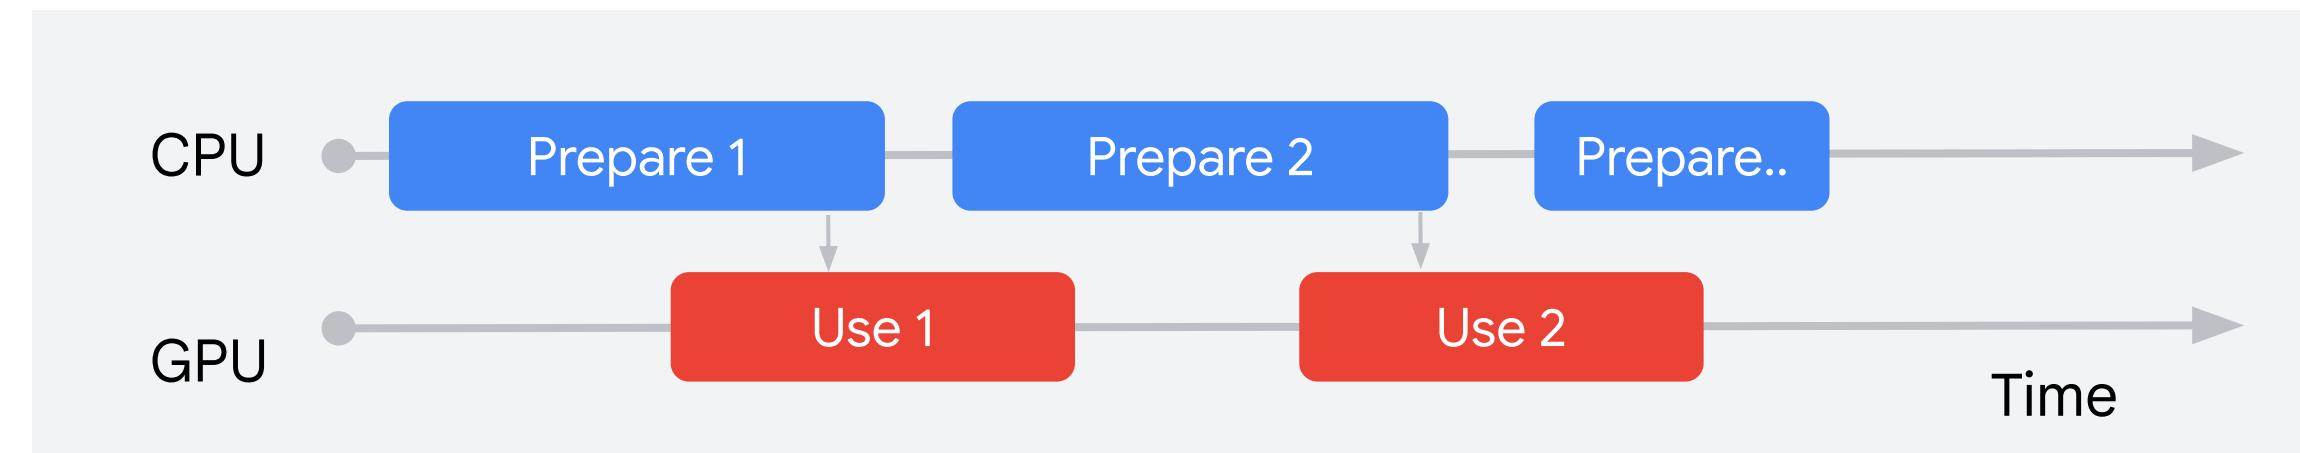
```
def parse_row(row):
    cols = tf.decode_csv(row, record_defaults=[[0],['house'],[0]])
    features = {'sq_footage': cols[0], 'type': cols[1]}
    label = cols[2] # price
    return features, label

def create_dataset(path):
    dataset = tf.data.Dataset.list_files(path) \
        .flat_map(tf.data.TextLineDataset) \
        .map(parse_row)

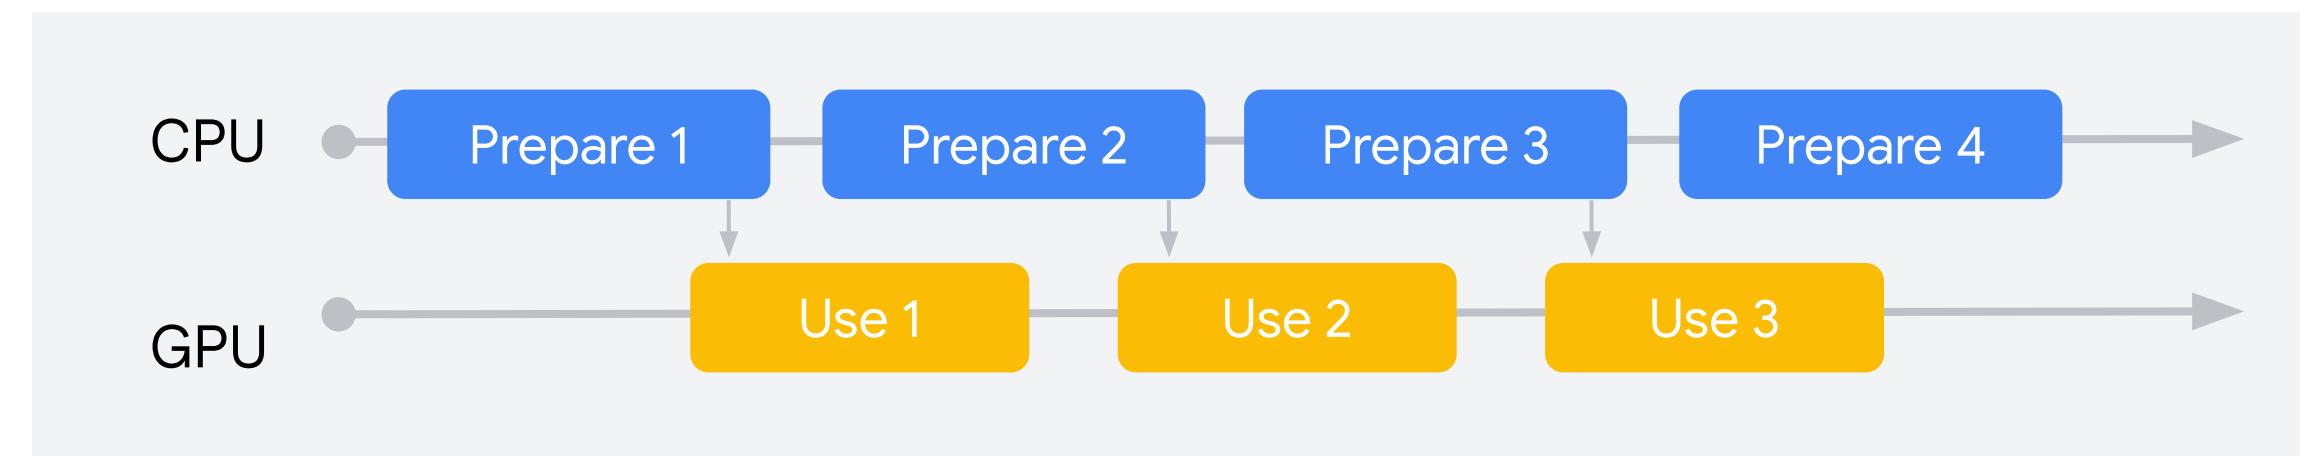
    dataset = dataset.shuffle(1000) \
        .repeat(15) \
        .batch(128)
    return dataset
```



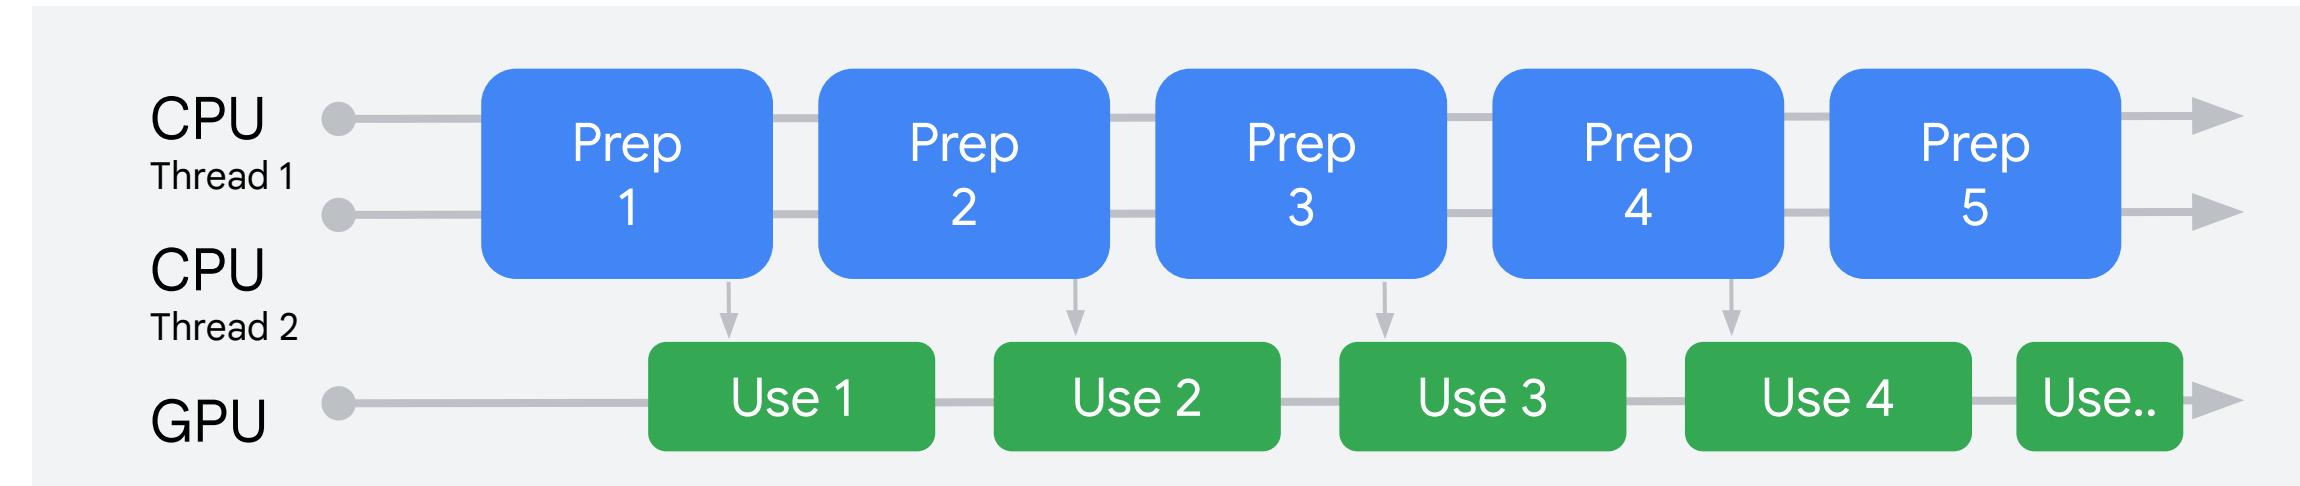
Without prefetching



With prefetching



With prefetching + multithreaded loading & preprocessing



The real benefit of Dataset is that you can do more than just ingest data

```
dataset = tf.data.TextLineDataset(filename)\\
    .skip(num_header_lines)\\
    .map(add_key)\\
    .map(decode_csv)\\
    .map(lambda feats, labels: preproc(feats), labels)
    .filter(is_valid)\\
    .cache()
```

Feature columns tell the model what inputs to expect

```
import tensorflow as tf  
  
featcols = [  
    tf.feature_column.numeric_column("sq_footage"),  
    tf.feature_column.categorical_column_with_vocabulary_list("type",  
        ["house", "apt"])  
]
```

...

Features

Under the hood:

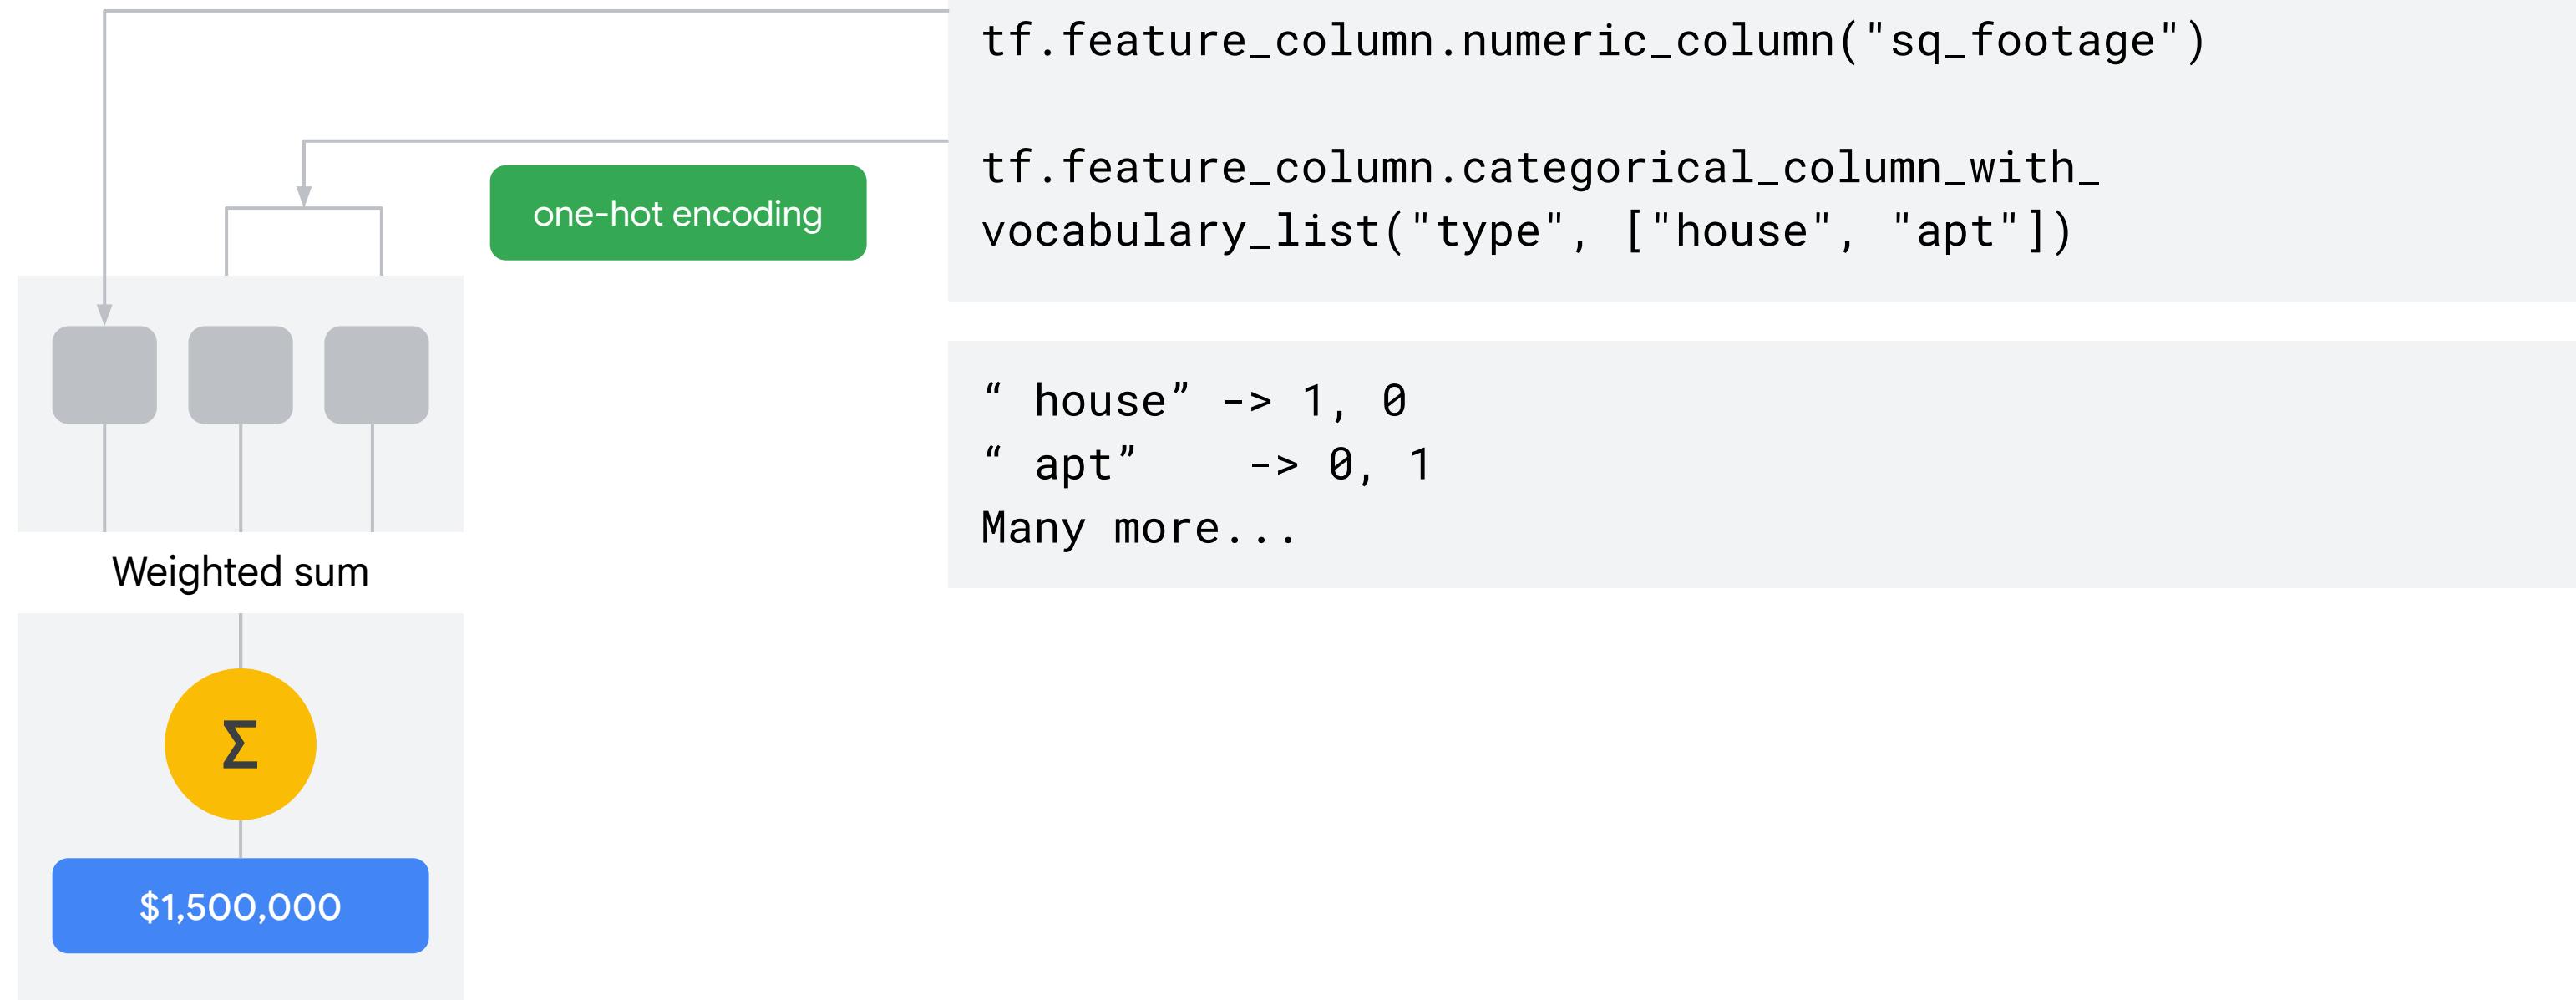
**Feature columns take
care of packing the
inputs into the input
vector of the model**

```
tf.feature_column.numeric_column("sq_footage")  
  
tf.feature_column.categorical_column_with_  
vocabulary_list("type", ["house", "apt"])
```

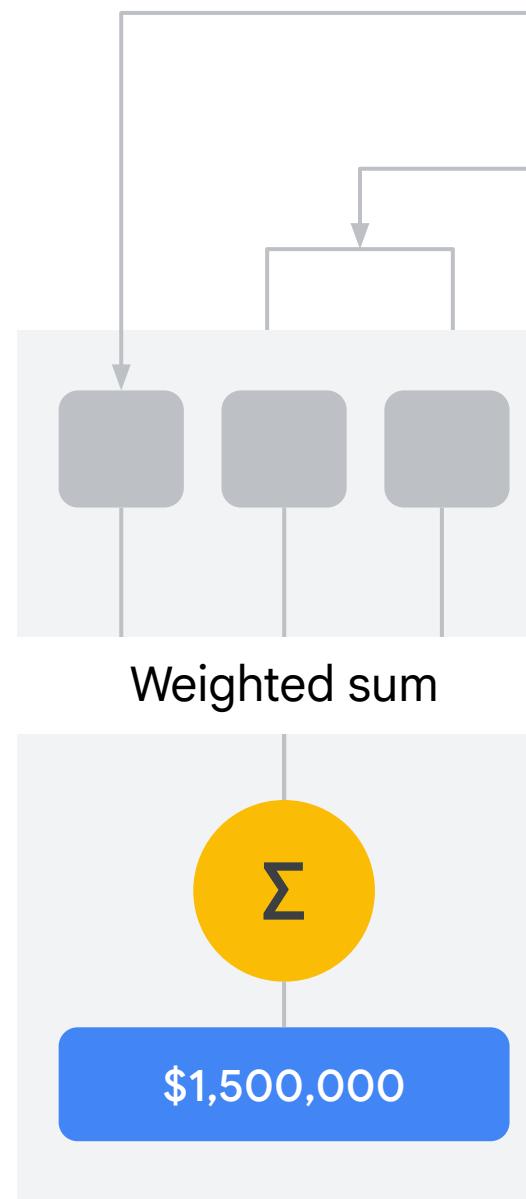
Under the hood: Feature columns take care of packing the inputs into the input vector of the model



Under the hood: Feature columns take care of packing the inputs into the input vector of the model



Under the hood: Feature columns take care of packing the inputs into the input vector of the model



```
tf.feature_column.numeric_column("sq_footage")
```

```
tf.feature_column.categorical_column_with_
vocabulary_list("type", ["house", "apt"])
```

“ house” -> 1, 0
“ apt” -> 0, 1
Many more...

```
tf.feature_column.bucketized_column(..)
tf.feature_column.embedding_column(..)
tf.feature_column.crossed_column(..)
tf.feature_column.categorical_column_with_hash_bucket(..)
...
```

fc.bucketized_column splits a numeric feature into categories based on numeric ranges

```
NBUCKETS = 16  
latbuckets = np.linspace(start=38.0, stop=42.0, num=NBUCKETS).tolist()  
lonbuckets = np.linspace(start=-76.0, stop=-72.0, num=NBUCKETS).tolist()
```

set up numeric ranges

```
fc_bucketized_plat = fc.bucketized_column(  
    source_column=fc.numeric_column("pickup_longitude"),  
    boundaries=lonbuckets)
```

create bucketized columns for pickup latitude and pickup longitude

```
fc_bucketized_plon = fc.bucketized_column(  
    source_column=fc.numeric_column("pickup_latitude"),  
    boundaries=latbuckets)
```

...

Representing feature columns as sparse vectors

These are all different ways to create a categorical column.

If you know the keys beforehand:

```
tf.feature_column.categorical_column_with_vocabulary_list('zipcode',  
    vocabulary_list = ['12345', '45678', '78900', '98723', '23451']),
```

If your data is already indexed; i.e., has integers in [0-N]:

```
tf.feature_column.categorical_column_with_identity('schoolsRatings',  
    num_buckets = 2)
```

If you don't have a vocabulary of all possible values:

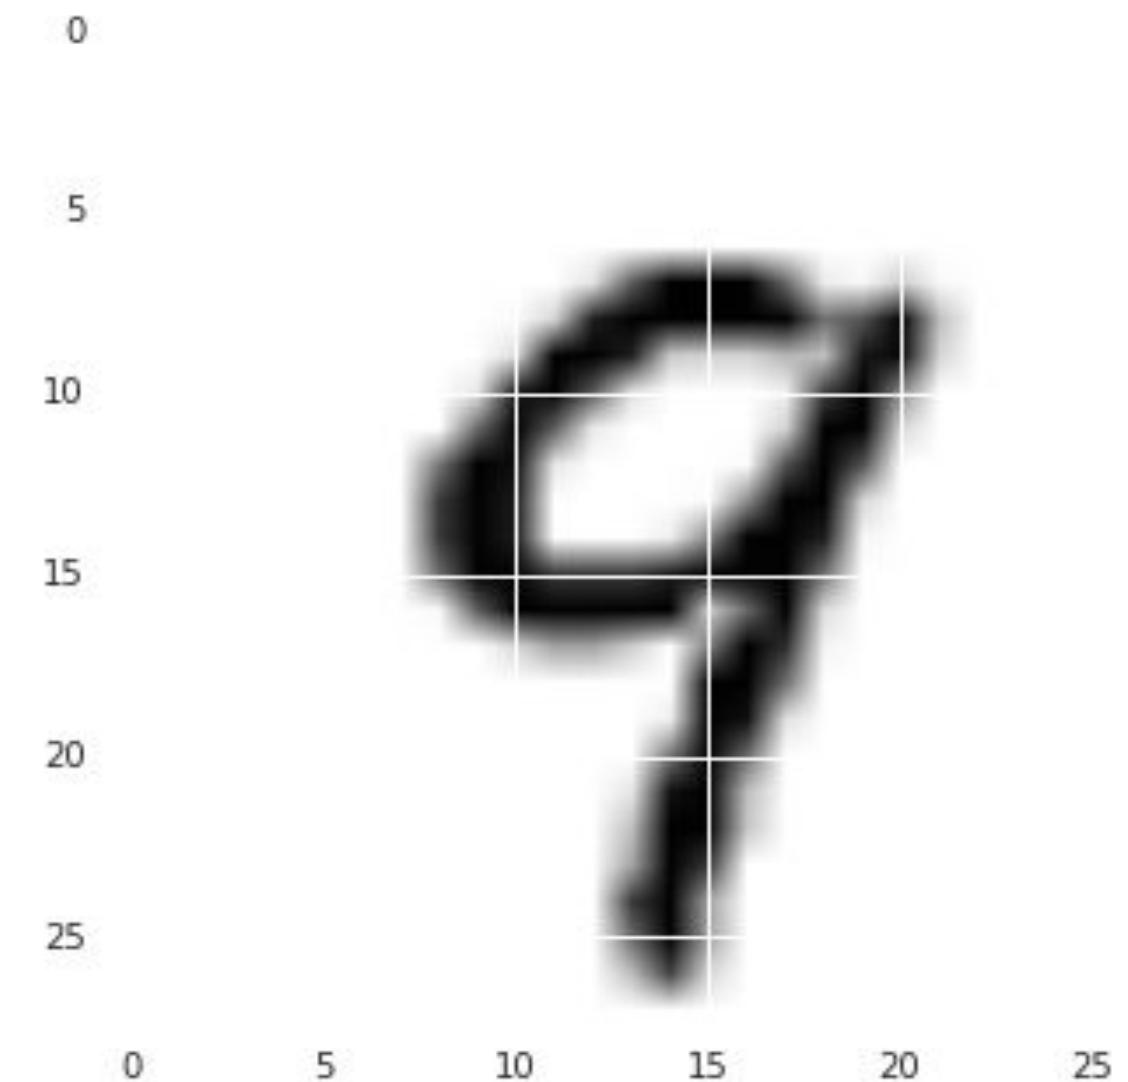
```
tf.feature_column.categorical_column_with_hash_bucket('nearStoreID',  
    hash_bucket_size = 500)
```

`fc.embedding_column` represents data as a lower-dimensional, dense vector

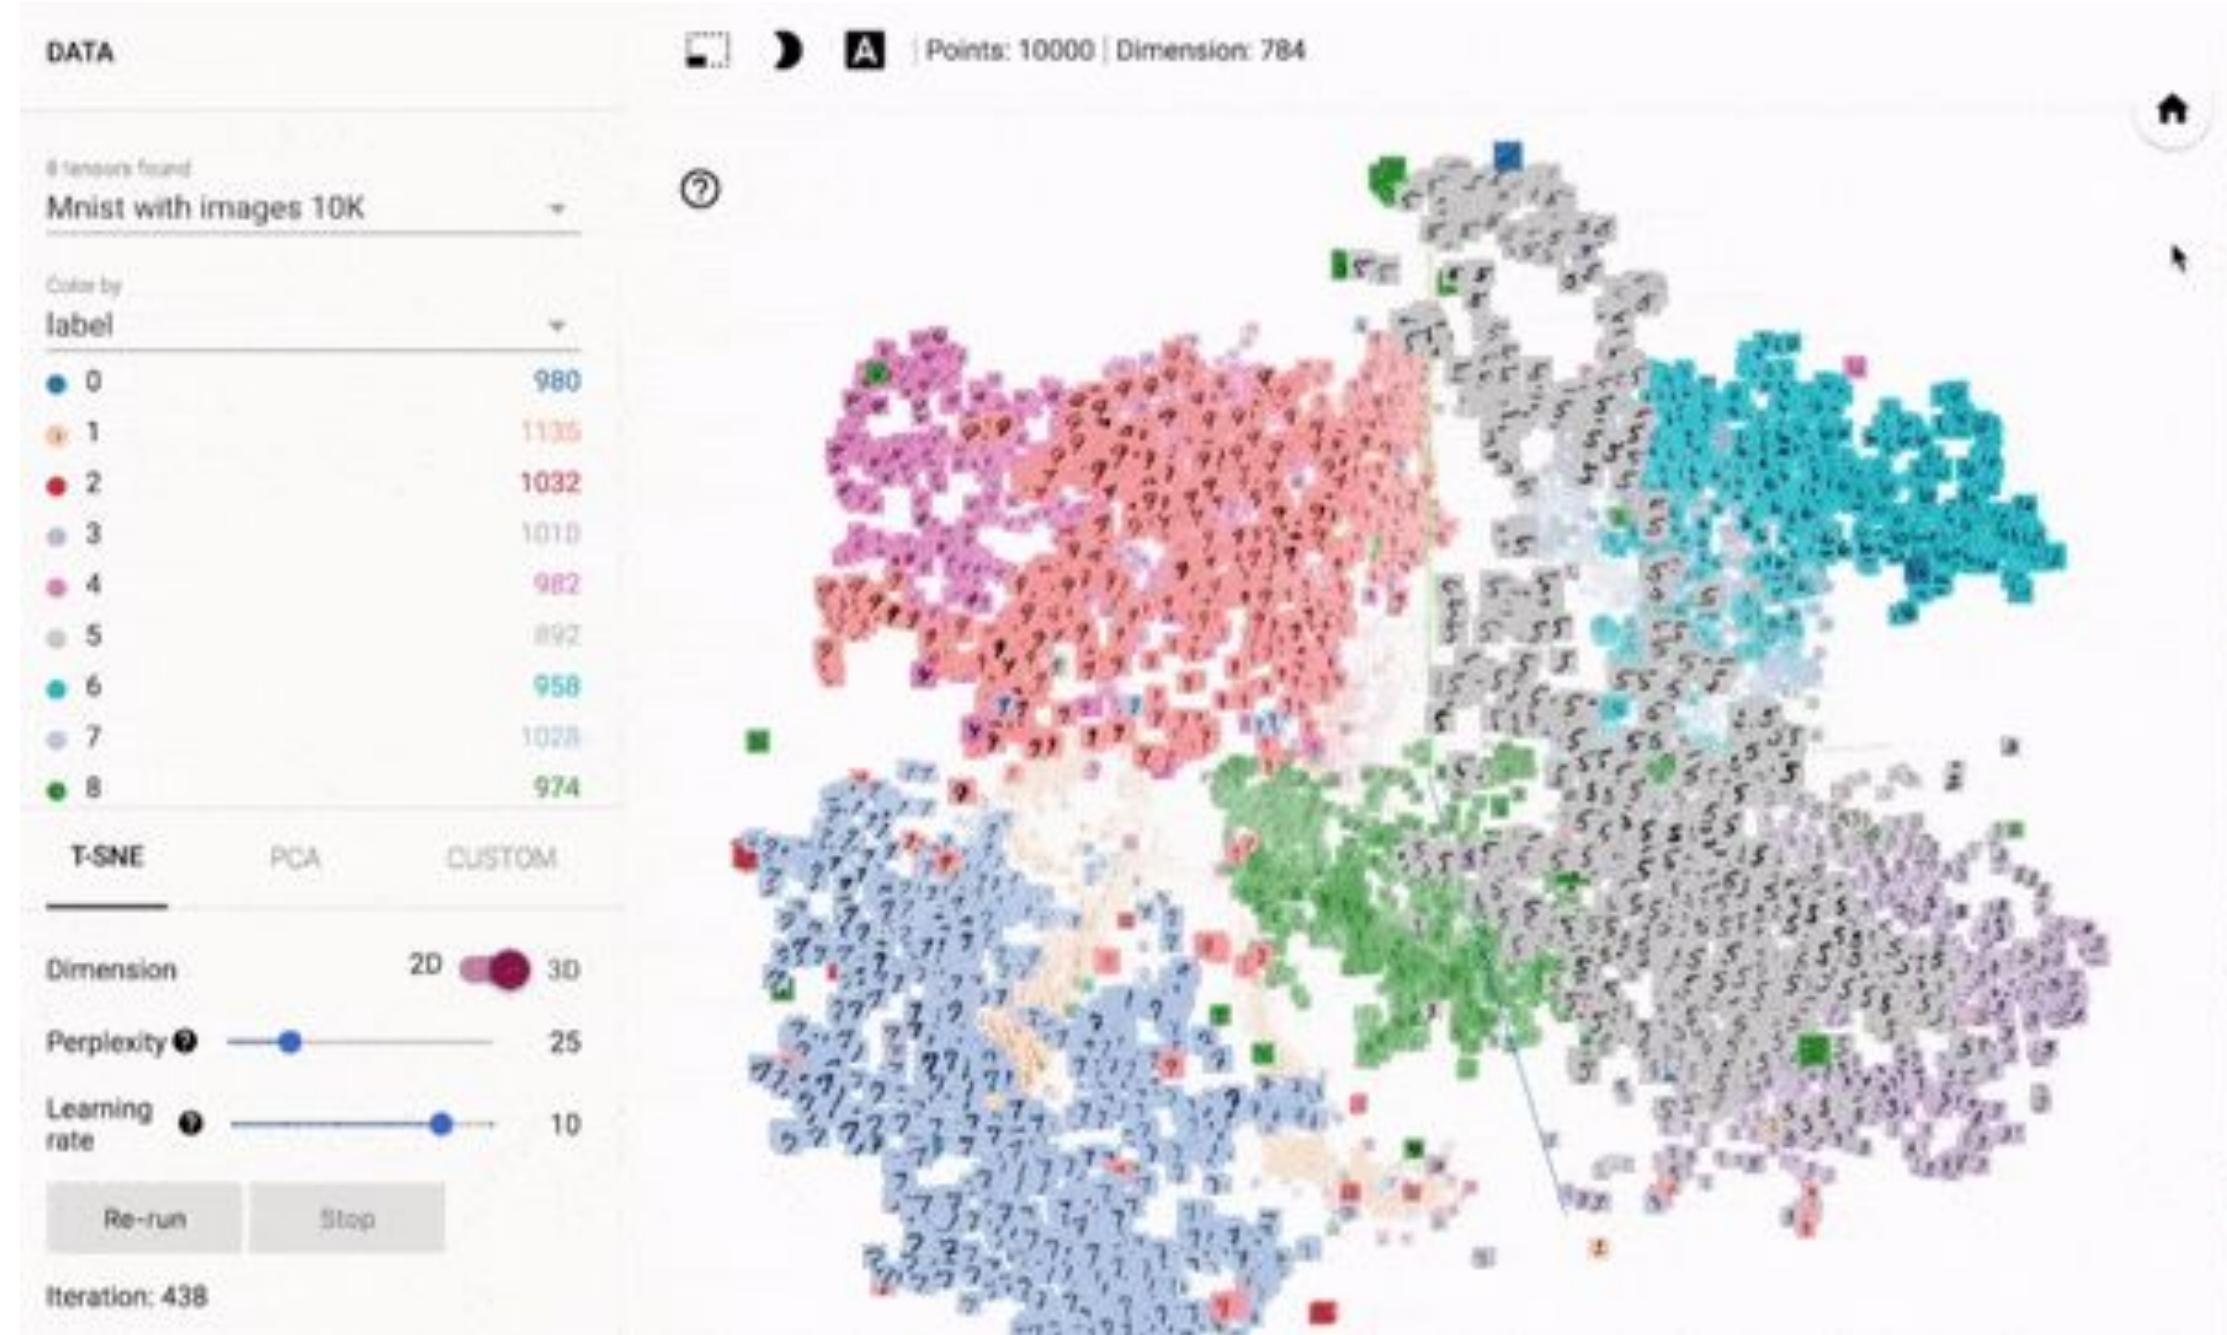
```
fc_ploc = fc.embedding_column(categorical_column=fc_crossed_ploc,  
                               dimension=3)
```

lower dimensional, dense vector in which each cell contains a number, not just 0 or 1

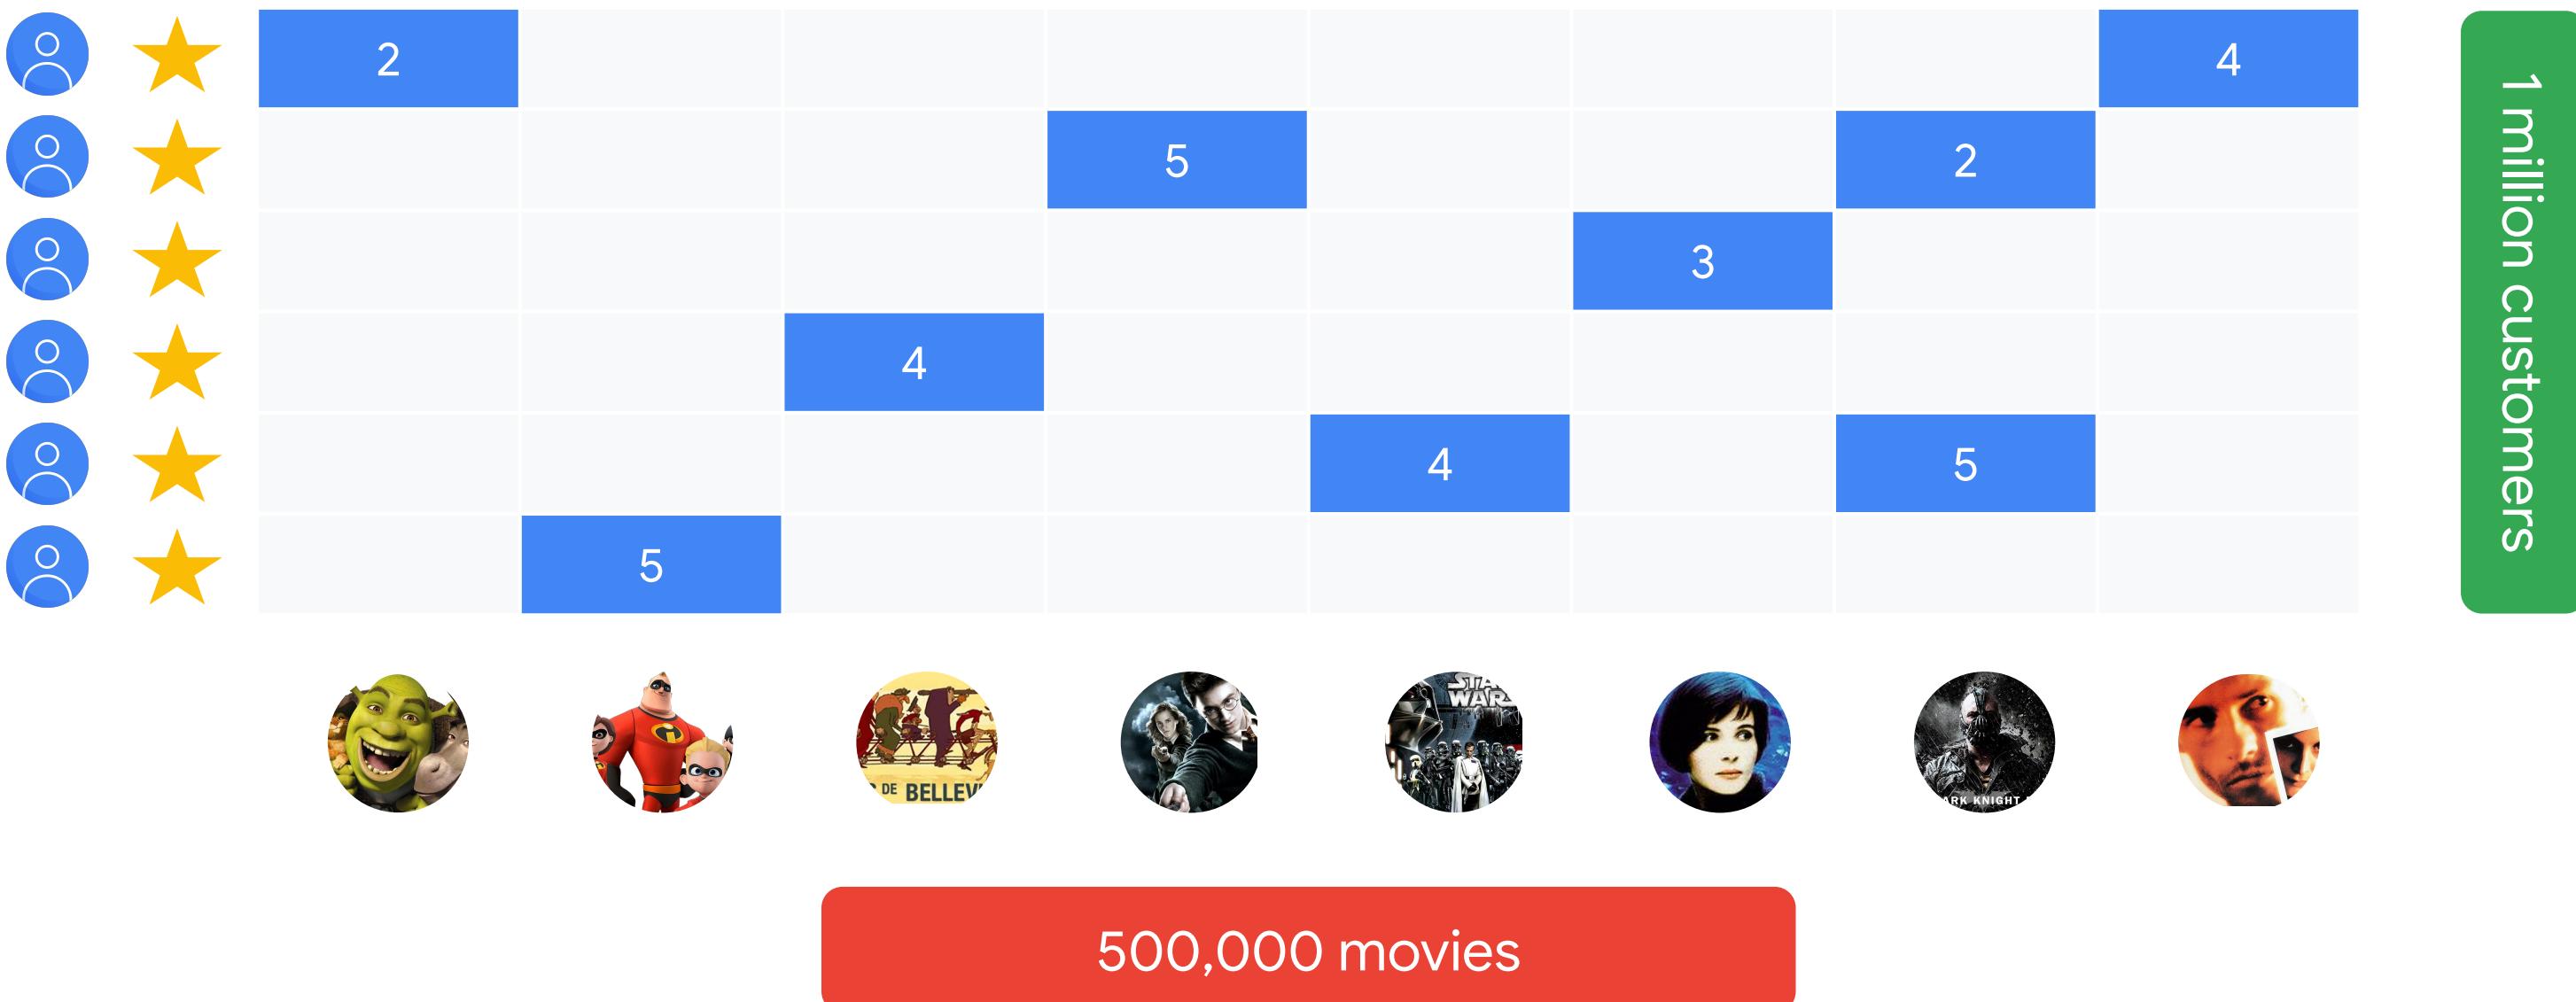
**How can we visually cluster
10,000 variations of
handwritten digits to look
for similarities?
Embeddings!**



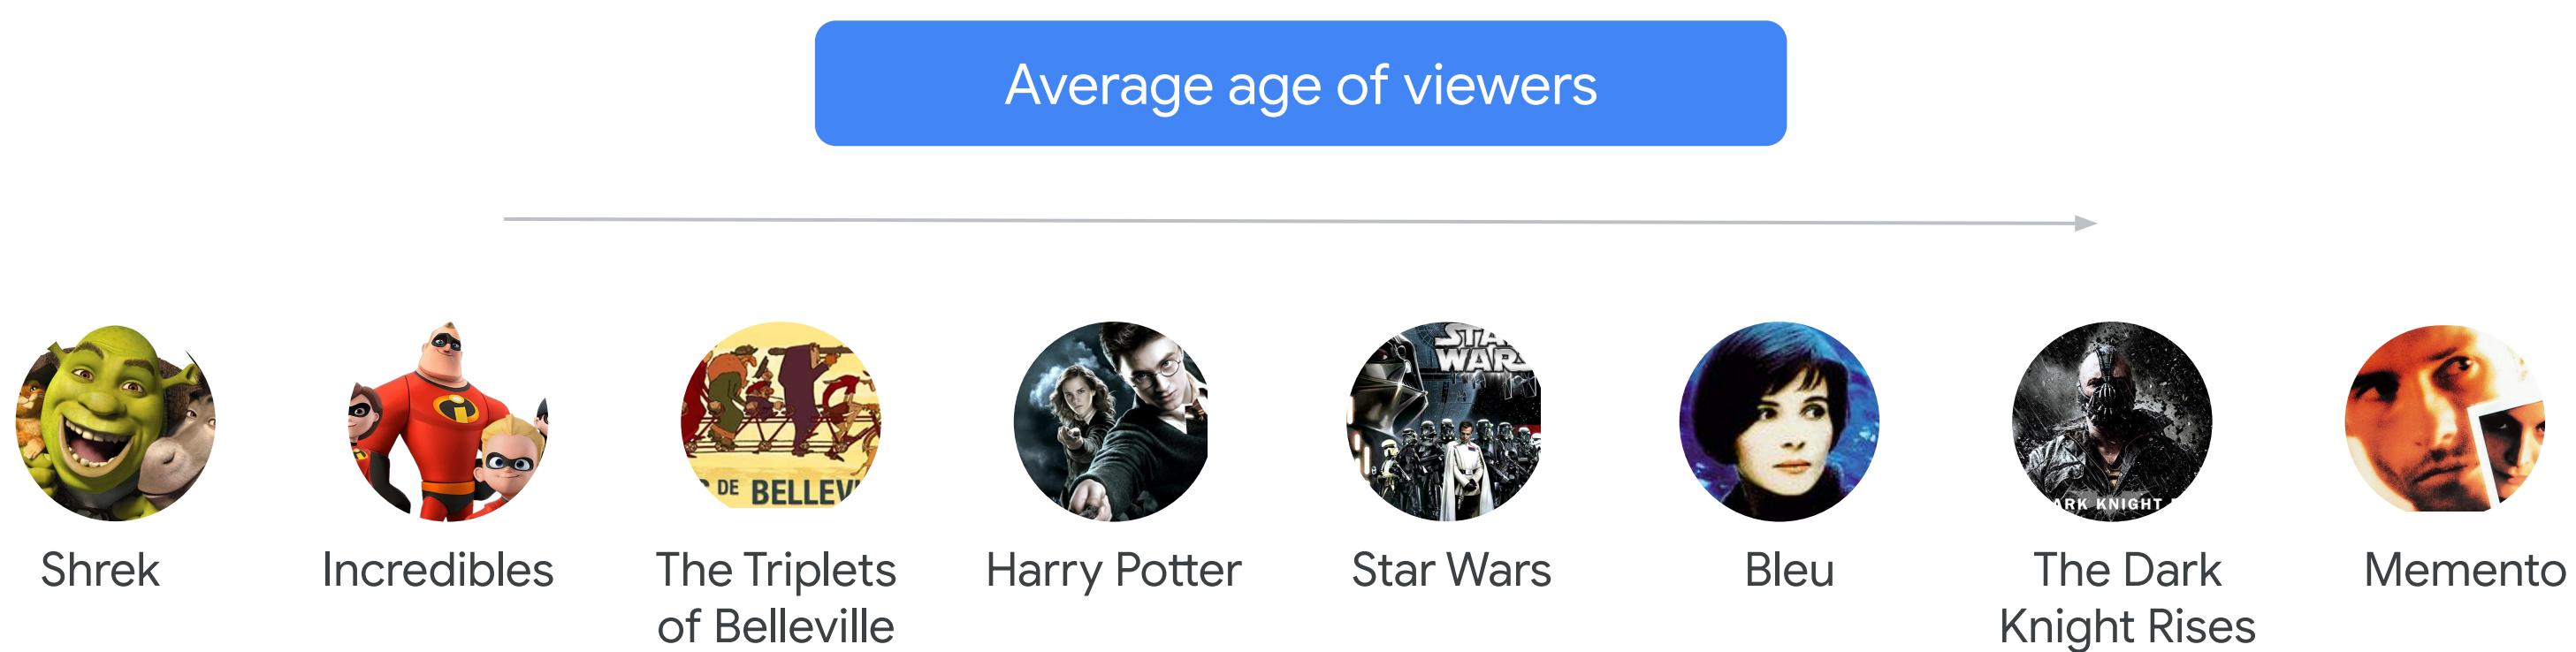
Embeddings are everywhere in modern machine learning



How do you recommend movies to customers?



One approach is to organize movies by similarity (1D)



Using a second dimension gives us more freedom in organizing movies by similarity



Shrek



Incredibles



The Triplets
of Belleville



Harry Potter

↑



Star Wars



The Dark
Knight Rises



Bleu

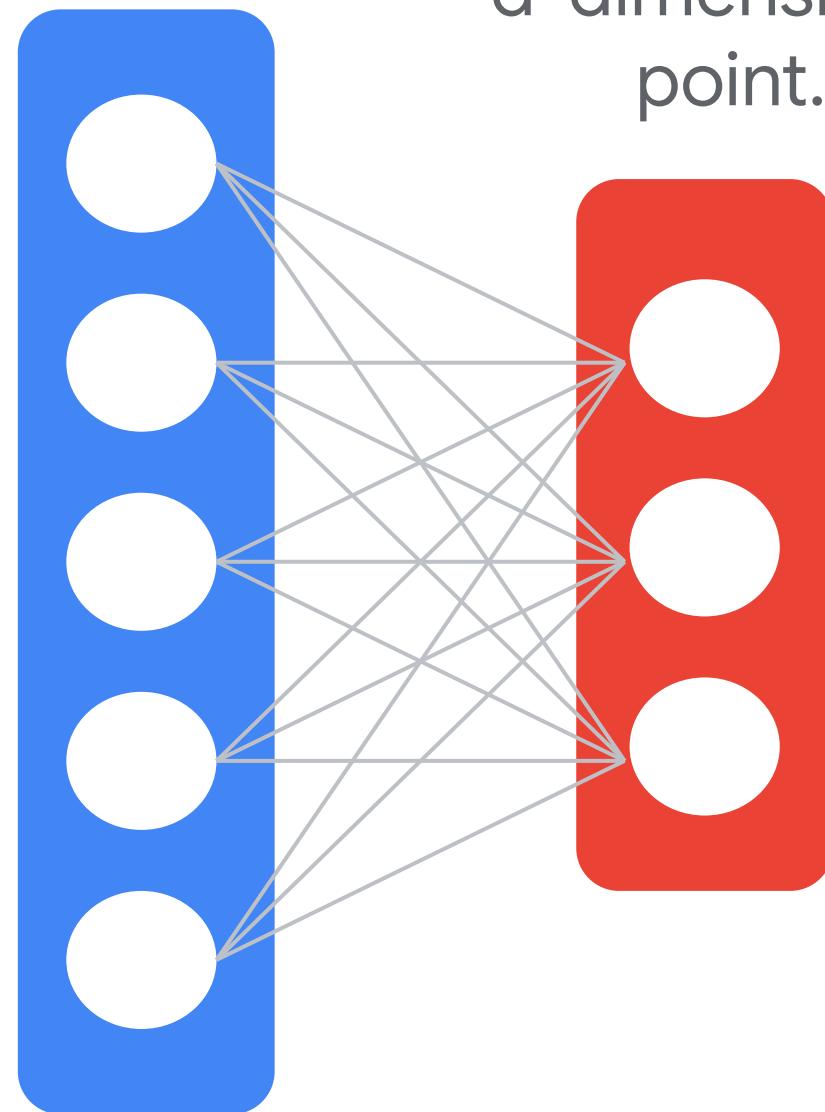


Memento

Gross ticket sales

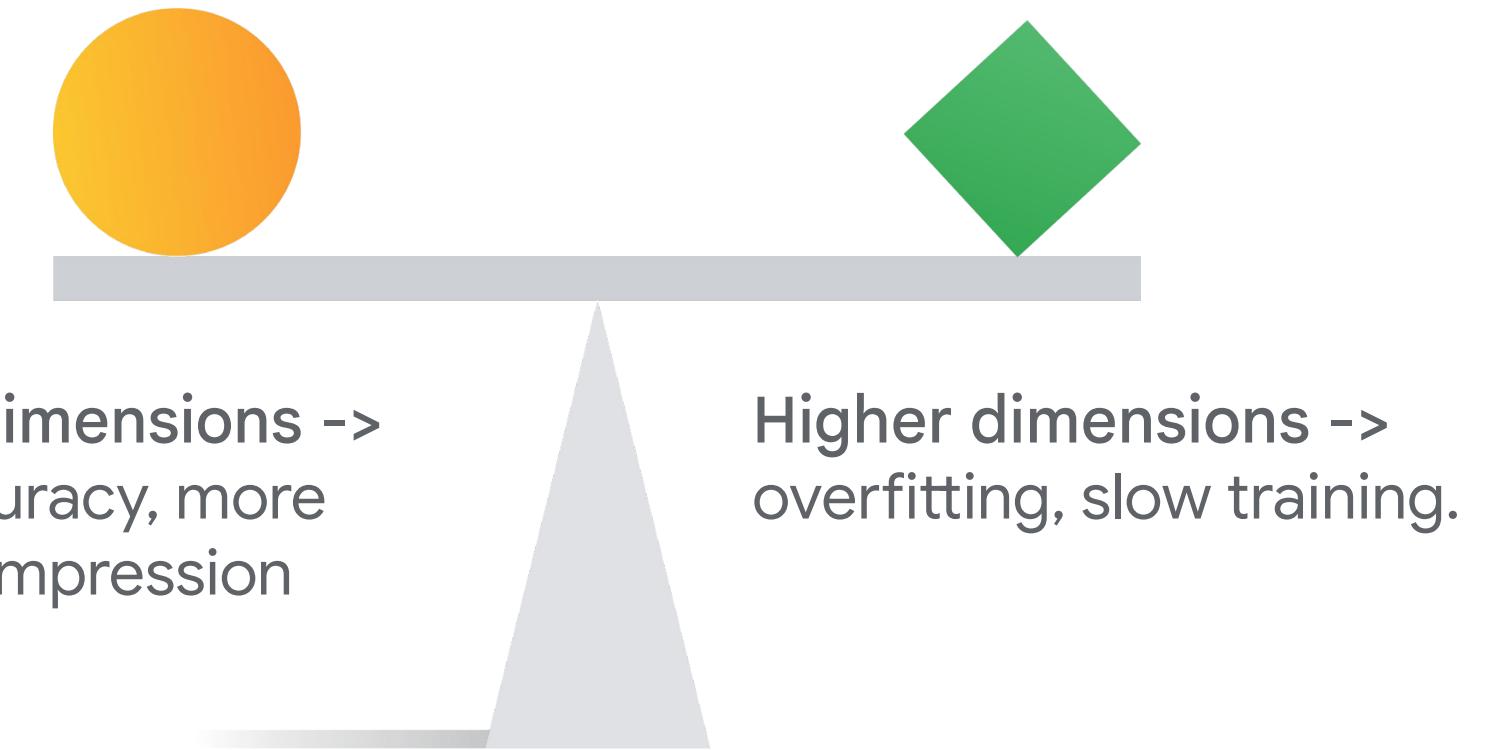
**A d-dimensional
embedding assumes that
user interest in movies can
be approximated by
d aspects**

Input has N
dimensions.

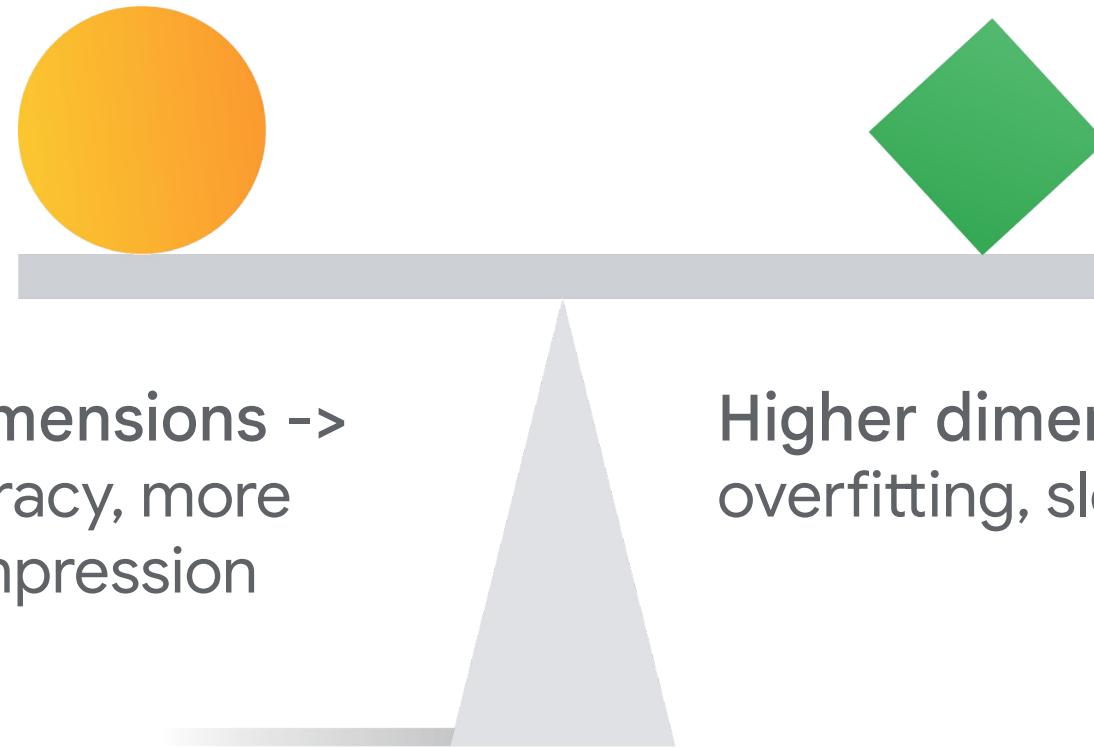


Each input is
reduced to a
d-dimensional
point.

A good starting
point for number
of embedding
dimensions



A good starting point for number of embedding dimensions



Lower dimensions ->
less accuracy, more
lossy compression

Higher dimensions ->
overfitting, slow training.

$$\text{Dimensions} \approx \sqrt[4]{\text{Possible values}}$$

Empirical tradeoff.

fc.crossed_column enables a model to learn separate for combination of features

```
fc_crossed_ploc = fc.crossed_column([fc_bucketized_plat, fc_bucketized_plon],  
                                     hash_bucket_size=NBUCKETS * NBUCKETS)
```

crossed_column is backed by a hashed_column, so you must set the size of the hash bucket

Training input data requires dictionary of features and a label

```
def features_and_label():
    # sq_footage and type
    features = {"sq_footage": [ 1000,      2000,      3000,      1000,      2000,      3000],
                 "type":          ["house", "house", "house", "apt", "apt", "apt"]}

    # prices in thousands
    labels = [ 500,      1000,      1500,      700,      1300,      1900]

    return features, labels
```

Create input pipeline using tf.data

```
def create_dataset(pattern, batch_size=1, mode=tf.estimator.ModeKeys.EVAL):
    dataset = tf.data.experimental.make_csv_dataset(
        pattern, batch_size, CSV_COLUMNS, DEFAULTS)
    dataset = dataset.map(features_and_labels)
    if mode == tf.estimator.ModeKeys.TRAIN:
        dataset = dataset.shuffle(buffer_size=1000).repeat()
    # take advantage of multi-threading; 1=AUTOTUNE
    dataset = dataset.prefetch(1)

    return dataset
```

Use DenseFeatures layer to input feature columns to the Keras model

```
feature_columns = [...]  
  
feature_layer = tf.keras.layers.DenseFeatures(feature_columns)  
  
model = tf.keras.Sequential([  
    feature_layer,  
    layers.Dense(128, activation='relu'),  
    layers.Dense(128, activation='relu'),  
    layers.Dense(1, activation='linear')  
])  
  
...
```

What about compiling and training the Keras model?

After your dataset is created, passing it into the Keras model for training is simple:

model.fit()

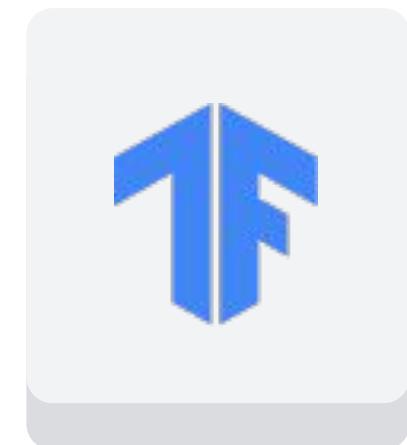
You will learn and practice this later after first mastering dataset manipulation!

Data is the a crucial component
of a machine learning model

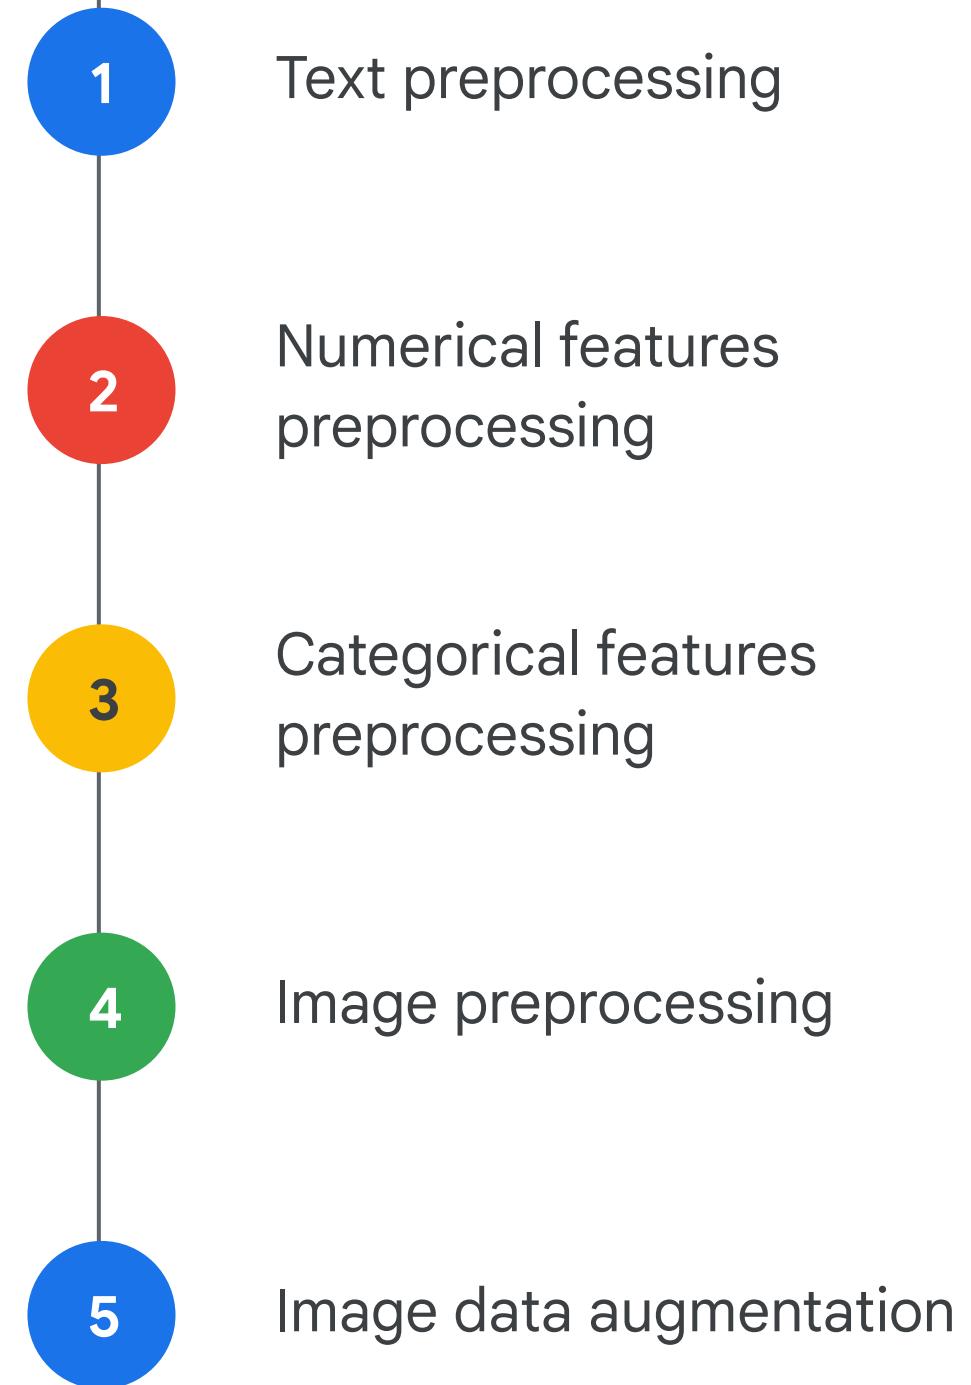
Scaling data preprocessing with `tf.data` and Keras preprocessing layers

Data preprocessing

- You can build and export end-to-end models that accept raw images or raw structured data as input.
- Models handle feature normalization or feature value indexing on their own.



Keras preprocessing layers



Text features preprocessing

```
tf.keras.layers.TextVectorization(  
    max_tokens=None,  
    standardize="lower_and_strip_punctuation",  
    split="whitespace",  
    ngrams=None,  
    output_mode="int",  
    output_sequence_length=None,  
    pad_to_max_tokens=False,  
    vocabulary=None,  
    **kwargs  
)
```

Text vectorization layer

[tf.keras.layers.TextVectorization:](#)

turns raw strings into an encoded representation that can be read by an [Embedding](#) layer or [Dense](#) layer.

Numerical features preprocessing

```
tf.keras.layers.Normalization(axis=-1, mean=None, variance=None, **kwargs)
```

Normalization class

Feature-wise normalization of the data

[tf.keras.layers.Normalization:](#)

performs feature-wise
normalization of input features.

Numerical preprocessing

```
tf.keras.layers.Discretization(  
    bin_boundaries=None, num_bins=None, epsilon=0.01, **kwargs  
)
```

Discretization class

Buckets data into discrete ranges.

[tf.keras.layers.Discretization:](#)

turns continuous numerical features into bucket data with discrete ranges.

Categorical features preprocessing

[Tf.keras.layers. CategoryEncoding](#)

Turns integer categorical features into one-hot, multi-hot, or count dense representations.

[Tf.keras.layers. Hashing](#)

Performs categorical feature hashing, also known as the "hashing trick."

[Tf.keras.layers. StringLookup](#)

Turns string categorical values into an encoded representation that can be read by an Embedding layer or Dense layer.

[Tf.keras.layers. IntegerLookup](#)

Turns integer categorical values into an encoded representation that can be read by an Embedding layer or Dense layer.

Categorical features preprocessing

[Tf.keras.layers. CategoryEncoding](#)

Turns integer categorical features into one-hot, multi-hot, or count dense representations.

[Tf.keras.layers. Hashing](#)

Performs categorical feature hashing, also known as the "hashing trick."

[Tf.keras.layers. StringLookup](#)

Turns string categorical values into an encoded representation that can be read by an Embedding layer or Dense layer.

[Tf.keras.layers. IntegerLookup](#)

Turns integer categorical values into an encoded representation that can be read by an Embedding layer or Dense layer.

Categorical features preprocessing

[Tf.keras.layers. CategoryEncoding](#)

Turns integer categorical features into one-hot, multi-hot, or count dense representations.

[Tf.keras.layers. Hashing](#)

Performs categorical feature hashing, also known as the "hashing trick."

[Tf.keras.layers. StringLookup](#)

Turns string categorical values into an encoded representation that can be read by an Embedding layer or Dense layer.

[Tf.keras.layers. IntegerLookup](#)

Turns integer categorical values into an encoded representation that can be read by an Embedding layer or Dense layer.

Categorical features preprocessing

[Tf.keras.layers. CategoryEncoding](#)

Turns integer categorical features into one-hot, multi-hot, or count dense representations.

[Tf.keras.layers. Hashing](#)

Performs categorical feature hashing, also known as the "hashing trick."

[Tf.keras.layers. StringLookup](#)

Turns string categorical values into an encoded representation that can be read by an Embedding layer or Dense layer.

[Tf.keras.layers. IntegerLookup](#)

Turns integer categorical values into an encoded representation that can be read by an Embedding layer or Dense layer.

The adapt() method

Stateful preprocessing layers that compute based on training data:/

01

TextVectorization:

Holds a mapping between string tokens and integer indices.

02

StringLookup and IntegerLookup:

Hold a mapping between input values and integer indices.

03

Normalization:

Holds the mean and standard deviation of the features.

04

Discretization:

Holds information about value bucket boundaries.

Important note: These layers are non-trainable. Their state is not set during training; it must be set before training.

PetFinder Dataset

Column	Description	Feature Type	Data Type
Type	Type of animal (Dog, Cat)	Categorical	string
Age	Age of the pet	Numerical	integer
Breed1	Primary breed of the pet	Categorical	string
Color1	Color 1 of pet	Categorical	string
Color2	Color 2 of pet	Categorical	string
MaturitySize	Size at maturity	Categorical	string
FurLength	Fur length	Categorical	string
Vaccinated	Pet has been vaccinated	Categorical	string
Sterilized	Pet has been sterilized	Categorical	string
Health	Health Condition	Categorical	string
Fee	Adoption Fee	Numerical	integer
Description	Profile write-up for this pet	Text	string
PhotoAmt	Total uploaded photos for this pet	Numerical	integer
AdoptionSpeed	Speed of adoption	Classification	integer

Classify structured data using Keras preprocessing layers:

Dataset

Column	Description	Feature Type	Data Type
Type	Type of animal (Dog, Cat)	Categorical	string
Age	Age of the pet	Numerical	integer
Breed1	Primary breed of the pet	Categorical	string
Color1	Color 1 of pet	Categorical	string
Color2	Color 2 of pet	Categorical	string
MaturitySize	Size at maturity	Categorical	string
FurLength	Fur length	Categorical	string
Vaccinated	Pet has been vaccinated	Categorical	string
Sterilized	Pet has been sterilized	Categorical	string
Health	Health Condition	Categorical	string
Fee	Adoption Fee	Numerical	integer
Description	Profile write-up for this pet	Text	string
PhotoAmt	Total uploaded photos for this pet	Numerical	integer
AdoptionSpeed	Speed of adoption	Classification	integer

Import the preprocessing library

```
[ ] !pip install -q sklearn  
  
[ ] import numpy as np  
import pandas as pd  
import tensorflow as tf  
  
from sklearn.model_selection import train_test_split  
from tensorflow.keras import layers  
from tensorflow.keras.layers.experimental import preprocessing
```

Categorical columns

Column	Description	Feature Type	Data Type
Type	Type of animal (Dog, Cat)	Categorical	string

```
def get_category_encoding_layer(name, dataset, dtype, max_tokens=None):
    # Create a StringLookup layer which will turn strings into integer indices
    if dtype == 'string':
        index = preprocessing.StringLookup(max_tokens=max_tokens)
    else:
        index = preprocessing.IntegerLookup(max_tokens=max_tokens)

    # Prepare a Dataset that only yields our feature
    feature_ds = dataset.map(lambda x, y: x[name])

    # Learn the set of possible values and assign them a fixed integer index.
    index.adapt(feature_ds)

    # Create a Discretization for our integer indices.
    encoder = preprocessing.CategoryEncoding(num_tokens=index.vocabulary_size())

    # Apply one-hot encoding to our indices. The lambda function captures the
    # layer so we can use them, or include them in the functional model later.
    return lambda feature: encoder(index(feature))
```

Categorical columns

tf.keras.layers.StringLookup:
Turns string categorical values into an encoded representation that can be read by an Embedding layer or Dense layer.

In this dataset, Type is represented as a string (e.g., 'Dog' or 'Cat'). You cannot feed strings directly to a model. The preprocessing layer takes care of representing strings as a one-hot vector.

Column	Description	Feature Type	Data Type
Type	Type of animal (Dog, Cat)	Categorical	string

```
def get_category_encoding_layer(name, dataset, dtype, max_tokens=None):
    # Create a StringLookup layer which will turn strings into integer indices
    if dtype == 'string':
        index = preprocessing.StringLookup(max_tokens=max_tokens)
    else:
        index = preprocessing.IntegerLookup(max_tokens=max_tokens)

    # Prepare a Dataset that only yields our feature
    feature_ds = dataset.map(lambda x, y: x[name])

    # Learn the set of possible values and assign them a fixed integer index.
    index.adapt(feature_ds)

    # Create a Discretization for our integer indices.
    encoder = preprocessing.CategoryEncoding(num_tokens=index.vocabulary_size())

    # Apply one-hot encoding to our indices. The lambda function captures the
    # layer so we can use them, or include them in the functional model later.
    return lambda feature: encoder(index(feature))
```

Categorical columns

tf.keras.layers.IntegerLookup:

Turns integer categorical values into an encoded representation that can be read by an Embedding layer or Dense layer.

```
def get_category_encoding_layer(name, dataset, dtype, max_tokens=None):
    # Create a StringLookup layer which will turn strings into integer indices
    if dtype == 'string':
        index = preprocessing.StringLookup(max_tokens=max_tokens)
    else:
        index = preprocessing.IntegerLookup(max_tokens=max_tokens)

    # Prepare a Dataset that only yields our feature
    feature_ds = dataset.map(lambda x, y: x[name])

    # Learn the set of possible values and assign them a fixed integer index.
    index.adapt(feature_ds)

    # Create a Discretization for our integer indices.
    encoder = preprocessing.CategoryEncoding(num_tokens=index.vocabulary_size())

    # Apply one-hot encoding to our indices. The lambda function captures the
    # layer so we can use them, or include them in the functional model later.
    return lambda feature: encoder(index(feature))
```

Numeric columns

For each of the Numeric features, you will use a `Normalization()` layer to ensure that the mean of each feature is 0 and its standard deviation is 1.

Recall that the `get_normalization_layer` function returns a layer that applies feature-wise normalization to numerical features.

```
def get_normalization_layer(name, dataset):
    # Create a Normalization layer for our feature.
    normalizer = preprocessing.Normalization(axis=None)

    # Prepare a Dataset that only yields our feature.
    feature_ds = dataset.map(lambda x, y: x[name])

    # Learn the statistics of the data.
    normalizer.adapt(feature_ds)

    return normalizer
```

Numeric columns

For each of the Numeric features, you will use a `Normalization()` layer to ensure that the mean of each feature is 0 and its standard deviation is 1.

Recall that the `get_normalization_layer` function returns a layer that applies feature-wise normalization to numerical features.

```
def get_normalization_layer(name, dataset):
    # Create a Normalization layer for our feature.
    normalizer = preprocessing.Normalization(axis=None)

    # Prepare a Dataset that only yields our feature.
    feature_ds = dataset.map(lambda x, y: x[name])

    # Learn the statistics of the data.
    normalizer.adapt(feature_ds)

    return normalizer
```

the `adapt()` method

Simple normalization example

```
#Load some data
(x_train, y_train), _ = keras.datasets.cifar10.load_data()
x_train = x_train.reshape((len(x_train), -1))
input_shape = x_train.shape[1:]
classes = 10

#Create a Normalization layer and set its internal state using the training data
normalizer = layers.Normalization()
normalizer.adapt(x_train)

#Create a model that include the normalization layer
inputs = keras.Input(shape=input_shape)
x = normalizer(inputs)
outputs = layers.Dense(classes, activation="softmax")(x)
model = keras.Model(inputs, outputs)

#Train the model
model.compile(optimizer="adam", loss="sparse_categorical_crossentropy")
model.fit(x_train, y_train)
```

Create a normalization layer

```
#Load some data
(x_train, y_train), _ = keras.datasets.cifar10.load_data()
x_train = x_train.reshape((len(x_train), -1))
input_shape = x_train.shape[1:]
classes = 10

#Create a Normalization layer and set its internal state using the training data
normalizer = layers.Normalization()
normalizer.adapt(x_train)
```

```
#Create a model that include the normalization layer
inputs = keras.Input(shape=input_shape)
x = normalizer(inputs)
outputs = layers.Dense(classes, activation="softmax")(x)
model = keras.Model(inputs, outputs)

#Train the model
model.compile(optimizer="adam", loss="sparse_categorical_crossentropy")
model.fit(x_train, y_train)
```

Create a model with a normalization layer

```
#Load some data
(x_train, y_train), _ = keras.datasets.cifar10.load_data()
x_train = x_train.reshape((len(x_train), -1))
input_shape = x_train.shape[1:]
classes = 10

#Create a Normalization layer and set its internal state using the training data
normalizer = layers.Normalization()
normalizer.adapt(x_train)

#Create a model that include the normalization layer
inputs = keras.Input(shape=input_shape)
x = normalizer(inputs)
outputs = layers.Dense(classes, activation="softmax")(x)
model = keras.Model(inputs, outputs)

#Train the model
model.compile(optimizer="adam", loss="sparse_categorical_crossentropy")
model.fit(x_train, y_train)
```

Model layers versus preprocessing dataset

(Preprocessing data before the model or inside the model)

Option 1

```
inputs = kera.Input(shape=input_shape)
x = preprocessing_layer(inputs)
outputs = rest_of_the_model(x)
model = keras.Model(inputs, outputs)
```

Option 2

```
dataset = dataset.map(
    lambda x, y: (preprocessing_layer(x),y))
```

Preprocessing as part of the model

Option 1

```
inputs = kera.Input(shape=input_shape)
x = preprocessing_layer(inputs)
outputs = rest_of_the_model(x)
model = keras.Model(inputs, outputs)
```

Option 2

```
dataset = dataset.map(
    lambda x, y: (preprocessing_layer(x),y))
```

Option 1:

Make them part
of the model,
like this

Preprocessing will happen on the device
synchronously with the rest of the
model execution.

```
inputs = keras.Input(shape=input_shape)
x = preprocessing_layer(inputs)
outputs = rest_of_the_model(x)
model = keras.Model(inputs, outputs)
```

Preprocessing before the model

Option 1

```
inputs = kera.Input(shape=input_shape)
x = preprocessing_layer(inputs)
outputs = rest_of_the_model(x)
model = keras.Model(inputs, outputs)
```

Option 2

```
dataset = dataset.map(
    lambda x, y: (preprocessing_layer(x),y))
```

Preprocessing will happen on the CPU asynchronously and will be buffered before going into the model.

Option 2: Apply it to your [tf.data.Dataset](#)

```
dataset = dataset.map(lambda x,y: (preprocessing_layer(x),y))
dataset = dataset.prefetch(tf.data.AUTOTUNE)
model.fit(dataset,...)
```

When running on a TPU, you should always place preprocessing layers in the [tf.data](#) pipeline (with the exception of [Normalization](#) and [Rescaling](#), which run well on TPU and are commonly used as the first layer in an image model).

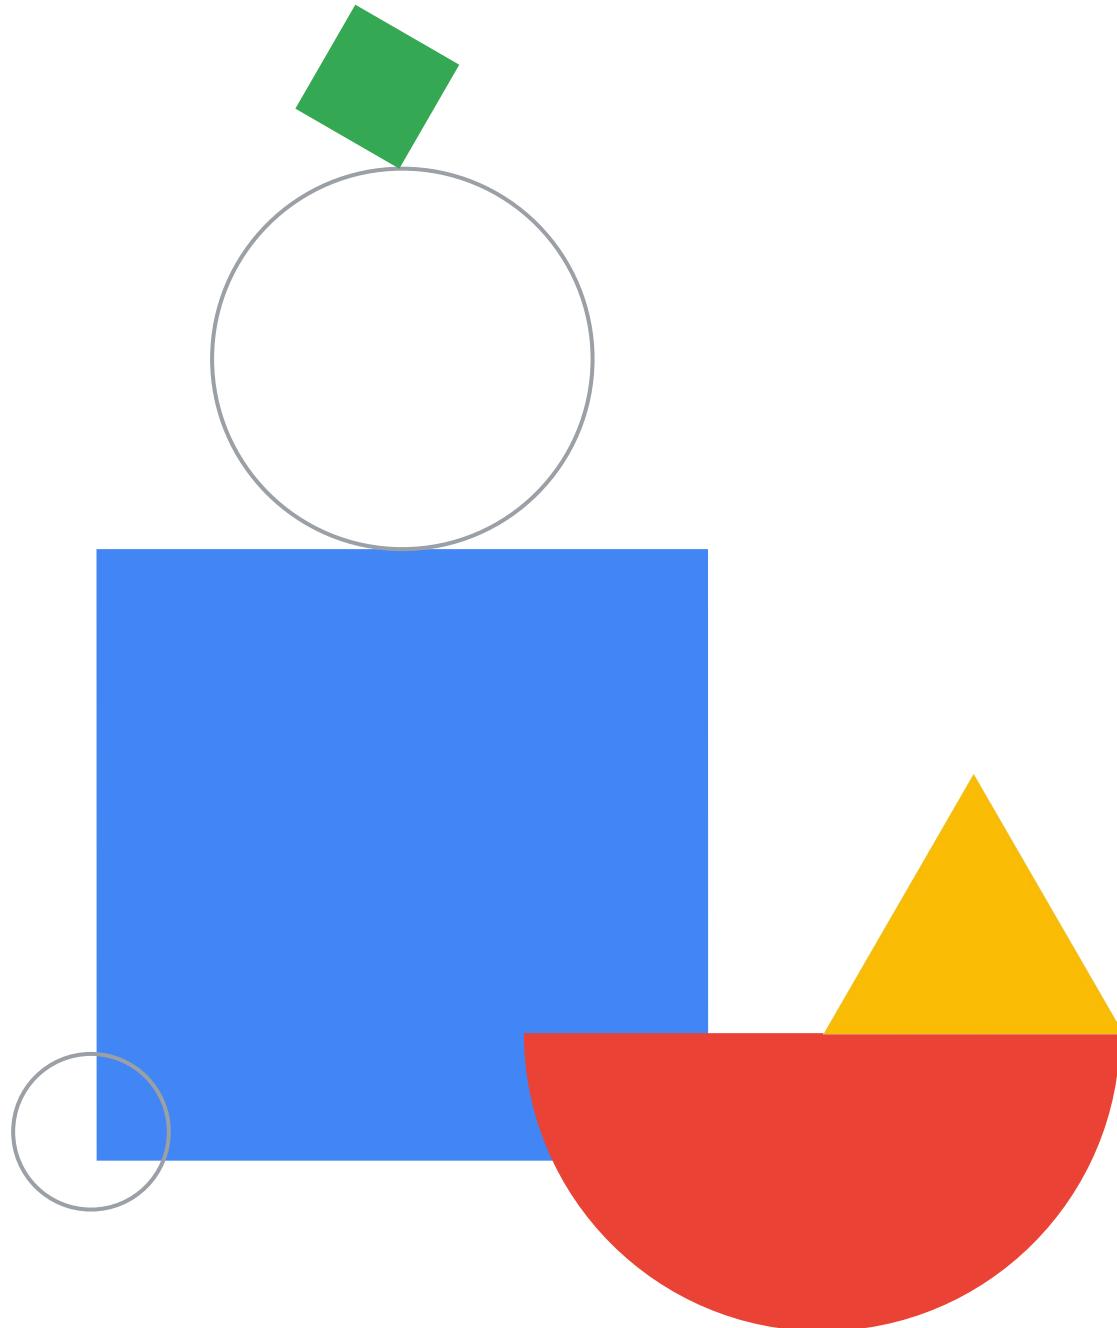
Benefits of doing preprocessing inside the model at inference time

Even if you choose option 2, you may later want to export an inference-only end-to-end model that will include the preprocessing layers.

The key benefits:

- Makes your model portable.
- Helps reduce the **training-serving skew**.





Building Neural Networks with the TensorFlow and Keras API

In this module, you learn to ...

01

Describe activation functions, loss, and optimizers

02

Build a DNN model using the Keras Sequential and Functional APIs

03

Use Keras preprocessing layers

04

Save/load and deploy a Keras model

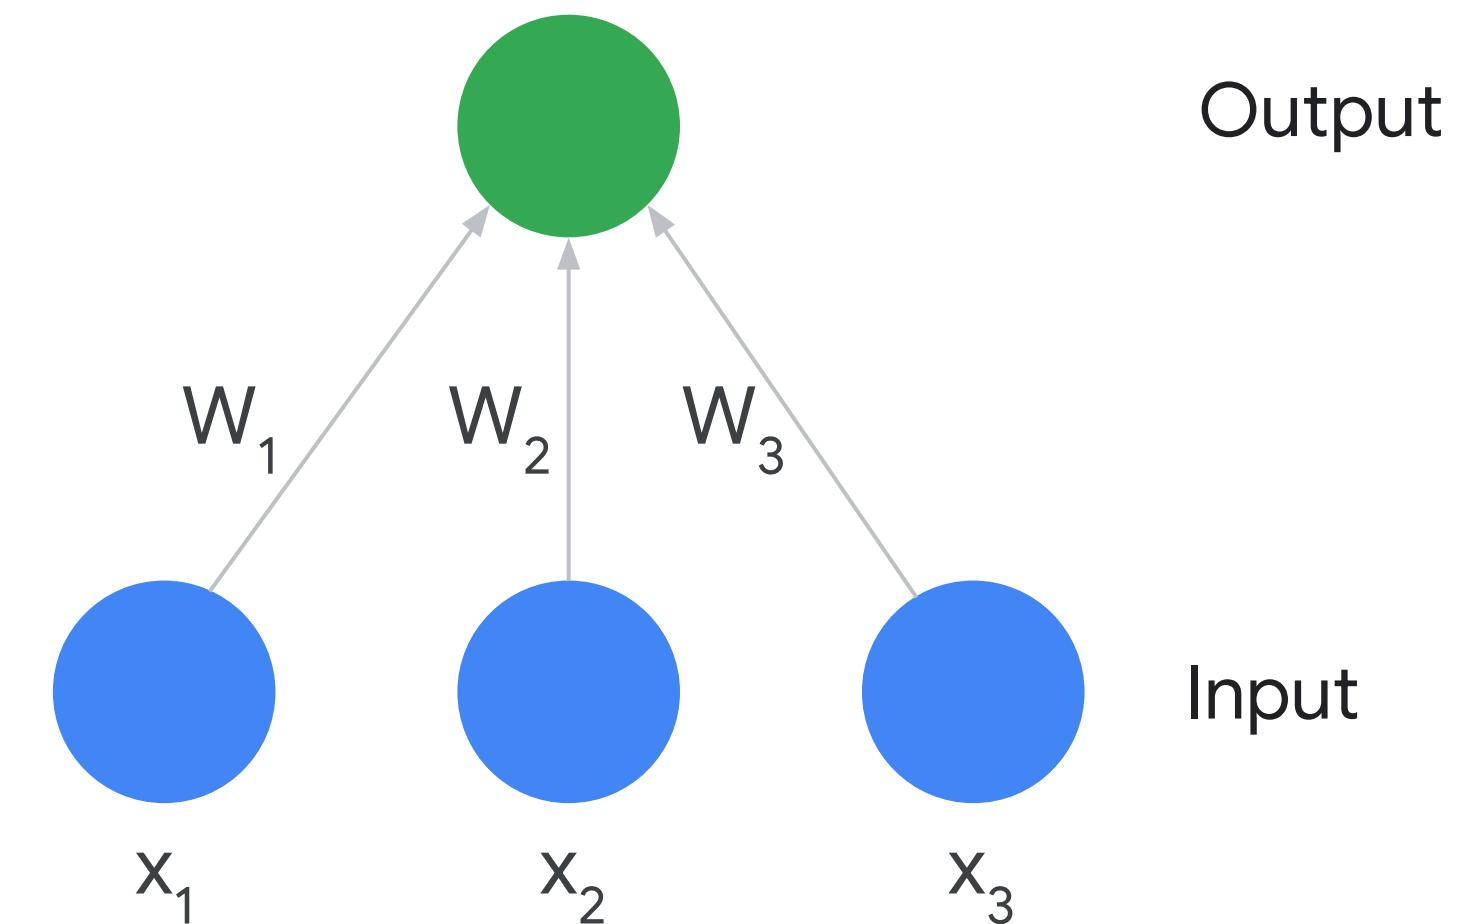
05

Describe model subclassing



A Linear Model can
be represented as
nodes and edges

$$\text{Output} = W_1 * x_1 + W_2 * x_2 + W_3 * x_3$$

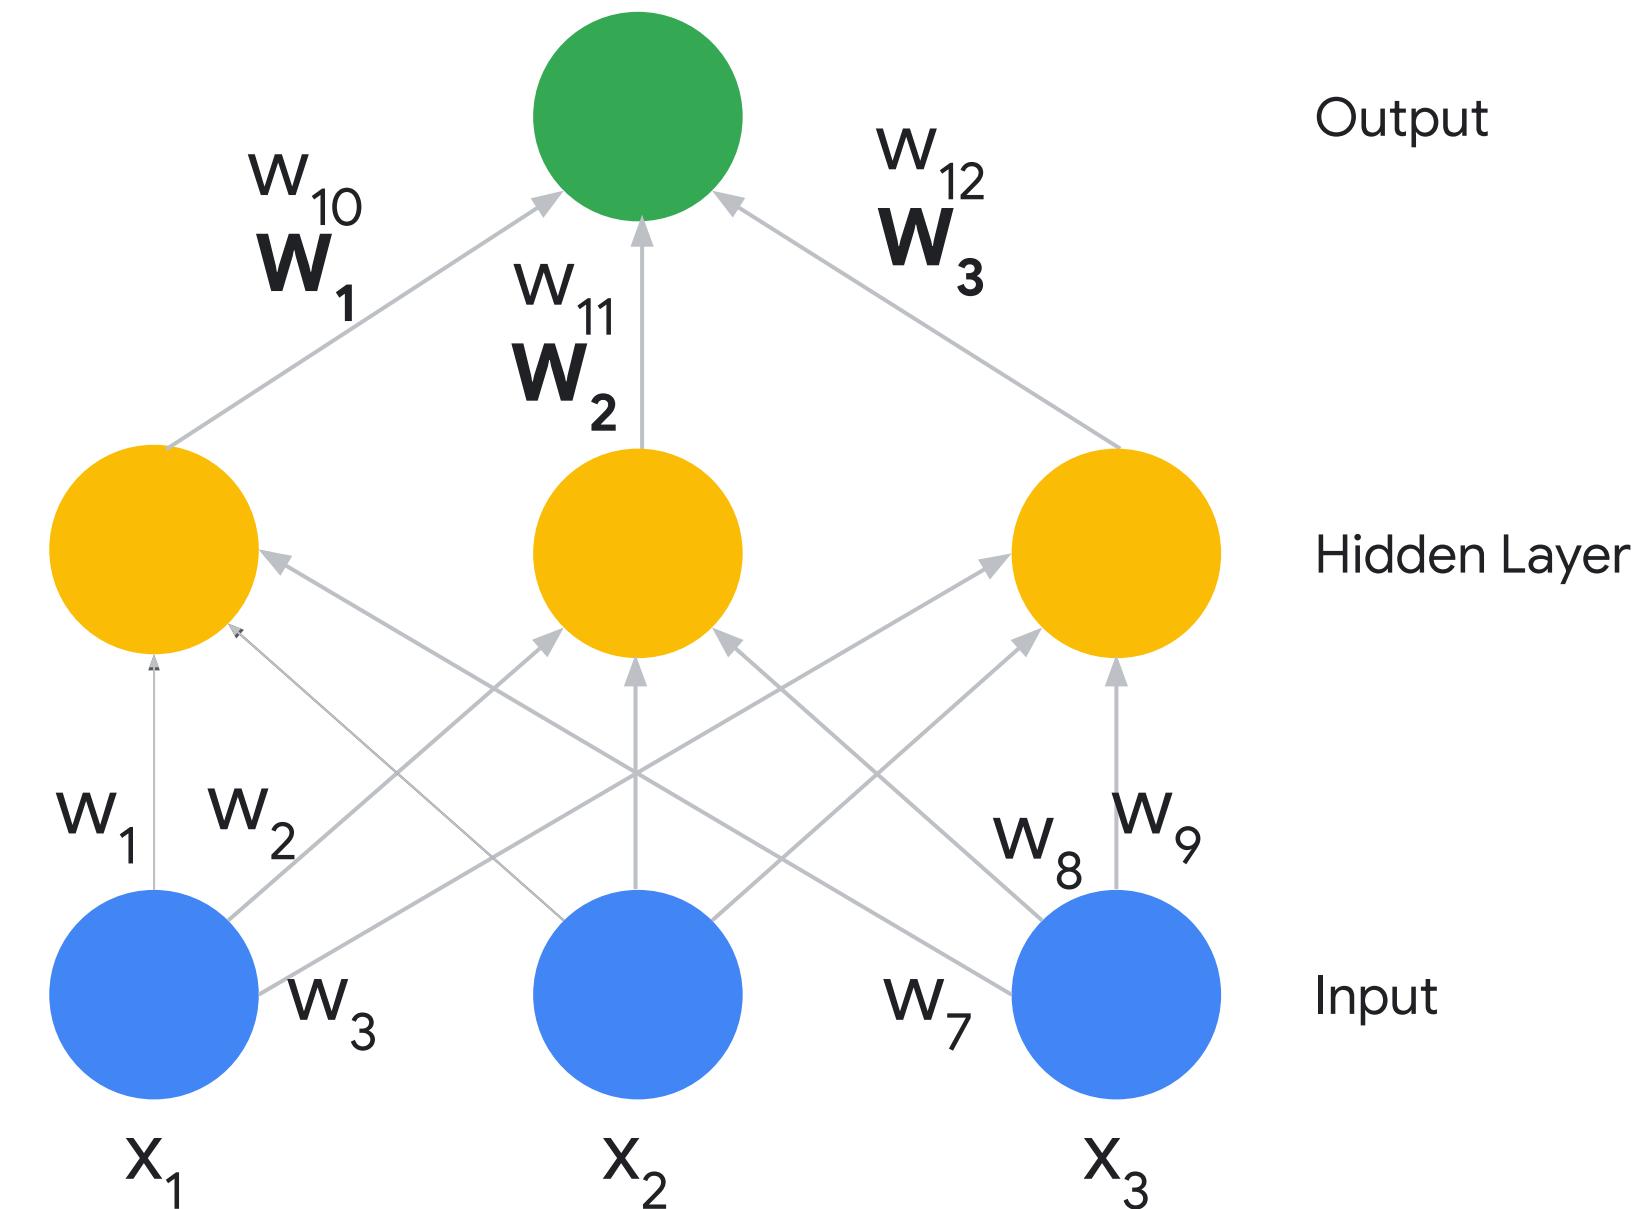


Add Complexity: Non-Linear?

$$\text{Output} = w_{10} * h_1 + w_{11} * h_2 + w_{12} * h_3$$

$$\begin{aligned} & ((w_{10} * w_1) + (w_{10} * w_4) + (w_{10} * w_7)) * x_1 \\ & + ((w_{11} * w_1) + (w_{11} * w_4) + (w_{11} * w_7)) * x_2 \\ & + ((w_{12} * w_1) + (w_{12} * w_4) + (w_{12} * w_7)) * x_3 \end{aligned}$$

$$= W_1 * x_1 + W_2 * x_2 + W_3 * x_3$$

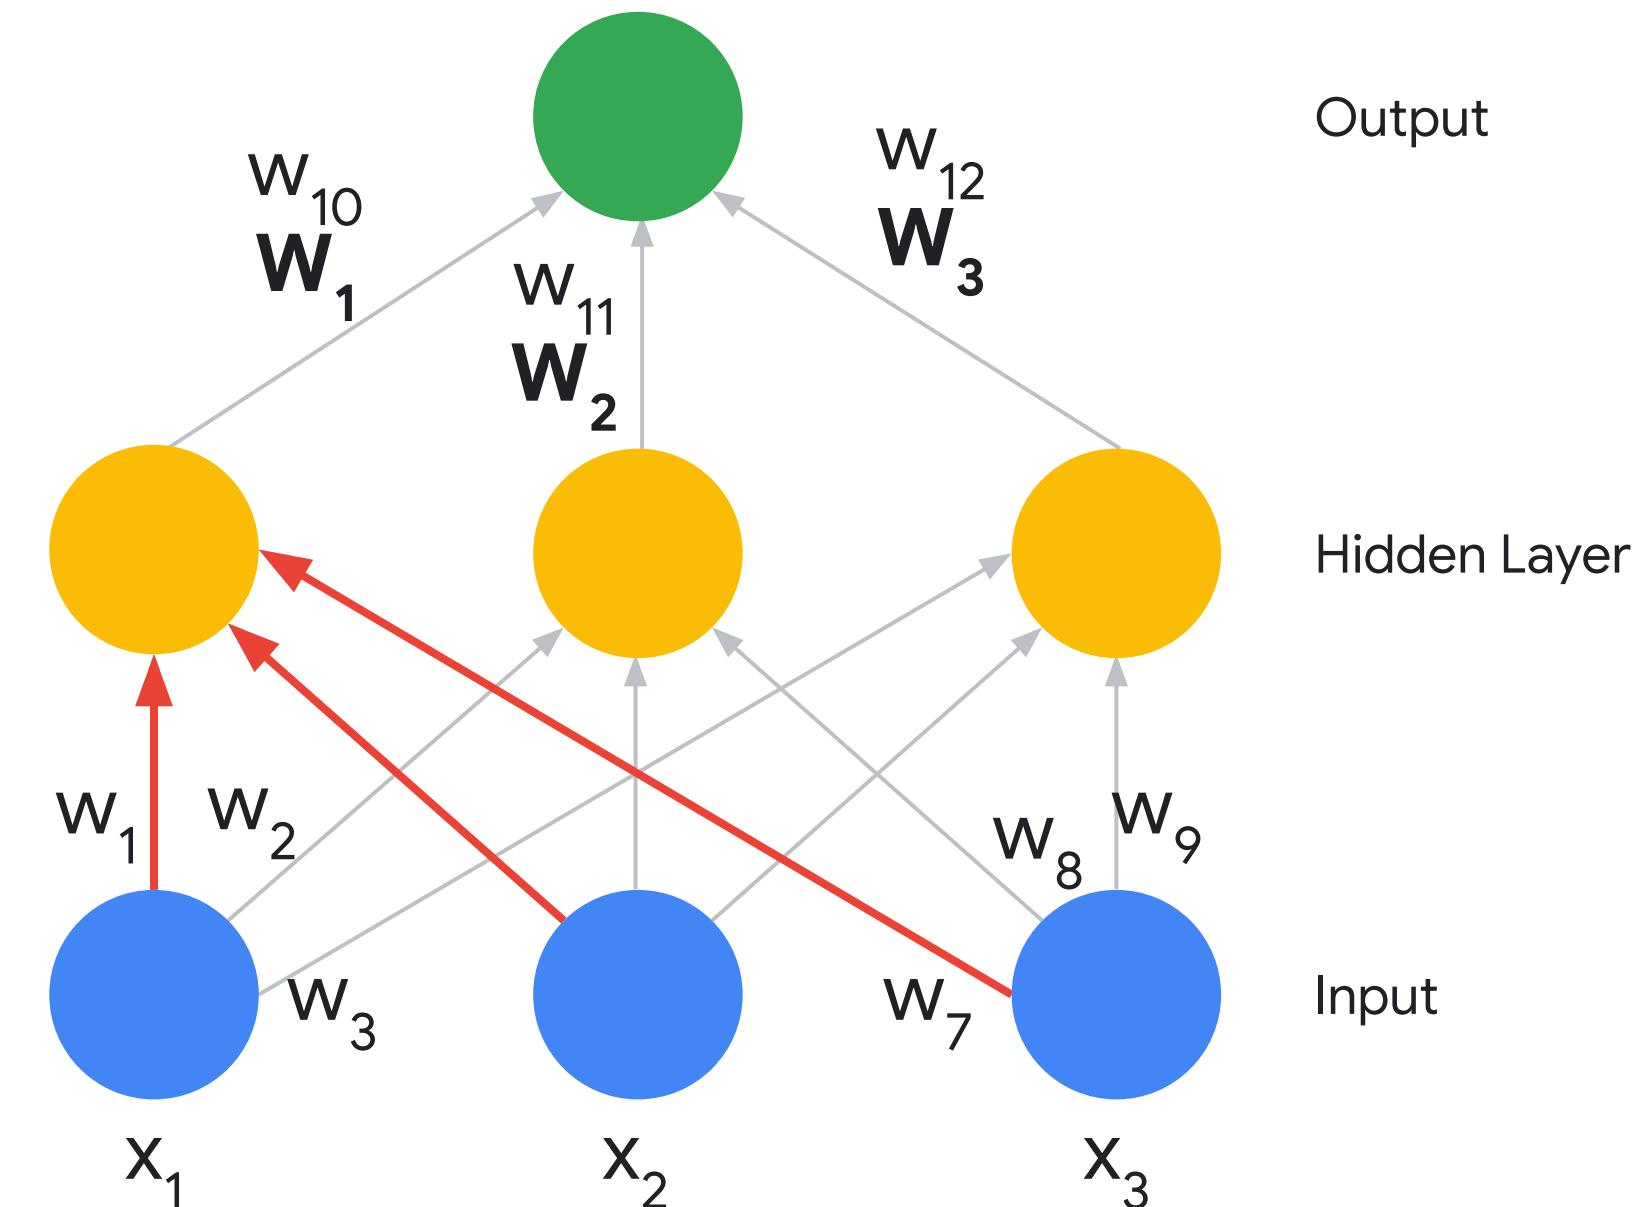


Add Complexity: Non-Linear?

$$\text{Output} = w_{10} * h_1 + w_{11} * h_2 + w_{12} * h_3$$

$$\begin{aligned} & ((w_{10} * w_1) + (w_{10} * w_4) + (w_{10} * w_7)) * x_1 \\ & + ((w_{11} * w_1) + (w_{11} * w_4) + (w_{11} * w_7)) * x_2 \\ & + ((w_{12} * w_1) + (w_{12} * w_4) + (w_{12} * w_7)) * x_3 \end{aligned}$$

$$= W_1 * x_1 + W_2 * x_2 + W_3 * x_3$$

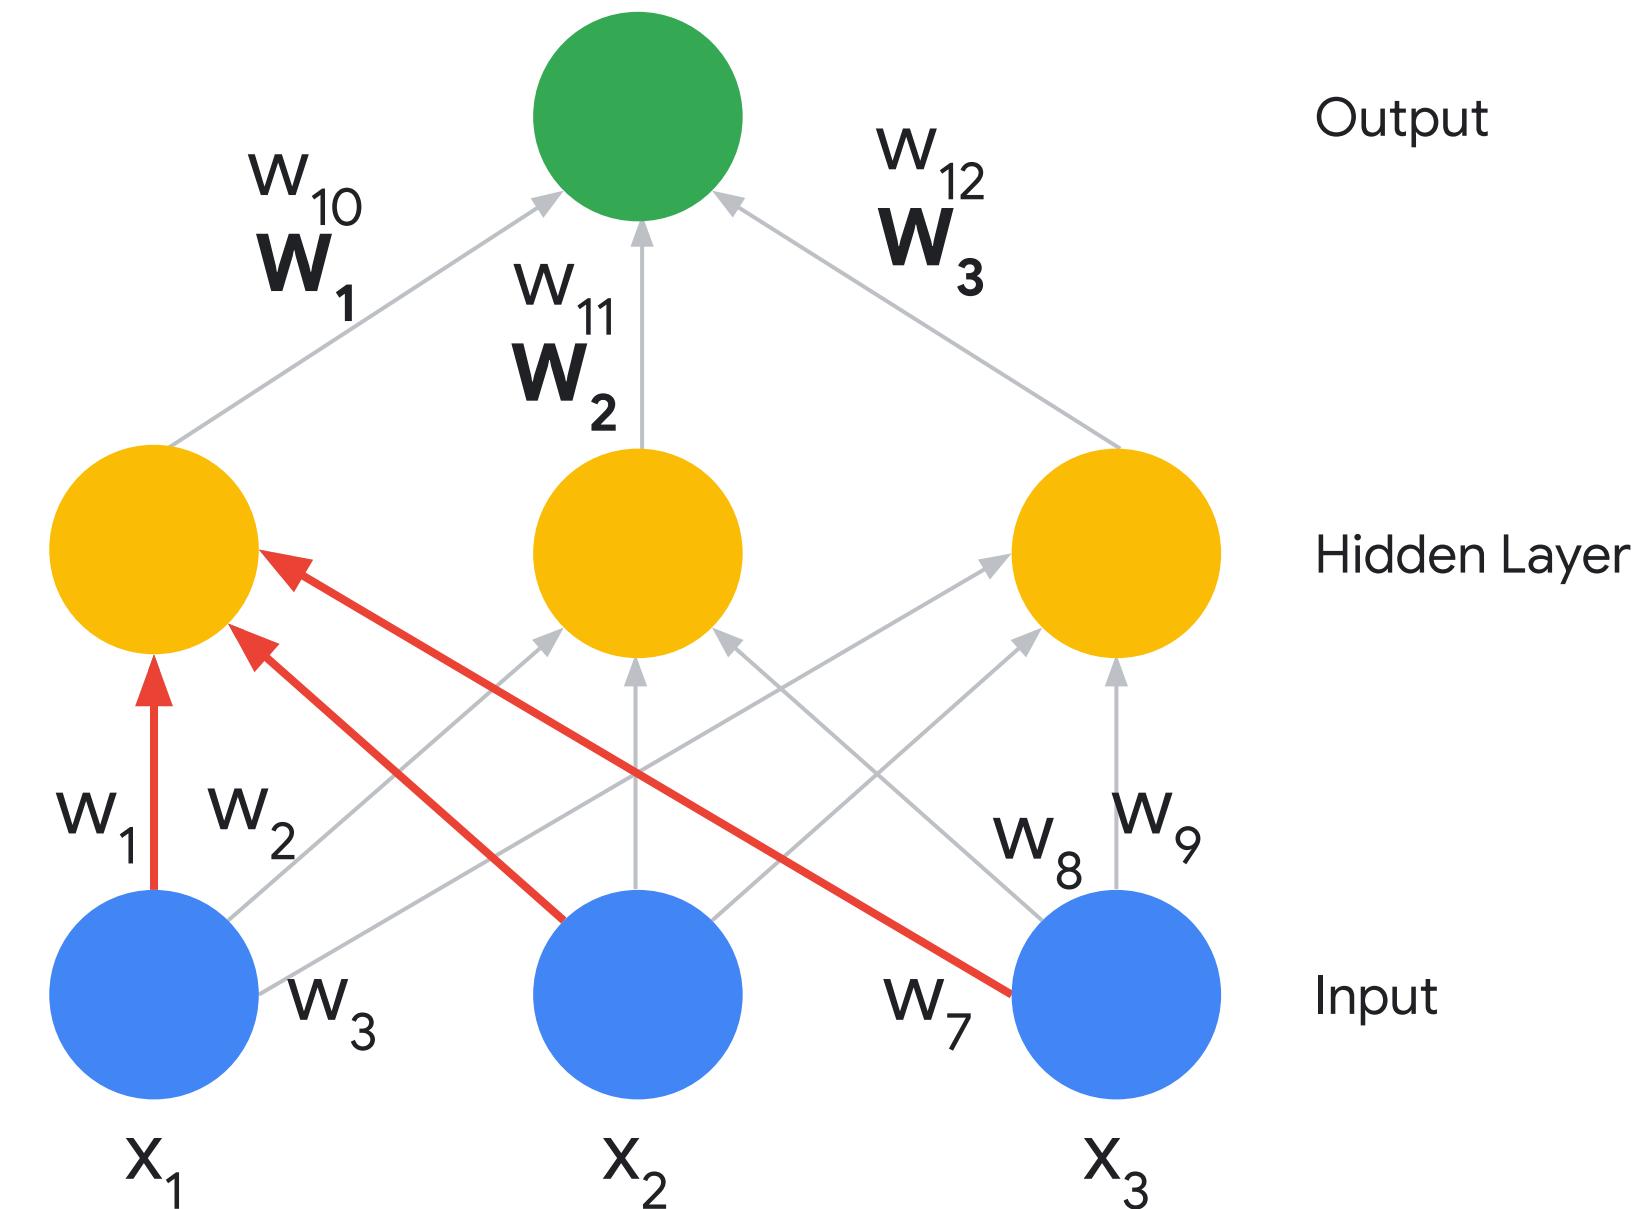


Add Complexity: Non-Linear?

$$\text{Output} = w_{10} * h_1 + w_{11} * h_2 + w_{12} * h_3$$

$$\begin{aligned} & ((w_{10} * w_1) + (w_{10} * w_4) + (w_{10} * w_7)) * x_1 \\ & + ((w_{11} * w_1) + (w_{11} * w_4) + (w_{11} * w_7)) * x_2 \\ & + ((w_{12} * w_1) + (w_{12} * w_4) + (w_{12} * w_7)) * x_3 \end{aligned}$$

$$= W_1 * x_1 + W_2 * x_2 + W_3 * x_3$$

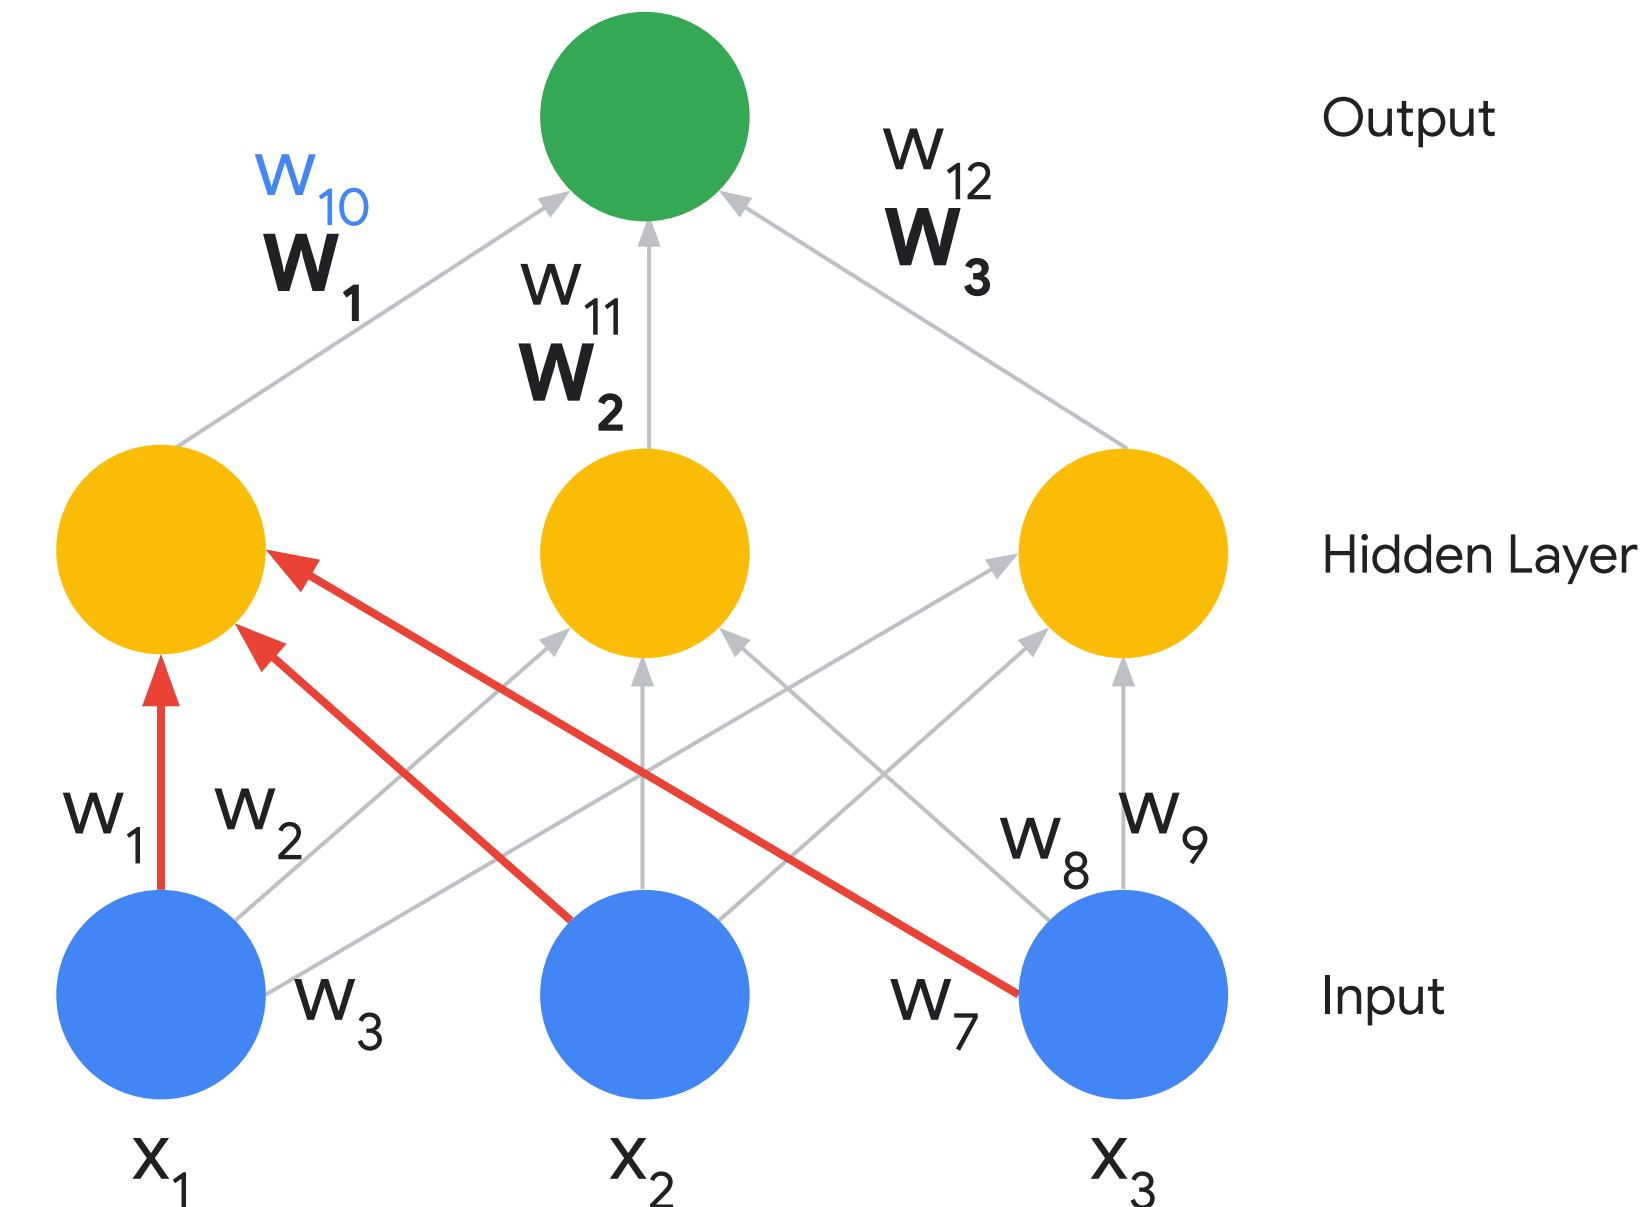


Add Complexity: Non-Linear?

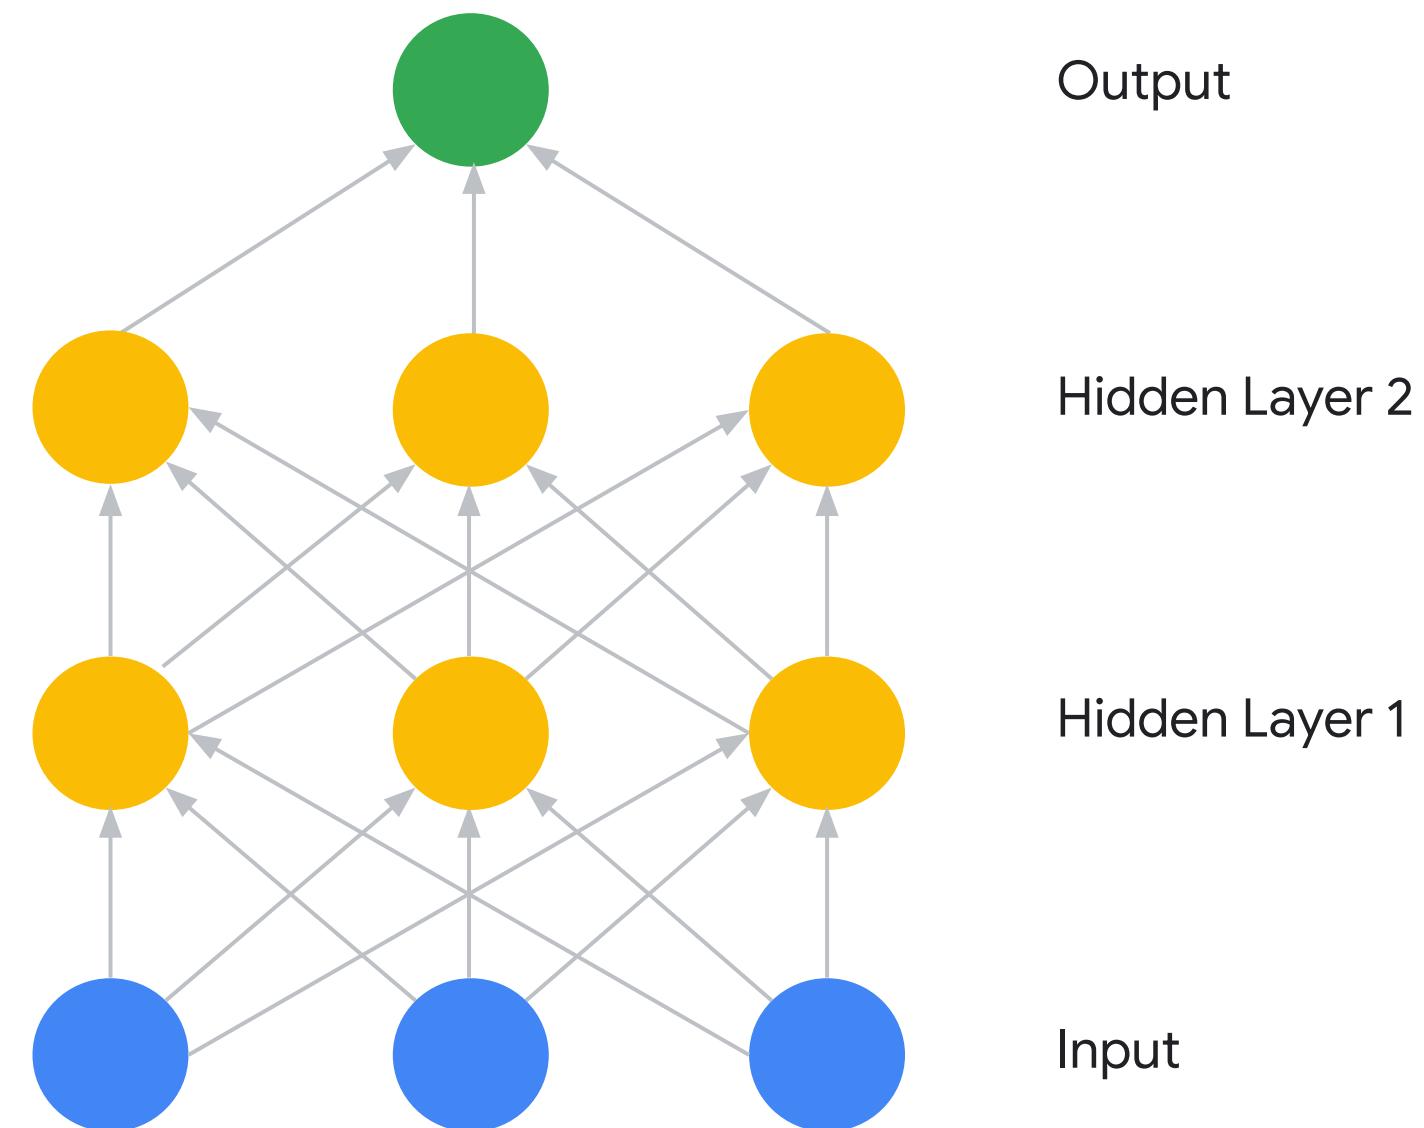
$$\text{Output} = w_{10} * h_1 + w_{11} * h_2 + w_{12} * h_3$$

$$\begin{aligned} & ((w_{10} * w_1) + (w_{10} * w_4) + (w_{10} * w_7)) * x_1 \\ & + ((w_{11} * w_1) + (w_{11} * w_4) + (w_{11} * w_7)) * x_2 \\ & + ((w_{12} * w_1) + (w_{12} * w_4) + (w_{12} * w_7)) * x_3 \end{aligned}$$

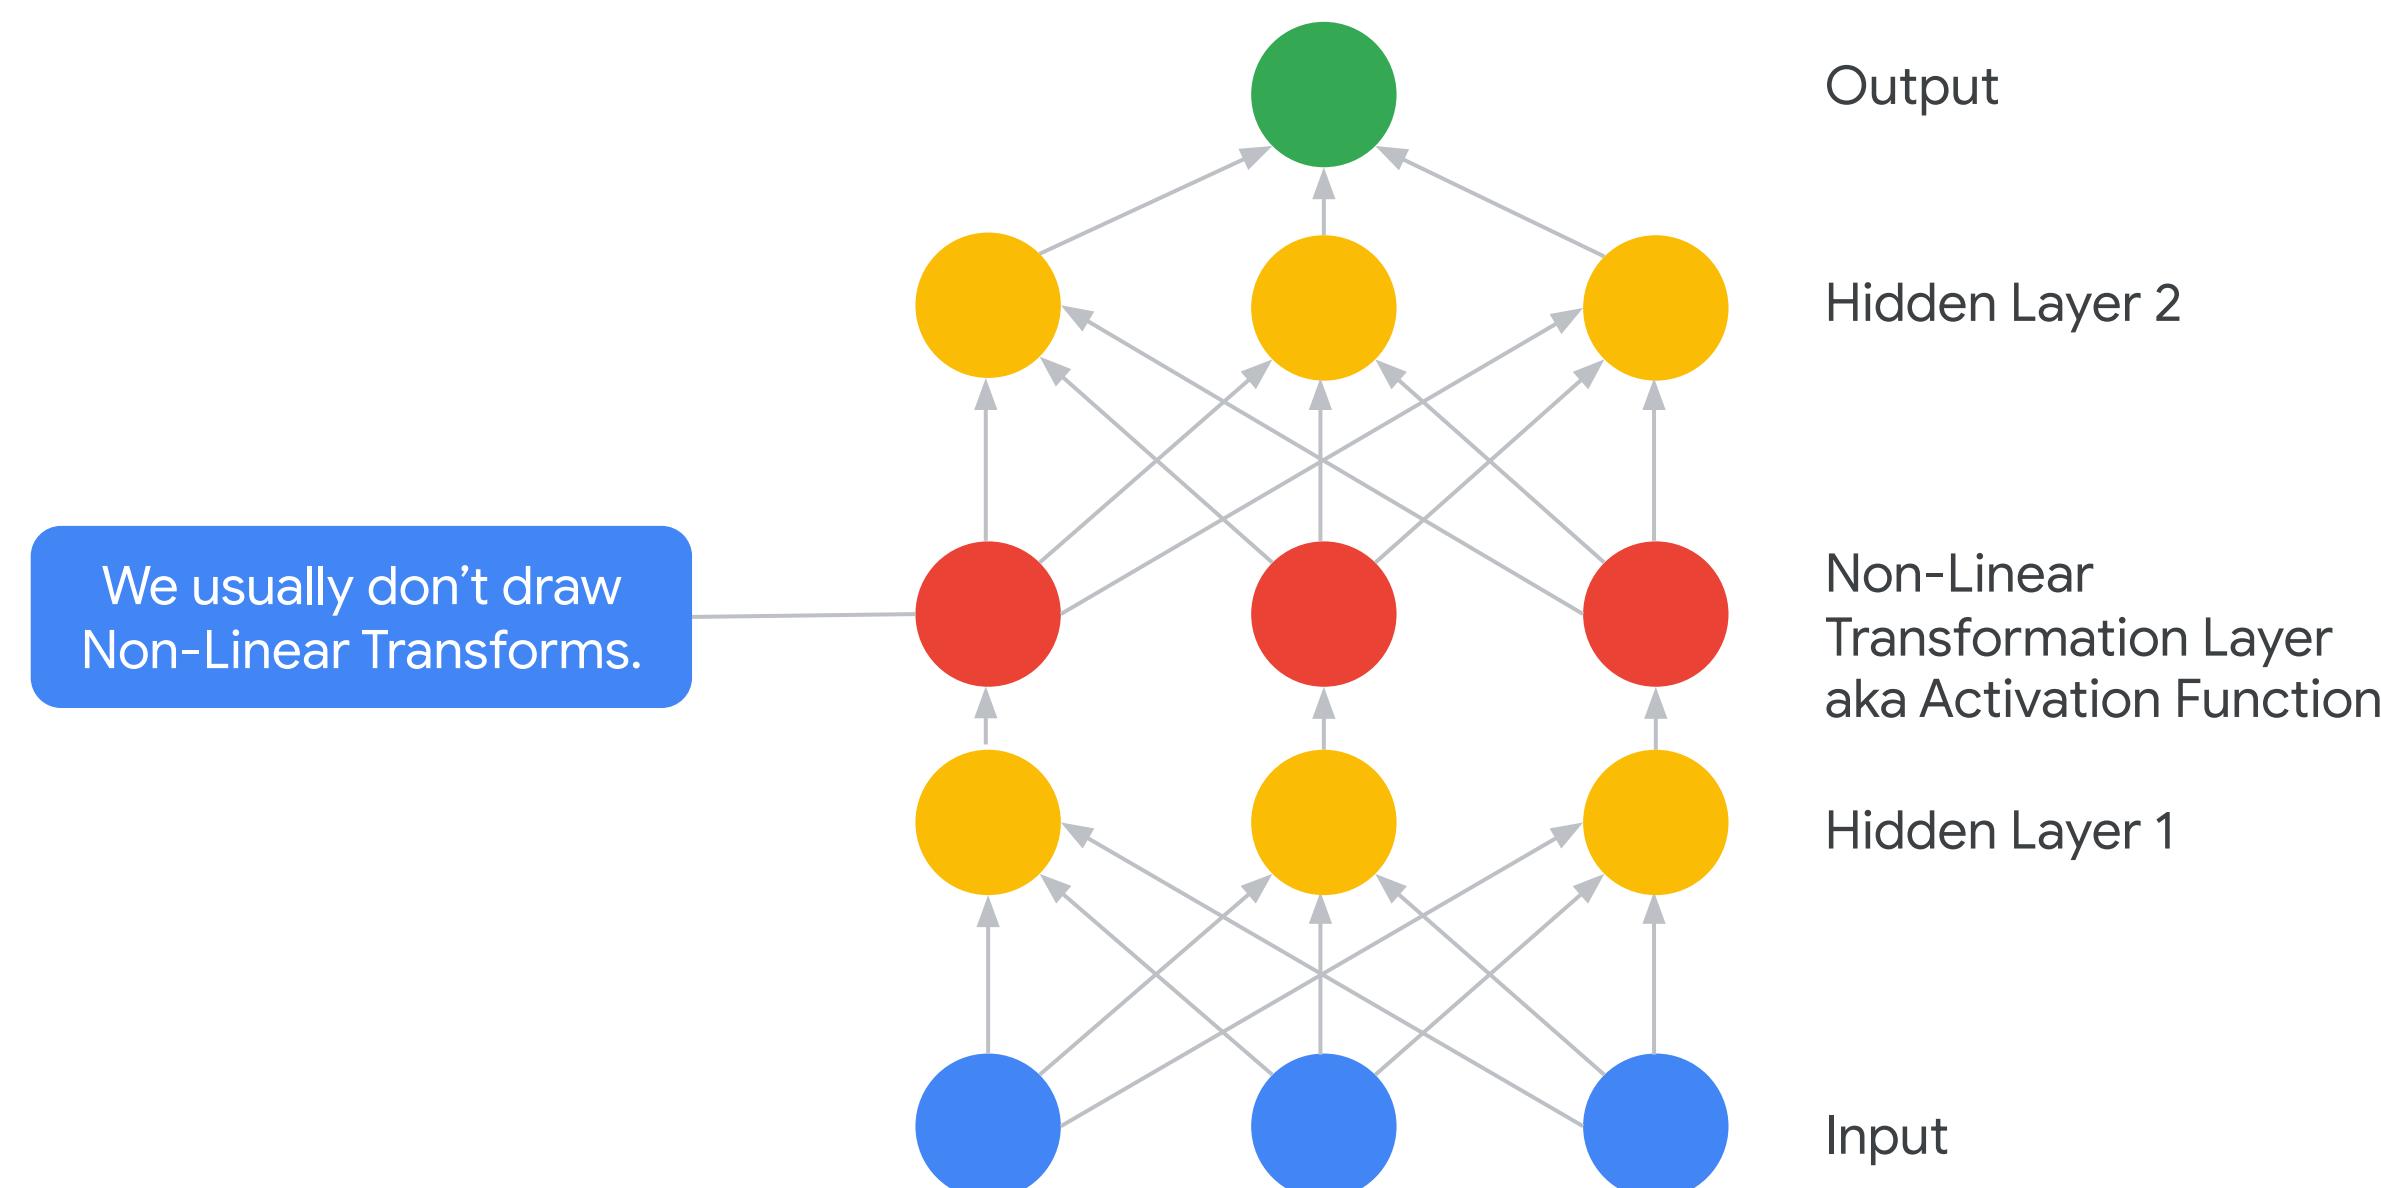
$$= W_1 * x_1 + W_2 * x_2 + W_3 * x_3$$



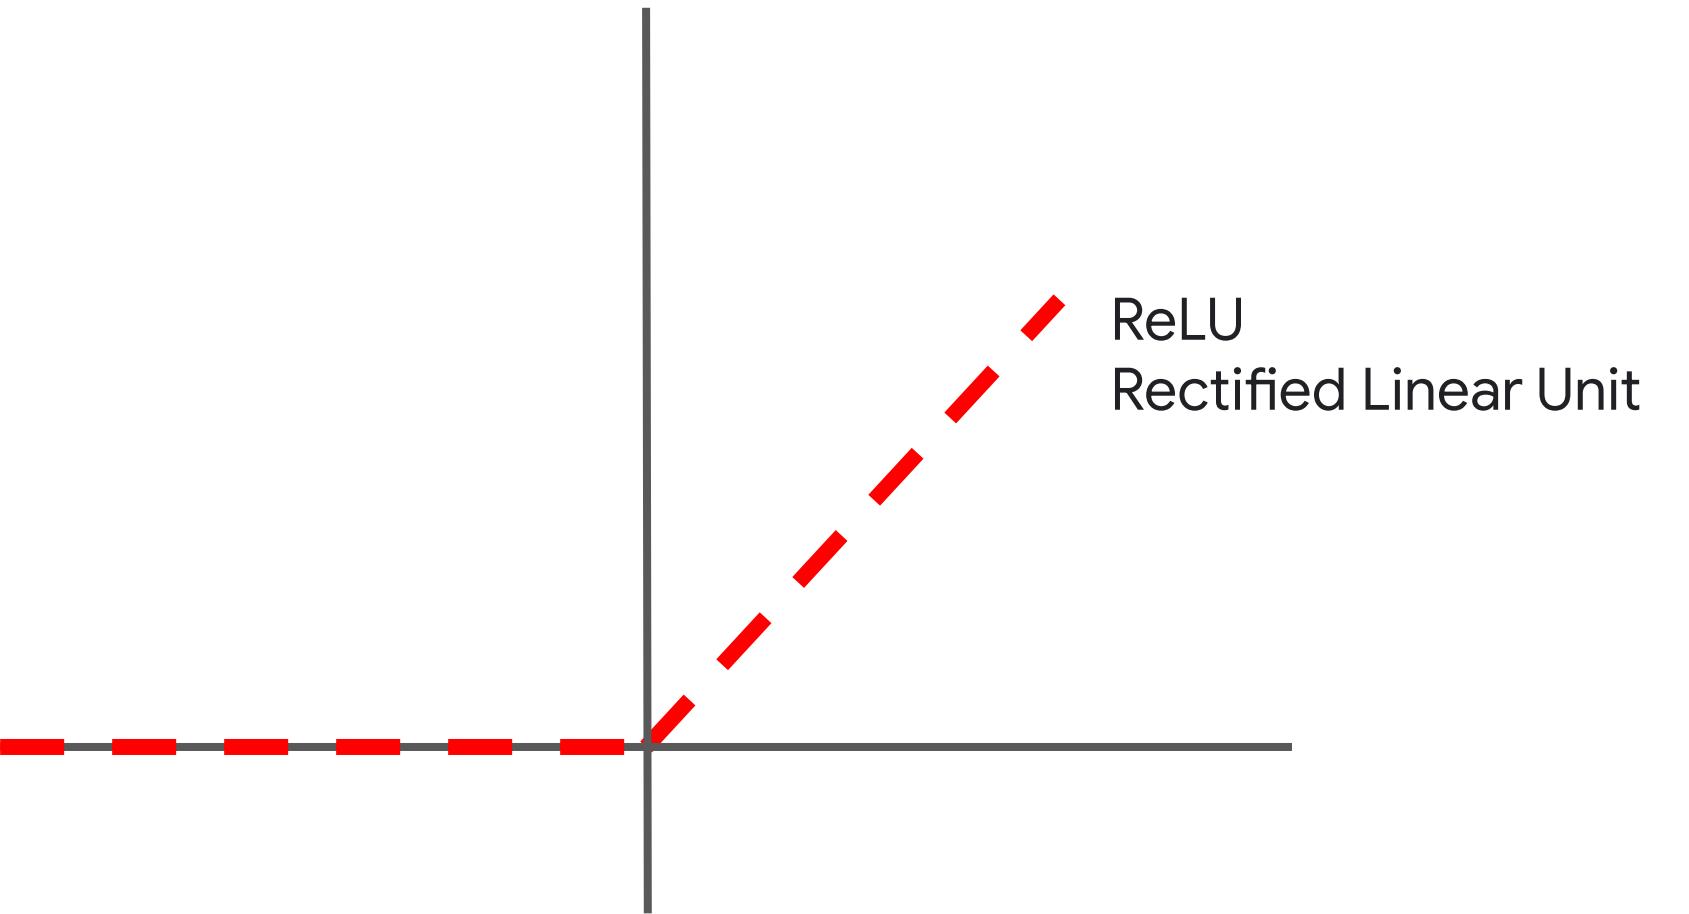
Add Complexity: Non-Linear?



Adding a Non-Linearity



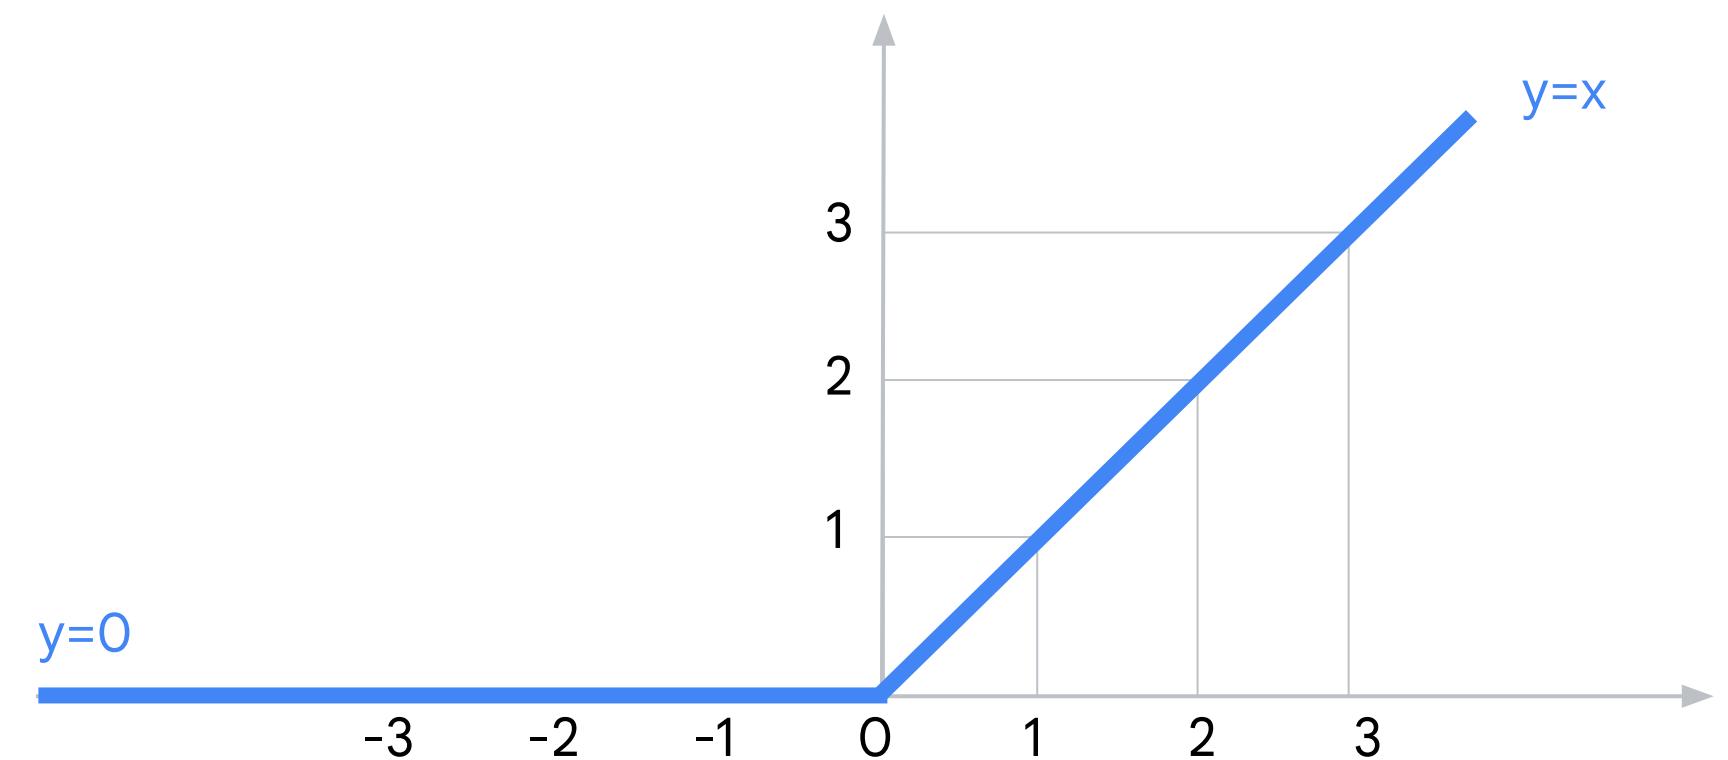
Our favorite
non-linearity is
the Rectified
Linear Unit



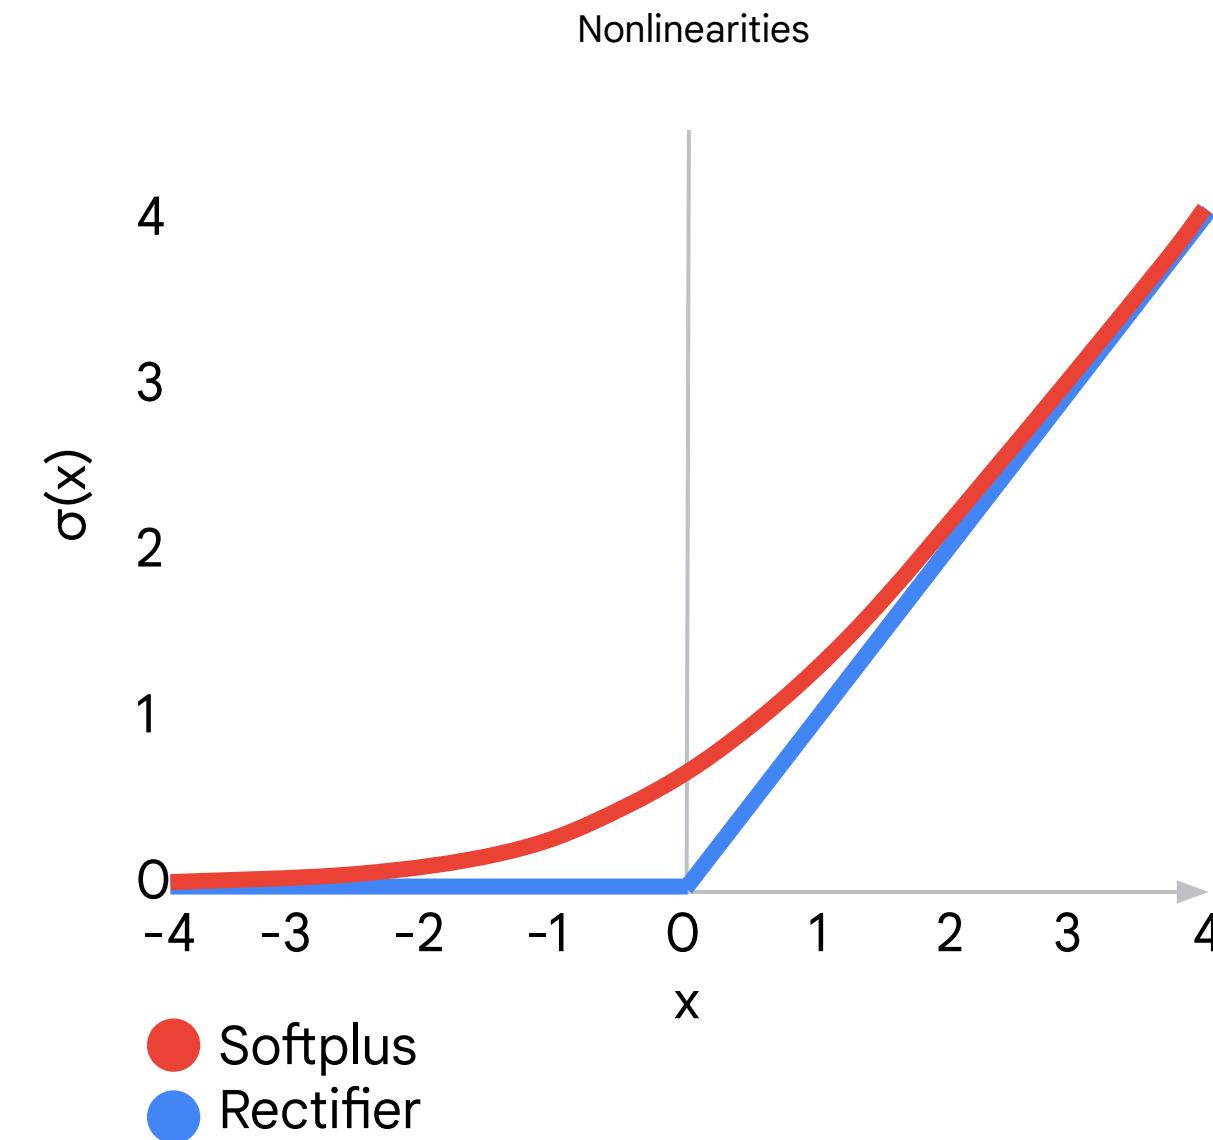
$$f(x) = \max(0, x)$$

Normal ReLU activation function

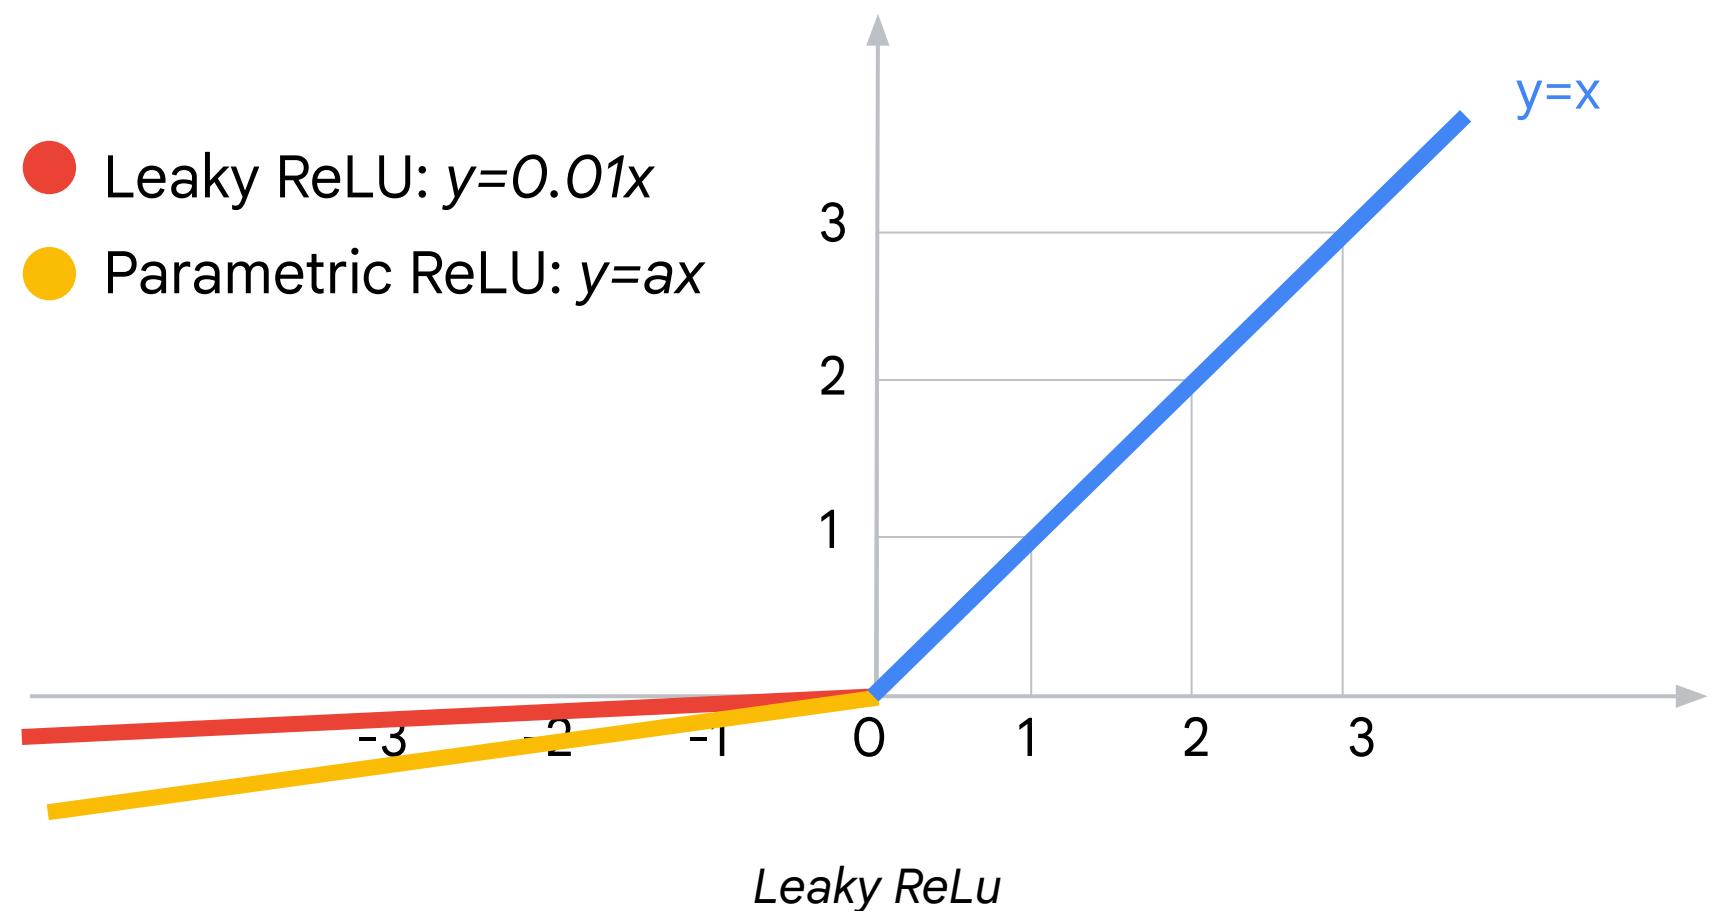
There are many
different ReLU variants



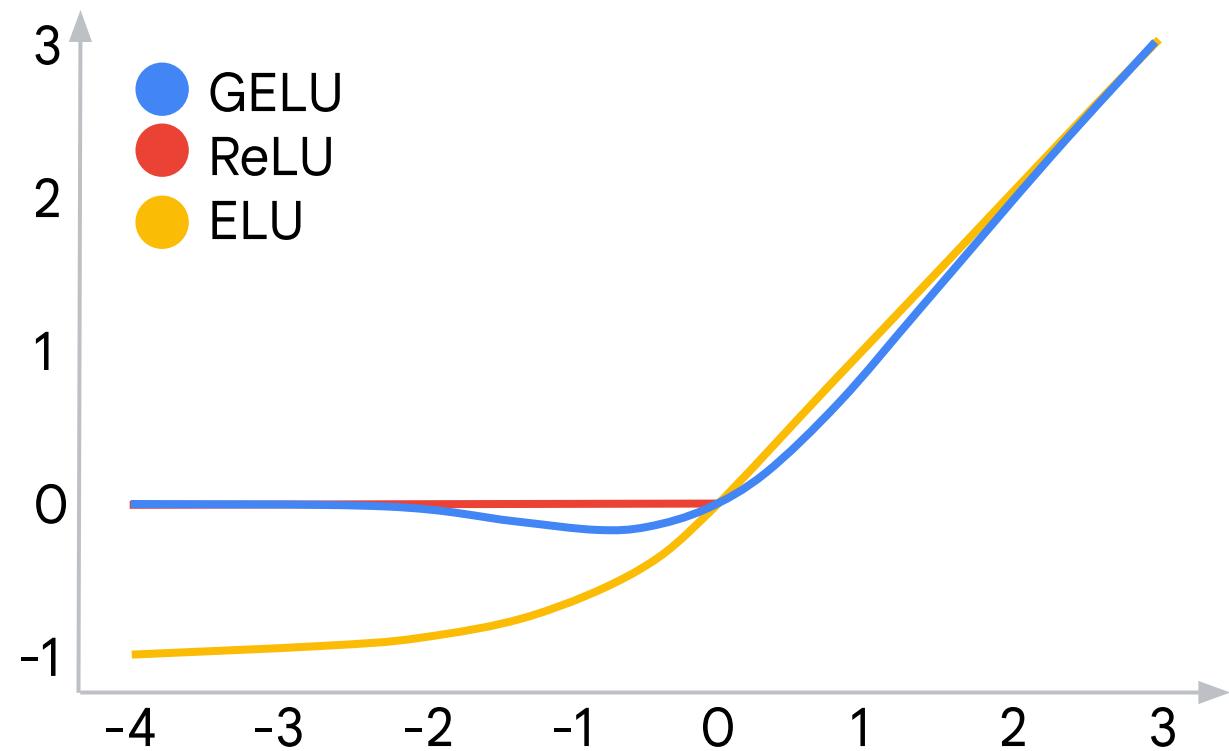
There are many
different **ReLU** variants



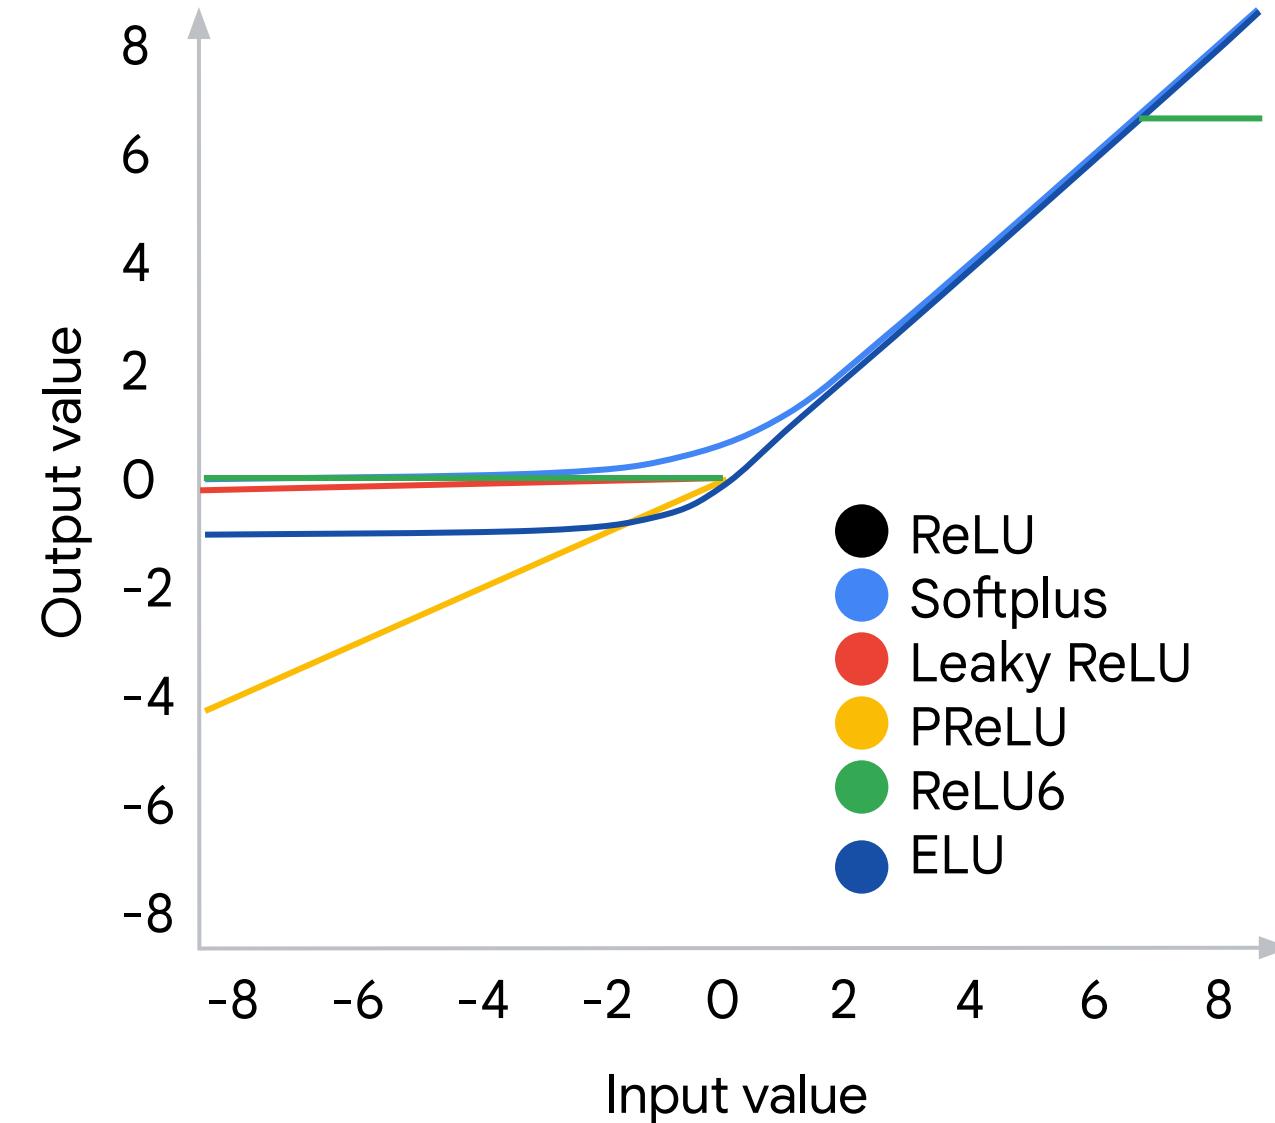
There are many
different **ReLU variants**



There are many
different **ReLU** variants



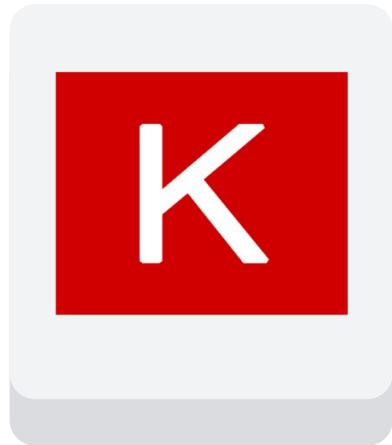
There are many
different ReLU variants



Keras is built-in to TF 2.x



Keras is built-in to TF 2.x



- User-friendly
- Modular and composable
- Easy to extend

Stacking layers with Keras Sequential model

```
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Dense  
  
model = Sequential([  
    Input(shape=(64,)),  
    Dense(units=32, activation="relu", name="hidden1"),  
    Dense(units=8, activation="relu", name="hidden2"),  
    Dense(units=1, activation="linear", name="output")  
])
```

The Keras sequential model stacks layers on the top of each other.

The batch size is omitted. Here the model expects batches of vectors with 64 components.

```
%tensorflow_version 2.x
import tensorflow as tf
from tensorflow import keras
(x_train, y_train), (x_test, y_test) =
keras.datasets.mnist.load_data()

# Define your model
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(10, activation='softmax')
])
```

A linear model (multiclass logistic regression)

```
%tensorflow_version 2.x
import tensorflow as tf
from tensorflow import keras
(x_train, y_train), (x_test, y_test) =
keras.datasets.mnist.load_data()

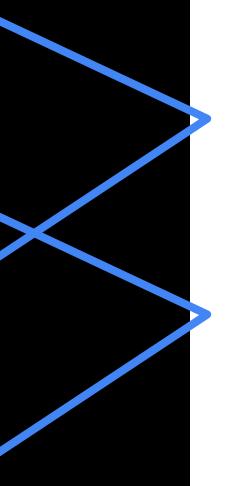
# Define your model
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(10, activation='softmax')
])
```

A linear model (multiclass logistic regression)

A linear model (a single Dense layer) aka multiclass logistic regression

```
%tensorflow_version 2.x
import tensorflow as tf
from tensorflow import keras
(x_train, y_train), (x_test, y_test) =
keras.datasets.mnist.load_data()

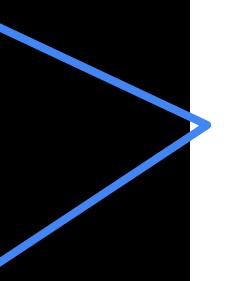
# Define your model
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
```



A neural network with one hidden layer
A neural network with one hidden layer

```
%tensorflow_version 2.x
import tensorflow as tf
from tensorflow import keras
(x_train, y_train), (x_test, y_test) =
keras.datasets.mnist.load_data()

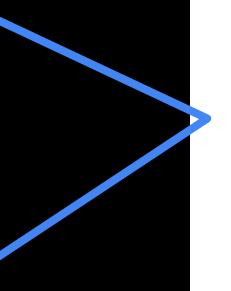
# Define your model
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
```



A neural network with multiple hidden layers (a deep neural network)

```
%tensorflow_version 2.x
import tensorflow as tf
from tensorflow import keras
(x_train, y_train), (x_test, y_test) =
keras.datasets.mnist.load_data()

# Define your model
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
```



A deeper neural network

Compiling a Keras model

```
def rmse(y_true, y_pred):  
    return tf.sqrt(tf.reduce_mean(tf.square(y_pred - y_true)))  
  
model.compile(optimizer="adam", loss="mse", metrics=[rmse, "mse"])
```

Custom Metric

Loss function

Compiling a Keras model

```
def rmse(y_true, y_pred):  
    return tf.sqrt(tf.reduce_mean(tf.square(y_pred - y_true)))  
  
model.compile(optimizer="adam", loss="mse", metrics=[rmse, "mse"])
```

Optimizer

Compiling a Keras model

```
def rmse(y_true, y_pred):  
    return tf.sqrt(tf.reduce_mean(tf.square(y_pred - y_true)))  
  
model.compile(optimizer="adam", loss="mse", metrics=[rmse, "mse"])
```

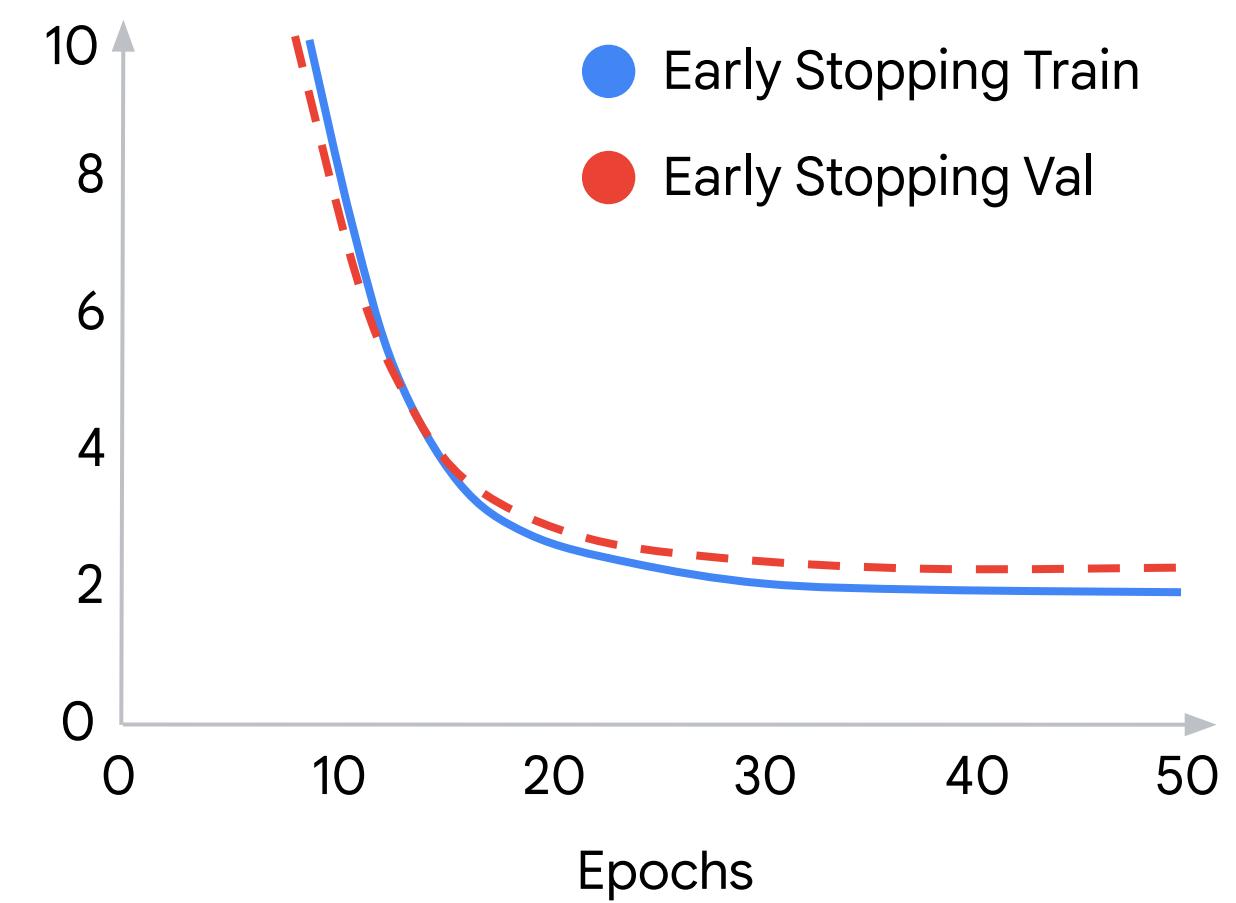
Training a Keras model

```
from tensorflow.keras.callbacks import TensorBoard  
  
steps_per_epoch = NUM_TRAIN_EXAMPLES // (TRAIN_BATCH_SIZE * NUM_EVALS)  
  
history = model.fit(  
    x=trains,  
    steps_per_epoch=steps_per_epoch,  
    epochs=NUM_EVALS,  
    validation_data=evals,  
    callbacks=[TensorBoard(LOGDIR)]  
)
```

This is a trick so that we have control on the total number of examples the model trains on (NUM_TRAIN_EXAMPLES) and the total number of evaluation we want to have during training (NUM_EVALS).

Training a Keras model

```
from tensorflow.keras.callbacks import TensorBoard  
  
steps_per_epoch = NUM_TRAIN_EXAMPLES //  
(TRAIN_BATCH_SIZE * NUM_EVALS)  
  
  
history = model.fit(  
    x=trains,  
    steps_per_epoch=steps_per_epoch,  
    epochs=NUM_EVALS,  
    validation_data=evals,  
    callbacks=[TensorBoard(LOGDIR)]  
)
```



```
%tensorflow_version 2.x
import tensorflow as tf
from tensorflow import keras
(x_train, y_train), (x_test, y_test) =
keras.datasets.mnist.load_data()
# Define your model
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
# Configure and train
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

A deeper neural network

Once trained, the model can be used for prediction

```
predictions = model.predict(input_samples, steps=1)
```

returns a Numpy array of predictions

With the predict() method you can pass

- a Dataset instance
- Numpy array
- a TensorFlow tensor, or list of tensors
- a generator of input samples

steps determines the total number of steps before declaring the prediction round finished. Here, since we have just one example, steps=1.

To serve our model for others to use, we
export the model file and deploy
the **model as a service**.

SavedModel is the universal serialization format for TensorFlow models

```
OUTPUT_DIR = "./export/savedmodel"  
shutil.rmtree(OUTPUT_DIR, ignore_errors=True)  
  
EXPORT_PATH = os.path.join(OUTPUT_DIR,  
    datetime.datetime.now().strftime("%Y%m%d%H%M%S"))  
  
tf.saved_model.save(model, EXPORT_PATH)
```

the directory in which to write the SavedModel

exports a model object to a SavedModel format

a trackable object such as a trained keras model

Create a model object in Vertex AI

```
MODEL_NAME=propertyprice
VERSION_NAME=dnn

if [[ $(gcloud ai-platform models list --format='value(name)' | grep $MODEL_NAME) ]];
then
    echo "$MODEL_NAME already exists"
else
    echo "Creating $MODEL_NAME"
    gcloud ai-platform models create --regions=$REGION $MODEL_NAME
fi
...

```

create the model in AI Platform

Create a version of the model in Vertex AI

```
MODEL_NAME=propertyprice
VERSION_NAME=dnn

...
if [[ $(gcloud ai-platform versions list --model $MODEL_NAME --format='value(name)' |
grep $VERSION_NAME) ]]; then
    echo "Deleting already existing $MODEL_NAME:$VERSION_NAME ... "
    echo yes | gcloud ai-platform versions delete --model=$MODEL_NAME $VERSION_NAME
    echo "Please run this cell again if you don't see a Creating message ... "
    sleep 2
fi
```

create the model version in AI Platform

Deploy SavedModel using gcloud ai-platform

```
gcloud ai-platform versions create \  
  --model=$MODEL_NAME $VERSION_NAME \  
  --framework=tensorflow \  
  --python-version=3.5 \  
  --runtime-version=2.1 \  
  --origin=$EXPORT_PATH \  
  --staging-bucket=gs://$BUCKET
```

specify the model name
and version

EXPORT_PATH denotes
location of SavedModel
directory

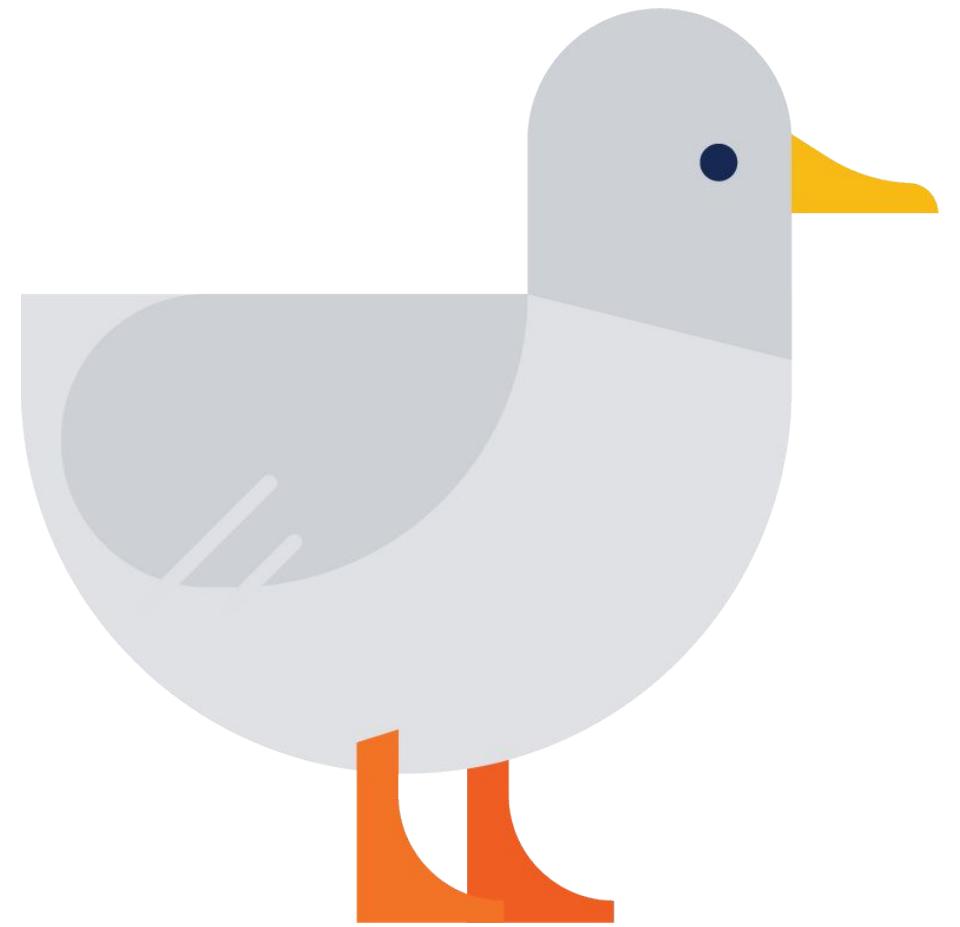
Make predictions using gcloud ai-platform

```
input.json = {"sq_footage": 3140,  
             "type": 'house'}
```

```
gcloud ai-platform predict \  
  --model propertyprice \  
  --version dnn \  
  --json-instances input.json
```

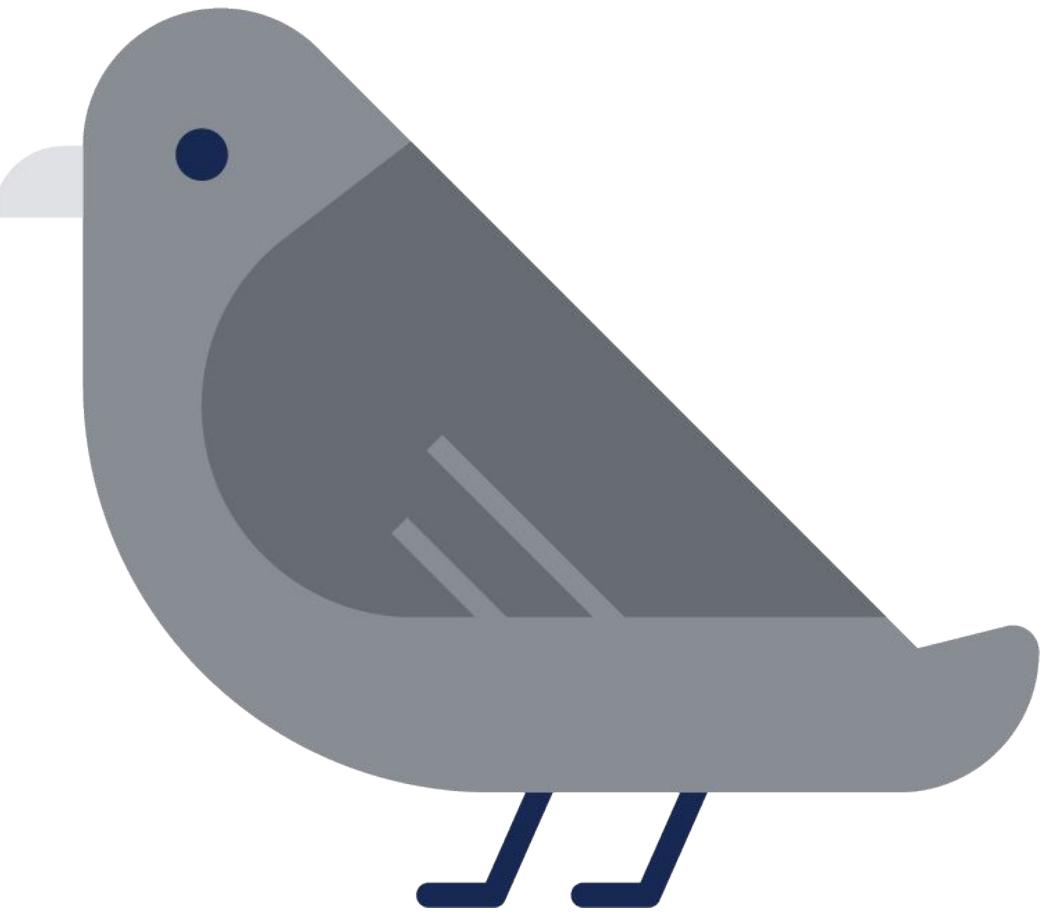
specify the name and version
of the deployed model

json file for prediction

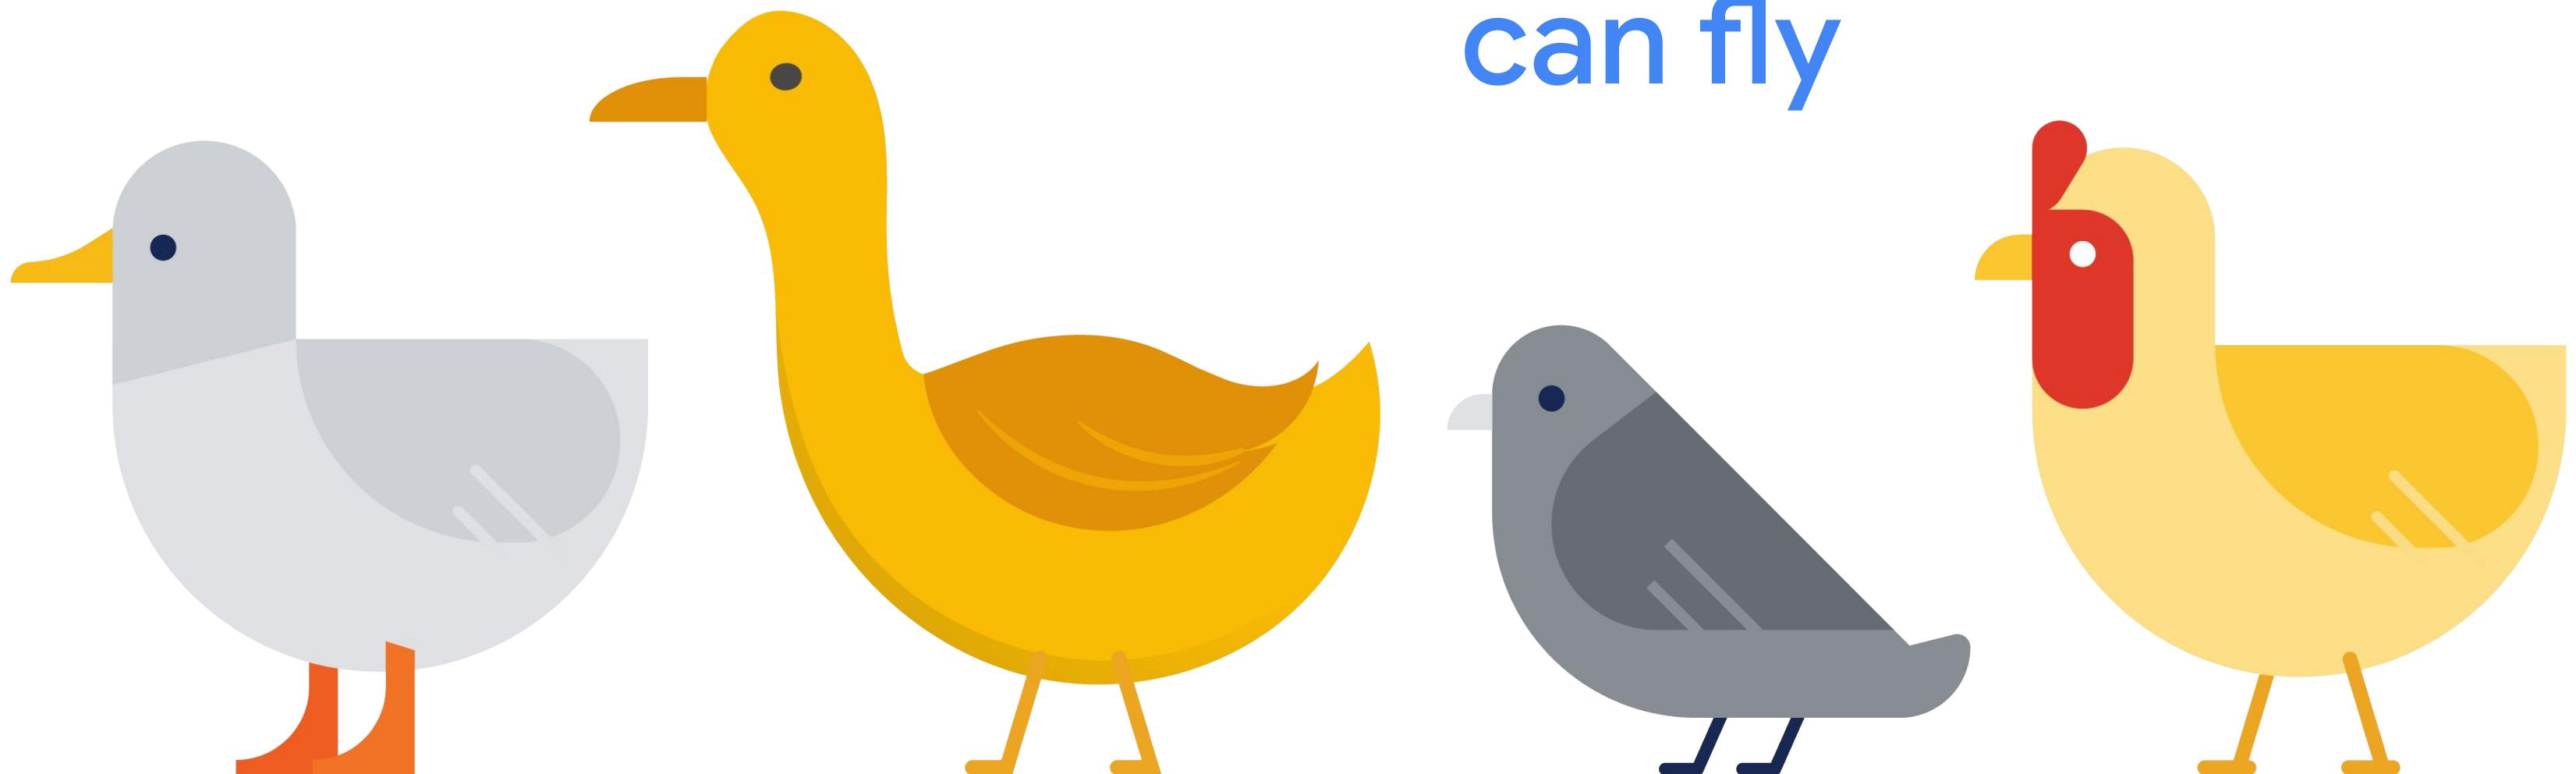


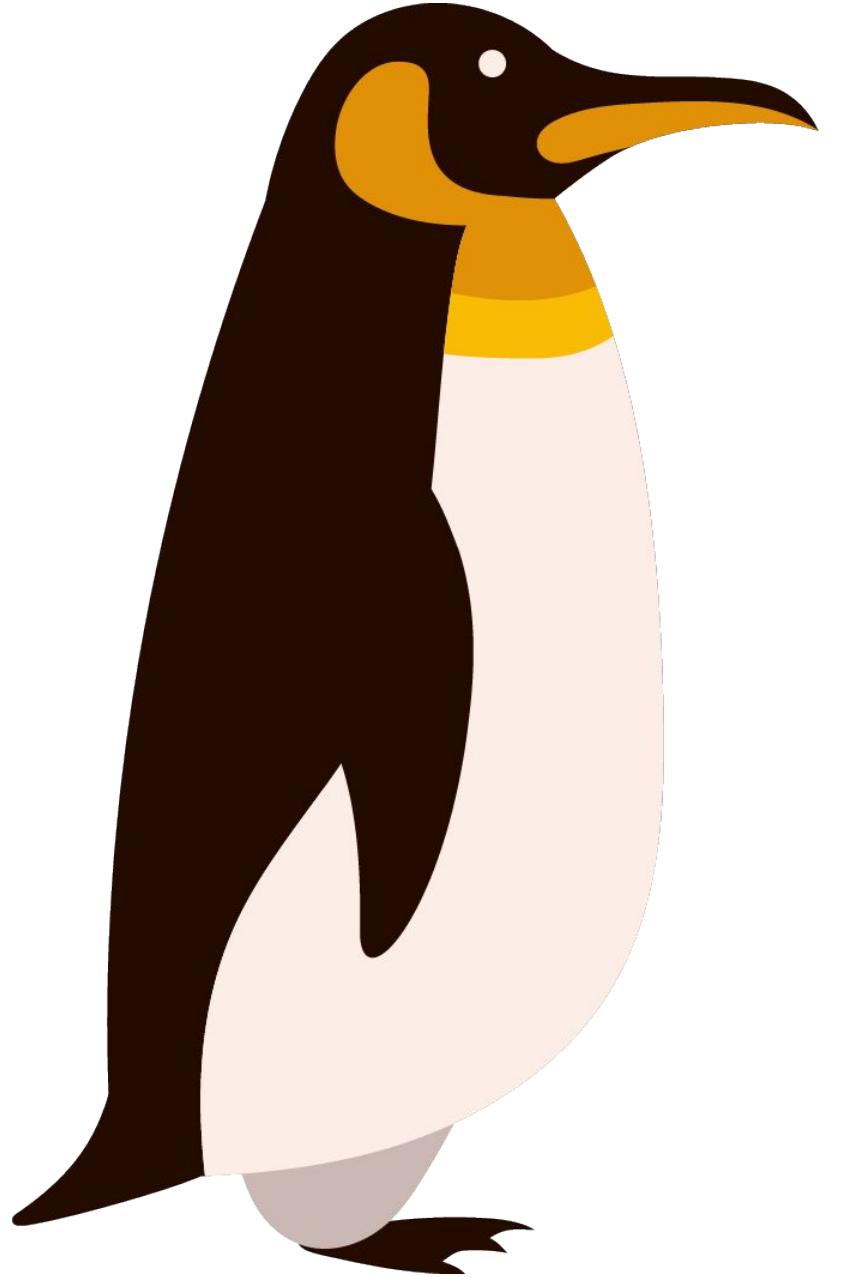
Seagulls can fly

Pigeons **can** fly



Animals with wings
can fly

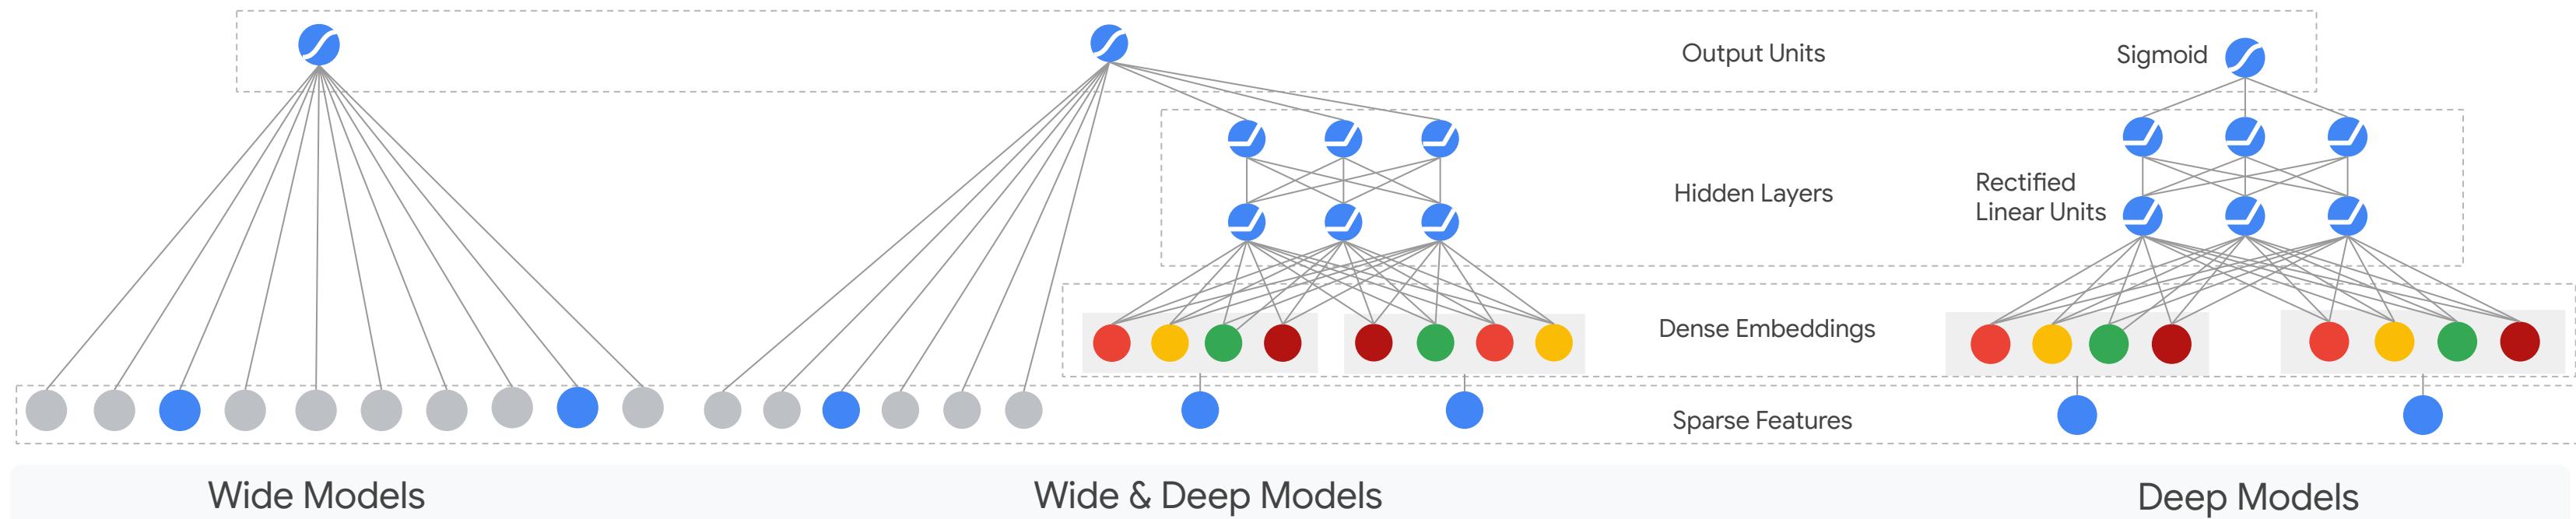




Penguins...

Using Wide and Deep learning

“ Combine the power of memorization and generalization
on one unified machine learning model. ”



Memorization + Generalization

- Memorization: “Seagulls can fly.” “Pigeons can fly.”
- Generalization: “**Animals with wings** can fly.”
- Generalization + memorizing exceptions: “**Animals with wings** can fly, but penguins cannot fly.”



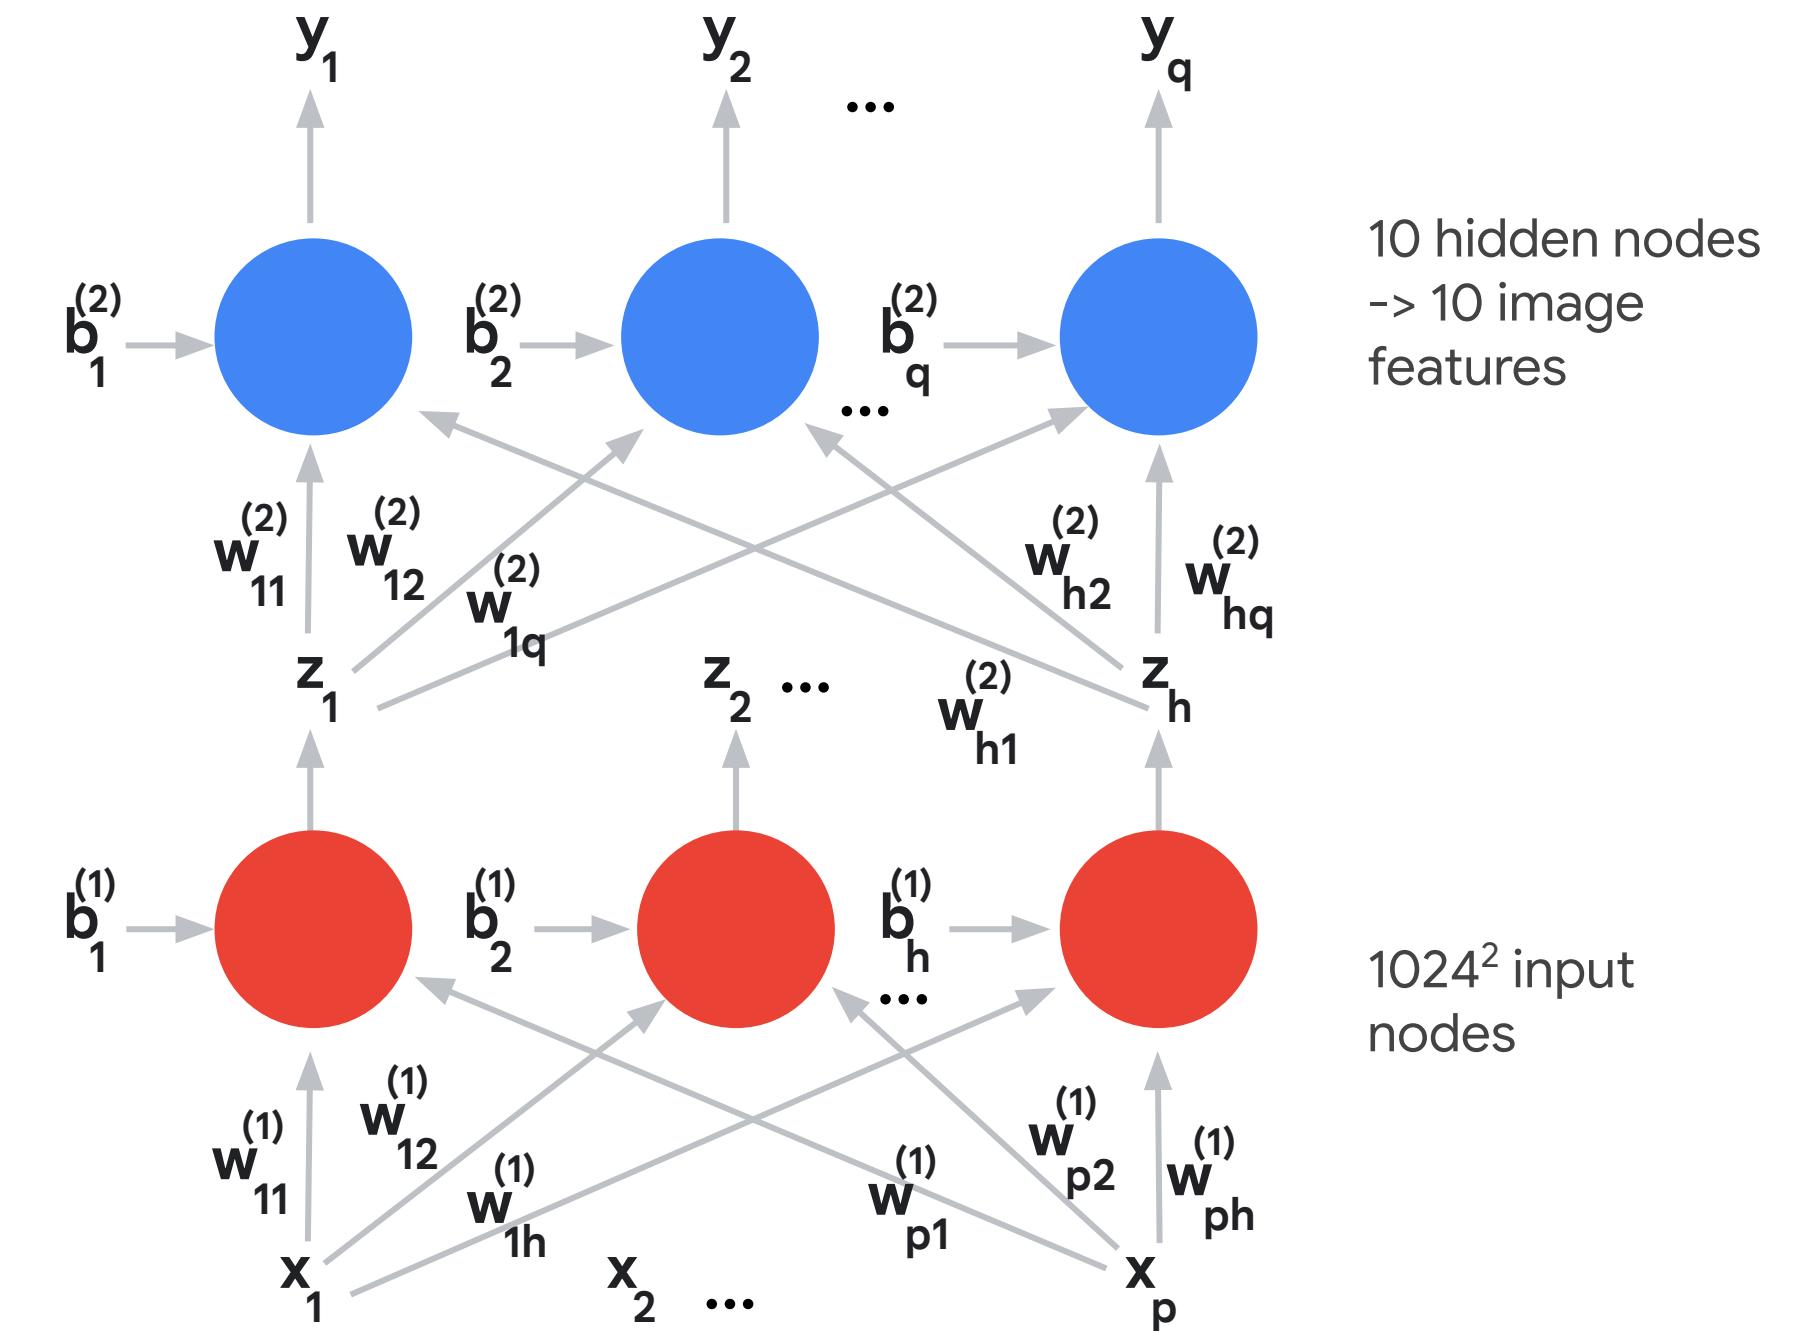
**Linear models are
good for sparse,
independent features**

```
[ 0.  1.  0.  0.  0.  0.  0.  0.  0.  0.]  
[ 0.  0.  0.  0.  0.  0.  0.  1.  0.  0.]  
[ 0.  0.  0.  0.  0.  0.  0.  0.  0.  1.]  
[ 0.  1.  0.  0.  0.  0.  0.  0.  0.  0.]  
[ 0.  0.  0.  0.  0.  0.  0.  1.  0.  0.]  
[ 0.  0.  0.  1.  0.  0.  0.  0.  0.  0.]  
[ 0.  0.  0.  0.  0.  0.  1.  0.  0.  0.]  
[ 0.  1.  0.  0.  0.  0.  0.  0.  0.  0.]  
[ 0.  0.  0.  0.  1.  0.  0.  0.  0.  0.]  
[ 0.  1.  0.  0.  0.  0.  0.  0.  0.  0.]
```

DNNs are good for dense,
highly correlated
features

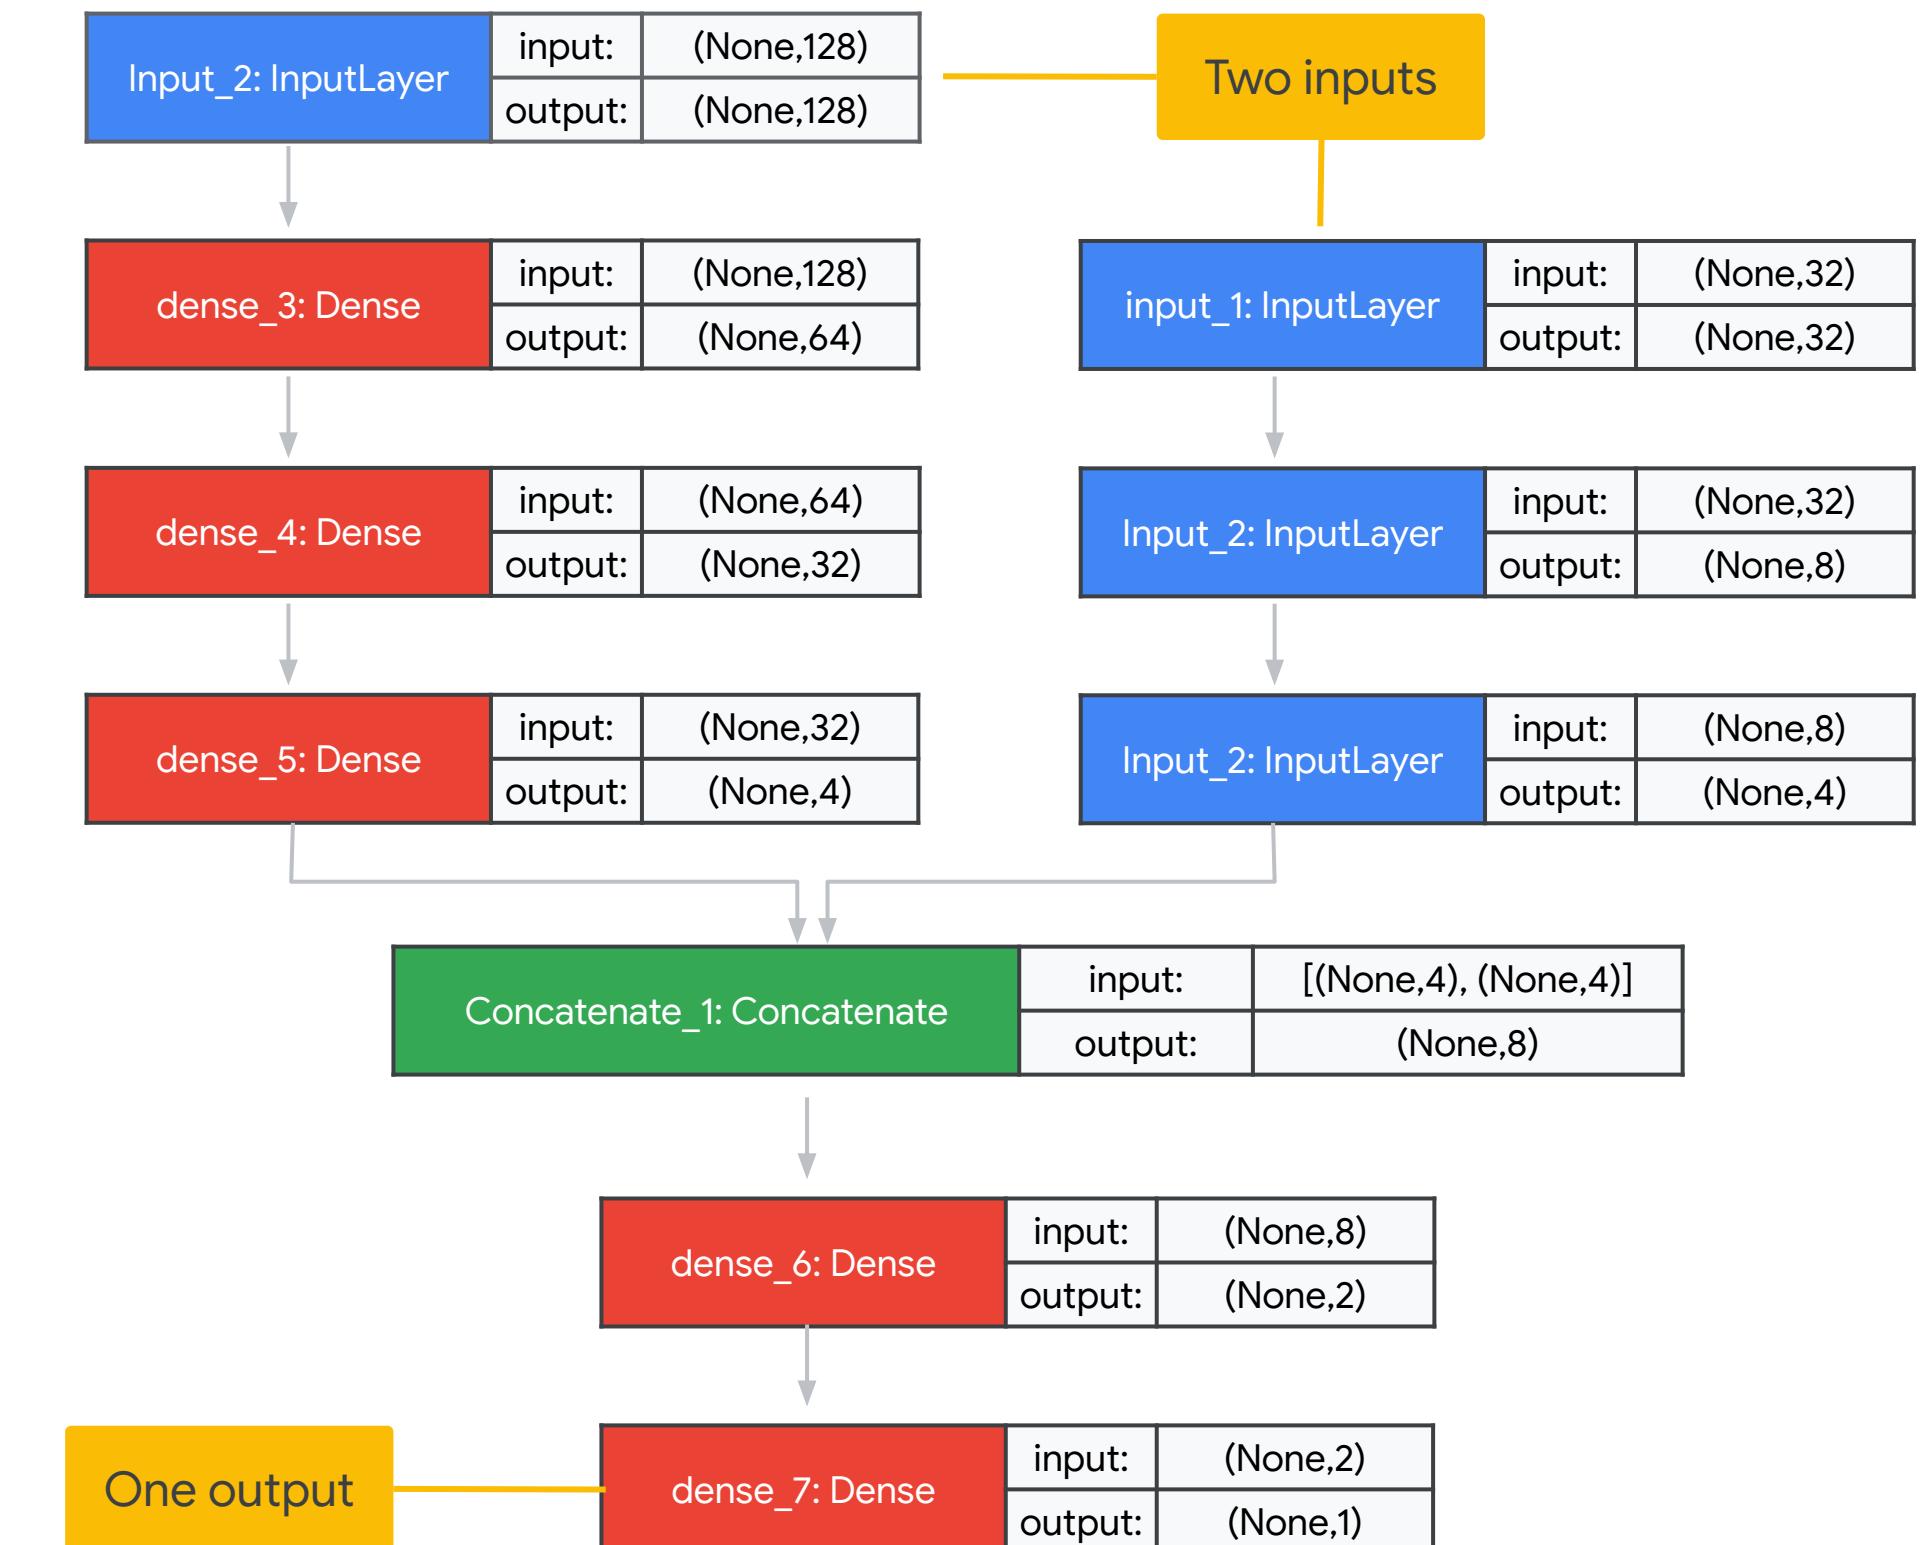


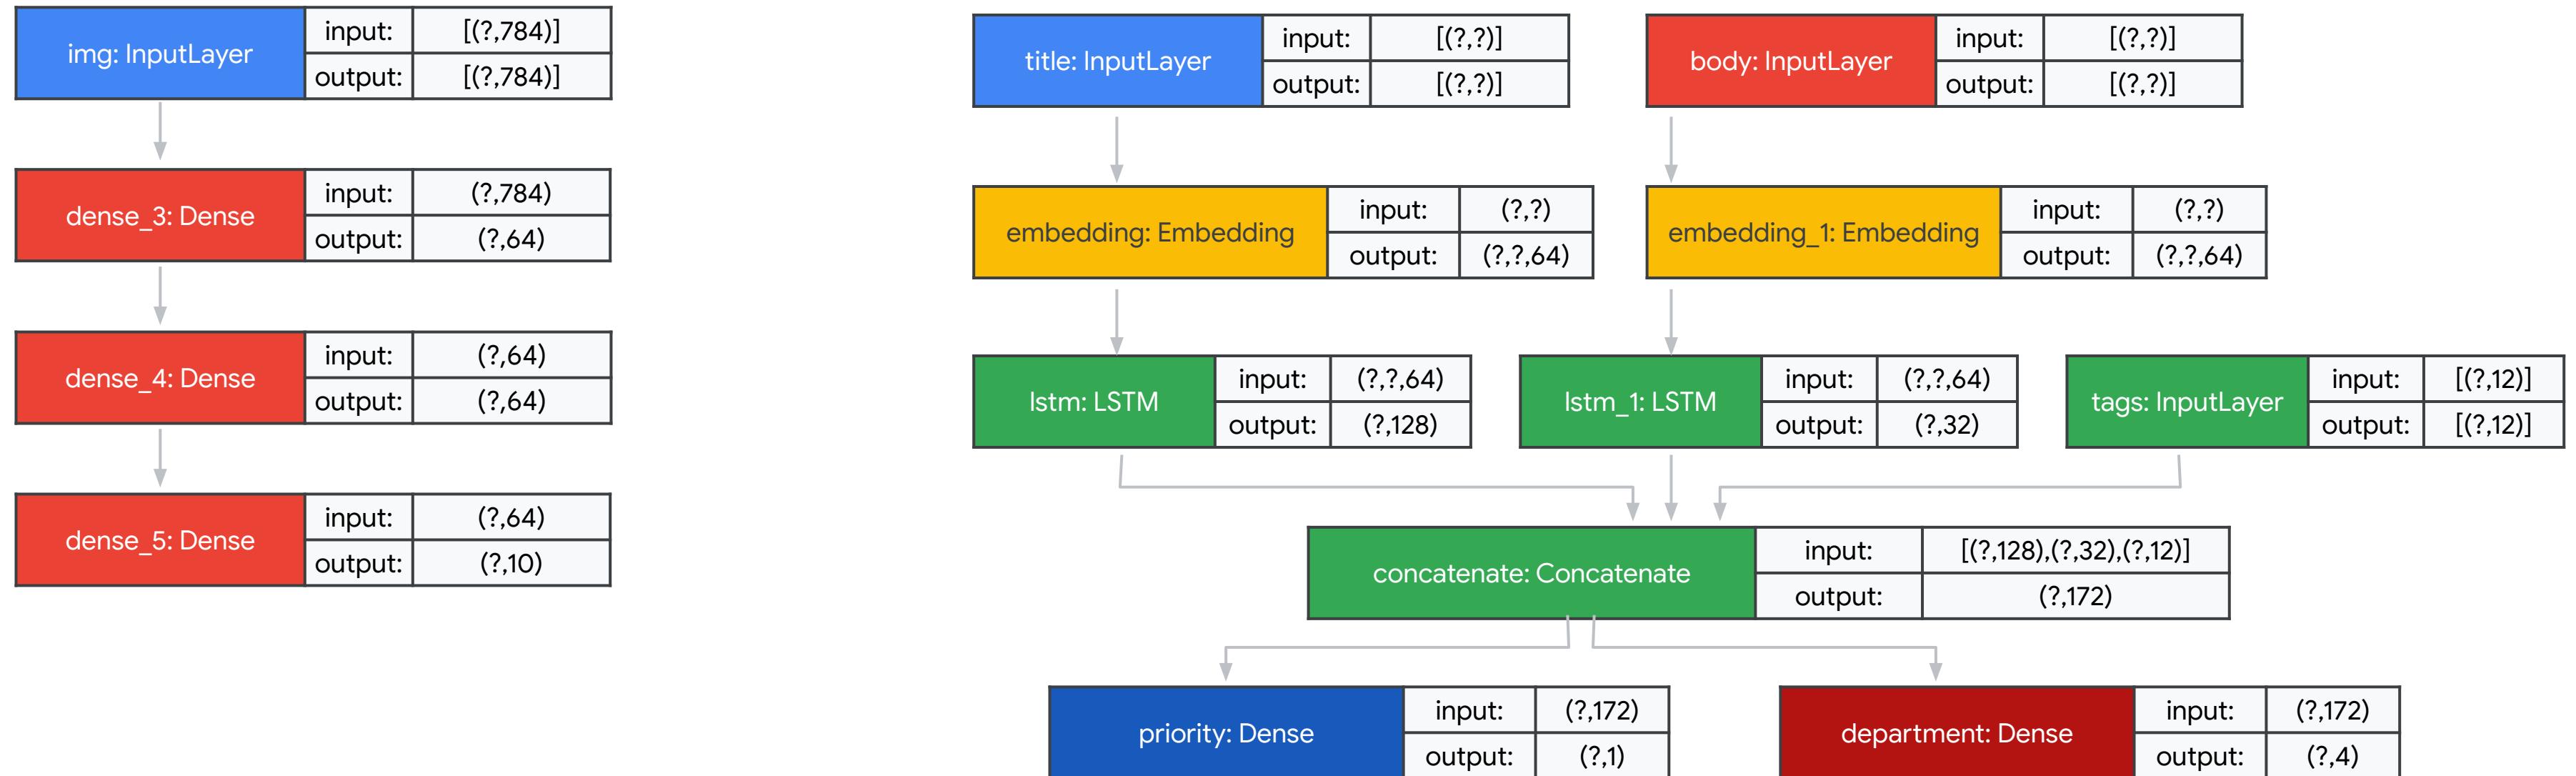
pixel_values (



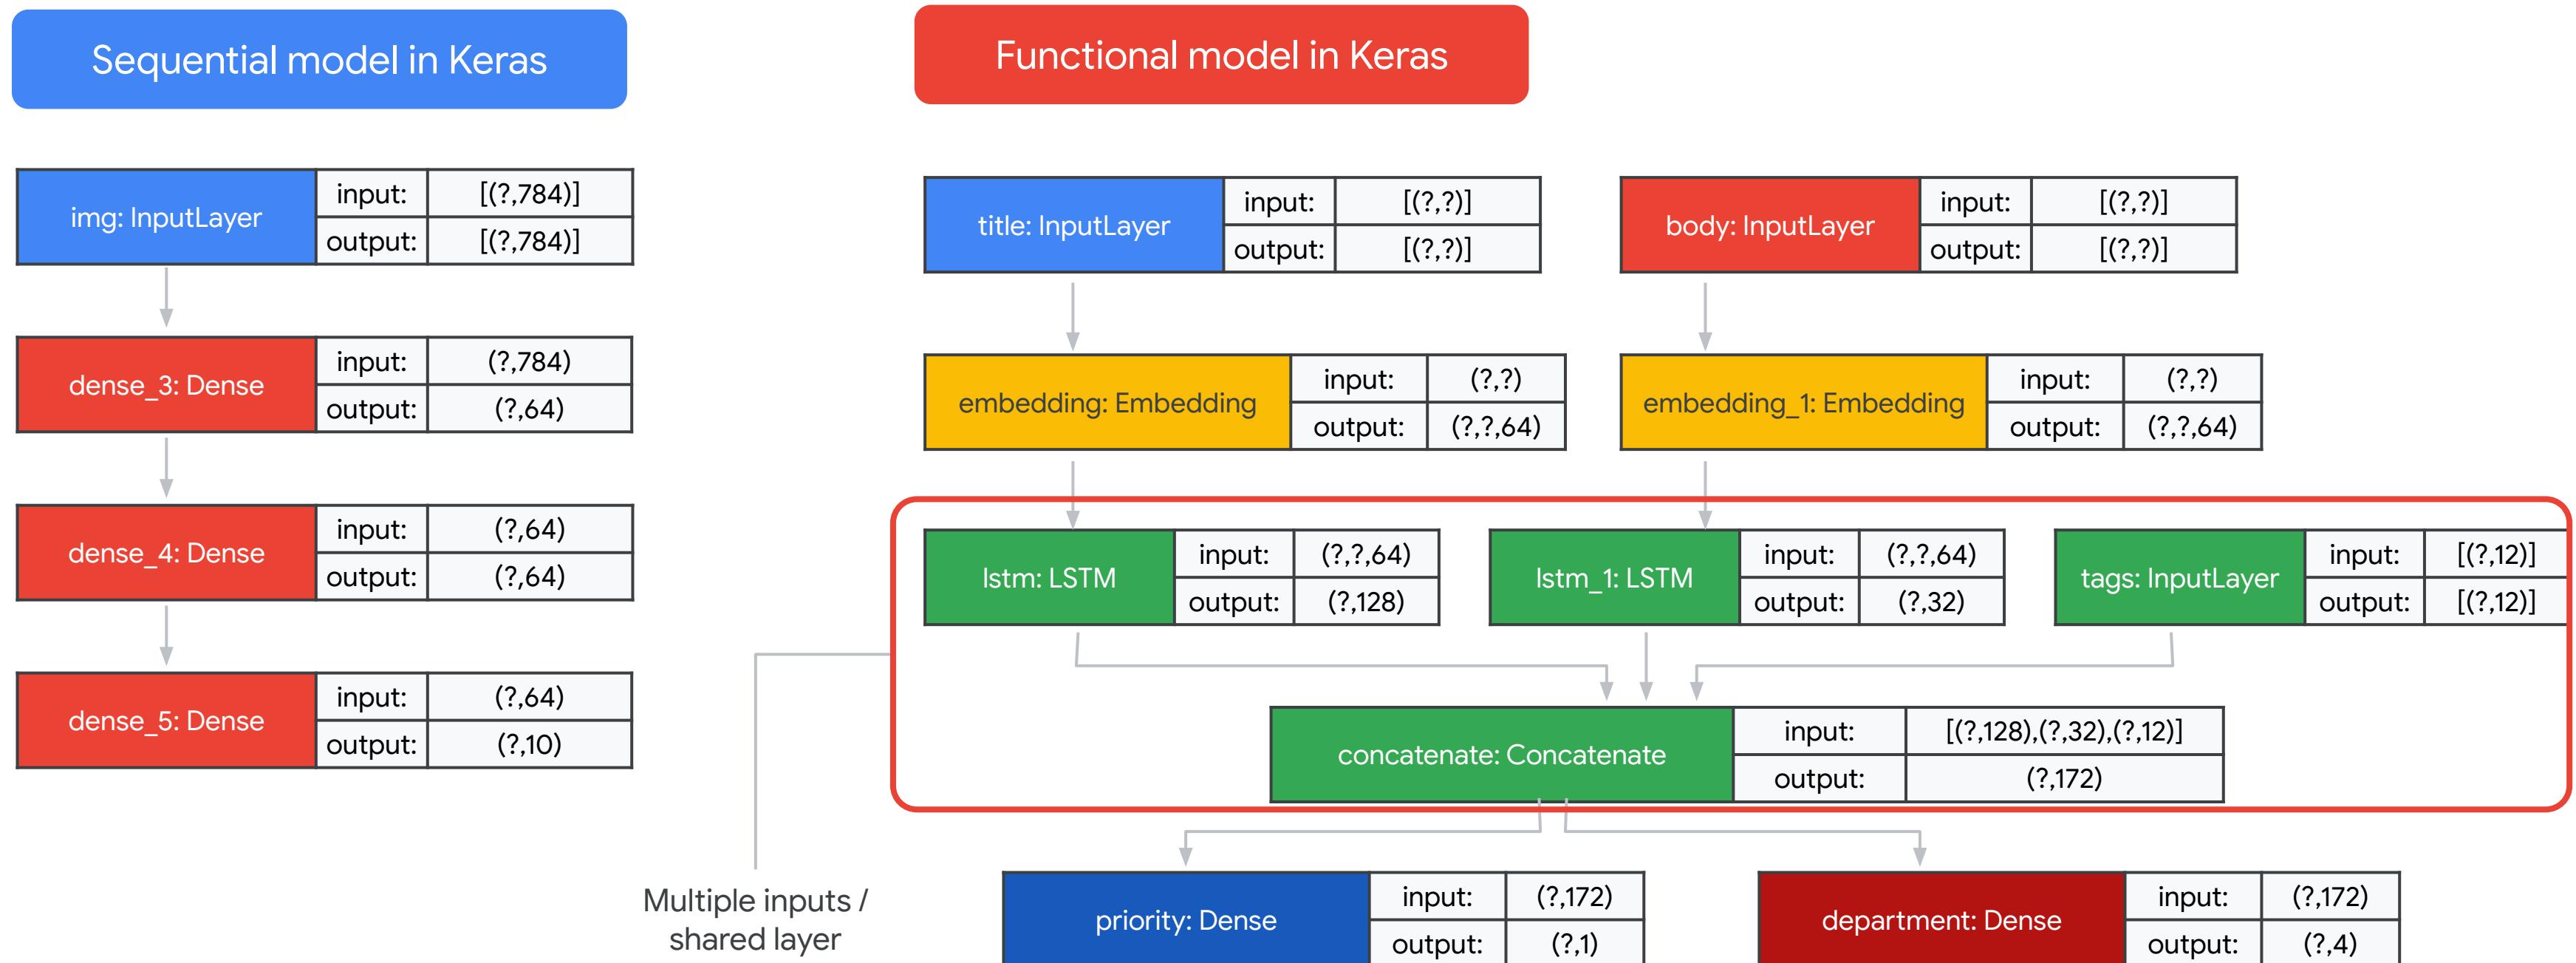
Wide and deep networks using Keras

Functional API



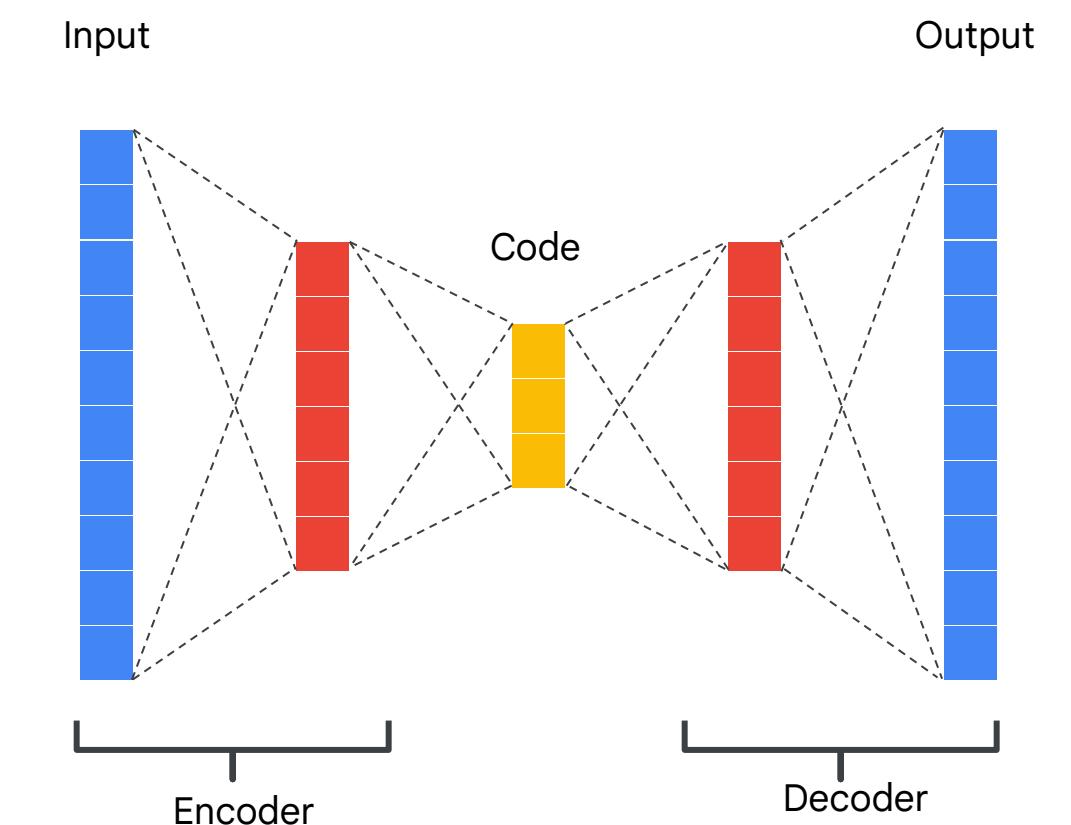


Functional API is more flexible than Sequential API



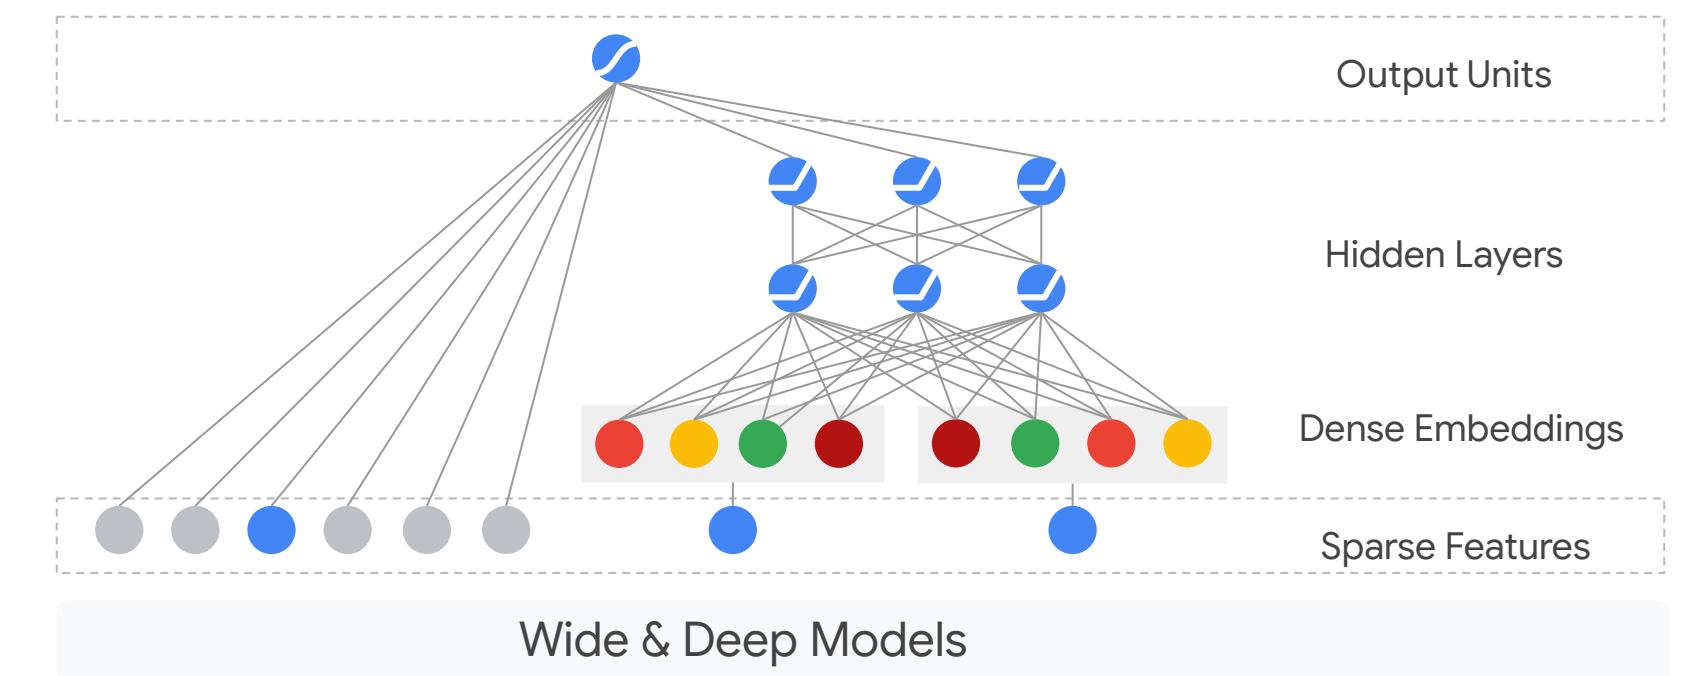
Models are created by specifying their inputs and outputs in a graph of layers

```
encoder_input = keras.Input(shape=(28, 28, 1), name='img')  
x = layers.Dense(16, activation='relu')(encoder_input)  
x = layers.Dense(10, activation='relu')(x)  
x = layers.Dense(5, activation='relu')(x)  
encoder_output = layers.Dense(3, activation='relu')(x)  
  
encoder = keras.Model(encoder_input, encoder_output, name='encoder')  
  
x = layers.Dense(5, activation='relu')(encoder_output)  
x = layers.Dense(10, activation='relu')(x)  
x = layers.Dense(16, activation='relu')(x)  
decoder_output = layers.Dense(28, activation='linear')(x)  
  
autoencoder = keras.Model(encoder_input, decoder_output, name='autoencoder')
```



Creating a Wide and Deep model in Keras

```
INPUT_COLS = [  
    'pickup_longitude',  
    'pickup_latitude',  
    'dropoff_longitude',  
    'dropoff_latitude',  
    'passenger_count'  
]  
  
# Prepare input feature columns  
inputs = {colname : layers.Input(name=colname,  
shape=(), dtype='float32')  
          for colname in INPUT_COLS  
}  
  
...
```



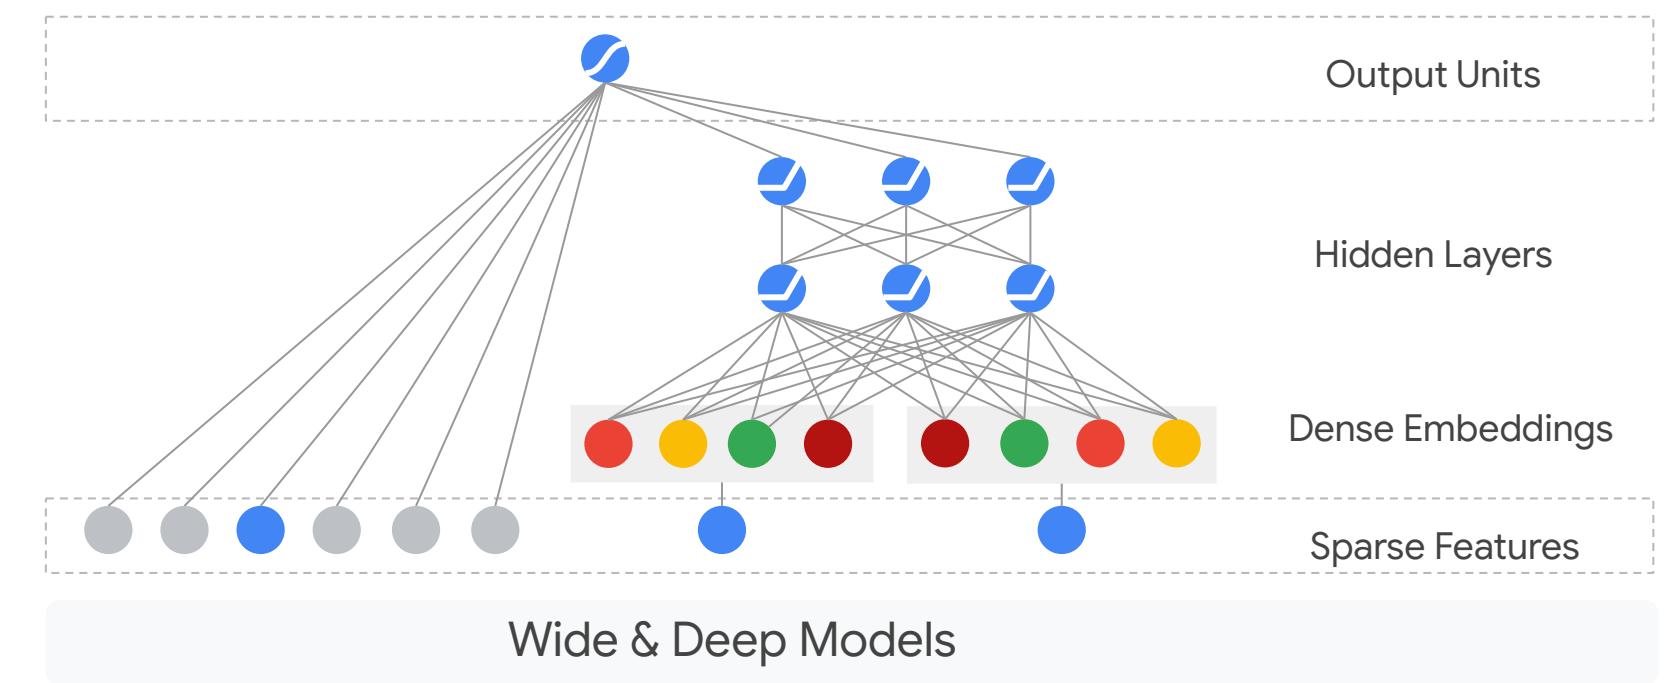
Creating a Wide and Deep model in Keras

```
# Create deep columns
deep_columns = [
    # Embedding_column to "group" together ...
    fc.embedding_column(fc_crossed_pd_pair, 10),

    # Numeric columns
    fc.numeric_column("pickup_latitude"),
    fc.numeric_column("pickup_longitude"),
    fc.numeric_column("dropoff_longitude"),
    fc.numeric_column("dropoff_latitude")]

# Create the deep part of model
deep_inputs = layers.DenseFeatures(
    (deep_columns, name='deep_inputs'))(inputs)
x = layers.Dense(30, activation='relu')(deep_inputs)
x = layers.Dense(20, activation='relu')(x)

deep_output = layers.Dense(10, activation='relu')(x)
```

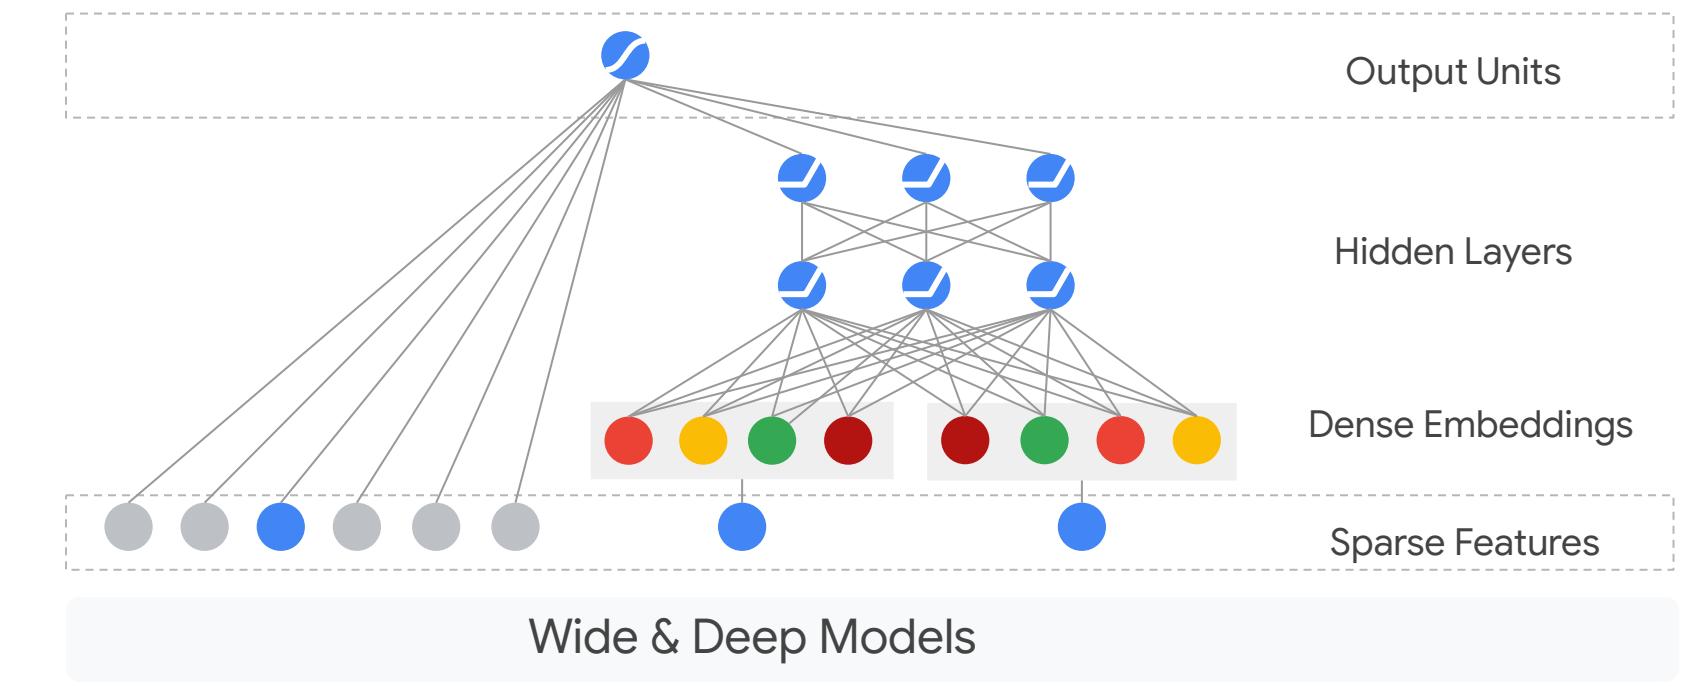


created in the previous slide

Creating a Wide and Deep model in Keras

```
# Create wide columns
wide_columns = [
    # One-hot encoded feature crosses
    fc.indicator_column(fc_crossed_dloc),
    fc.indicator_column(fc_crossed_ploc),
    fc.indicator_column(fc_crossed_pd_pair)
]

# Create the wide part of model
wide = layers.DenseFeatures(wide_columns,
name='wide_inputs')(inputs)
```

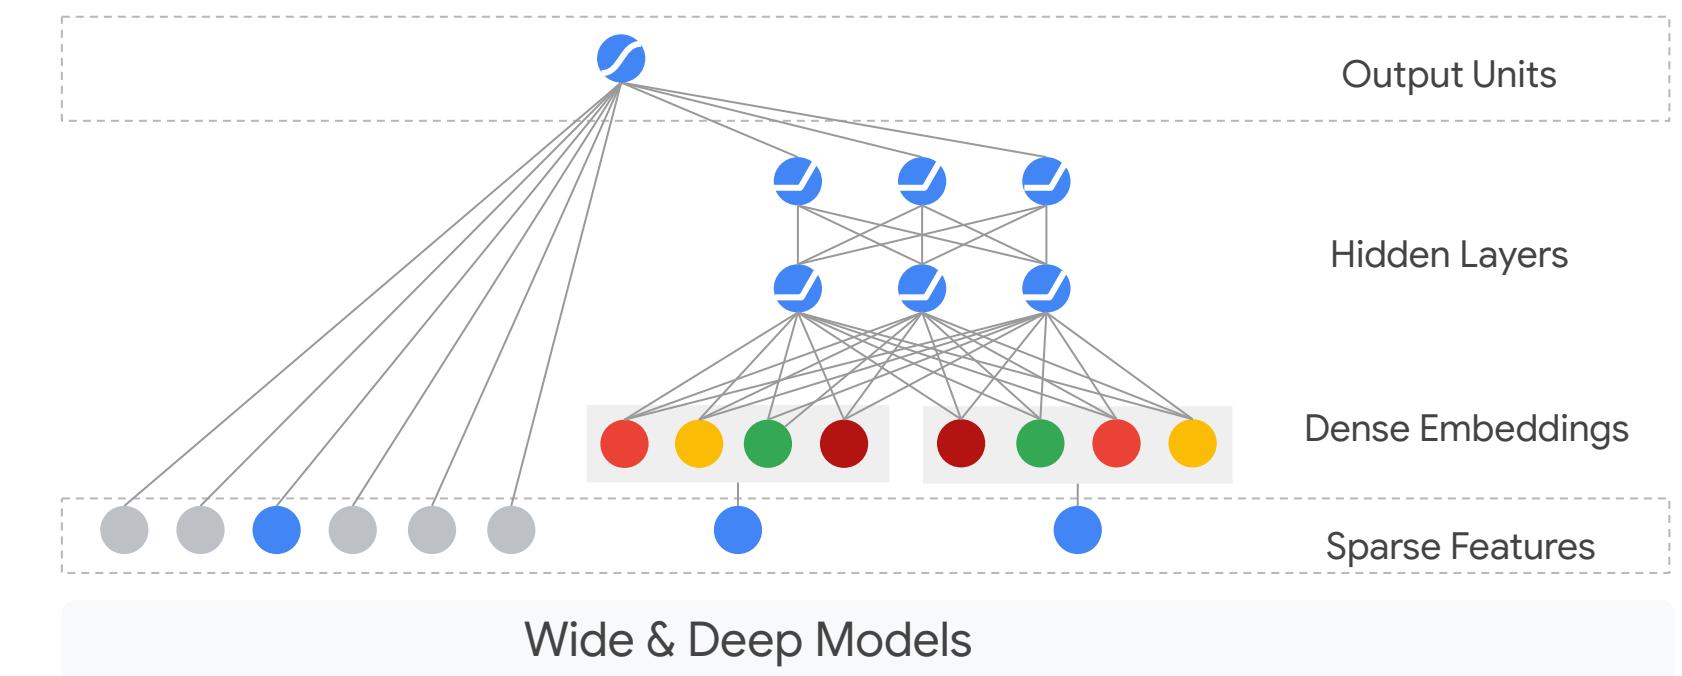


created in the previous slide

Creating a Wide and Deep model in Keras

```
# Combine outputs
combined = concatenate(inputs=[deep, wide],
                       name='combined')
output = layers.Dense(1,
                      activation=None,
                      name='prediction')(combined)

# Finalize model
model = keras.Model(inputs=list(inputs.values()),
                     outputs=output,
                     name='wide_and_deep')
model.compile(optimizer="adam",
              loss="mse",
              metrics=[rmse, "mse"])
```



To finalize the model, specify
the inputs and outputs

Strengths and weaknesses of the Functional API

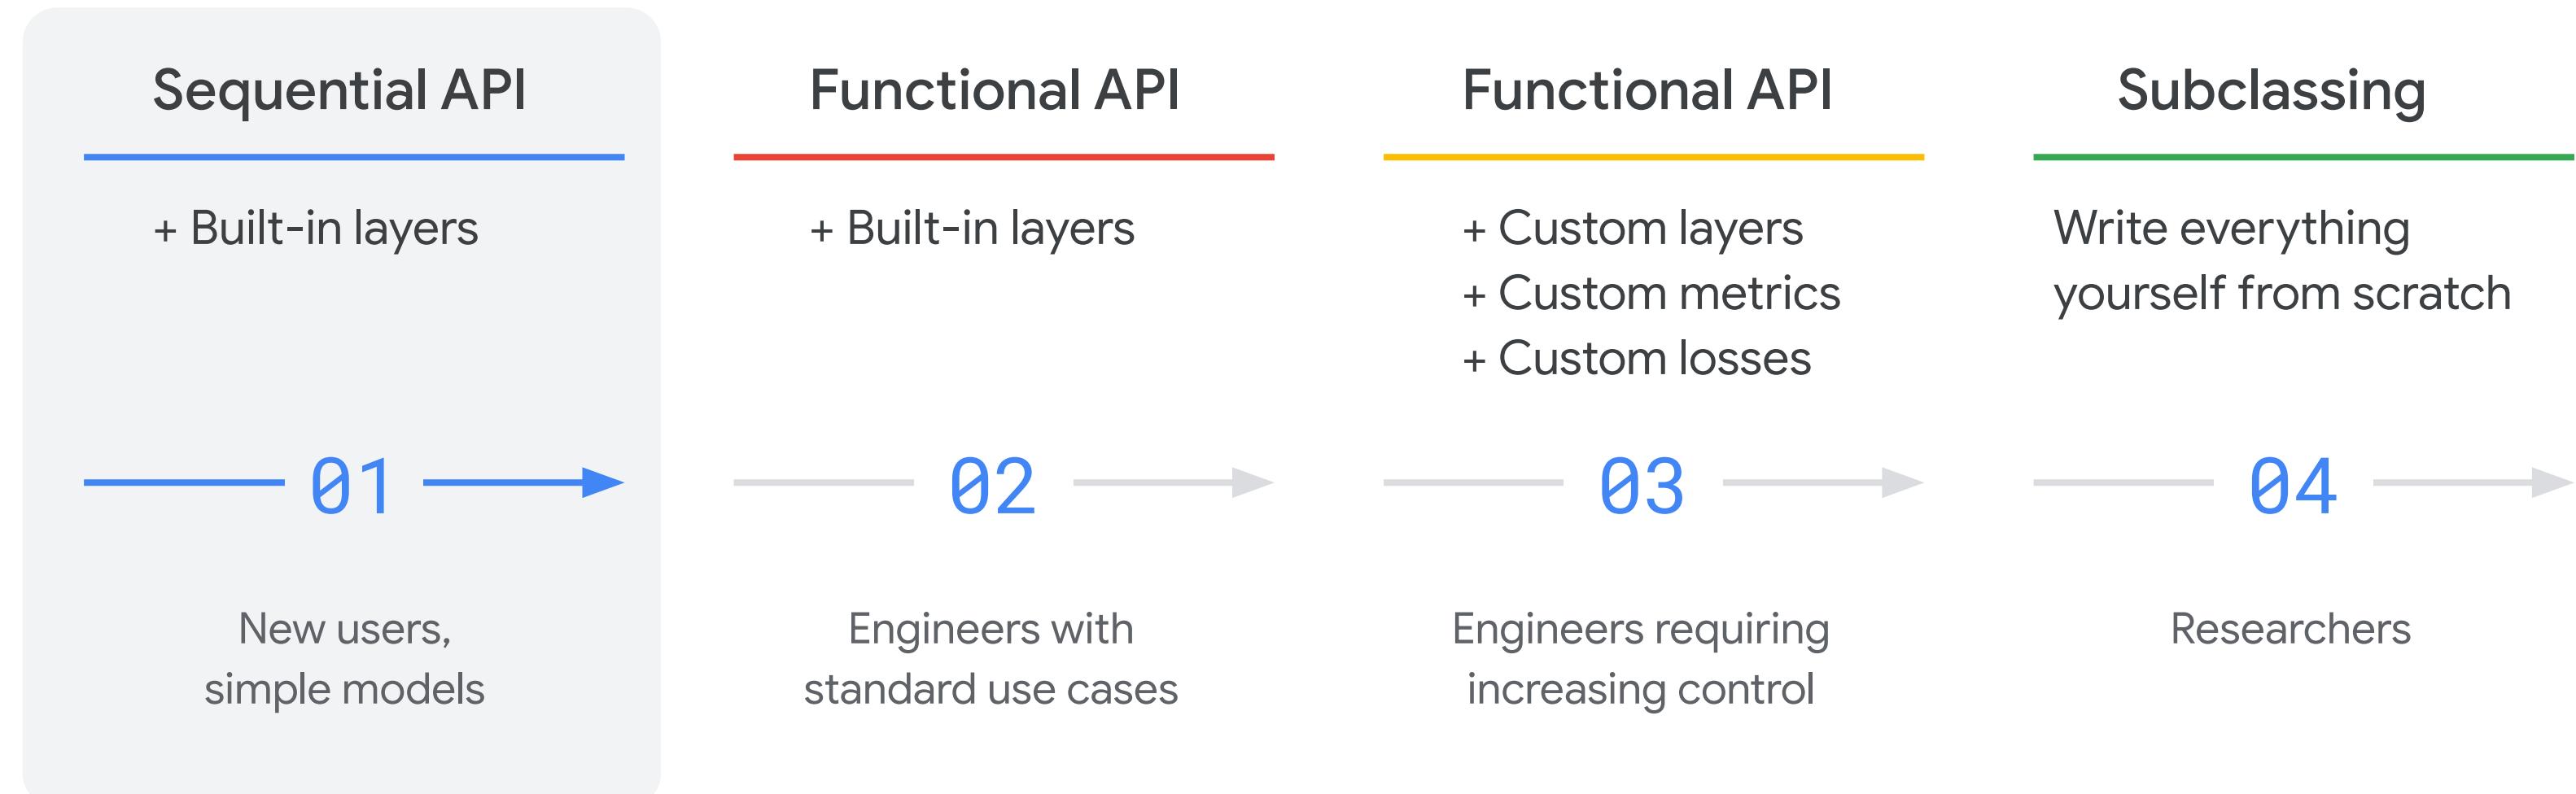
Strengths

- Less verbose than using keras.Model subclasses
- Validates your model while you're defining it
- Your model is plottable and inspectable
- Your model can be serialized or cloned

Weaknesses

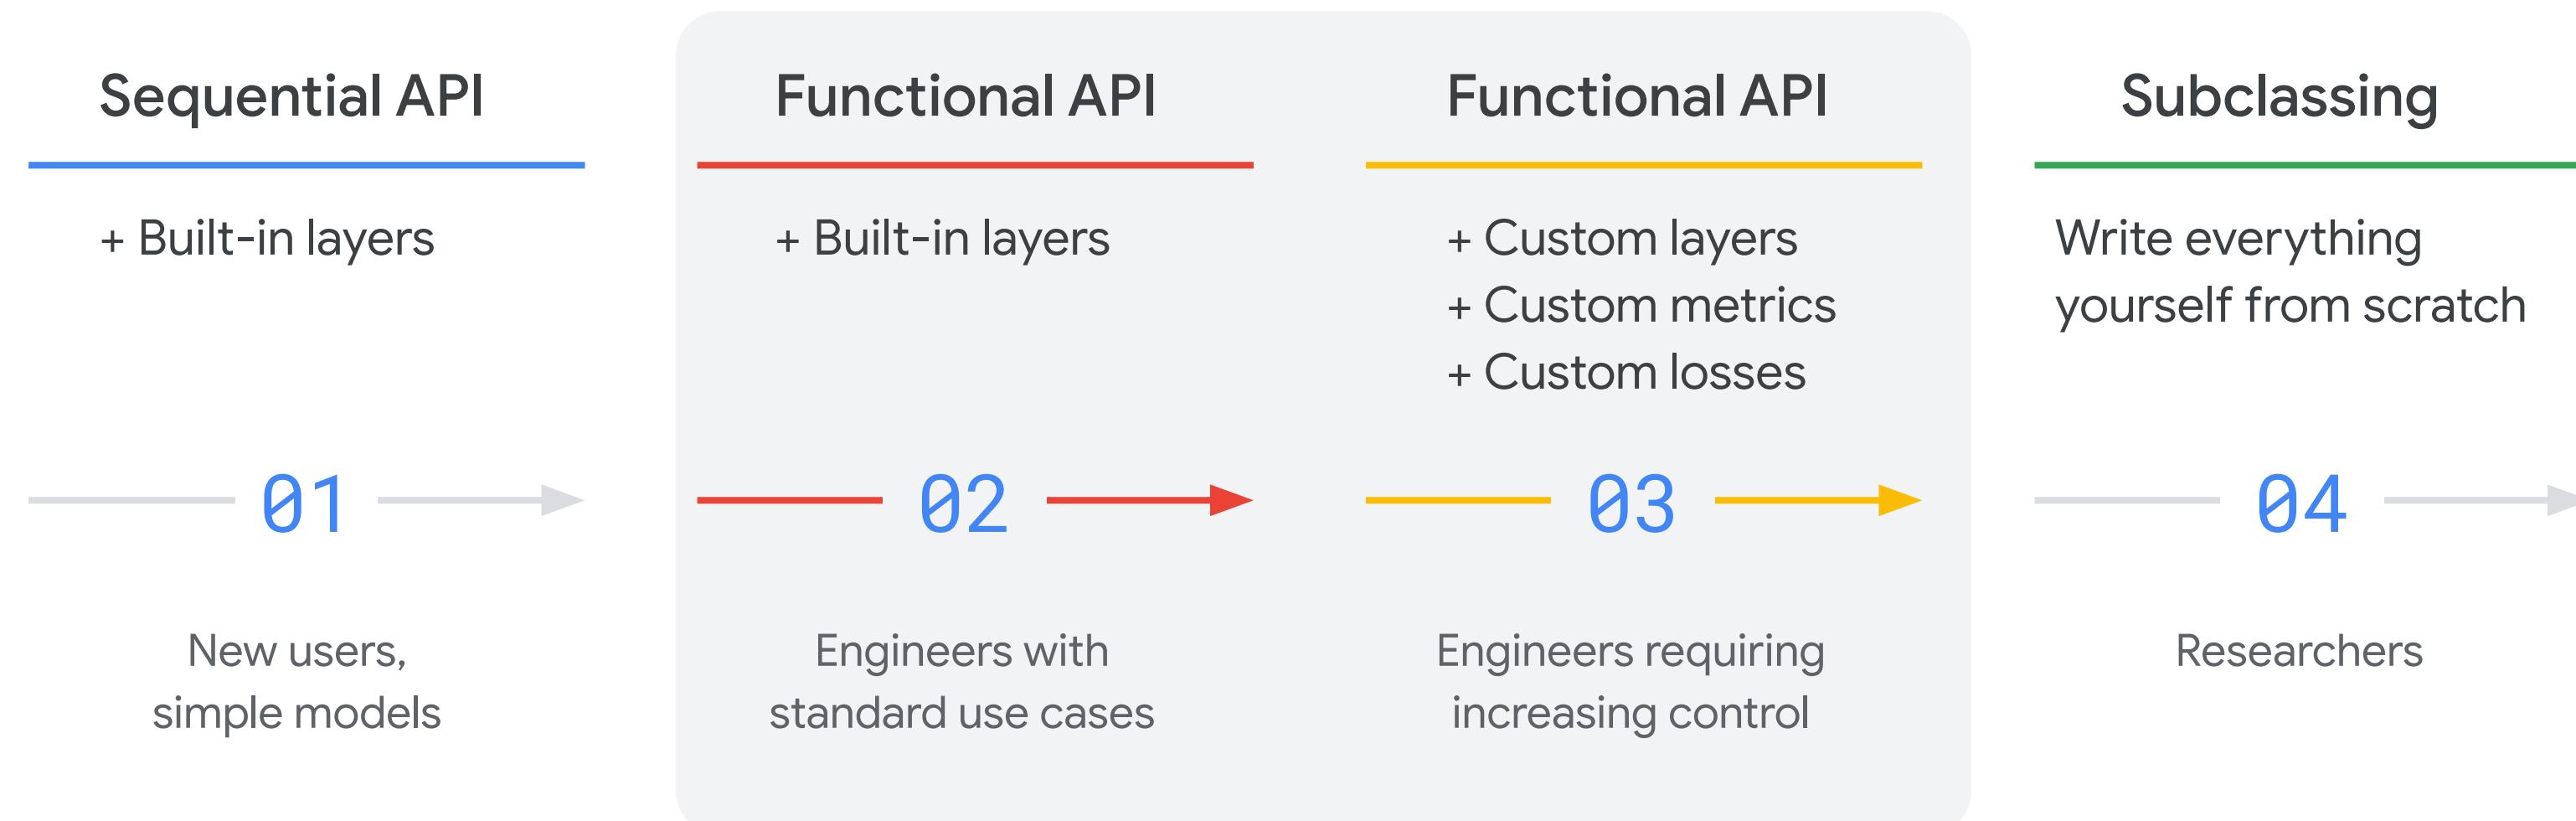
- Doesn't support dynamic architectures
- Sometimes you have to write from scratch and you need to build subclasses, e.g. custom training or inference layers

Model building: From simple to arbitrarily flexible



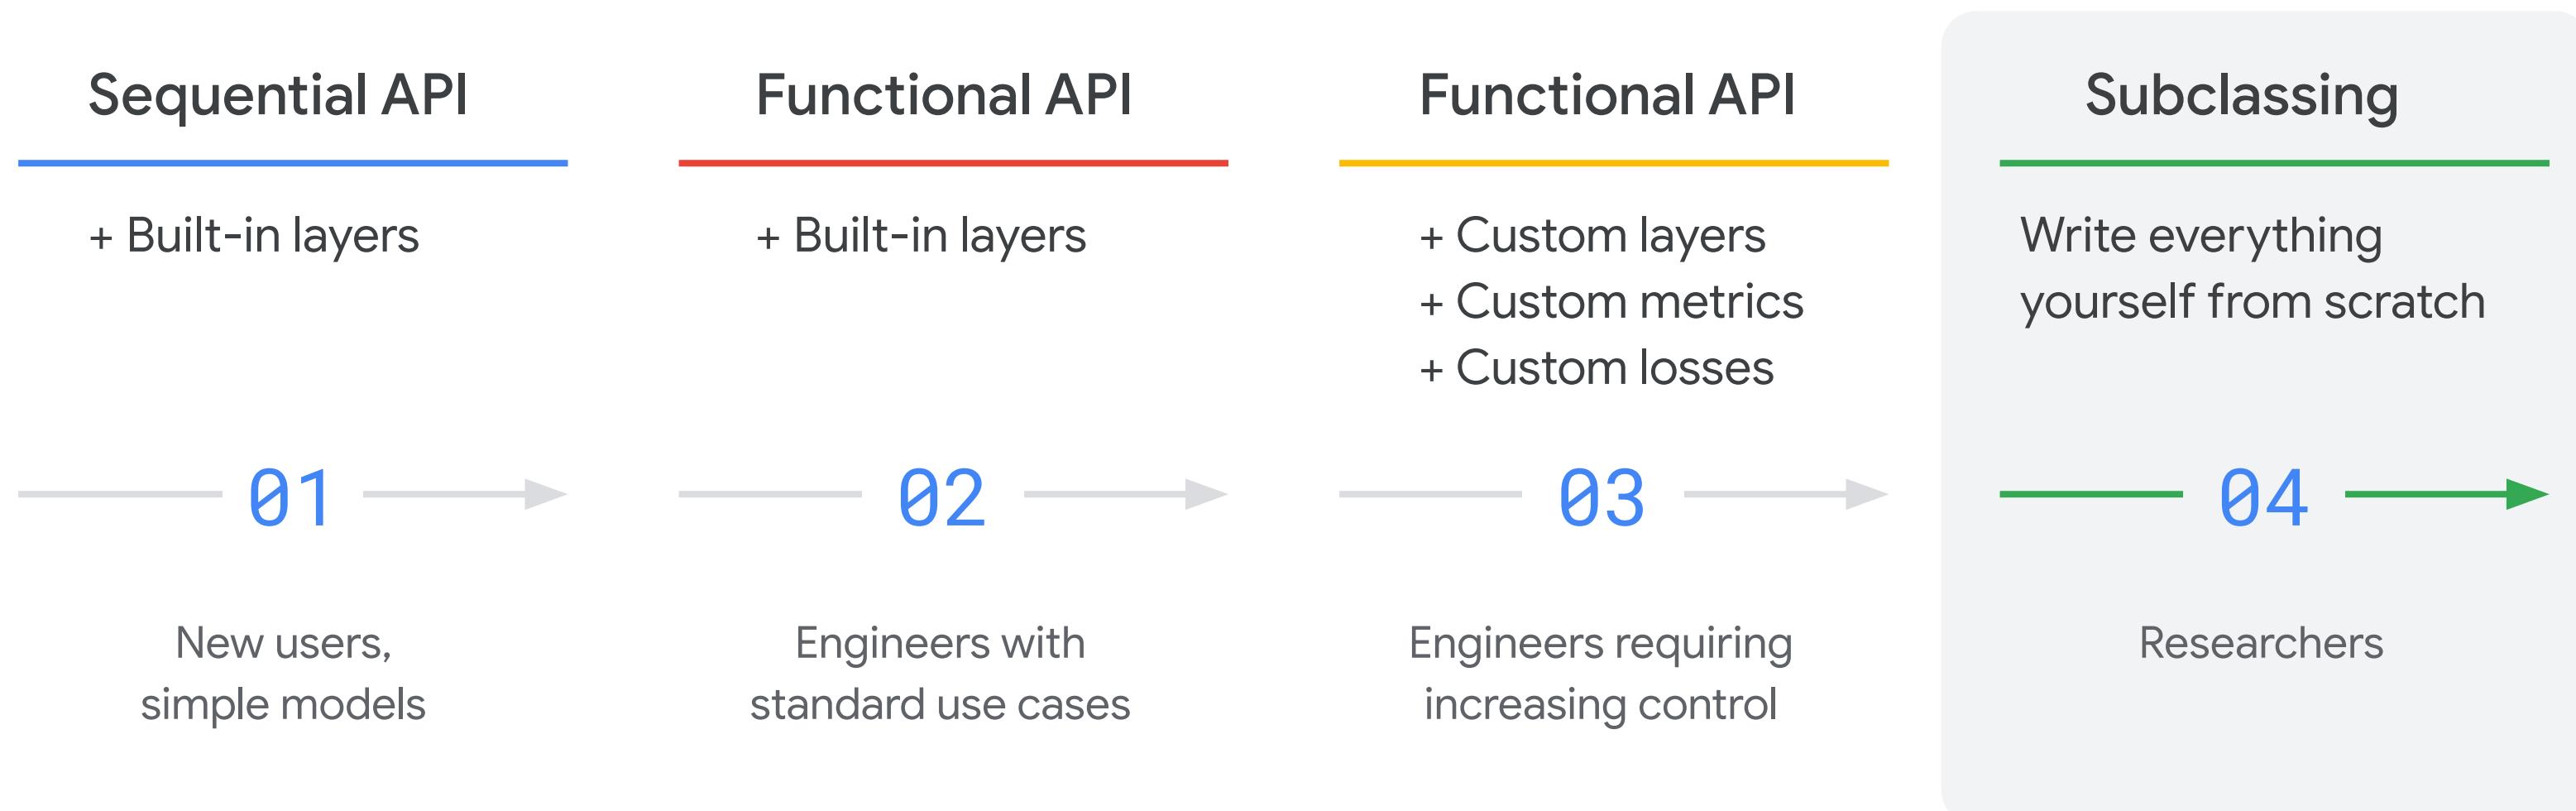
Progressive disclosure of complexity

Model building: From simple to arbitrarily flexible



Progressive disclosure of complexity

Model building: From simple to arbitrarily flexible



Progressive disclosure of complexity

Define layers

```
import tensorflow as tf

class MyModel(tf.keras.Model):

    def __init__(self):
        super(myModel, self).__init__()
        self.dense1 = tf.keras.layers.Dense(4, activation=tf.nn.relu)
        self.dense2 = tf.keras.layers.Dense(5, activation=tf.nn.softmax)

    def call(self, inputs):
        x = self.dense1(inputs)
        return self.dense2(x)

model = MyModel()
```

Implement the forward pass

```
import tensorflow as tf

class MyModel(tf.keras.Model):

    def __init__(self):
        super(myModel, self).__init__()
        self.dense1 = tf.keras.layers.Dense(4, activation=tf.nn.relu)
        self.dense2 = tf.keras.layers.Dense(5, activation=tf.nn.softmax)

    def call(self, inputs):
        x = self.dense1(inputs)
        return self.dense2(x)

model = MyModel()
```

Constructor method

```
import tensorflow as tf

class MyModel(tf.keras.Model):

    def __init__(self):
        super(myModel, self).__init__()
        self.dense1 = tf.keras.layers.Dense(4, activation=tf.nn.relu)
        self.dense2 = tf.keras.layers.Dense(5, activation=tf.nn.softmax)
```

Constructor

Call method

```
def call(self, inputs):  
    x = self.dense1(inputs)  
    return self.dense2(x)  
  
model = MyModel()
```

Call method

The “constructor” and “call method”

```
import tensorflow as tf

class MyModel(tf.keras.Model):
    def __init__(self):
        super(myModel, self).__init__()
        self.dense1 = tf.keras.layers.Dense(4, activation=tf.nn.relu)
        self.dense2 = tf.keras.layers.Dense(5, activation=tf.nn.softmax)

    def call(self, inputs):
        x = self.dense1(inputs)
        return self.dense2(x)

model = MyModel()
```

Subclass the model directly

Constructor

Call method

The “constructor” and “call method”

```
class MyModel(tf.keras.Model):
    def __init__(self, num_classes=10):
        super(myModel, self).__init__(name='my_model')
        # Define your layers here.
        self.dense_1 = layers.Dense(32, activation='relu')
        self.dense_2 = layers.Dense(num_classes, activation='softmax')

    def call(self, inputs):
        # Define your forward pass here
        # using layers you previously defined in '__init__'
        x = self.dense_1(inputs)
        return self.dense_2(x)
```

Subclass the model directly

Constructor

Call method

The “constructor” and “call method”

```
class MyModel(tf.keras.Model):
    def __init__(self, num_classes=10):
        super(myModel, self).__init__(name='my_model')
        # Define your layers here.
        self.dense_1 = layers.Dense(32, activation='relu')
        self.dense_2 = layers.Dense(num_classes, activation='softmax')

    def call(self, inputs):
        # Define your forward pass here
        # using layers you previously defined in '__init__'
        x = self.dense_1(inputs)
        return self.dense_2(x)
```

Subclass the model directly

Constructor

Call method

The “constructor” and “call method”

```
class MyModel(tf.keras.Model):
    def __init__(self, num_classes=10):
        super(myModel, self).__init__(name='my_model')
        # Define your layers here.
        self.dense_1 = layers.Dense(32, activation='relu')
        self.dense_2 = layers.Dense(num_classes, activation='softmax')

    def call(self, inputs):
        # Define your forward pass here
        # using layers you previously defined in '__init__'
        x = self.dense_1(inputs)
        return self.dense_2(x)
```

Subclass the model directly

Constructor

Call method

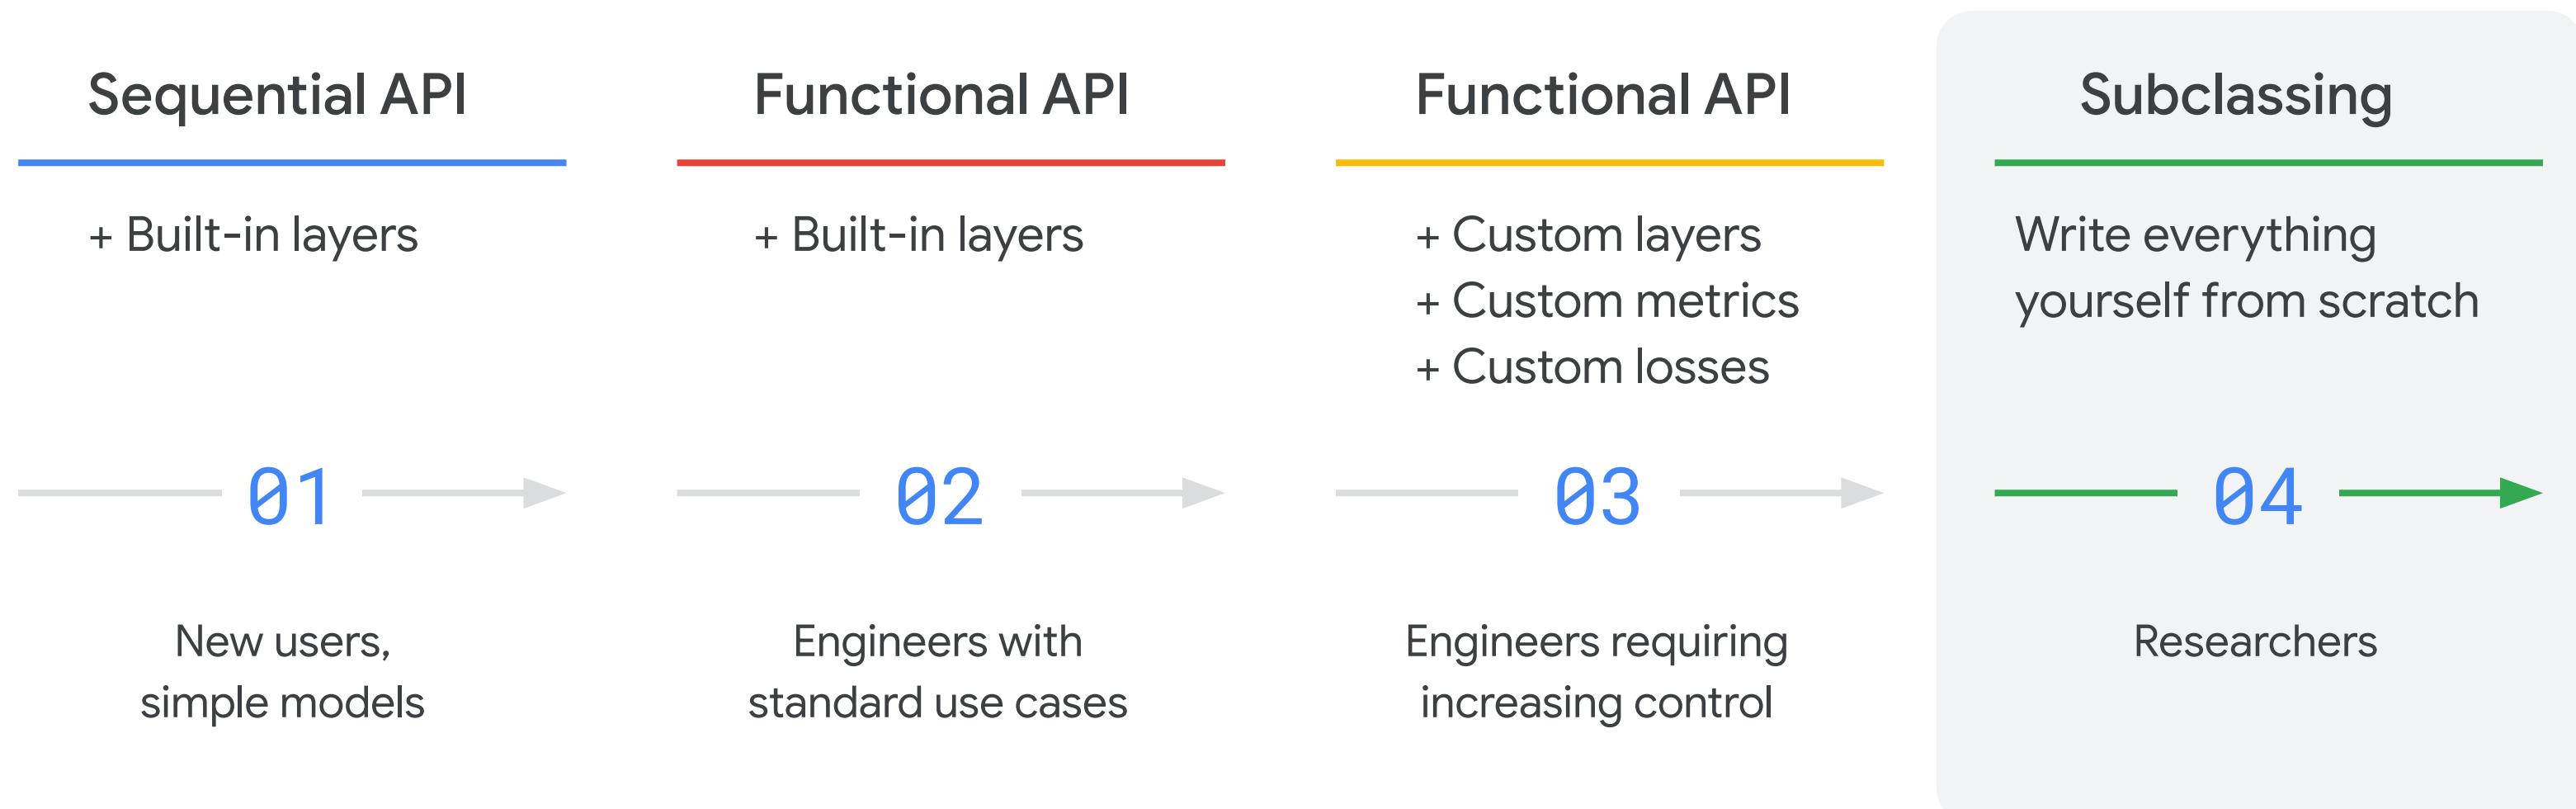
Custom training loop

```
model = MyModel()
with tf.GradientTape() as tape:
    logits = model(images, training=True)
    loss_value = loss(logits, labels)
grads = tape.gradient(loss_value, model.variables)
optimizer.apply_gradients(zip(grads, model.variables))
```

Custom training loop

```
model = MyModel()
with tf.GradientTape() as tape:
    logits = model(images, training=True)
    loss_value = loss(logits, labels)
grads = tape.gradient(loss_value, model.variables)
optimizer.apply_gradients(zip(grads, model.variables))
```

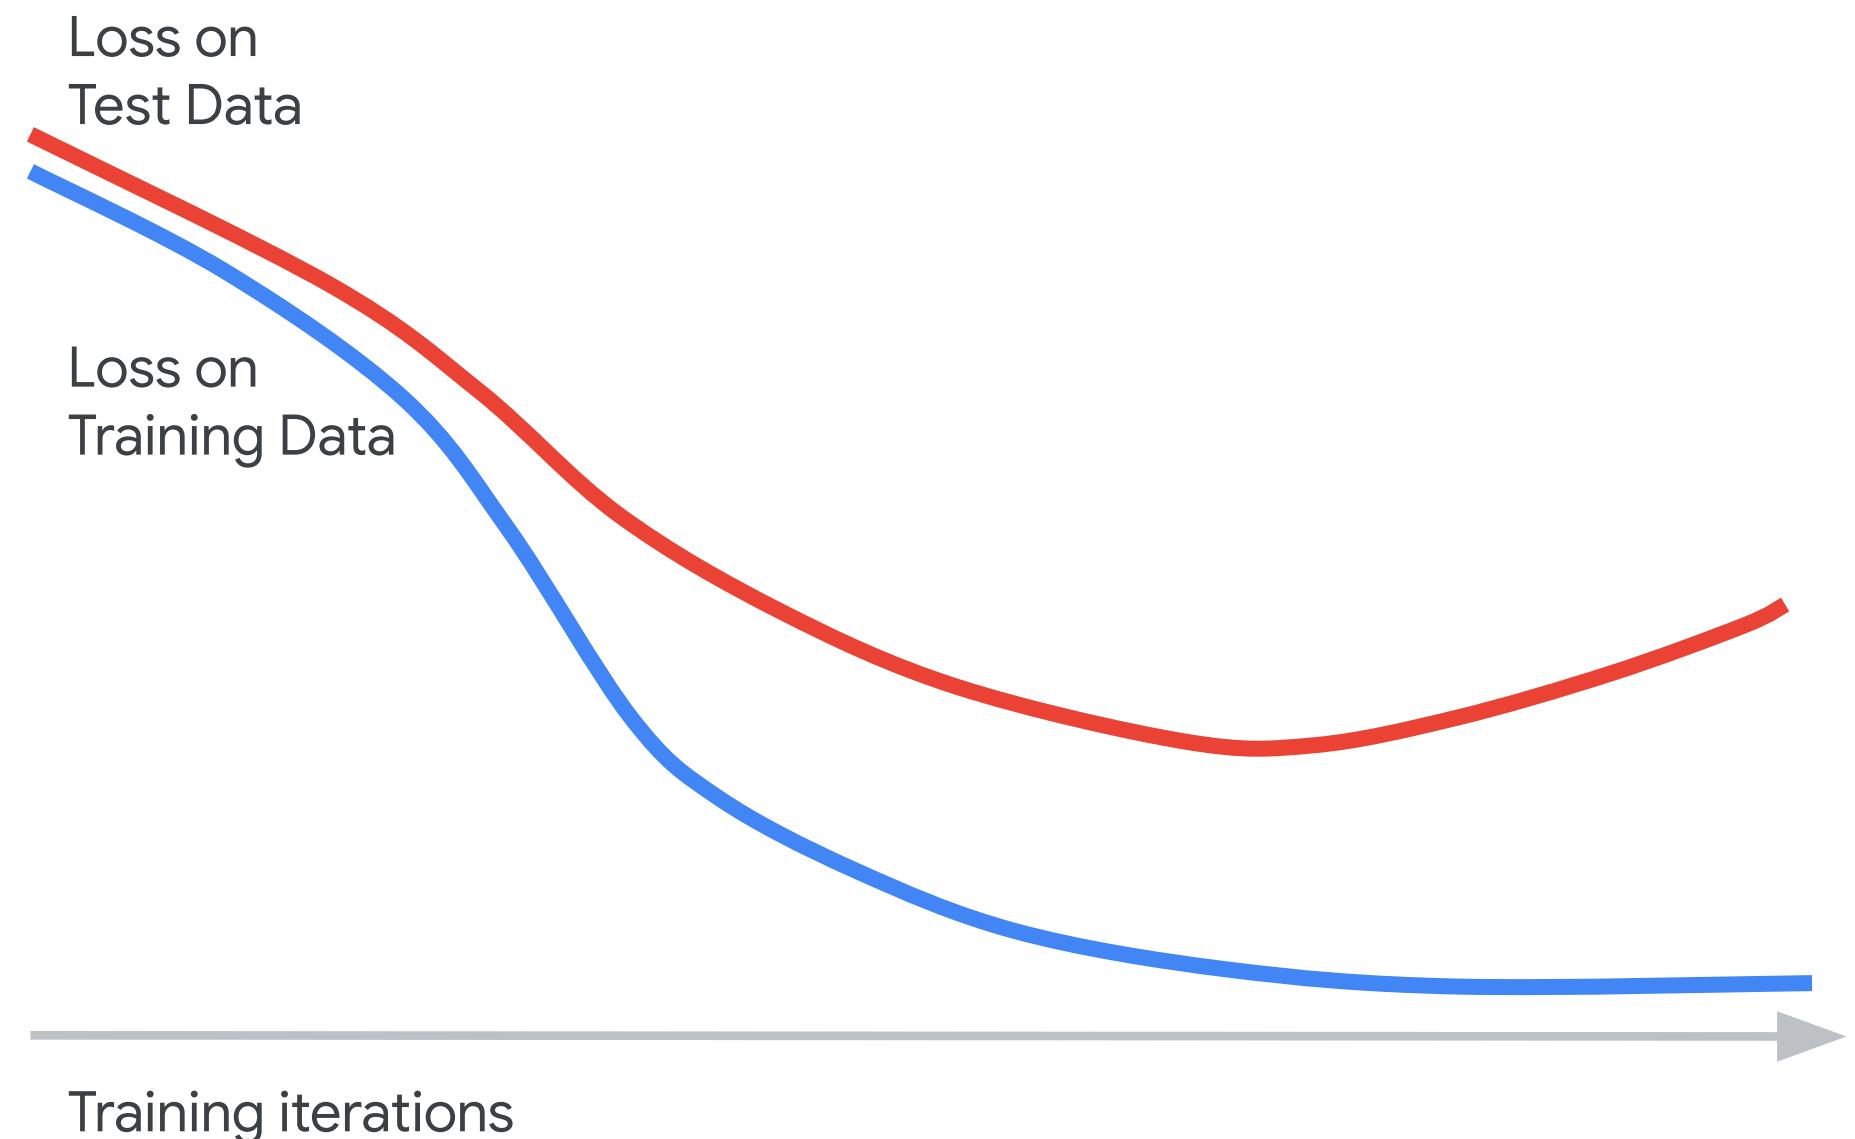
Model building: From simple to arbitrarily flexible



Progressive disclosure of complexity

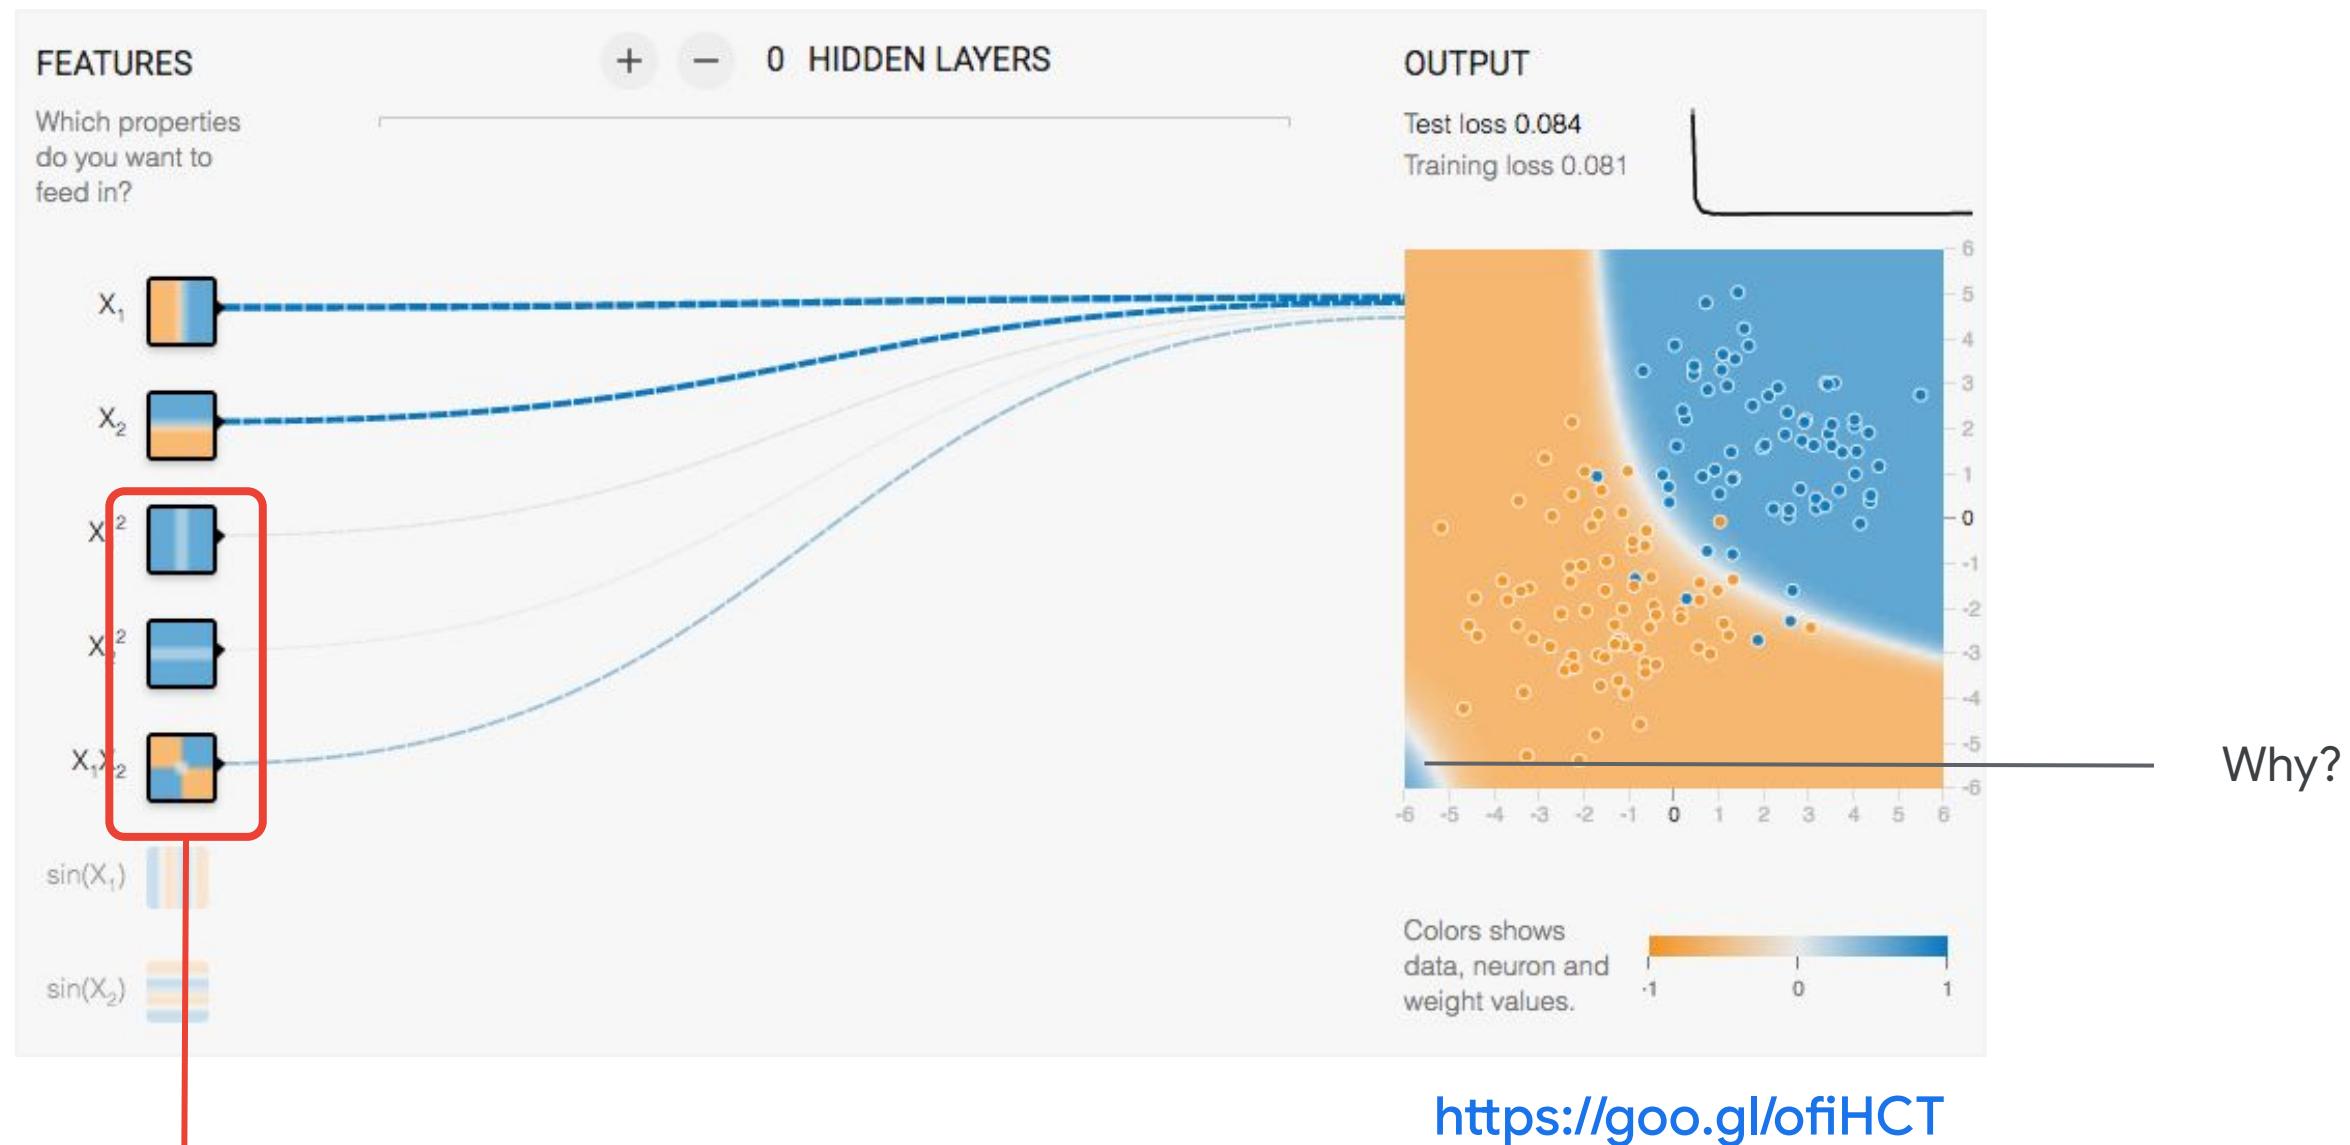
**What is
happening here?**

**How can we
address this?**



Remember the splotch of blue?

Why does it happen?



Is the model behavior surprising? What's the issue?
Try removing cross-product features. Does performance improve?

Occam's razor

When presented with competing hypothetical answers to a problem, one should select the one that makes the fewest assumptions.

The idea is attributed to William of Ockham (c. 1287–1347).



Factor in model complexity when calculating error

Minimize: $\text{loss}(\text{Data}|\text{Model}) + \text{complexity}(\text{Model})$

aim for low
training error

...but balance against
complexity

Optimal model complexity is data-dependent, so requires hyperparameter tuning.

Regularization is a major field of ML research

Early Stopping

Parameter Norm Penalties

L1 regularization

L2 regularization

Max-norm regularization

Dataset Augmentation

Noise Robustness

Sparse Representations

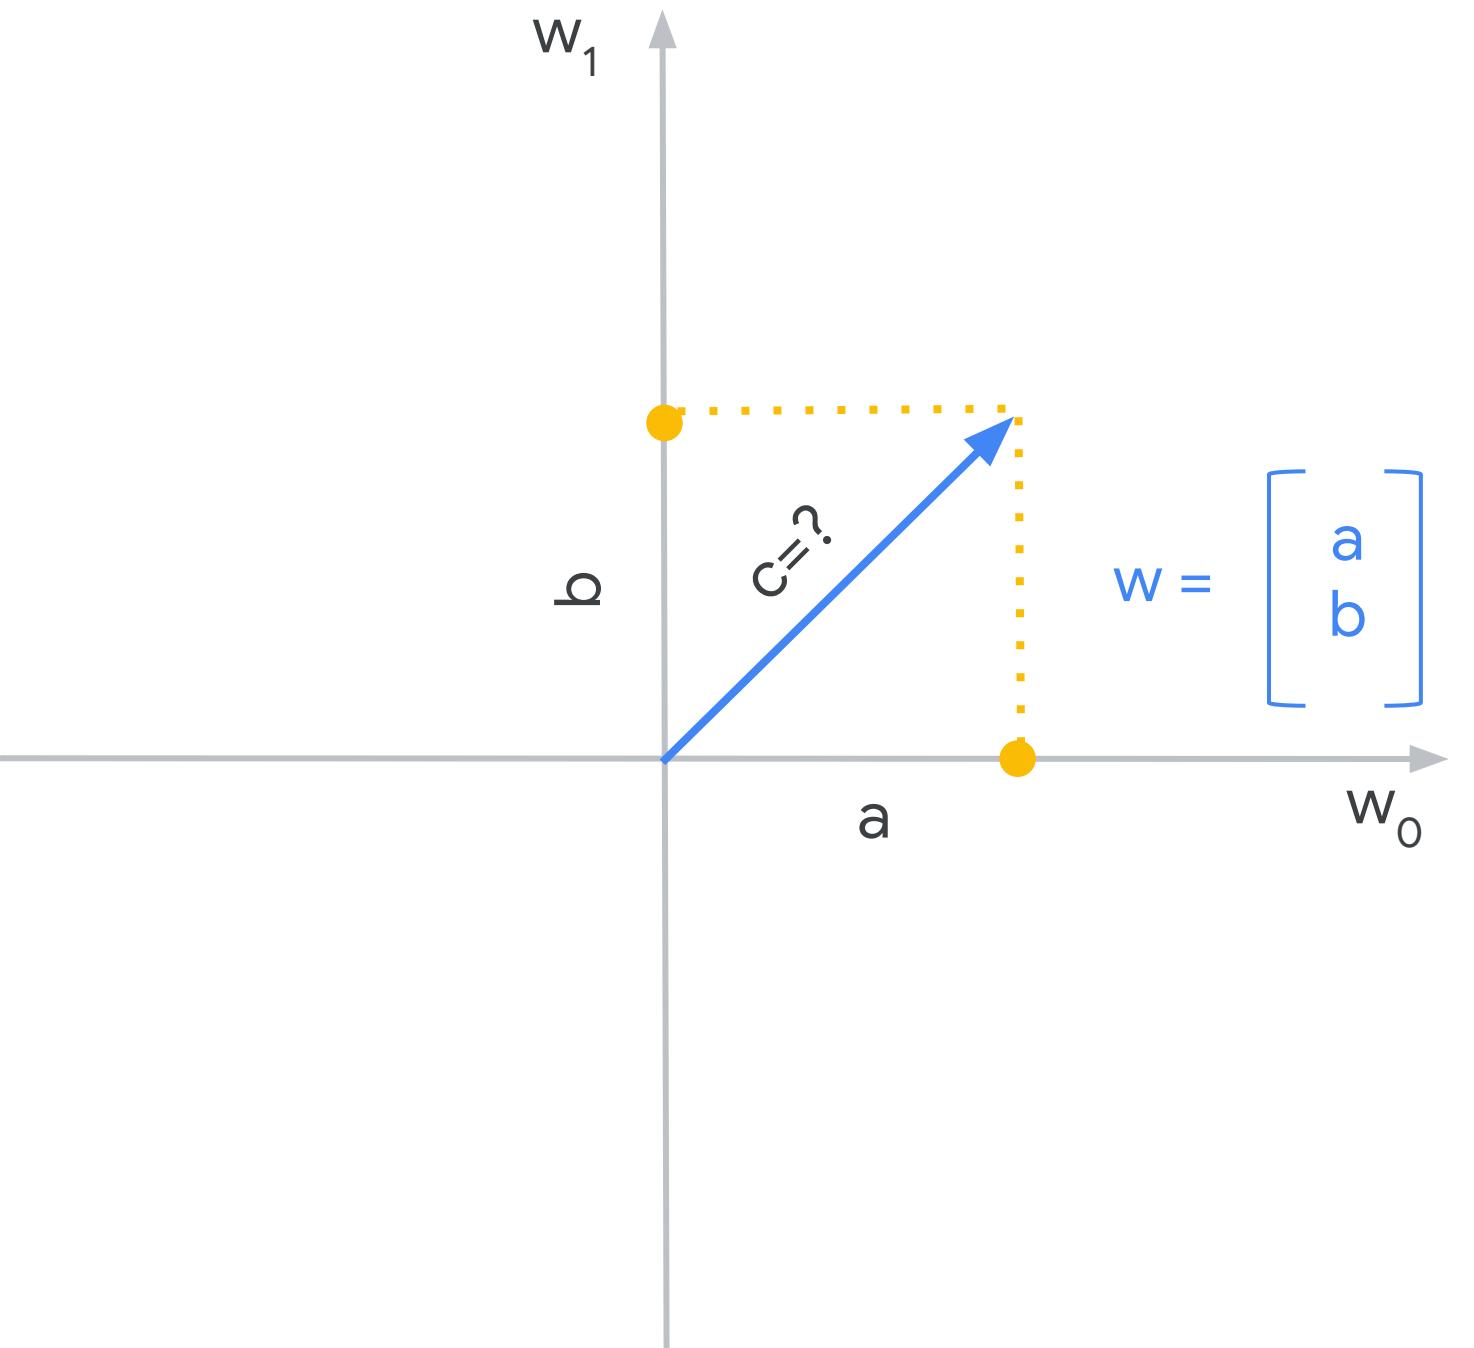
· · ·

We will look into
these methods.

How can we measure
model complexity?

L2 vs. L1 Norm

$$w = [w_0, w_1, \dots, w_n]^T$$

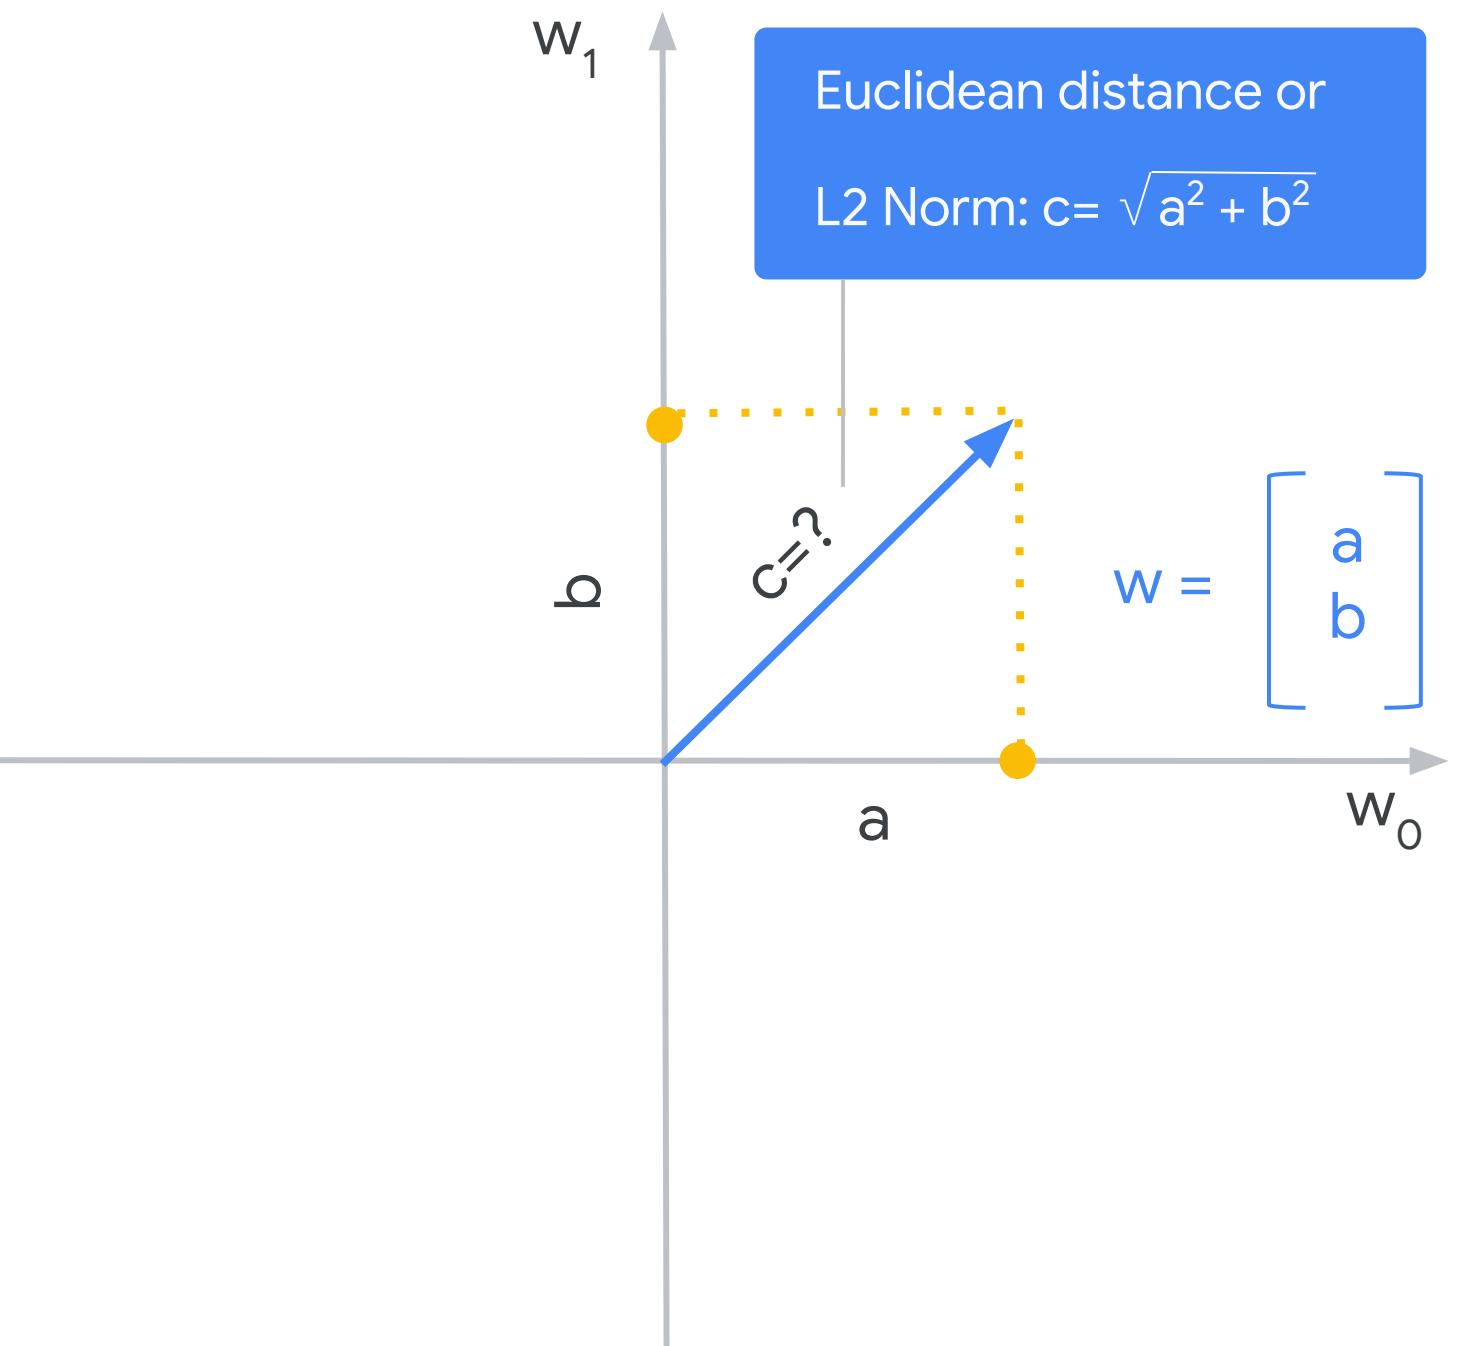


L2 vs. L1 Norm

$$w = [w_0, w_1, \dots, w_n]^T$$

$$\|w\|_2 = [w_0^2 + w_1^2 + \dots + w_n^2]^{1/2}$$

$$\|w\|_1 = (|w_0| + |w_1| + \dots + |w_n|)$$



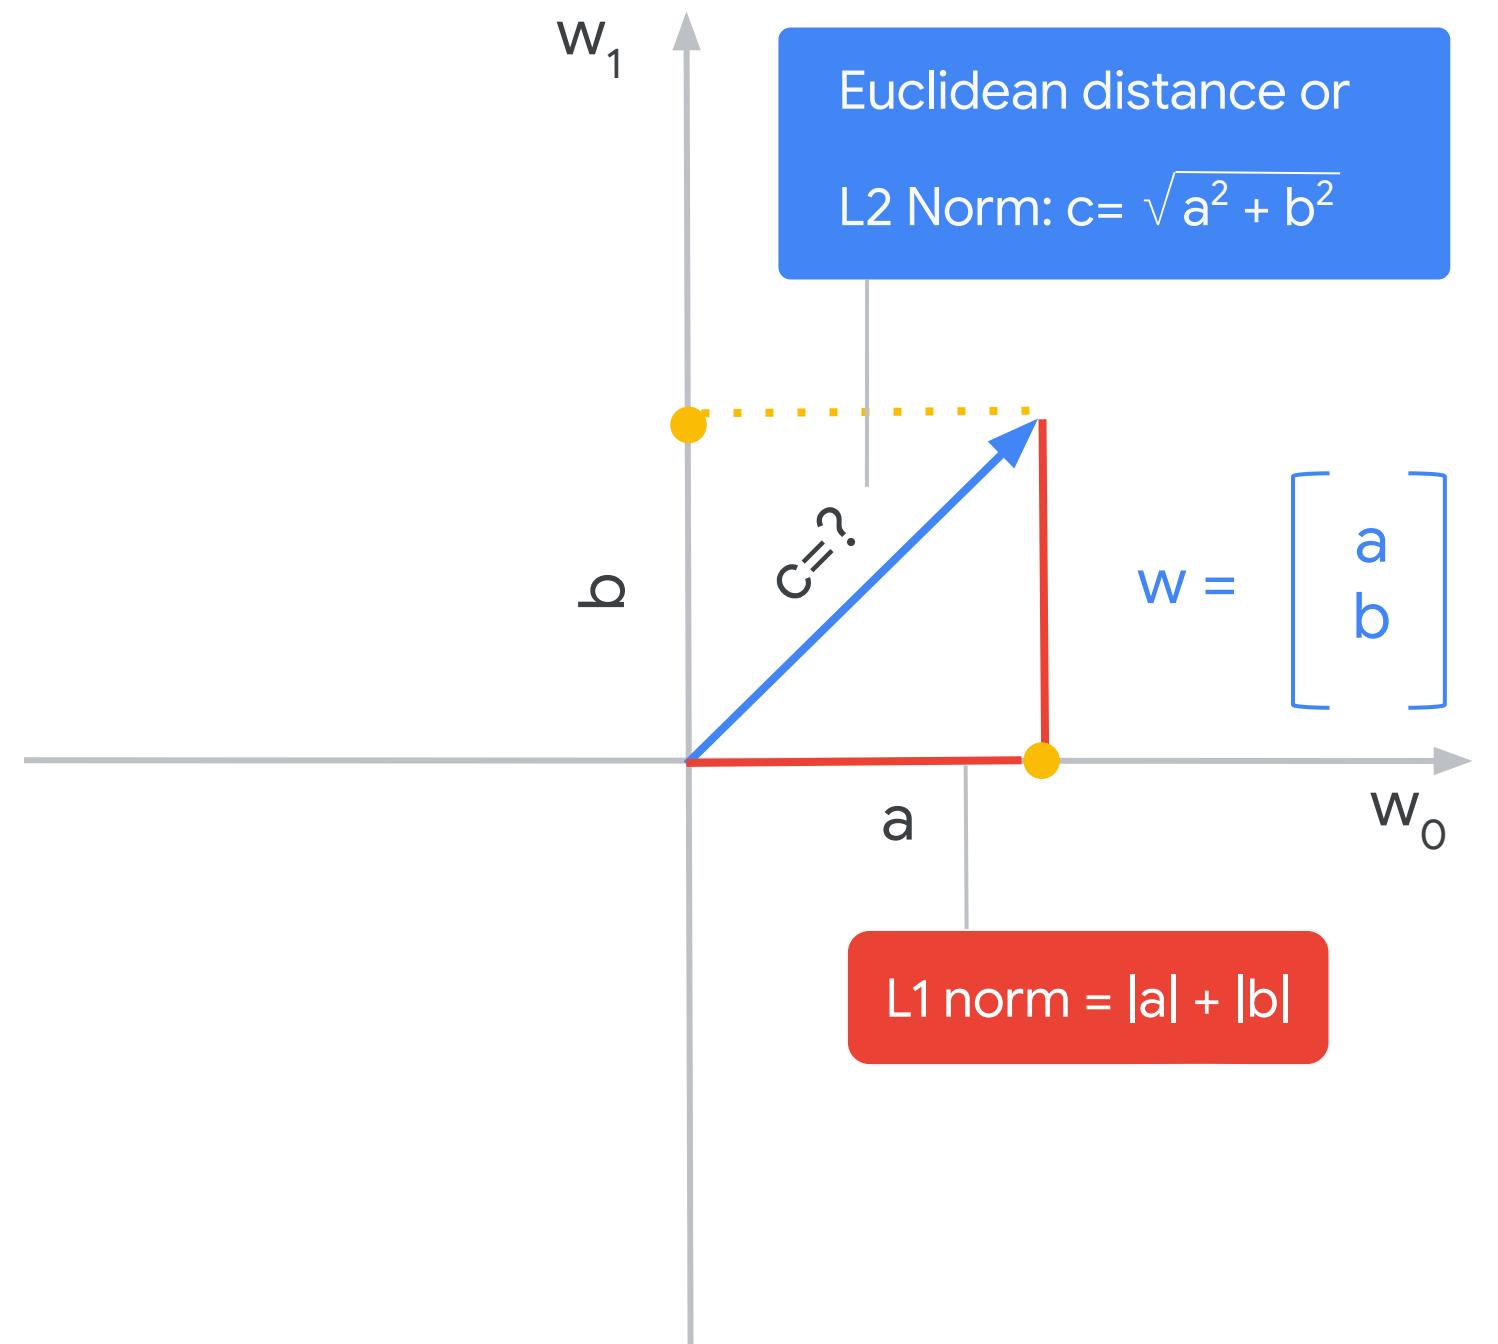
L2 vs. L1 Norm

$$\mathbf{w} = [w_0, w_1, \dots, w_n]^T$$

$$\|\mathbf{w}\|_2 = [w_0^2 + w_1^2 + \dots + w_n^2]^{1/2}$$

L2 norm

$$\|\mathbf{w}\|_1 = (|w_0| + |w_1| + \dots + |w_n|)$$



L2 vs. L1 Norm

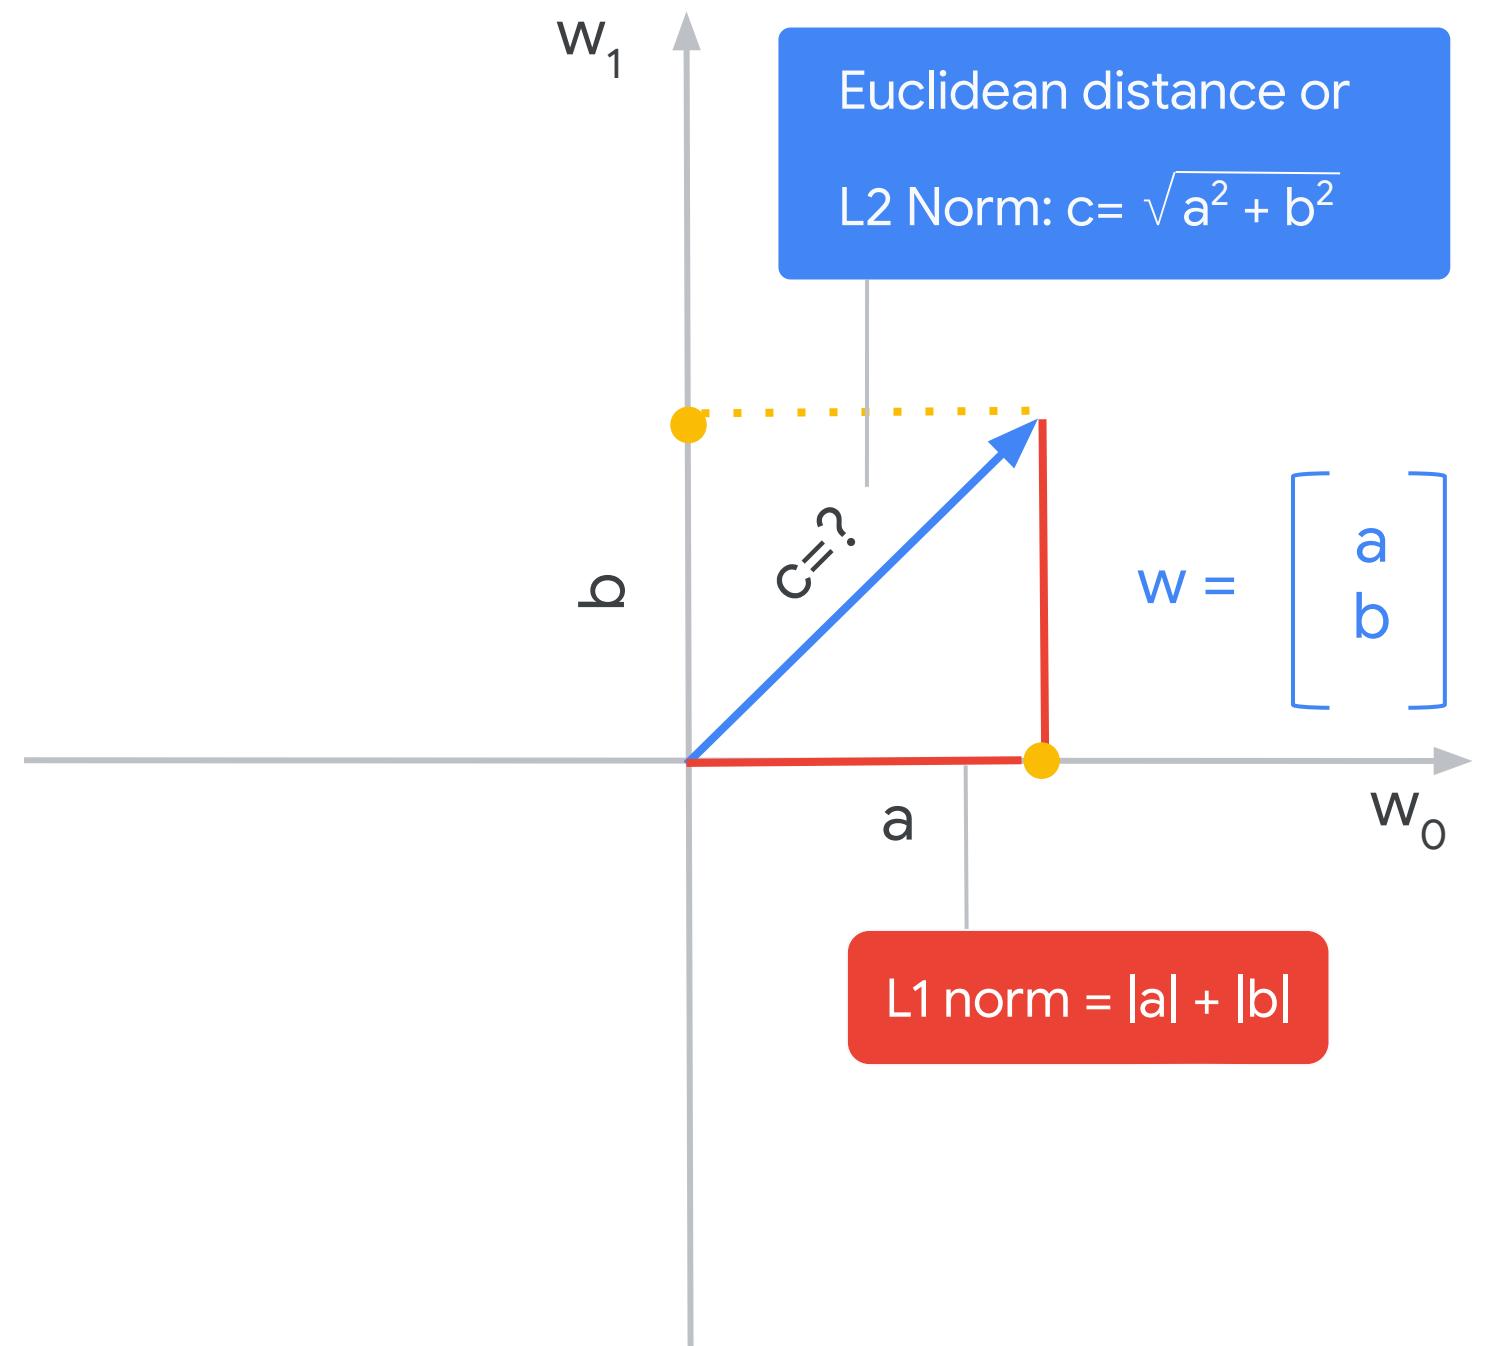
$$w = [w_0, w_1, \dots, w_n]^T$$

$$\|w\|_2 = [w_0^2 + w_1^2 + \dots + w_n^2]^{1/2}$$

L2 norm

$$\|w\|_1 = (|w_0| + |w_1| + \dots + |w_n|)$$

L1 norm



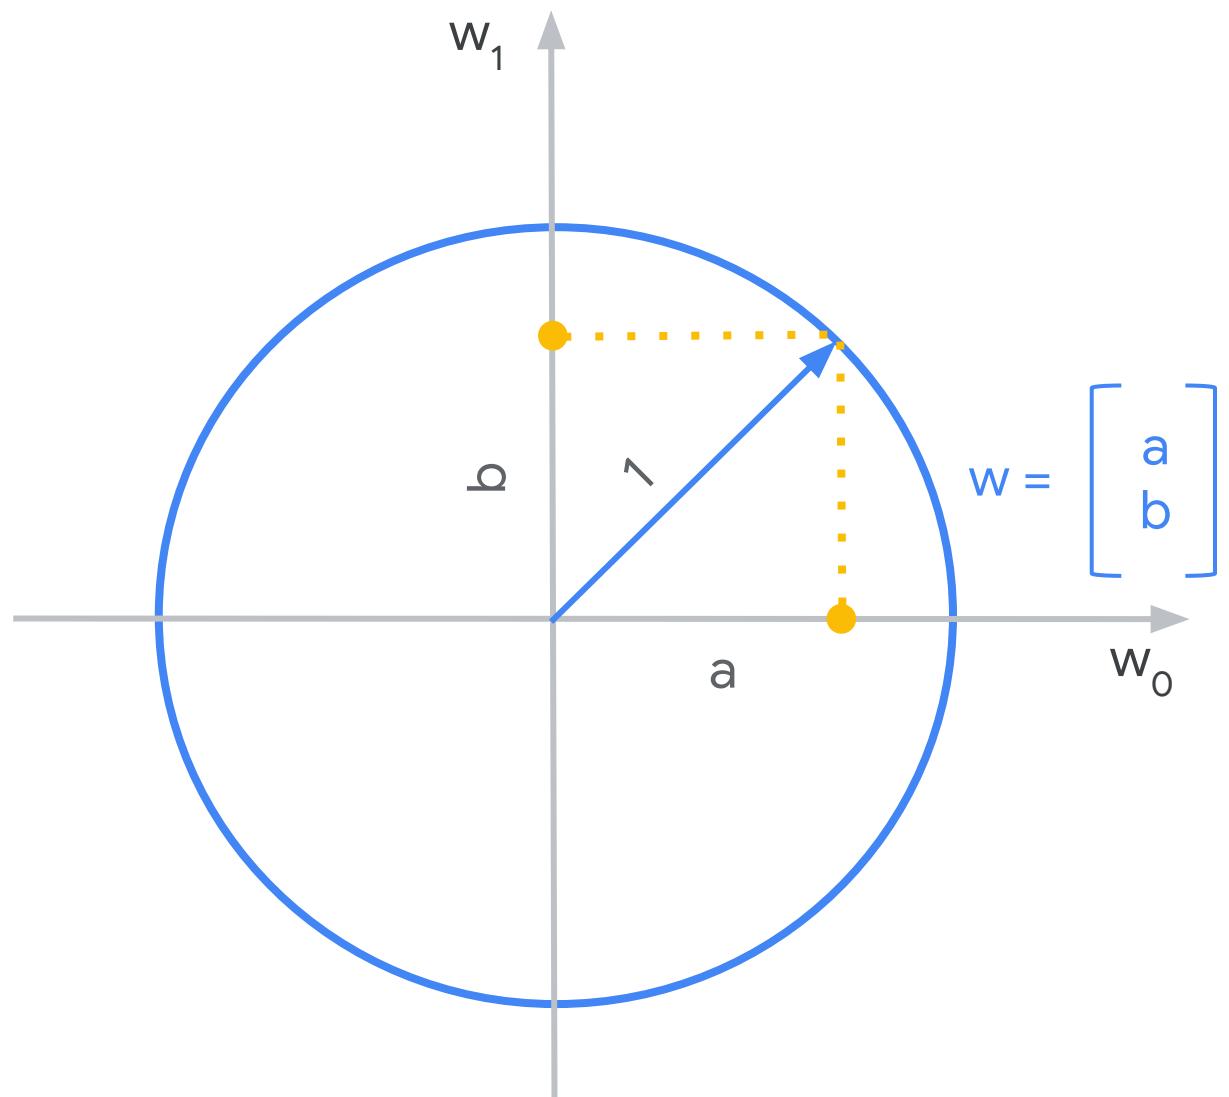
L2 vs. L1 Norm

$$\mathbf{w} = [w_0, w_1, \dots, w_n]^T$$

$$\|\mathbf{w}\|_2 = [w_0^2 + w_1^2 + \dots + w_n^2]^{1/2}$$

L2 norm

$$\|\mathbf{w}\|_1 = (|w_0| + |w_1| + \dots + |w_n|)$$



L2 vs. L1 Norm

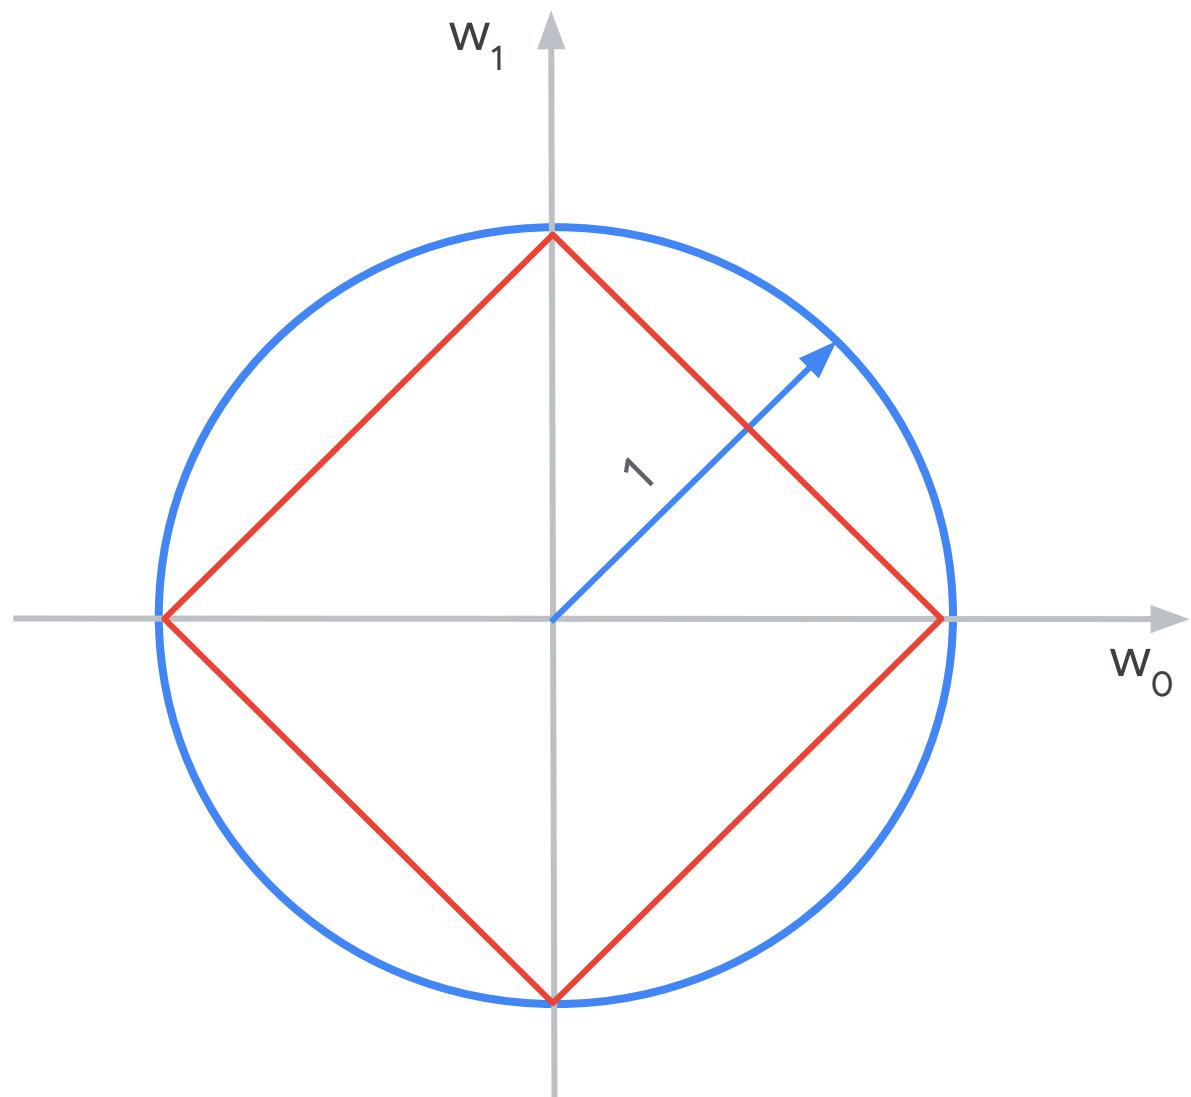
$$\mathbf{w} = [w_0, w_1, \dots, w_n]^T$$

$$\|\mathbf{w}\|_2 = [w_0^2 + w_1^2 + \dots + w_n^2]^{1/2}$$

L2 norm

$$\|\mathbf{w}\|_1 = (|w_0| + |w_1| + \dots + |w_n|)$$

L1 norm



In **L2 regularization**,
complexity of model is
defined by the **L2 norm**
of the weight vector

Aim for low
training error

...but balance against
complexity

$$L(w, D) + \lambda \|w\|_2$$

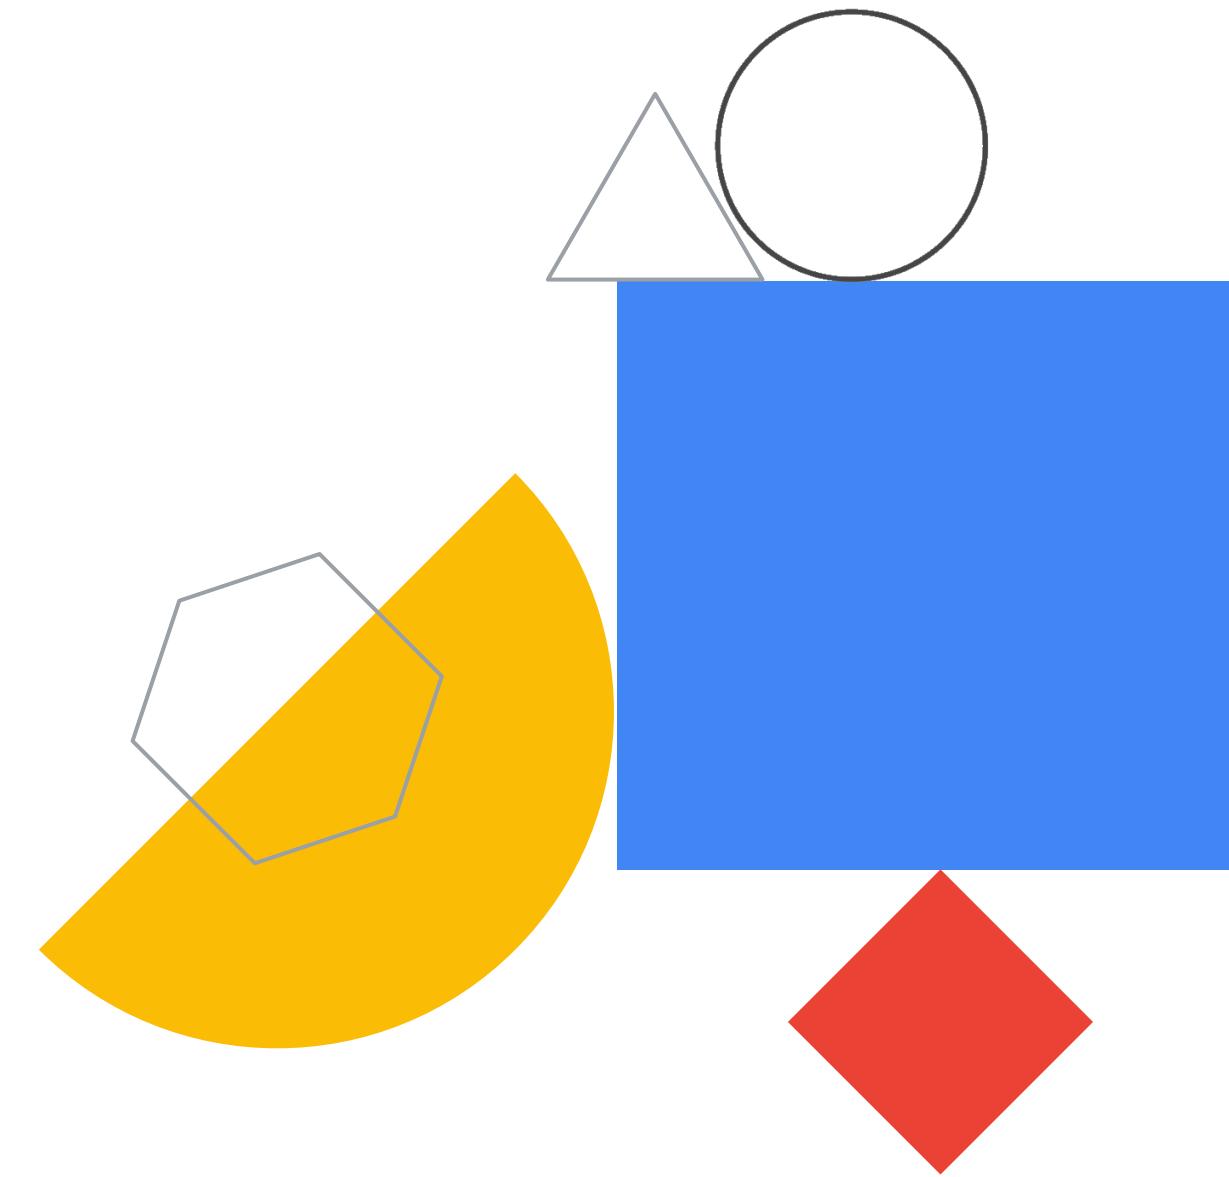
Lambda controls how
these are balanced

In L1 regularization,
complexity of model is
defined by the L1 norm
of the weight vector

$$L(w, D) + \lambda \|w\|_1$$

L1 regularization can be used as a
feature selection mechanism.

Training at Scale with Vertex AI



In this module, you learn to ...

01

Use TensorFlow to create a training job

02

Package up a TensorFlow model as a Python package

03

Configure, start, and monitor a Vertex AI training job



We will use distributed TensorFlow on Vertex AI

tf.keras	High-level APIs for distributed training
tf.losses, tf.metrics	Components useful when building custom NN models
Core TensorFlow (Python)	Python API gives you full control
Core TensorFlow (C++)	C++ API is quite low level
CPU GPU TPU Android	TF runs on different hardware



Vertex AI

Run TF at scale with Vertex AI.

We will use
distributed
TensorFlow on
Vertex AI

- 
- 1 Use TensorFlow to create your Keras model
 - 2 Package your trainer application
 - 3 Configure and start a Vertex AI training job

Create task.py to parse command-line parameters and send along to train_and_evaluate

model.py

```
def train_and_evaluate(hparams):
    # [...] Obtain param values from hparam

    model = build_dnn_model(nbuckets, nnsize, lr)
    trainds = create_train_dataset(train_data_path, batch_size)
    evalds = create_eval_dataset(eval_data_path, batch_size)

    # [...] create the callbacks
    steps_per_epoch = num_examples_to_train_on // (batch_size * num_evals)

    history = model.fit(
        trainds,
        validation_data=evalds,
        epochs=num_evals,
        steps_per_epoch=max(1, steps_per_epoch),
        verbose=2,  # 0=silent, 1=progress bar, 2=one line per epoch
        callbacks=[checkpoint_cb, tensorboard_cb]
    )
    tf.saved_model.save(model, model_export_path)
    return history
```

task.py

```
parser.add_argument(
    '--train_data_paths', required=True)
parser.add_argument(
    '--train_steps', ...)
```

Package up TensorFlow model as Python package

Python modules need to contain an `__init__.py` in every folder.

`taxifare/trainer/__init__.py`

`taxifare/trainer/task.py`

`taxifare/trainer/model.py`

Test your code locally first

```
EVAL_DATA_PATH=./taxifare/tests/data/taxi-valid*
TRAIN_DATA_PATH=./taxifare/tests/data/taxi-train*
OUTPUT_DIR=./taxifare-model

test ${OUTPUT_DIR} && rm -rf ${OUTPUT_DIR}
export PYTHONPATH=${PYTHONPATH}:$(PWD)/taxifare

python3 -m trainer.task \
--eval_data_path $EVAL_DATA_PATH \
--output_dir $OUTPUT_DIR \
--train_data_path $TRAIN_DATA_PATH \
--batch_size 5 \
--num_examples_to_train_on 100 \
--num_evals 1 \
--nbuckets 10 \
--lr 0.001 \
--nnsize "32 8"
```

To make our code compatible with Vertex AI

01

Upload data to Google
Cloud Storage

02

Move code into a
trainer Python package

03

Submit training job with
gcloud to train on
Vertex AI

02

Move code to trainer Python package

We package our code as a
source distribution using
setup.py and setuptools.

```
taxifare/trainer/__init__.py  
taxifare/trainer/task.py  
taxifare/trainer/model.py  
taxifare/setup.py
```

02

Move code to trainer Python package

We package our code as a
source distribution using
setup.py and setuptools.

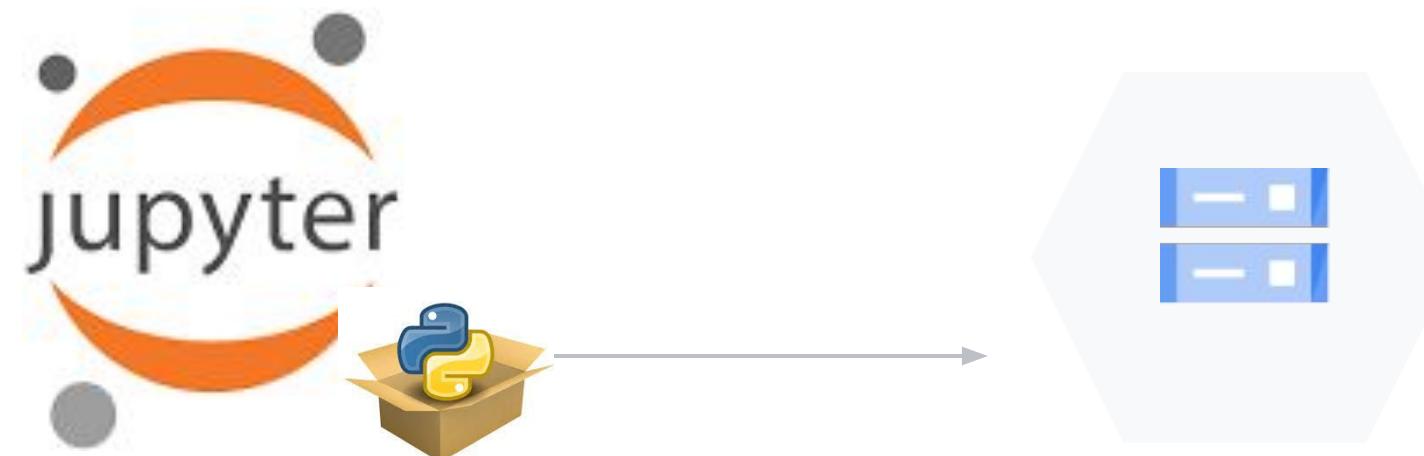
```
taxifare/trainer/__init__.py  
taxifare/trainer/task.py  
taxifare/trainer/model.py  
taxifare/setup.py
```

```
python ./setup.py sdist --formats=gztar
```

02

Copy that Python package to your GCS bucket

```
gsutil cp taxifare/dist/taxifare_trainer-0.1.tar.gz  
gs://{$BUCKET}/taxifare/
```



03

Submitting a job with gcloud ai custom-jobs create

using a pre-built container....

```
gcloud ai custom-jobs create \  
  --region=$REGION \  
  --display-name=$JOB_NAME \  
  --python-package-uris=$PYTHON_PACKAGE_URIS \  
  --worker-pool-spec=machine-type=$MACHINE_TYPE, \  
    replica-count=$REPLICA_COUNT, \  
    executor-image-uri=$PYTHON_PACKAGE_EXECUTOR_IMAGE_URI, \  
    python-module=$PYTHON_MODULE
```

using a custom container....

```
gcloud ai custom-jobs create \  
  --region=$LOCATION \  
  --display-name=$JOB_NAME \  
  --worker-pool-spec=machine-type=$MACHINE_TYPE, \  
    replica-count=$REPLICA_COUNT, \  
    container-image-uri=$CUSTOM_CONTAINER_IMAGE_URI
```

03

For distributed training, specify multiple worker-pool-spec

using a pre-built container....

```
gcloud ai custom-jobs create \
--region=$REGION \
--display-name=$JOB_NAME \
--python-package-uris=$PYTHON_PACKAGE_URIS \
--worker-pool-spec(machine-type=$MACHINE_TYPE, \
replica-count=$REPLICA_COUNT, \
executor-image-uri=$PYTHON_PACKAGE_EXECUTOR_IMAGE_URI, \
python-module=$PYTHON_MODULE \
--worker-pool-spec(machine-type=$SECOND_MACHINE_TYPE, \
replica-count=$SECOND_REPLICA_COUNT, \
executor-image-uri=$SECOND_PYTHON_PACKAGE_EXECUTOR_IMAGE_URI, \
python-module=$SECOND_POOL_PYTHON_MODULE
```

03

For our taxifare
model, we'll specify
the following
components...

using a pre-built container....

```
gcloud ai custom-jobs create \  
  --region=$REGION \  
  --display-name=$JOB_NAME \  
  --python-package-uris=$PYTHON_PACKAGE_URIS \  
  --worker-pool-spec=$WORKER_POOL_SPEC \  
  --args="$ARGS"
```

```
TIMESTAMP=$(date -u +%Y%m%d_%H%M%S)  
JOB_NAME=taxifare_$TIMESTAMP
```

03

For our taxifare
model, we'll specify
the following
components...

using a pre-built container....

```
gcloud ai custom-jobs create \  
  --region=$REGION \  
  --display-name=$JOB_NAME \  
  --python-package-uris=$PYTHON_PACKAGE_URIS \  
  --worker-pool-spec=$WORKER_POOL_SPEC \  
  --args="$ARGS"
```

PYTHON_PACKAGE_URIS=gs://\$BUCKET/taxifare/taxifar
e_trainer-0.1.tar.gz

03

For our taxifare model, we'll specify the following components...

```
gcloud ai custom-jobs create \  
  --region=$REGION \  
  --display-name=$JOB_NAME \  
  --python-package-uris=$PYTHON_PACKAGE_URIS \  
  --worker-pool-spec=$WORKER_POOL_SPEC \  
  --args="--eval_data_path=$EVAL_DATA_PATH, \  
         --output_dir=$OUTDIR, \  
         --train_data_path=$TRAIN_DATA_PATH, \  
         --batch_size=$BATCH_SIZE, \  
         --num_examples_to_train_on=$NUM_EXAMPLES_TO_TRAIN_ON, \  
         --num_evals=$NUM_EVALS, \  
         --nbuckets=$NBUCKETS, \  
         --lr=$LR, \  
         --nnsize=$NNSIZE"
```

using a pre-built container....

03

You can also use config.yaml to control training parameters

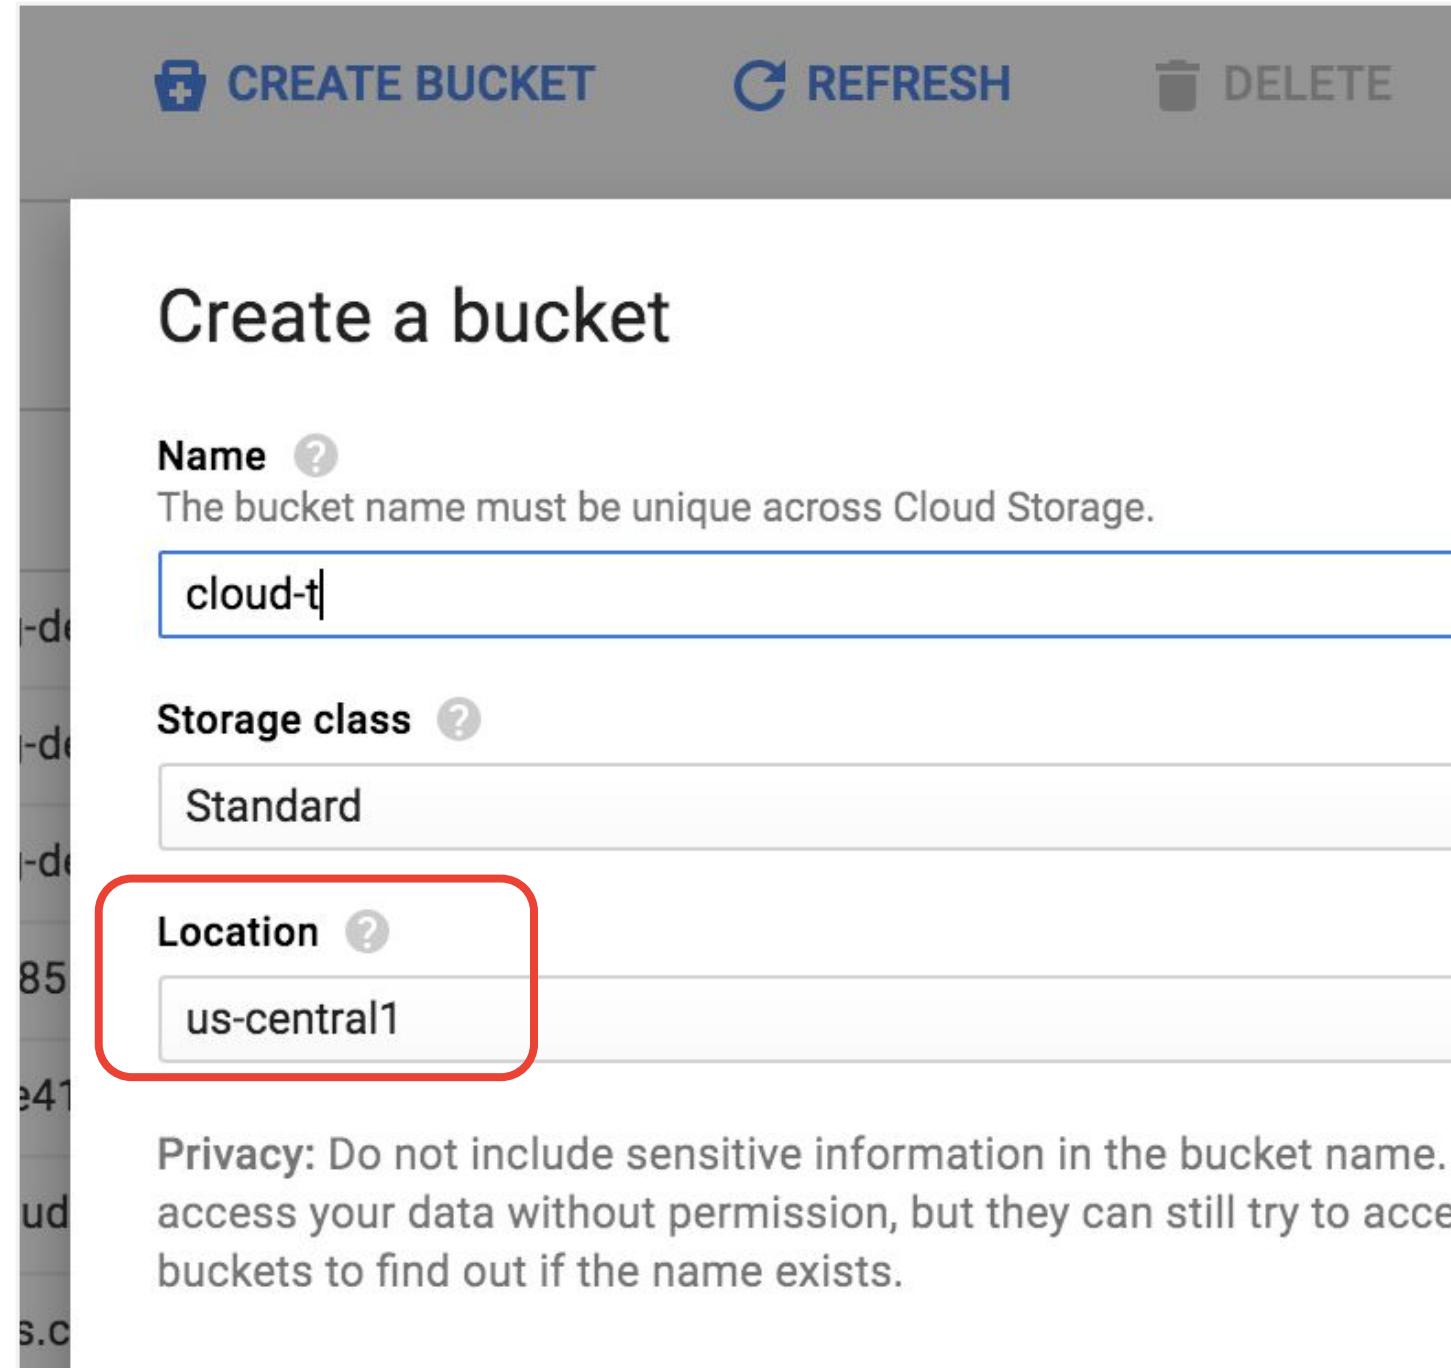
using a config.yaml file....

```
gcloud ai custom-jobs create \  
  --region=$LOCATION \  
  --display-name=$JOB_NAME \  
  --config=config.yaml
```

config.yaml

```
workerPoolSpecs:  
  machineSpec:  
    machineType: n1-highmem-2  
  replicaCount: 1  
  containerSpec:  
    imageUri:  
      gcr.io/ucaip-test/ucaip-training-test  
    args:  
      - port=8500  
    command:  
      - start
```

Tip:
**Use single-region
bucket for ML**



Monitor training jobs with GCP console

You can also view CPU and Memory utilization charts for this training job with StackDriver Monitoring.

The screenshot shows the Google Cloud Platform Vertex AI Training interface. On the left, there is a sidebar with the following navigation items:

- Vertex AI
- Dashboard
- Datasets
- Features
- Labeling tasks
- Notebooks
- Pipelines
- Training** (selected)
- Experiments
- Models
- Endpoints
- Batch predictions
- Metadata

The main content area has the following tabs at the top:

- TRAINING PIPELINES
- CUSTOM JOBS** (highlighted with a red box)
- HYPERPARAMETER TUNING JOBS

Below the tabs, there is a descriptive text about Custom jobs:

Custom jobs specify how Vertex AI runs your custom training code, including worker pools, machine types, and settings related to your Python training application and custom container. Custom jobs are only used by custom-trained models and not AutoML models. [Learn More](#)

A dropdown menu for the Region is set to "us-central1 (Iowa)".

A "Filter" input field is present with the placeholder "Enter a property name".

A table lists six training jobs:

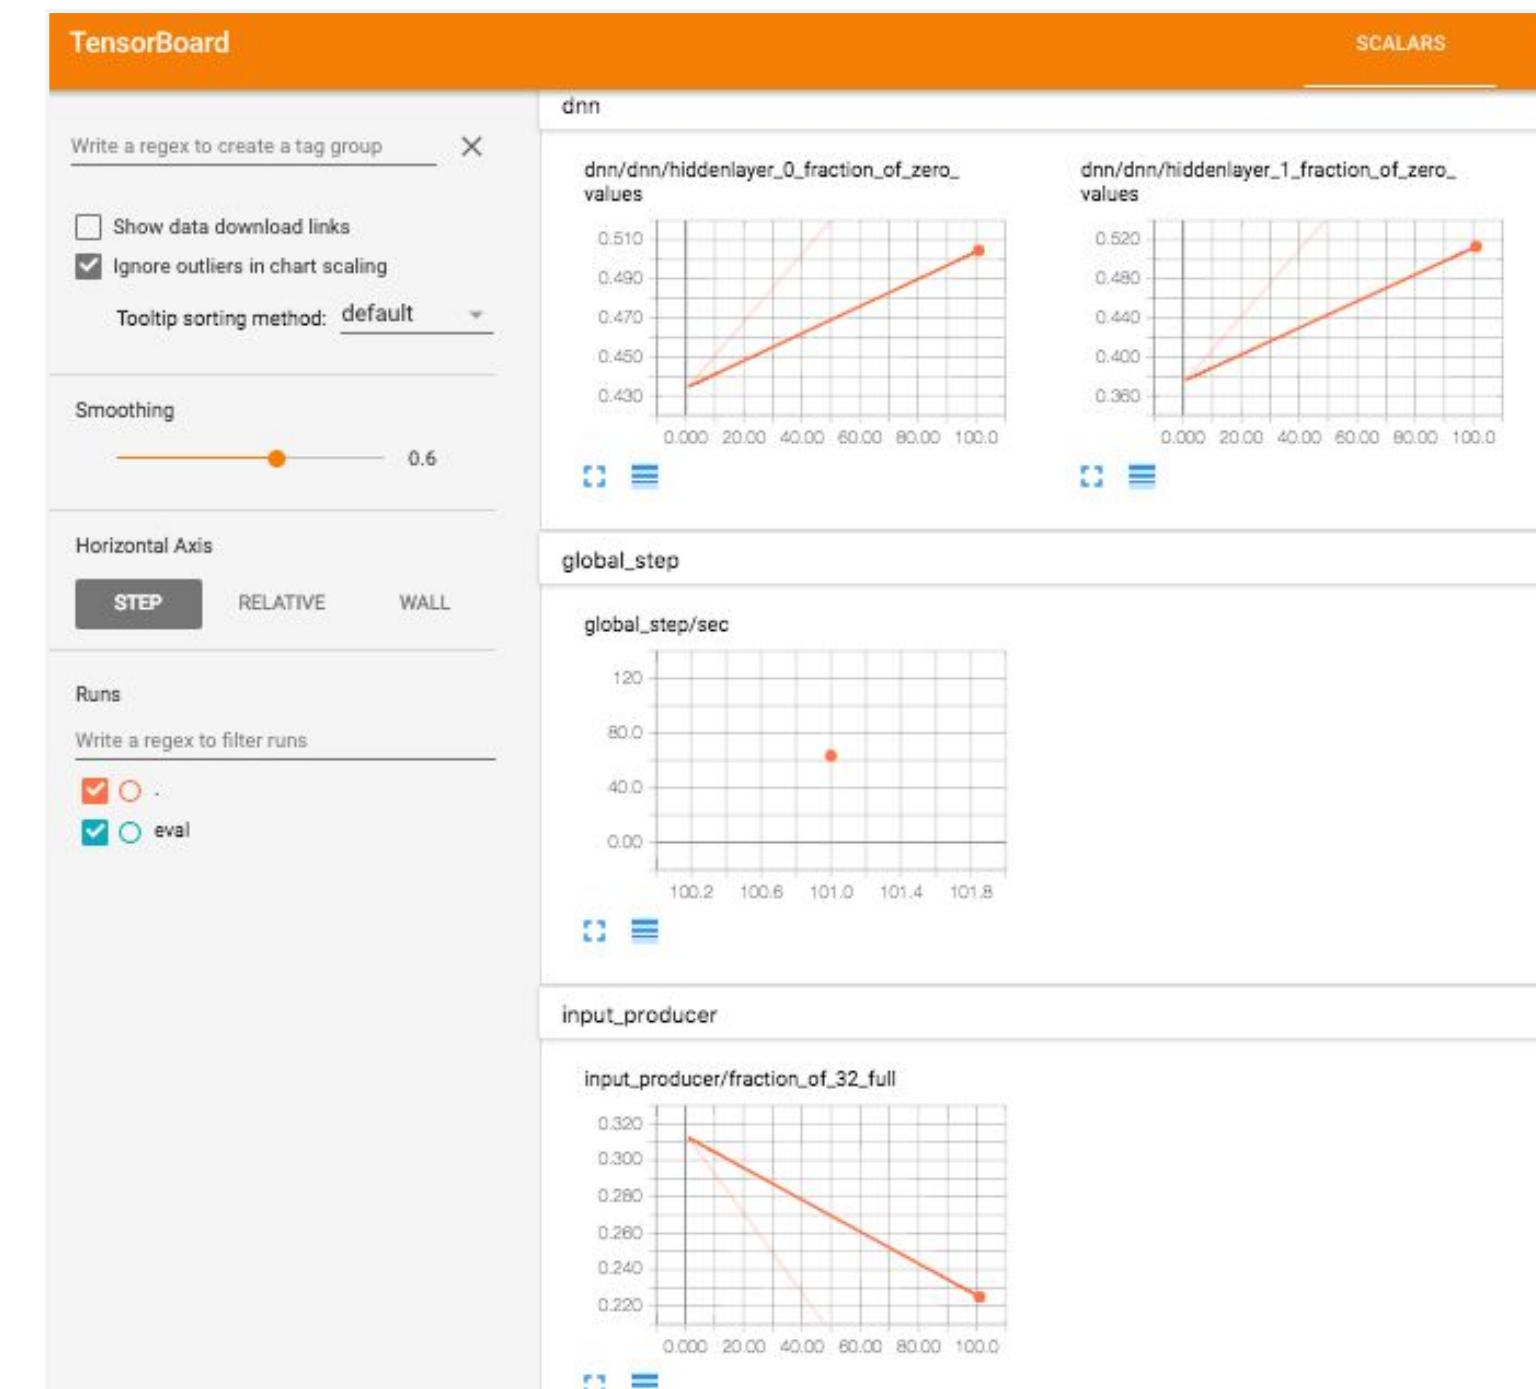
Name	ID	Job type	Model
mnist_cnn_20210617_152600	6634123308512051200	Custom job	-
babyweight_tuned_20210616_213056	5391974236287926272	Custom job	-
babyweight_20210616_212749	3993606551989387264	Custom job	-
taxifare_container_20210614_174703	4832014954494164992	Custom job	-
taxifare_20210614_162401	1869772299591221248	Custom job	-

Monitor training jobs with TensorBoard

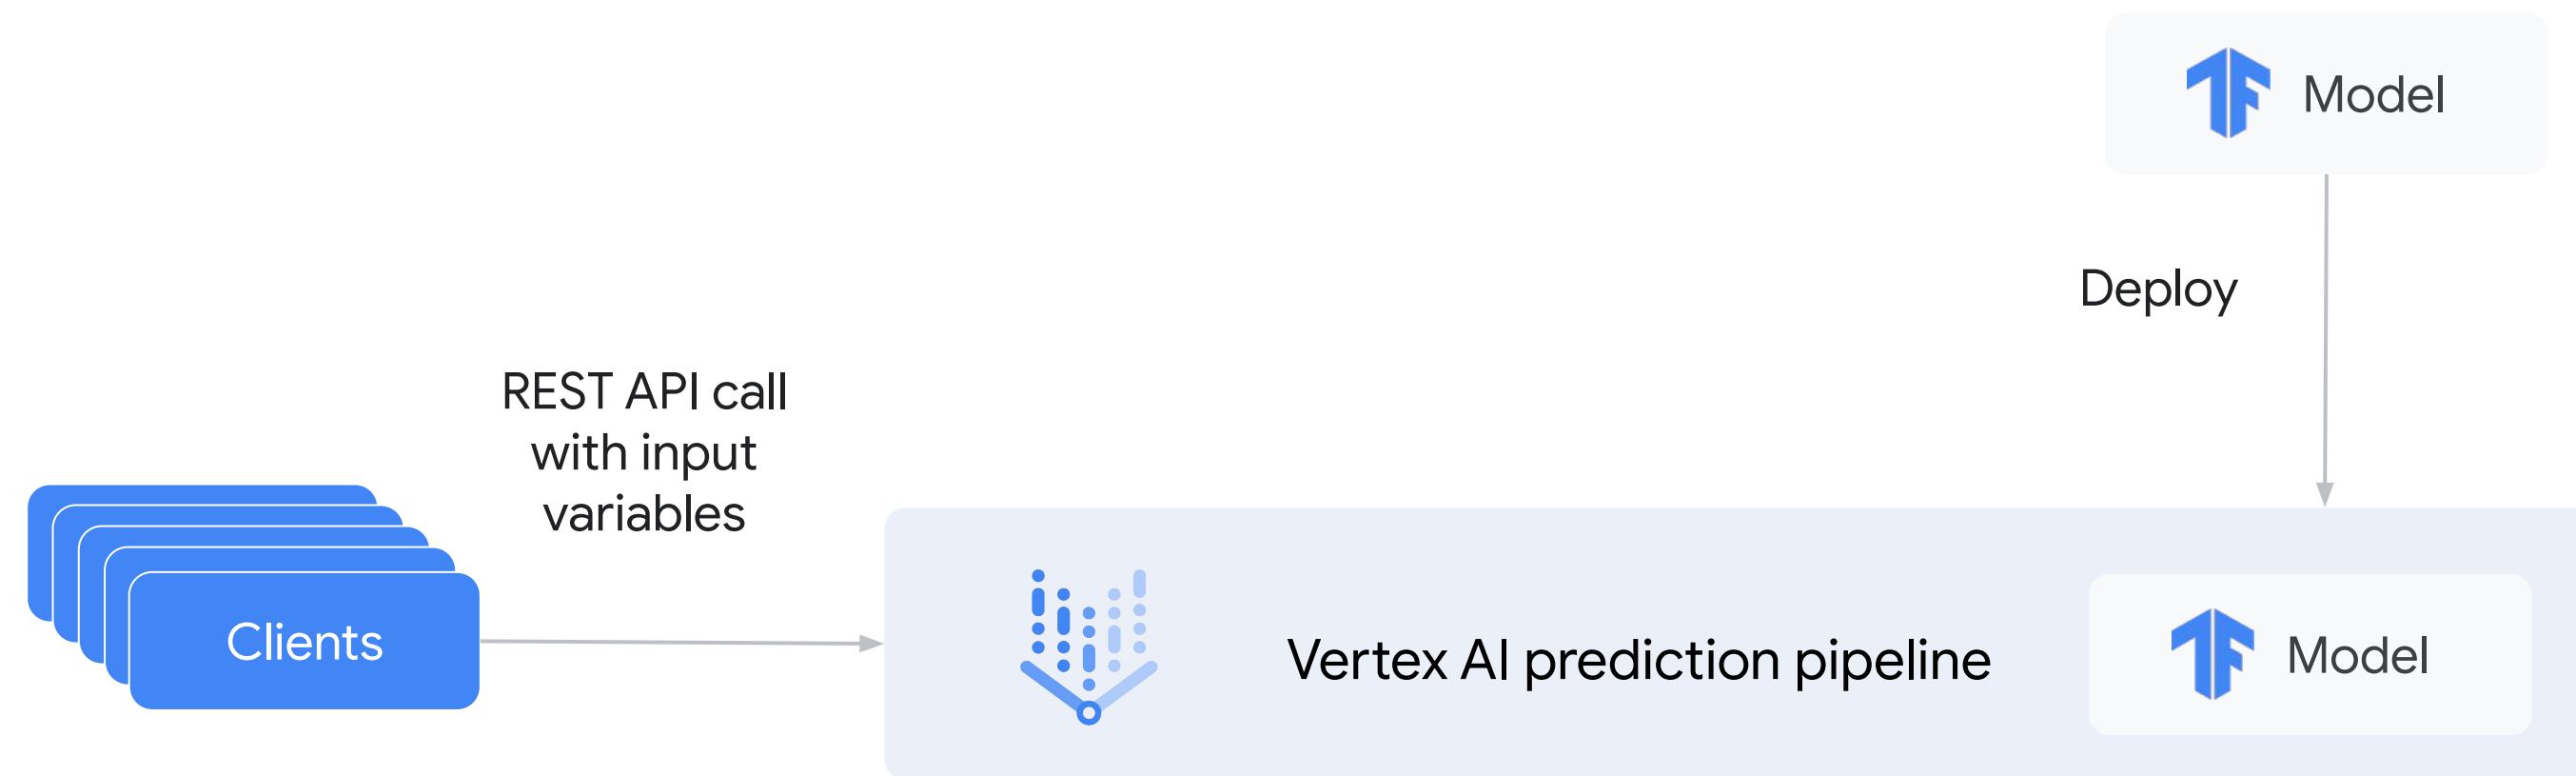
Configure your trainer to save summary data that you can examine and visualize using TensorBoard.

```
tensorboard_cb = callbacks.TensorBoard(tensorboard_path)

history = model.fit(
    trainds,
    validation_data=evalds,
    epochs=num_evals,
    steps_per_epoch=max(1, steps_per_epoch),
    verbose=2, # 0=silent, 1=progress bar, 2=one line per epoch
    callbacks=[checkpoint_cb, tensorboard_cb]
)
```



Vertex AI Prediction service makes deploying models and scaling the infrastructure easy



We'll see how to deploy with models and endpoints and make predictions in the later labs